

Advertisement

STARTUPS



A former OpenAI
engineer describes what
it's really like to work
there

OpenAI

Julie Bort — 1:58 PM PDT · July 15, 2025

Three weeks ago, an engineer named Calvin French-Owen, who worked on one of OpenAI's most promising new products, resigned from the company.

He just published a [fascinating blog post](#) on what it was like to work there for a year, including the sleepless sprint to build Codex. That's [OpenAI's new coding agent](#) that competes with [tools like Cursor](#) and Anthropic's [Claude Code](#).

French-Owen said he didn't leave because of any "drama," but because he wants to get back to being a startup founder. He was a co-founder of customer data startup Segment, [which was bought by Twilio in 2020 for \\$3.2 billion](#).

Some of what he revealed about the OpenAI culture would surprise no one, but other observations combat some misconceptions about the company. (He could not be immediately reached for comment.)

Fast growth: OpenAI grew from 1,000 to 3,000 people in the year he was there, he wrote.

The LLM model maker certainly has reasons for such hiring. It is the fastest-growing consumer product ever, and its competitors are also growing fast. In March, it said that [ChatGPT had over 500 million active users](#) and climbing quickly.


TE DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion

Chaos: “Everything breaks when you scale that quickly: how to communicate as a company, the reporting structures, how to ship product, how to manage and organize people, the hiring processes, etc.,” French-Owen wrote.

- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

[REGISTER NOW](#)

Like a small startup, people there are still empowered to act on their ideas with little to no red tape. But that also means that multiple teams are duplicating efforts. “I must’ve seen half a dozen libraries for things like queue management or agent loops,” he offered as examples.

Coding skill varies, too, from seasoned Google engineers who write code that can handle a billion users, to newly minted PhDs who do not. This, coupled with the flexible Python language, means that the central code repository, aka “the back-end monolith,” is “a bit of a dumping ground,” he described.

Stuff frequently breaks or can take excessive time to run. But top engineering managers are aware of

this and are working on improvements, he wrote.

“Launching spirit”: OpenAI doesn’t seem to know yet that it’s a giant company, right down to running entirely on Slack. It feels very much like move-fast-and-break-things Meta in its early Facebook years, he observed. The company is also full of hires from Meta.

French-Owen described how his senior team of around eight engineers, four researchers, two designers, two go-to-market staff, and a product manager built and launched Codex in only seven weeks, start to finish, with almost no sleep.

But launching it was magic. Just by turning it on, they got users. “I’ve never seen a product get so much immediate uptick just from appearing in a left-hand sidebar, but that’s the power of ChatGPT.”

Secretive fishbowl: ChatGPT is a highly scrutinized company. This led to a culture of secrecy in an attempt to clamp down on leaks to the public. At the same time, the company watches X. If a post goes viral there, OpenAI will see it and, possibly, respond to it. “A friend of mine joked, ‘this company runs on twitter vibes,’” he wrote.

Biggest misconception: French-Owen implied that the biggest misconception about OpenAI is that it isn’t as concerned about safety as it should be. Certainly a lot of AI safety folks, including former [OpenAI employees, have criticized its processes.](#)

While there are doomsayers worrying about theoretic risks to humanity, internally there’s more focus on practical safety like “hate speech, abuse, manipulating political biases, crafting bio-weapons, self-harm, prompt injection,” he wrote.

OpenAI isn’t ignoring the long-term potential impacts, he wrote. There are researchers looking at

them, and it's aware that hundreds of millions of people are using its LLMs today for everything from medical advice to therapy.

Governments are watching. Competitors are watching (and OpenAI is watching competitors in return). "The stakes feel really high."

Topics: [AI](#) [OpenAI](#) [Startups](#) [tech culture](#) [tech jobs](#)



Julie Bort

Venture Editor |

View Bio >

———— KEEP READING ————

AI



OpenAI's new GPT-4.1 AI models focus on coding

OpenAI on Monday launched a new family of models called GPT-4.1. Yes, “4.1” — as if the company’s nomenclature wasn’t confusing enough already.

There’s GPT-4.1, GPT-4.1 mini, and GPT-4.1 nano, all of which OpenAI says “excel” at coding and instruction following. Available through OpenAI’s API but not [ChatGPT](#), the multimodal models have a 1-million-token context window, meaning they can take in roughly 750,000 words in one go (longer than “War and Peace”).

GPT-4.1 arrives as OpenAI rivals like Google and Anthropic ratchet up efforts to build sophisticated programming models. Google’s recently released [Gemini 2.5 Pro](#), which also has a 1-million-token context window, ranks highly on popular coding benchmarks. So do Anthropic’s [Claude 3.7 Sonnet](#) and Chinese AI startup [DeepSeek’s upgraded V3](#).

It’s the goal of many tech giants, including OpenAI, to train AI coding models capable of performing complex software engineering tasks. OpenAI’s grand ambition is to create an “agentic software engineer,” as [CFO Sarah Friar put it](#) during a tech summit in London last month. The company asserts its future models will be able to program entire apps end-to-end, handling aspects such as quality assurance, bug testing, and documentation writing.

GPT-4.1 is a step in this direction.

Advertisement

Newsletters

[See More ↗](#)

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

“We’ve optimized GPT-4.1 for real-world use based on direct feedback to improve in areas that developers care most about: frontend coding, making fewer extraneous edits, following formats reliably, adhering to response structure and ordering, consistent tool usage, and more,” an OpenAI spokesperson told TechCrunch via email. “These improvements enable developers to build agents that are considerably better at real-world software engineering tasks.”

OpenAI claims the full GPT-4.1 model outperforms its [GPT-4o and GPT-4o mini](#) models on coding benchmarks, including SWE-bench. GPT-4.1 mini and nano are said to be more efficient and faster at the cost of some accuracy, with OpenAI saying GPT-4.1 nano is its speediest — and cheapest — model ever.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

GPT-4.1 costs \$2 per million input tokens and \$8 per million output tokens. GPT-4.1 mini is \$0.40/million input tokens and \$1.60/million output tokens, and

GPT-4.1 nano is \$0.10/million input tokens and \$0.40/million output tokens.

According to OpenAI's internal testing, GPT-4.1, which can generate more tokens at once than GPT-4o (32,768 versus 16,384), scored between 52% and 54.6% on SWE-bench Verified, a human-validated subset of SWE-bench. (OpenAI noted in a blog post that some solutions to SWE-bench Verified problems couldn't run on its infrastructure, hence the range of scores.) Those figures are slightly under the scores reported by Google and Anthropic for Gemini 2.5 Pro (63.8%) and Claude 3.7 Sonnet (62.3%), respectively, on the same benchmark.

In a separate evaluation, OpenAI probed GPT-4.1 using Video-MME, which is designed to measure the ability of a model to "understand" content in videos. GPT-4.1 reached a chart-topping 72% accuracy on the "long, no subtitles" video category, claims OpenAI.

While GPT-4.1 scores reasonably well on benchmarks and has a more recent "knowledge cutoff," giving it a better frame of reference for current events (up to June 2024), it's important to keep in mind that even some of the best models today struggle with tasks that wouldn't trip up experts. For example, [many studies](#) have [shown](#) that code-generating models often fail to fix, and even introduce, security vulnerabilities and bugs.

OpenAI acknowledges, too, that GPT-4.1 becomes less reliable (i.e., likelier to make mistakes) the more input tokens it has to deal with. On one of the company's own tests, OpenAI-MRCR, the model's accuracy decreased from around 84% with 8,000 tokens to 50% with 1 million tokens. GPT-4.1 also tended to be more "literal" than GPT-4o, says the company, sometimes necessitating more specific, explicit prompts.



Kyle Wiggers

AI Editor | [X](#) [🦋](#) [@](#) [in](#)

Kyle Wiggers was TechCrunch’s AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

[View Bio](#) >

KEEP READING



Latest

Startups

Venture

Apple

Security

AI

Apps

Events

Podcasts

Newsletters

🔍

☰



Kyle Wiggers

12:44 PM PDT · July 6, 2023

OpenAI today [announced](#) the general availability of [GPT-4](#), its latest text-generating model, through its API.

Starting this afternoon, all existing OpenAI API developers “with a history of successful payments” can access GPT-4. The company plans to open up access to new developers by the end of this month, and then start raising availability limits after that “depending on compute availability.”

“Millions of developers have requested access to the GPT-4 API since March, and the range of innovative products leveraging GPT-4 is growing every day,” OpenAI wrote in a [blog post](#). “We envision a future where chat-based models can support any use case. ”

GPT-4 can generate text (including code) and accept image and text inputs — an improvement over [GPT-3.5](#), its predecessor, which only accepted text — and performs at “human level” on various professional and academic benchmarks. Like previous GPT models from OpenAI, GPT-4 was trained using publicly available data, including from public web pages, as well as data that OpenAI licensed.

The image-understanding capability isn’t available to all OpenAI customers just yet. OpenAI’s testing it with a single partner, [Be My Eyes](#), to start with. But it hasn’t indicated when it’ll open it up to the wider customer base.

It’s worth noting that, as with even the best generative AI models today, GPT-4 isn’t perfect. It “hallucinates” facts and makes reasoning errors, sometimes with confidence. And it doesn’t learn from its experience, failing at hard problems such as introducing security vulnerabilities into code it generates.

Latest in AI

[See More](#)

AI



Google denies AI search features are killing website traffic

Sarah Perez - 2 hours ago

AI



First impressions of Alexa+, Amazon’s upgraded, AI-powered digital assistant

Sarah Perez - 3 hours ago

AI



Google takes on ChatGPT’s Study Mode with new ‘Guided Learning’ tool in Gemini

Aisha Malik - 3 hours ago

In the future, OpenAI says that it'll allow developers to fine-tune GPT-4 and [GPT-3.5 Turbo](#), one of its other recent but less capable text-generating models (and one of the original models powering [ChatGPT](#)), with their own data, as has long been possible with several of OpenAI's other text-generating models. That capability should arrive later this year, according to OpenAI.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Since the reveal of GPT-4 in March, the generative AI competition has grown fiercer. Recently, Anthropic expanded the context window for [Claude](#) — its flagship text-generating AI model, still in preview — from 9,000 tokens to 100,000 tokens. (Context window refers to the text the model considers before generating additional text, while tokens represent raw text — e.g. the word “fantastic” would be split into the tokens “fan,” “tas” and “tic.”)

[GPT-4](#) held the previous crown in terms of context window, weighing in at 32,000 tokens on the high

end. Generally speaking, models with small context windows tend to “forget” the content of even very recent conversations, leading them to veer off topic.

In a related announcement today, OpenAI said that it’s making its DALL-E 2 and Whisper APIs generally available — [DALL-E 2](#) being OpenAI’s image-generating model and “[Whisper](#)” referring to the company’s speech-to-text model. The company also said that it plans to deprecate old models available through its API to “optimize [its] compute capacity.” (Over the past few months, thanks in large part to the exploding popularity of ChatGPT, OpenAI has [struggled](#) to keep up with demand for its generative models.)

Starting January 4, 2024, certain older OpenAI models — specifically GPT-3 and its derivatives — will no longer be available, and will be replaced with new “base GPT-3” models that one would presume are more compute efficient. Developers using the old models will have to manually upgrade their integrations by January 4, and those who wish to continue using fine-tuned old models beyond January 4 will need to fine-tune replacements atop the new base GPT-3 models.

“We will be providing support to users who previously fine-tuned models to make this transition as smooth as possible,” OpenAI wrote. “In the coming weeks, we will reach out to developers who have recently used these older models, and will provide more information once the new completion models are ready for early testing.”

Topics: [AI](#) [api](#) [ChatGPT](#) [Generative AI](#) [gpt-4](#)

[OpenAI](#)



Kyle Wiggers

AI Editor | X reddit @ in

Kyle Wiggers was TechCrunch’s AI Editor until June 2025.
His writing has appeared in VentureBeat and Digital...

View Bio >

KEEP READING

AI

OpenAI intros new generative text features while reducing pricing

Kyle Wiggers — 11:12 AM PDT · June 13, 2023

IMAGE CREDITS:
STEFANI REYNOLDS/AFP / GETTY IMAGES

As the competition in the generative AI space grows fiercer, OpenAI is upgrading its text-

generating models while reducing pricing.

Today, OpenAI [announced](#) the release of new versions of [GPT-3.5-turbo](#) and [GPT-4](#), the latter being its latest text-generating AI, with a capability called function calling. As OpenAI explains in a [blog post](#), function calling allows developers to describe programming functions to GPT-3.5-turbo and GPT-4 and have the models create code to execute those functions.

For example, function calling can help to create chatbots that answer questions by calling external tools, convert natural language into database queries and extract structured data from text. “These models have been fine-tuned to both detect when a function needs to be called ... and to respond with [JSON](#) that adheres to the function signature,” OpenAI writes. “Function calling allows developers to more reliably get structured data back from the model.”

Beyond function calling, OpenAI is introducing a flavor of GPT-3.5-turbo with a greatly expanded context window. Context window, measured in [tokens](#), or raw bits of text, refers to the text the model considers before generating any additional text. Models with small context windows tend to “forget” the content of even very recent conversations, leading them to veer off topic — often in problematic ways.

The new GPT-3.5-turbo offers four times the context length (16,000 tokens) of the vanilla GPT-3.5-turbo at twice the price — \$0.003 per 1,000 input tokens (i.e. tokens fed into the model) and \$0.004 per 1,000 output tokens (tokens the model generates). OpenAI says that it can ingest around 20 pages of text in a single go — short of the hundreds of pages that AI startup Anthropic’s flagship model can [process](#), notably. (OpenAI is [testing](#) a version of

Advertisement

DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government’s website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue’s AI-generated ad isn’t just about fashion
- Lina Khan points to Figma IPO as vindication of M&A scrutiny

GPT-4 with a 32,000-token context window, but only in limited release.)

On the plus side, OpenAI says that it's reducing pricing for GPT-3.5-turbo — the original, not the version with the expanded context window — by 25%. Developers can now use the model for \$0.0015 per 1,000 input tokens and \$0.002 per 1,000 output tokens, which equates to roughly 700 pages per dollar.

Pricing is also being reduced for text-embedding-ada-002, one of OpenAI's more popular text embedding models. Text embeddings measure the relatedness of text strings, and are commonly used for search (where results are ranked by relevance to a query string) and recommendations (where items with related text strings are recommended).

- Your public ChatGPT queries are getting indexed by Google and other search engines

Advertisement

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Text-embedding-ada-002 now costs \$0.0001 per 1,000 tokens, a 75% reduction from the previous price. OpenAI says the reduction was made possible by increased efficiency in its systems — a

key area of focus for the startup, no doubt, as it spends [hundreds of millions](#) of dollars on R&D and infrastructure.

OpenAI has signaled that incremental updates to existing models — not massive new from-scratch models — are its MO following the release of GPT-4 in early March. At a recent [conference](#) hosted by Economic Times, CEO Sam Altman reaffirmed that OpenAI hasn't begun training the successor to GPT-4, indicating that the company “has a lot of work to do” before it starts that model.

Topics: [AI](#) [Generative AI](#) [gpt-4](#) [OpenAI](#)

Advertisement



Kyle Wiggers

AI Editor |

Kyle Wiggers was TechCrunch's AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

[View Bio](#) >

Loading the next article





[TechCrunch Staff](#)

[Contact Us](#)

[Advertise](#)

[Crunchboard Jobs](#)

[Site Map](#)

[Terms of Service](#)

[Privacy Policy](#)

[RSS Terms of Use](#)

[OpenAI](#)

[Perplexity](#)

[Cisco](#)

[DeepMind](#)

[Windsurf](#)

[Tech Layoffs](#)

[ChatGPT](#)