

Advertisement

AI



A new AI coding
challenge just published
its first results — and
they aren't pretty

A new AI coding challenge has revealed its first winner — and set a new bar for AI-powered software engineers.

On Wednesday at 5 p.m. PT, the nonprofit Laude Institute announced the first winner of the K Prize, a multi-round AI [coding challenge](#) launched by Databricks and Perplexity co-founder Andy Konwinski. The winner was a Brazilian prompt engineer named Eduardo Rocha de Andrade, who will receive \$50,000 for the prize. But more surprising than the win was his final score: He won with correct answers to just 7.5% of the questions on the test.

IMAGE CREDITS: SASHKINW / GETTY IMAGES

“We’re glad we built a benchmark that is actually hard,” said Konwinski. “Benchmarks should be hard if they’re going to matter,” he continued, adding: “Scores would be different if the big labs had entered with their biggest models. But that’s kind of the point. K Prize runs offline with limited compute, so it favors smaller and open models. I love that. It levels the playing field.”

Konwinski has pledged \$1 million to the first open source model that can score higher than 90% on the test.

Similar to the well-known SWE-Bench system, the K Prize tests models against flagged issues from GitHub as a test of how well models can deal with real-world programming problems. But while SWE-

DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government’s website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue’s AI-generated ad isn’t just about fashion

Bench is based on a fixed set of problems that models can train against, the K Prize is designed as a “contamination-free version of SWE-Bench,” using a timed entry system to guard against any benchmark-specific training. For round one, models were due by March 12. The K Prize organizers then built the test using only GitHub issues flagged after that date.

The 7.5% top score stands in marked contrast to SWE-Bench itself, which currently shows a 75% top score on its easier “Verified” test and 34% on its harder “Full” test. Konwinski still isn’t sure whether the disparity is due to contamination on SWE-Bench or just the challenge of collecting new issues from GitHub, but he expects the K Prize project to answer the question soon.

“As we get more runs of the thing, we’ll have a better sense,” he told TechCrunch, “because we expect people to adapt to the dynamics of competing on this every few months.”

- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

[REGISTER NOW](#)

It might seem like an odd place to fall short, given the wide range of AI coding tools already publicly available — but with benchmarks becoming too easy, many critics see projects like the K Prize as a necessary step toward solving [AI’s growing evaluation problem](#).

“I’m quite bullish about building new tests for existing benchmarks,” says Princeton researcher Sayash Kapoor, who put forward a similar idea [in a recent paper](#). “Without such experiments, we can’t actually tell if the issue is contamination, or even just targeting the SWE-Bench leaderboard with a human in the loop.”

For Konwinski, it’s not just a better benchmark, but an open challenge to the rest of the industry. “If you listen to the hype, it’s like we should be seeing AI doctors and AI lawyers and AI software engineers, and that’s just not true,” he says. “If we can’t even get more than 10% on a contamination-free SWE-Bench, that’s the reality check for me.”

Topics: [AI](#) [Andy Konwinski](#) [K Prize](#) [Laude Institute](#)

[f](#) [X](#) [in](#) [📷](#) [✉](#) [🔗](#)

Russell Brandom

✕|

Russell is a freelance writer based in New York.

[View Bio](#) >

———— KEEP READING ————



Databricks, Perplexity co-founder pledges \$100M on new fund for AI researchers

Julie Bort — 1:30 PM PDT · June 23, 2025

IMAGE CREDITS: GETTY IMAGES

Andy Konwinski, computer scientist and co-founder of Databricks and Perplexity, announced Monday that his company, Laude, is forming a new AI research institute backed with \$100 million of his own money.

The [Laude Institute](#) functions less as an AI research lab and more like a fund looking to make investments structured similar to grants. In addition to Konwinski, the institute's board includes UC Berkeley professor Dave Patterson (known for a string of award-winning research), Jeff Dean (known as Google's chief scientist), and Joelle Pineau (Meta's vice president of AI Research).

Konwinski announced the institute's first and "flagship" grant of \$3 million annually for five years, which will anchor the new AI Systems Lab at UC

Advertisement

Newsletters

See More ↗

Subscribe for the industry's biggest tech news

Berkeley. This is a new lab led by one of Berkeley's most celebrated researchers, Ion Stoica, current director of the Sky Computing Lab. Stoica is also a co-founder of startup Anyscale (an AI and Python platform) and AI big data company Databricks, both built from tech developed in Berkeley's lab system.

The new AI Systems Lab is set to open in 2027 and in addition to Stoica will include a number of other well-known researchers.

In his [blog post](#) announcing the institute, Konwinski described its mission as "built by and for computer science researchers ... We exist to catalyze work that doesn't just push the field forward but guides it towards more beneficial outcomes."

That's not necessarily a direct dig at OpenAI, which started out as an AI research facility and is now, arguably, consumed by its enormous commercial side. But other researchers have fallen prey to the lure of money as well.

For instance, popular AI researcher Epoch [faced controversy](#) when it revealed that OpenAI supported the creation of one of its AI benchmarks that was then used to unveil its new o3 model. Epoch's founder also launched a startup with the controversial mission to [replace all human workers everywhere with AI agents](#).

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch's coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).



the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Like other AI research organizations with commercial ambitions, Konwinski has structured his institute across boundaries: as a nonprofit with a public benefit corporation operating arm.

He's dividing his research investments into two buckets that he calls "Slingshots and Moonshots." Slingshots are for early-stage research that can benefit from grants and hands-on help. Moonshots are, as the name implies, for "long-horizon labs tackling species-level challenges like AI for scientific discovery, civic discourse, healthcare, and workforce re-skilling."

His lab has, for instance, collaborated with "terminal-bench," a Stanford-led benchmark for how well AI agents handle tasks, [used by Anthropic.](#)

One thing to note, Konwinski's endeavors under the name "Laude" aren't solely a grant-writing research institute. He also [co-founded a for-profit venture fund](#) launched in 2024. The fund's co-founder is former NEA VC Pete Sonsini, and a Laude spokesperson tells us it has more than 50 leading researchers as LPs in the fund. As TechCrunch previously reported, [Laude led a \\$12 million investment in AI agent infrastructure startup Arcade.](#) It has quietly backed other startups, too.

A Laude spokesperson tells us that while Konwinski has pledged \$100 million, he's also looking for, and open to, investment from other successful technologists. As to how Konwinski amassed a

fortune big enough to guarantee \$100 million for this new endeavor: [Databricks closed a \\$15.3 billion](#) funding round in January that valued the company at \$62 billion. [Perplexity last month secured](#) a \$14 billion valuation, too.

Does the world really need yet more AI “good for humanity” research or another AI company with a murky nonprofit/commercial structure? No, and yes.

AI research has become [increasingly muddled](#). For instance, AI benchmarks designed to prove that a particular vendor’s model works best have become plentiful these days. (Even Salesforce [has its own LLM benchmark for CRMs](#).)



Latest

Startups

Venture

Apple

Security

AI

Apps

Events

Podcasts

Newsletters

Q

≡

[f](#) [X](#) [in](#) [tumblr](#) [email](#) [rss](#)



Julie Bort

Venture Editor | [X](#)

[View Bio](#) >

———— KEEP READING ————



IMAGE CREDITS: POKÉMON

AI



Debates over AI benchmarking have reached Pokémon

Kyle Wiggers — 3:27 PM PDT · April 14, 2025

Not even Pokémon is safe from AI benchmarking controversy.

Last week, a [post on X](#) went viral, claiming that Google's latest Gemini model surpassed Anthropic's flagship Claude model in the original Pokémon video game trilogy. Reportedly, Gemini had reached Lavender Town in a developer's Twitch stream; Claude was [stuck at Mount Moon](#) as of late February.

Gemini is literally ahead of Claude atm in pokemon after reaching Lavender Town

Latest in AI

[See More](#)

AI



Google denies AI search features are killing website traffic

Sarah Perez - 10 hours ago

AI



First impressions of Alexa+, Amazon's upgraded, AI-powered digital assistant

119 live views only btw, incredibly underrated stream pic.twitter.com/8AvSovAl4x

— Jush (@Jush21e8) [April 10, 2025](#)

But what the post failed to mention is that Gemini had an advantage.

As [users on Reddit](#) pointed out, the developer who maintains the Gemini stream built a custom minimap that helps the model identify “tiles” in the game like cuttable trees. This reduces the need for Gemini to analyze screenshots before it makes gameplay decisions.

Now, Pokémon is a semi-serious AI benchmark at best — few would argue it’s a very informative test of a model’s capabilities. But it *is* an instructive example of how different implementations of a benchmark can influence the results.

For example, Anthropic [reported](#) two scores for its recent Anthropic 3.7 Sonnet model on the benchmark SWE-bench Verified, which is designed to evaluate a model’s coding abilities. Claude 3.7 Sonnet achieved 62.3% accuracy on SWE-bench Verified, but 70.3% with a “custom scaffold” that Anthropic developed.

Sarah Perez - 11 hours ago

AI



Google takes on ChatGPT’s Study Mode with new ‘Guided Learning’ tool in Gemini

Aisha Malik - 11 hours ago

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your

ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

More recently, Meta [fine-tuned](#) a version of one of its newer models, Llama 4 Maverick, to perform well on a particular benchmark, LM Arena. The [vanilla version](#) of the model scores significantly worse on the same evaluation.

Given that AI benchmarks — Pokémon included — are [imperfect measures](#) to begin with, custom and non-standard implementations threaten to muddy the waters even further. That is to say, it doesn't seem likely that it'll get any easier to compare models as they're released.

Topics: [AI](#) [Gaming](#) [pokemon](#)



Kyle Wiggers

AI Editor |

Kyle Wiggers was TechCrunch's AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

[View Bio](#) >

KEEP READING

OpenAI's new GPT-4.1 AI models focus on coding

Kyle Wiggers — 10:00 AM PDT · April 14, 2025

OpenAI on Monday launched a new family of models called GPT-4.1. Yes, "4.1" — as if the company's nomenclature wasn't confusing enough already.

IMAGE CREDITS:
JAKUB PORZYCKI/NURPHOTO / GETTY IMAGES

There's GPT-4.1, GPT-4.1 mini, and GPT-4.1 nano, all of which OpenAI says "excel" at coding and instruction following. Available through OpenAI's API but not [ChatGPT](#), the multimodal models have a 1-million-token context window, meaning they can take in roughly 750,000 words in one go (longer than "War and Peace").

GPT-4.1 arrives as OpenAI rivals like Google and Anthropic ratchet up efforts to build sophisticated programming models. Google's recently released [Gemini 2.5 Pro](#), which also has a 1-million-token context window, ranks highly on popular coding


DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website

benchmarks. So do Anthropic's [Claude 3.7 Sonnet](#) and Chinese AI startup [DeepSeek's upgraded V3](#).

It's the goal of many tech giants, including OpenAI, to train AI coding models capable of performing complex software engineering tasks. OpenAI's grand ambition is to create an "agentic software engineer," as [CFO Sarah Friar put it](#) during a tech summit in London last month. The company asserts its future models will be able to program entire apps end-to-end, handling aspects such as quality assurance, bug testing, and documentation writing.

GPT-4.1 is a step in this direction.

"We've optimized GPT-4.1 for real-world use based on direct feedback to improve in areas that developers care most about: frontend coding, making fewer extraneous edits, following formats reliably, adhering to response structure and ordering, consistent tool usage, and more," an OpenAI spokesperson told TechCrunch via email. "These improvements enable developers to build agents that are considerably better at real-world software engineering tasks."

OpenAI claims the full GPT-4.1 model outperforms its [GPT-4o and GPT-4o mini](#) models on coding benchmarks, including SWE-bench. GPT-4.1 mini and nano are said to be more efficient and faster at the cost of some accuracy, with OpenAI saying GPT-4.1 nano is its speediest — and cheapest — model ever.

- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion
- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and

sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

GPT-4.1 costs \$2 per million input tokens and \$8 per million output tokens. GPT-4.1 mini is \$0.40/million input tokens and \$1.60/million output tokens, and GPT-4.1 nano is \$0.10/million input tokens and \$0.40/million output tokens.

According to OpenAI's internal testing, GPT-4.1, which can generate more tokens at once than GPT-4o (32,768 versus 16,384), scored between 52% and 54.6% on SWE-bench Verified, a human-validated subset of SWE-bench. (OpenAI noted in a blog post that some solutions to SWE-bench Verified problems couldn't run on its infrastructure, hence the range of scores.) Those figures are slightly under the scores reported by Google and Anthropic for Gemini 2.5 Pro (63.8%) and Claude 3.7 Sonnet (62.3%), respectively, on the same benchmark.

In a separate evaluation, OpenAI probed GPT-4.1 using Video-MME, which is designed to measure the ability of a model to "understand" content in videos. GPT-4.1 reached a chart-topping 72% accuracy on the "long, no subtitles" video category, claims OpenAI.

While GPT-4.1 scores reasonably well on benchmarks and has a more recent "knowledge cutoff," giving it a better frame of reference for current events (up to June 2024), it's important to keep in mind that even some of the best models

today struggle with tasks that wouldn't trip up experts. For example, [many studies](#) have [shown](#) that code-generating models often fail to fix, and even introduce, security vulnerabilities and bugs.

OpenAI acknowledges, too, that GPT-4.1 becomes less reliable (i.e., likelier to make mistakes) the more input tokens it has to deal with. On one of the company's own tests, OpenAI-MRCR, the model's accuracy decreased from around 84% with 8,000 tokens to 50% with 1 million tokens. GPT-4.1 also tended to be more "literal" than GPT-4o, says the company, sometimes necessitating more specific, explicit prompts.

Topics: [AI](#) [gpt4.1](#) [OpenAI](#)



Kyle Wiggers

AI Editor |

Kyle Wiggers was TechCrunch's AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

[View Bio](#) >

Loading the next article



[TechCrunch Staff](#)

[Contact Us](#)

[Advertise](#)

[Crunchboard Jobs](#)

[Terms of Service](#)

[Privacy Policy](#)

[RSS Terms of Use](#)

[OpenAI](#)
[Perplexity](#)

[Cisco](#)
[DeepMind](#)

[Windsurf](#)

[Tech Layoffs](#)

[ChatGPT](#)

