

Advertisement

TRANSPORTATION



Aeva strikes lidar chip-making deal with LG subsidiary

[LG Innotek](#), the components and materials subsidiary of South Korea's [LG Group](#) has struck a manufacturing partnership with [Aeva Technologies](#), which makes 4D lidar sensing systems.

The partnership will see LG Innotek manufacturing and supplying Aeva's [Atlas Ultra](#) 4D lidar sensor for automobiles, and eventually expanding the technology for use in consumer electronics, robotics, and industrial automation. As part of the deal, LG Innotek will invest up to \$50 million in Aeva, acquiring an equity stake of about 6% in the U.S. company.

 SPONSORED

IMAGE CREDITS: AEVA



Design agents are here — Lovart launches the first end-to-end AI canvas for design

Sponsored by Lovart

“The first part of the partnership is focused on the automotive sector,” Aeva’s co-founder and CEO Soroush Salehian told TechCrunch. “LG Innotek will act as our manufacturing partner for some of the top 10 global passenger vehicle OEMs.”

Founded in 2017, Aeva says it has poured nearly half a billion dollars into developing what it calls “4D lidar on a chip.” Unlike conventional lidar systems that rely on time-of-flight measurements to estimate distance, Aeva’s frequency modulated continuous wave-based perception technology


DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government’s website
- Instagram takes on Snapchat with new ‘Instagram Map’
- Google says hackers stole its customers’ data by breaching its Salesforce database
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Spotify raises subscription prices

measures both distance and velocity for every pixel in real time, Salehian said.

“It’s like going from black-and-white to color camera. We get this velocity information as a new dimension,” he explained.

Aeva says it has integrated the entire lidar system, including optics, into a silicon photonics module. This miniaturized, chip-scale design enables more efficient production and integration, particularly for scaling in markets such as automotive, robotics, and consumer electronics, Salehian said.

The company plans to put the capital from the deal toward product development and augmenting its team. “Over the past eight years, we’ve made [the hardware] super small, and we’re sticking to that trajectory. We’re going to make it down to a monolithic, single-chip,” Salehian said.

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.


San Francisco | October 27-29, 2025

REGISTER NOW

Aeva is also targeting the manufacturing sector and has struck partnerships Nikon in Japan and

- The uproar over Vogue’s AI-generated ad isn’t just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines

Sponsored by Lovart



SPONSORED

Lovart’s design agent hints at the future: AI-native brands built in minutes

Sick AG in Germany. The company is also expanding into smart infrastructure and transportation, providing security and monitoring solutions at major U.S. airports, including SFO and JFK Terminal One.

Topics: [Aeva](#) [Asia](#) [LG Innotek](#) [Lidar](#) [South Korea](#)

[Transportation](#) [United States](#)



Kate Park

Reporter, Asia |

Kate Park is a reporter at TechCrunch, with a focus on technology, startups and venture capital in Asia. She...


[View Bio](#) >

———— KEEP READING ————

TRANSPORTATION



Tesla shuts down Dojo,
the AI training
supercomputer that
Musk said would be key
to full self-driving



Rebecca Bellan — 3:19 PM PDT · August 7, 2025

Tesla is breaking up the team behind its [Dojo supercomputer](#), ending the automaker's play at developing in-house chips for driverless technology, according to [Bloomberg](#).

Dojo's lead, Peter Bannon, is leaving the company, and the remaining team members will be reassigned to other data center and compute projects within Tesla, per Bloomberg's reporting, which cited anonymous sources.

The disbanding of Tesla's Dojo efforts follows the departure of around 20 workers, who left the automaker to start their own [AI company called DensityAI](#). The new startup is reportedly coming out of stealth soon and is building chips, hardware, and software that will power data centers for AI that are used in robotics, by AI agents, and in automotive applications. DensityAI was founded by former Dojo head [Ganesh Venkataramanan](#) and ex-Tesla employees Bill Chang and Ben Floering.

It also comes at a crucial time for Tesla.

[CEO Elon Musk has pushed](#) to get shareholders to view Tesla as an AI and robotics company, despite a limited robotaxi launch in Austin this past June that featured Model Y vehicles with a human in the front passenger seat and resulted in a number of reported incidents of the vehicles exhibiting problematic driving behavior.

Advertisement

Newsletters

[See More ↗](#)

Subscribe for the industry's biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch's coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

Tesla’s decision to shut down Dojo, which Musk has been talking about since 2019, is a major shift in strategy. Musk has said that Dojo would be the cornerstone of Tesla’s AI ambitions and its goal to reach full self-driving due to its ability to “process truly vast amounts of video data.” He talked about Dojo, albeit briefly, as recently as the company’s second-quarter earnings call.

In 2023, Morgan Stanley predicted Dojo could add \$500 billion to the company’s market value by unlocking new revenue streams in the form of robotaxis and software services. Just last year, Musk noted that Tesla’s AI team would “[double down](#)” on Dojo in the lead-up to [Tesla’s robotaxi reveal](#), which happened in October.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

But talk about Dojo halted around August 2024, when Musk began touting Cortex instead, Tesla’s “giant new AI training supercluster being built at Tesla HQ in Austin to solve real-world AI.”

The Dojo project was one part supercomputer, one part in-house chip-making. Tesla unveiled its D1 chip when it formally announced Dojo at its first AI Day in 2021. [Venkataramanan](#) presented the chip, which Tesla said would be used alongside Nvidia's GPU to power the Dojo supercomputer. The automaker also said it was working on a next-gen D2 chip that would solve any information flow bottlenecks of its predecessor.

Sources told Bloomberg that now Tesla plans to increase its reliance on Nvidia, as well as other external tech partners like AMD for compute and Samsung for chip manufacturing. Tesla last month signed a [\\$16.5 billion deal with Samsung](#) to make its AI6 inference chips, a chip design that promises to scale from powering FSD and Tesla's Optimus humanoid robots all the way to high-performance AI training in data centers.

During Tesla's second-quarter earnings call, Musk hinted at potential redundancies.

"Thinking about Dojo 3 and the AI6 inference chip, it seems like intuitively, we want to try to find convergence there, where it's basically the same chip," Musk said.

The news comes as [Tesla's board offers Musk a \\$29 billion pay package](#) to keep him at Tesla and help push the company's AI efforts forward, rather than getting too sidetracked by his other companies, including the more pure-play AI startup xAI.

TechCrunch has reached out to Tesla for more information.

Have a sensitive tip or confidential documents? We're reporting on the inner workings of the AI industry — from the companies shaping its future to the people impacted by their decisions. Reach out to Rebecca Bellan at

rebecca.bellan@techcrunch.com and Maxwell Zeff
at maxwell.zeff@techcrunch.com. For secure
communication, you can contact us via Signal at
@rebeccabellan.491 and @mzeff.88.

Topics: [AI](#) [autonomous vehicles](#) [Dojo](#) [Elon Musk](#)

[robotaxis](#) [Tesla](#) [Transportation](#)

Latest Security
Startups AI
Venture Apps
Apple



Events
Podcasts
Newsletters

SPACE



Rocket Lab eyes big defense opportunities with new acquisition

Rocket Lab is signaling to investors, yet again, that it's more than "just" a rocket company.

Rocket Lab's second-quarter results, which were posted Thursday, show revenues continue to be driven by its space systems business rather than launch. The results also highlight the company's acquisition strategy and how its purchase of a new optical payloads company will make it more competitive for lucrative government contracts.

The company's space systems brought in \$97.9 million of the \$144.5 million in total revenues for the second quarter. Rocket Lab's total revenue, its highest quarterly revenue in the company's history, jumped 36% from a year ago. The company's net loss widened to \$66.4 million.

Rocket Lab founder and CEO Peter Beck highlighted the "busy quarter of M&A activity" as the company gets close to closing its deal to buy [Geost](#), a company that builds optical payloads. Through that acquisition, which will close at \$275 million in cash and equity, Rocket Lab is opening a new business unit called Optical Systems to scale electro-optical and infrared sensor manufacturing.

Those sensors are used in missile warning, tracking, and space domain awareness. The acquisition is also one part of a larger play from Rocket Lab — spelled out quite explicitly in the company's earnings presentation — to bid for

Latest in Space

[See More](#)

SPACE



Rocket Lab eyes big defense opportunities with new acquisition

Aria Alamalhodaie

- 7 hours ago

SPACE



SpaceX is building a water pipeline to Starbase — but access comes with some conditions

Aria Alamalhodaie

- 11 hours ago

SPACE



Hubble Network plans massive satellite upgrade to create global Bluetooth layer

Aria Alamalhodaie

- 18 hours ago

multibillion-dollar [DOD initiatives like Golden Dome](#).

Rocket Lab has already won some major defense contracts, including a \$515 million, 18-satellite build contract to support the Space Development Agency's constellation of missile-tracking satellites. The company has moved into production on those satellites, following confirmation that they meet the DOD's mission requirements, Rocket Lab said.

Down on Earth, Rocket Lab says it's progressing well toward the first launch of its larger Neutron rocket. The rocket's launch complex in Virginia is expected to be complete in the third quarter of this year, with [hardware en route](#) and the new rocket engine, called Archimedes, undergoing multiple tests per day.

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

[REGISTER NOW](#)

The company is staying vague about the exact launch date, saying only in the presentation that it

is pushing an “all-out effort to get Neutron to the launch pad before the end of 2025.”

Rocket Lab ended the quarter with \$564 million in cash and cash equivalents and estimates revenues will hit between \$145 to \$155 million next quarter.

Topics: [earnings](#) [Mergers and Acquisitions](#)

[Rocket Lab](#) [Space](#)

[f](#) [X](#) [in](#) [t](#) [✉](#) [🔗](#)



Aria Alamalhodaei

Reporter, Space and Defense | [X](#)

Aria Alamalhodaei covers the space and defense industries at TechCrunch. Previously, she covered the...

[View Bio](#) >

———— KEEP READING ————

STARTUPS

[f](#) [X](#) [in](#) [t](#) [✉](#) [🔗](#)

High costs and thin margins threatening AI coding startups

In February, AI coding startup Windsurf was in talks to raise a big new round at a [\\$2.85 billion valuation](#) led by [Kleiner Perkins](#), at double the valuation it hit six months earlier, sources told TechCrunch at the time. That deal didn't happen, according to a source familiar with the matter. Instead, news broke in April that the startup planned to sell itself to OpenAI for roughly the same valuation: [\\$3 billion](#).

While that deal famously fell apart, one bigger question remains: If the startup was growing that fast and attracting VC interest, why would it sell at all?

IMAGE CREDITS:
YUICHIRO CHINO / GETTY IMAGES

Insiders tell TechCrunch that for all the popularity and hype around AI coding assistants, they can actually be massively money-losing businesses. Vibe coders generally, and Windsurf in particular, can have such expensive structures that their gross margins are “very negative,” one person close to Windsurf told TechCrunch. Meaning it cost more to run the product than the startup could charge for it.

This is due to the high costs of using large language models (LLMs), the person explained. AI coding assistants are particularly pressured to always offer the most recent, most advanced, and most expensive LLMs because model makers are particularly fine-tuning their latest models for

TC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Instagram takes on Snapchat with new 'Instagram Map'
- Google says hackers stole its customers' data by breaching its Salesforce database
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door

Improvements in coding and related tasks like debugging.

This is a challenge compounded by fierce competition in the vibe-coding and code-assist market. Rivals include companies that already have huge customer bases like Anysphere's Cursor and GitHub Copilot.

The most straightforward path to improving margins in this business involves the startups building their own models, thereby eliminating costs of paying suppliers like Anthropic and OpenAI.

"It's a very expensive business to run if you're not going to be in the model game," said the person.

- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

But that idea comes with its own risks. Windsurf's co-founder and CEO, Varun Mohan, ultimately decided against the company building its own model — an expensive undertaking, the person said.

In addition, model makers are already competing directly. Anthropic offers Claude Code and OpenAI offers Codex, for instance.

Selling the business was a strategic move to lock in a high return before it could be undermined by the very companies that supplied its AI, including OpenAI and Anthropic, which were also entering the AI coding market.

Multiple people believe that the same pressure on margins Windsurf faced could be impacting Anysphere, the maker of Cursor, as well as vibe coders like Lovable, Replit, and others.

“Margins on all of the ‘code gen’ products are either neutral or negative. They’re absolutely abysmal,” said Nicholas Charriere, founder of Mocha, a vibe-coding startup and back-end hosting solution serving small and medium businesses (SMBs). He added that he believes the variable costs for all the startups in the sector are very close, likely within 10% to 15% of one another.

Unlike Windsurf, Anysphere has been growing so fast that it intends to remain an independent company, having already [turned down acquisition offers](#), including, reports say, from OpenAI.

And Anysphere [announced in January](#) that it is attempting to build its own model, which could give it more control over its expenses. In July, the startup [hired two leaders from Anthropic’s Claude Code](#) team, the Information reported, but two weeks later, these employees [returned](#) to work at Anthropic.

In addition to building a model, Anysphere could expect the cost of LLMs to decrease over time.

“That’s what everyone’s banking on,” said Erik Nordlander, a general partner at Google Ventures.

“The inference cost today, that’s the most expensive it’s ever going to be.”

It’s not entirely clear how true that is. Rather than falling as expected, the [cost of some of the latest AI models has risen](#), as they use more time and computational resources to handle complicated, multistep tasks.

When that will change remains to be seen. On Thursday, for instance, [OpenAI introduced](#) a new flagship model, GPT-5, with fees that are significantly less than its competitor, Anthropic’s Claude Opus 4.1. And Anysphere [immediately offered](#) this model as a choice for Cursor users.

Anysphere has also recently changed its pricing structure to pass along the increased costs of running Anthropic’s latest Claude model, particularly to its most active users. The move caught some of [Cursor customers by surprise](#), since they didn’t expect additional charges on top of its \$20-per-month Pro plan. Anysphere CEO Michael Truell later apologized for unclear communication about the pricing change in a [blog post](#).

This is the rock and the hard place. Although Cursor is one of the most popular AI applications, having reached [\\$500 million in ARR in June](#), the company’s user base may not be so loyal to the product if another company develops a tool that is superior to Cursor, investors say.

Anysphere didn’t respond to a request for comment.

Given the competitive landscape and the costs, Windsurf’s decision to get out may prove to be understandable. After the OpenAI deal fell through, the founders and key employees left to join Google in a deal that led to a [\\$2.4 billion payout](#) to key

shareholders. The remaining business then [sold itself to Cognition](#).

While many, including prominent VCs, criticized Mohan for leaving approximately 200 employees without roles at Google, a source familiar with the deal insisted the acquisition actually maximized the outcomes for all employees.

Beyond Cursor, other AI coding tools are also among the fastest growing startups of the LLM generation, like Replit, Lovable, and Bolt, and all of them rely on model makers as well.

Additionally, if this extremely popular business sector, already generating hundreds of millions in revenue or more a year, has difficulty building on top of model makers, what might it mean for other, more nascent industries?

Topics: [AI](#) [coding assistants](#) [cursor](#) [Startups](#)

[Venture](#) [Windsurf](#)

Advertisement

[f](#) [X](#) [in](#) [🍷](#) [✉](#) [🔗](#)


Marina Temkin

Reporter, Venture | [X](#)

Marina Temkin is a venture capital and startups reporter at TechCrunch. Prior to joining TechCrunch, she wrote abo...

View Bio >

Loading the next article



TechCrunch
Staff

Contact Us

Advertise

Crunchboard
Jobs

Site Map

Terms of
Service

Privacy
Policy

RSS Terms of
Use

OpenAI
Perplexity

Cisco
DeepMind

Windsurf
Tech Layoffs
ChatGPT