


Advertisement

AI



Anthropic unveils new rate limits to curb Claude Code power users



Anthropic says its rolling out new weekly rate limits for Claude to curb usage among subscribers who are running its AI coding tool, Claude Code, “continuously in the background, 24/7.” Anthropic says the rate limits also aim to stop a handful of users who are violating Claude’s usage policy by sharing accounts and reselling access to Claude Code.

The new rate limits will go into effect August 28 for subscribers to Anthropic’s \$20-per-month Pro plan, as well as its \$100- and \$200-per-month Max plans, the company said Monday in an email to subscribers and a [post on X](#).

IMAGE CREDITS: ANTHROPIC

DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government’s website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue’s AI-generated ad isn’t just about fashion




Anthropic

@AnthropicAI · Follow



We're rolling out new weekly rate limits for Claude Pro and Max in late August. We estimate they'll apply to less than 5% of subscribers based on current usage.



6:23 PM · Jul 28, 2025



3K



Reply



Copy link

[Read 563 replies](#)

- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

Anthropic says its existing usage limits, which reset every five hours, will remain in place. The company is also introducing two new weekly rate limits that reset every seven days; one is an overall usage limit, whereas the other is specific to Anthropic's most advanced AI model, Claude Opus 4. Anthropic says Max subscribers can purchase additional usage, beyond what the rate limit provides, at standard API rates.

The announcement comes just weeks after [Anthropic quietly introduced rate limits for Claude Code](#). The company said at the time it was aware of the issues but declined to elaborate further. While Anthropic's AI coding tool has been a hit with developers, the company seems to be having a difficult time serving it broadly. Anthropic's [status page](#) shows that Claude Code has experienced a partial or major outage at least seven times in the

last month — perhaps because some power users seem to be running Claude Code nonstop.

“Claude Code has experienced unprecedented demand since launch,” said Anthropic spokesperson Amie Rotherham in an email to TechCrunch about the weekly rate limits. Rotherham notes that “most users won’t notice a difference,” and that this limit will affect less than 5% of subscribers, based on their current usage patterns.

Anthropic tells TechCrunch that most Pro users can expect 40 to 80 hours of Sonnet 4 through Claude Code within their weekly rate limits. Subscribers to Anthropic’s \$100-per-month Max plan can expect 140 to 280 hours of Sonnet 4 and 15 to 35 hours of Opus 4. And subscribers to Anthropic’s \$200-per-month Max plan can expect 240 to 480 hours of Sonnet 4 and 24 to 40 hours of Opus 4.

The company notes that usage may vary based on codebase size and other factors. Anthropic claims [that the \\$200 Max plan offers 20x more usage than the Pro plan](#) — the company says that rate limits are activated when users hit a certain number of tokens, rather than a certain time limit.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

[REGISTER NOW](#)

Anthropic has said before that [it's very constrained when it comes to computational resources](#), which seems to be the case for most AI model providers today. Most AI companies are racing to bring new AI data centers online to meet the massive demands of serving and training their AI models.

Several providers of AI coding tools are revisiting the pricing strategy around their products. In June, the company behind Cursor, Anysphere, changed the way it priced usage for its \$20-per-month Pro plan to limit power users from abusing the plan. However, Anysphere later [apologized for poorly communicating those changes](#), leading to some users paying more than they expected. Another AI coding tool provider, Replit, [made similar pricing changes in June](#) as well.

In an email to Claude subscribers, Anthropic says it's committed to "supporting long-running use cases through other options in the future." However, the company claims these rate limits will help them maintain reliable service broadly in the short term.

Update 7/29 10 a.m. PT: This story was updated to include comments from an Anthropic spokesperson and more details about Anthropic's weekly rate limits.

Topics: [AI](#) [AI coding tools](#) [Anthropic](#) [Claude](#)





Maxwell Zeff

Senior AI Reporter | [X](#)

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >

KEEP READING



AI



Anthropic rolls out a \$200-per-month Claude subscription

Maxwell Zeff — 10:00 AM PDT · April 9, 2025

IMAGE CREDITS:

BENJAMIN GIRETTE/BLOOMBERG / GETTY IMAGES

Anthropic announced on Wednesday that it's launching a new, very expensive subscription plan for its AI chatbot Claude: Max. An answer to [OpenAI's \\$200-a-month ChatGPT Pro](#) tier, Max

comes with higher usage limits than Anthropic’s \$20-per-month Claude Pro subscription, as well as priority access to the company’s newest AI models and features.

A bit confusingly, Max comes in two flavors with different price points and usage limits. There’s a \$100-per-month Max tier with 5x higher rate limits than Claude Pro and a \$200-per-month Max option with 20x higher rate limits.

Frontier AI model developers are looking for new ways to increase revenue — and offering power users an expensive AI subscription seems like a promising way to do it. Just two months after launching ChatGPT Pro, OpenAI reportedly told investors [its annualized revenue grew by \\$300 million](#). Should Anthropic’s new Max plan be similarly successful, it could give the company a [much-needed boost](#).

Anthropic’s product lead Scott White told TechCrunch in an interview that the company is not ruling out launching even pricier subscriptions someday. Notably, Anthropic still doesn’t have an unlimited usage plan, which OpenAI offers with ChatGPT Pro.

When asked whether Anthropic would ever consider launching a \$500-per-month Claude plan, White said, “We’ll always keep a number of exploratory options available to us.” User feedback will continue to guide Anthropic’s product roadmap, he added.

To meet the massive costs of developing frontier AI models, Anthropic is exploring various new revenue channels, one of which is [Claude for Education](#). Claude for Education targets university customers, providing specific capabilities and benefits for colleges.

Advertisement

Newsletters

See More ↗

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

Anthropic declined to tell TechCrunch how many Claude subscriptions it has sold to date. However, White said the company's [new Claude 3.7 Sonnet AI model](#) has created "a lot of demand" for its products. Claude 3.7 Sonnet is Anthropic's first reasoning model, a type of model that uses more computing power than traditional models to answer certain questions more reliably.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Topics: [AI](#) [Anthropic](#) [Claude](#) [Startups](#)




Maxwell Zeff

Senior AI Reporter |

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >



Anthropic on Thursday [launched](#) a new way to connect apps and tools to its AI chatbot Claude, as well as an expanded “deep research” capability that allows Claude to search the web, enterprise accounts, and more.

The new app connection feature, called Integrations, and expanded deep research tool, dubbed Advanced Research, are available in beta for subscribers to Anthropic’s Claude Max, Team, and Enterprise plans, and soon Pro. In related news, Anthropic has also upped the rate limits for its AI-

Latest in Apps

[See More](#)

APPS



The Browser Company launches a \$20 monthly subscription for its AI-powered browser

Ivan Mehta - 1 hour ago

powered coding tool, [Claude Code](#), for Max customers.

Integrations and Advanced Research are both a part of Anthropic's effort to keep pace with competing chatbots like Google's Gemini and OpenAI's ChatGPT. Anthropic [reportedly](#) hopes to reach \$34.5 billion in revenue in 2027, and while it's certainly [making inroads](#), it has a long way to go. Anthropic's annualized revenue was around \$1.4 billion in early March, [according to one report](#).

Integrations builds on the company's [MCP protocol](#), which lets AI models draw data from sources such as business tools to complete tasks, as well as from content repositories and app development environments. Integrations taps MCP to allow developers to create and host app servers that enhance Claude's capabilities and let users discover and connect these servers to Claude.

“When you connect your tools to Claude, it gains deep context about your work — understanding project histories, task statuses, and organizational knowledge — and can take actions across every surface,” [wrote Anthropic in a blog post](#).

To start, there are a number of Integrations from Anthropic partners, including Atlassian, Zapier, Cloudflare, Intercom, Square, and PayPal. Each expands what Claude can do. For example, the Atlassian Integration lets Claude summarize and create pages in Atlassian's Confluence workplace software, while the Zapier Integration allows Claude to connect to the former's app automation workflows.


AI



Google takes on ChatGPT's Study Mode with new 'Guided Learning' tool in Gemini

Aisha Malik - 14 hours ago

APPS



Instagram takes on Snapchat with new 'Instagram Map'

Aisha Malik - 18 hours ago



As for the other new Claude feature launching today, Advanced Research, it allows Claude to crawl “hundreds” of internal and external sources to deliver what Anthropic describes as “more comprehensive” reports on a topic in anywhere from five to 45 minutes. Advanced Research can take advantage of Claude’s newly expanded connectors, Anthropic says, to search across Integrations and — when using the Claude Desktop app on macOS or Windows — MCP-connected local drives.

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

“With its new ability to do more complex research, available when you toggle on the Research button, Claude breaks down your request into smaller parts, investigating each deeply before compiling a comprehensive report,” continued Anthropic in the blog post. “When Claude incorporates information from sources, it provides clear citations that link directly to the original material.”

There’s been a raft of deep research tools launched recently across chatbots such as Gemini, Microsoft’s Copilot, and xAI’s Grok. Driving them are “reasoning” AI models, which possess the ability to think through problems and fact-check themselves — skills arguably important for conducting in-depth research on a subject.

Anthropic’s Research tool previously was among the faster tools in the deep research category, taking around a minute to compile reports. But the results tended to be shallower, in part because Research didn’t employ a reasoning model.

Topics: [AI](#) [Anthropic](#) [Apps](#) [Claude](#)



Kyle Wiggers

AI Editor |

Kyle Wiggers was TechCrunch’s AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

[View Bio](#) >



IMAGE CREDITS: ANTHROPIC

AI



Anthropic launches a new AI model that ‘thinks’ as long as you want

Maxwell Zeff — 10:30 AM PST · February 24, 2025

Anthropic is releasing a new frontier AI model called Claude 3.7 Sonnet, which the company designed to “think” about questions for as long as users want it to.

Anthropic calls Claude 3.7 Sonnet the industry’s first “hybrid AI reasoning model,” because it’s a single model that can give both real-time answers and more considered, “thought-out” answers to questions. Users can choose whether to activate the AI model’s “reasoning” abilities, which prompt Claude 3.7 Sonnet to “think” for a short or long period of time.

TC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

The model represents Anthropic’s broader effort to simplify the user experience around its AI products. Most AI chatbots today have a daunting model picker that forces users to choose from several different options that vary in cost and capability. Labs like Anthropic would rather you not have to think about it — ideally, one model does all the work.

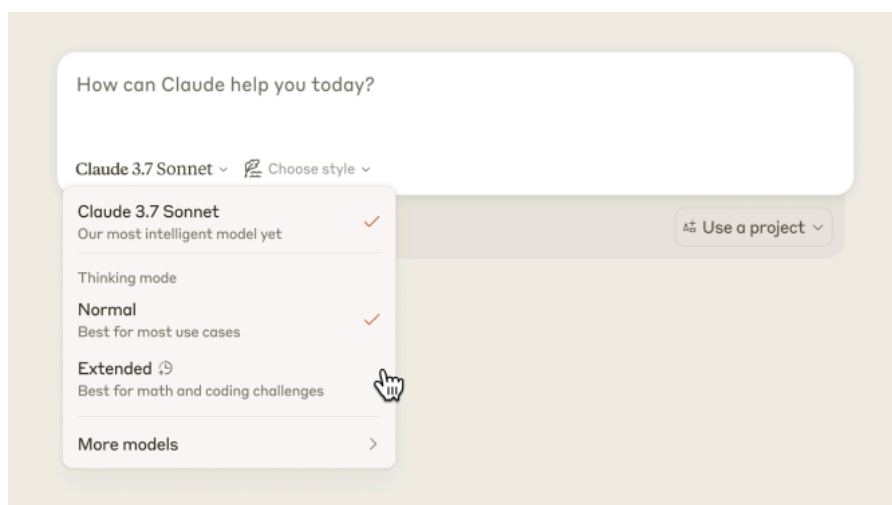
Claude 3.7 Sonnet is rolling out to all users and developers on Monday, Anthropic said, but only people who pay for Anthropic’s premium Claude chatbot plans will get access to the model’s reasoning features. Free Claude users will get the standard, non-reasoning version of Claude 3.7 Sonnet, which Anthropic claims outperforms its previous frontier AI model, [Claude 3.5 Sonnet](#). (Yes, the company skipped a number.)

Claude 3.7 Sonnet costs \$3 per million input tokens (meaning you could enter roughly 750,000 words, more words than the entire “Lord of the Rings” series, into Claude for \$3) and \$15 per million output tokens. That makes it more expensive than OpenAI’s o3-mini (\$1.10 per 1 million input tokens/\$4.40 per 1 million output tokens) and DeepSeek’s R1 (55 cents per 1 million input tokens/\$2.19 per 1 million output tokens), but keep in mind that o3-mini and R1 are strictly reasoning models — not hybrids like Claude 3.7 Sonnet.

Most Popular



- Key sections of the US Constitution deleted from government’s website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue’s AI-generated ad isn’t just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines
- The best dating apps aren’t even dating apps



Claude 3.7 Sonnet is Anthropic's first AI model that can “reason,” a technique [many AI labs have turned to as traditional methods of improving AI performance taper off.](#)

Reasoning models like o3-mini, R1, Google's Gemini 2.0 Flash Thinking, and xAI's Grok 3 (Think) use more time and computing power before answering questions. The models break problems down into smaller steps, which tends to improve the accuracy of the final answer. Reasoning models aren't thinking or reasoning like a human would, necessarily, but their process is modeled after deduction.

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025


[REGISTER NOW](#)

Eventually, Anthropic would like Claude to figure out how long it should “think” about questions on its own, without needing users to select controls in advance, Anthropic's product and research lead, Dianne Penn, told TechCrunch in an interview.

“Similar to how humans don't have two separate brains for questions that can be answered

immediately versus those that require thought,” Anthropic wrote in a [blog post](#) shared with TechCrunch, “we regard reasoning as simply one of the capabilities a frontier model should have, to be smoothly integrated with other capabilities, rather than something to be provided in a separate model.”

Anthropic says it’s allowing Claude 3.7 Sonnet to show its internal planning phase through a “visible scratch pad.” Penn told TechCrunch users will see Claude’s full thinking process for most prompts, but that some portions may be redacted for trust and safety purposes.



CLAUDE’S THINKING PROCESS IN THE CLAUDE APP
IMAGE CREDITS: ANTHROPIC

Anthropic says it optimized Claude’s thinking modes for real-world tasks, such as difficult coding problems or agentic tasks. Developers tapping Anthropic’s API can control the “budget” for thinking, trading speed, and cost for quality of answer.

On one test to measure real-world coding tasks, SWE-Bench, Claude 3.7 Sonnet was 62.3% accurate, compared to OpenAI’s o3-mini model which scored 49.3%. On another test to measure an AI model’s ability to interact with simulated users and external APIs in a retail setting, TAU-Bench, Claude 3.7 Sonnet scored 81.2%, compared to OpenAI’s o1 model which scored 73.5%.

Anthropic also says Claude 3.7 Sonnet will refuse to answer questions less often than its previous models, claiming the model is capable of making more nuanced distinctions between harmful and benign prompts. Anthropic says it reduced unnecessary refusals by 45% compared to Claude 3.5 Sonnet. This comes at a time when [some other AI labs are rethinking their approach to restricting their AI chatbot's answers](#).

In addition to Claude 3.7 Sonnet, Anthropic is also releasing an agentic coding tool called Claude Code. Launching as a research preview, the tool lets developers run specific tasks through Claude directly from their terminal.

In a demo, Anthropic employees showed how Claude Code can analyze a coding project with a simple command such as, “Explain this project structure.” Using plain English in the command line, a developer can modify a codebase. Claude Code will describe its edits as it makes changes, and even test a project for errors or push it to a GitHub repository.

Claude Code will initially be available to a limited number of users on a “first come, first serve” basis, an Anthropic spokesperson told TechCrunch.

Anthropic is releasing Claude 3.7 Sonnet at a time when AI labs are shipping new AI models at a breakneck pace. Anthropic has historically taken a more methodical, safety-focused approach. But this time, the company's looking to lead the pack.

For how long, though, is the question. [OpenAI may be close to releasing a hybrid AI model of its own](#); the company's CEO, Sam Altman, has said it'll arrive in “months.”

Topics: [AI](#) [Anthropic](#) [Claude](#)

Advertisement




Maxwell Zeff

Senior AI Reporter | X

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

View Bio >

Loading the next article



TechCrunch

Staff

Contact Us

Advertise

Crunchboard
Jobs

Site Map

Terms of
Service

Privacy
Policy

RSS Terms of
Use

OpenAI

Perplexity

Cisco

DeepMind

Windsurf

Tech Layoffs

ChatGPT