


Advertisement

AI



ElevenLabs launches an
AI music generator,
which it claims is cleared
for commercial use



The AI audio-generation [unicorn ElevenLabs](#) announced a new model on Tuesday that allows users to [generate music](#), which it claims is cleared for commercial use.

This move marks ElevenLabs' expansion beyond its main focus thus far in its three years of existence, which has been building AI audio tools. ElevenLabs is a leader among companies making text-to-speech AI products, and it has expanded into [conversational bots](#) and tools that translate speech into other languages.

Alongside the launch, ElevenLabs shared samples of its [AI-generated music](#).

IMAGE CREDITS: ELEVENLABS

One features a synthetic voice rapping about how it “came up through the cracks with ambition in my pocket” and left its hometown, traveling from “Compton to the Cosmos.” It’s unsettling to hear a computer reflect the influence and language of artists like Dr. Dre, N.W.A., and Kendrick Lamar, who actually lived the experiences that this technology is attempting to emulate.

Given these concerns around what material AI music generation tools are trained on, it’s not so straightforward for startups to delve into music generation.

Last year, Suno and Udio were [sued](#) by the Recording Industry Association of America (RIAA),


TEC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government’s website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue’s AI-generated ad isn’t just about fashion

the trade organization that represents the U.S. music industry. These lawsuits allege that Suno and Udio trained their music-generation models on copyrighted material. The companies are now [reportedly](#) discussing licensing deals with major record labels.

ElevenLabs also announced deals with Merlin Network and Kobalt Music Group, two digital publishing platforms for independent musicians, to use their materials for AI training.

According to [Merlin's website](#), the company represents major artists like Adele, Nirvana, Mitski, Carly Rae Jepsen, and Phoebe Bridgers; [Kobalt](#) represents stars like Beck, Bon Iver, and Childish Gambino.

A representative from Kobalt told TechCrunch that artists have to voluntarily opt-in for their music to be licensed for AI use.

“Our clients benefit directly from this agreement in several key ways: it opens a new revenue stream in a growing market, includes revenue sharing so they participate in the upside, provides strong safeguards against infringement and misuse, and offers favorable terms comparable to other publishing and recording rightsholders,” the Kobalt representative told TechCrunch.

Updated, 12:50 PM ET, with comment from Kobalt.

Topics: [AI](#) [ai music](#) [ElevenLabs](#)

- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines



Amanda Silberling is a senior writer at TechCrunch covering the intersection of technology and culture. She...

[View Bio](#) >

KEEP READING

Need some inspiration?

Check out a range of different genders, nationalities and moods.




IMAGE CREDITS: PODCASTLE

AI



Podcasting platform Podcastle launches a text-to-speech model with more than 450 AI voices

Ivan Mehta — 10:24 AM PST · March 3, 2025

Podcast recording and editing platform Podcastle is now joining other companies in the AI-powered, text-to-speech race by releasing its own AI model called [Asyncflow v1.0](#). An API for developers will

also be available, allowing them to directly integrate the text-to-speech model in their apps.

Thanks to the new model, the company is able to offer more than 450 AI voices that can narrate your text. The startup said that it developed the technology and model in such a way that its training and inference costs are low, giving it an advantage against competitors.

With the move, Podcastle joins a number of startups, including ElevenLabs, Speechify, and WellSaid, that have developed technology and AI models to convert any kind of text into a voice clip narrated by AI. This technology spans use cases like marketing, advertisement, content creation, education, and corporate training.

Podcastle’s founder, Arto Yeritsyan, told TechCrunch that the company had always wanted to build a text-to-speech model, but the cost of training and data requirements for that were very high.

“We wanted to build a robust text-to-speech model since our inception. However, the costs of development were very high. Thanks to recent large language model developments, we were able to reach a breakthrough last year to get to a place where we could build a high-quality voice model without needing a ton of data,” Yeritsyan said.

The company was also aided in its efforts by its [\\$13.5 million Series A fundraise last year](#).

Yeritsyan said that while Podcastle charges around \$40 per 500 minutes of text-to-speech conversion, ElevenLabs charges \$99 for the same.

TC

Techcrunch event
Tech and VC heavyweights join the Disrupt 2025 agenda

Advertisement

Newsletters

[See More ↗](#)

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

Podcastle's voice cloning feature is getting an upgrade as well to create a quicker process for training.

Earlier, the training process involved reading roughly 70 different sentences. Now it just needs a few seconds of recording from you to create a clone of your voice. The new process also used [Podcastle's Magic Dust AI, which was released last year, to improve audio recording quality.](#)

Advertisement




IMAGE CREDITS:PODCASTLE

In our testing, the voice created with the new process sounded a bit robotic, though it mimicked our tone. The company said that, over time, it will improve the feature. Plus, you can train different samples of your voice to get different results.

SPONSORED



Design agents are here — Lovart launches the first end-to-end AI canvas for design

Sponsored by Lovart

Podcastle said that apart from costs, having tools for audio, video, podcasts, and AI-powered narration under one redesigned site will give it an edge over competitors. Yeritsyan said that while the majority of the users use Podcastle to work on audio content, video is catching up to it as well.

Topics: [AI](#) [AI](#) [ElevenLabs](#) [Podcastle](#)

[text-to-speech](#)



Ivan Mehta

[X](#) [in](#)

Ivan covers global consumer tech developments at TechCrunch. He is based out of India and has previously...

View Bio >

KEEP READING




IMAGE CREDITS: ELEVENLABS

AI



ElevenLabs now lets authors create and publish audiobooks on its own platform


Ivan Mehta — 7:45 PM PST · February 25, 2025

Voice AI company [ElevenLabs](#) is now letting authors publish AI-generated [audiobooks on its own Reader app](#), TechCrunch has learned and the company confirmed. The announcement comes [days after the company partnered with Spotify](#) for AI-narrated audiobooks.

Latest in AI

See More

ElevenLabs, which raised [a \\$180 million mega-round last month](#), started inviting authors to try out their publishing program through their app on a trial basis last year, TechCrunch previously spotted. That program is newly open to all authors as of today.



AI



Google denies AI search features are killing website traffic

Sarah Perez - 13 hours ago

AI



First impressions of Alexa+, Amazon's upgraded, AI-powered digital assistant

Sarah Perez - 14 hours ago


AI



Google takes on ChatGPT's Study Mode with new 'Guided Learning' tool in Gemini

Aisha Malik - 14 hours ago

Sponsored by Lovart



Latest

Security


Startups

AI

Venture

Apps

Apple



Events

Podcasts

Newsletters

Sponsored by Lovart

The platform itself aims to compete with Audible, which ElevenLabs believes offers lower royalty rates for authors. Under its model, ElevenLabs' audiobooks will be offered within its own Reader app and the company will pay authors when users engage with their content.

Currently, it pays roughly \$1.10 to authors when listeners engage with an audiobook for 11 minutes or more.



ElevenLabs said the average user spent 19 minutes listening to the published books on its app during the testing phase. While the startup thinks that these rates are among the best in the industry, they could still change as the program scales.

At launch, the payout is offered to authors in the U.S. and for English-only titles. Later, it aims to extend payouts to titles in the 32 languages it supports for audiobooks.

TC

Techcrunch event

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th

anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

The company also plans to create a marketplace where authors can sell their content.

The bigger opportunity for ElevenLabs involves authors and publishers generating audiobooks using its AI tech by way of its paid plans ranging from \$11 to \$330 per month. This is less expensive than booking studio time and paying voice actors.

Notably, ElevenLabs has already powered other audio platforms like [Pocket FM and Kuku FM](#) to turn text into audio content.

The company's move to become a publishing and distribution surface to host more indie content is in line with ElevenLabs CEO [Mati Staniszewski's plans to expand into more consumer experiences](#).

Topics: [AI](#) [Apps](#) [Artificial Intelligence \(AI\)](#) [Audible](#)

[audiobooks](#) [ElevenLabs](#)



Ivan Mehta

X in

Ivan covers global consumer tech developments at TechCrunch. He is based out of India and has previously...

View Bio >

KEEP READING




IMAGE CREDITS: PHONIC

AI

Voice AI platform Phonic gets backing from Lux

Ivan Mehta — 8:00 AM PDT · April 3, 2025

The quality of AI-generated voices is good enough for things like creating audiobooks and podcasts, having articles read aloud to you, and basic customer support. But many businesses don't think AI voice tech [is quite reliable enough to deploy](#).

That's why two MIT grads, Moin Nadeem and Nikhil Murthy (pictured above), founded [Phonic](#), a company offering an end-to-end voice stack to increase synthetic voice reliability while decreasing latency.

★ SPONSORED



Just launched: Lovart, the AI design agent building entire campaigns in 3 minutes

Sponsored by Lovart

Nadeem and Murthy met at MIT, and have known each other for more than seven years. When the duo started building Phonic last year, they felt there weren't many companies crafting complete voice tech solutions.

"Voice AI is at a place where you tie up different parts, such as automatic voice recognition [and] text-to-speech, and [then integrate] intelligence," Murthy told TechCrunch. "However, when we talked to actual customers, we found that there is a lack of [solutions] that [are] reliable at scale."

Nadeem, who previously worked at MosaicML, [a company Databricks acquired for \\$1.3 billion in 2023](#), said that a lot of companies that are building in the voice AI space (e.g. [Vapi](#), [Rounded](#)) are creating workflows to piece together separate AI models.

Phonic takes a different approach: It trains its models in-house end-to-end. Murthy said that there are a few advantages to this.


TC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion
- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

“Owning the models allows us to deeply integrate some [...] reliability pieces into the [models themselves],” he said. “If you don’t own that layer [...] you’re just tying disparate pieces that don’t really fit seamlessly.”

TC

Techcrunch event

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

Tech and VC heavyweights join the Disrupt 2025 agenda



Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

Murthy added that Phonic’s method also allows the company to host and run models cost-efficiently. He claims that Phonic trains its models on a range of recordings, including recordings of accented

Sponsored by Lovart



★ SPONSORED

Lovart’s design agent hints at the future: AI-native brands built in minutes

and muffled speech, to make the models highly robust.

Phonic is currently working with a limited set of partners, including companies in the insurance and healthcare spaces, but plans to launch its product broadly in a few months. Soon, prospective clients will be able to try out Phonic’s tech from its website, Nadeem said.

Phonic has raised \$4 million in a seed round led by Lux with participation from Replit co-founder Amjad Masad, Hugging Face co-founder Clem Delangue, Applied Intuition co-founder Qasar Younis, and Modal Labs founder Erik Bernhardsson.

★ SPONSORED



Just launched: Lovart, the AI design agent building entire campaigns in 3 minutes

Sponsored by Lovart

Grace Isford, a partner at Lux Capital, said that the company’s in-house way of training models was appealing to the investment firm.

“We think both Moin and Nikhil are incredible technologists,” she said. “They founded [a] machine learning club at MIT. And they have worked on training models for a while now. Plus, their approach of combining diffusion and proprietary models in the voice AI sector is novel.”

Topics: [AI](#) [ElevenLabs](#) [Fundraising](#) [lux capital](#)

[MIT](#) [voice AI](#)





Ivan Mehta

Ivan covers global consumer tech developments at TechCrunch. He is based out of India and has previously...

[View Bio](#) >

Loading the next article



[TechCrunch](#)
[Staff](#)

[Contact Us](#)

[Advertise](#)

[Crunchboard](#)
[Jobs](#)

[Site Map](#)

[Terms of](#)
[Service](#)

[Privacy](#)
[Policy](#)

[RSS Terms of](#)
[Use](#)

[OpenAI](#)

[Perplexity](#)

[Cisco](#)

[DeepMind](#)

[Windsurf](#)

[Tech Layoffs](#)

[ChatGPT](#)