

Advertisement

ENTERPRISE



Enterprises prefer
Anthropic's AI models
over anyone else's,
including OpenAI's

AI research lab Anthropic's AI models are now the top choice for enterprises, surpassing OpenAI.

Anthropic now holds 32% of the enterprise large language model market share by usage, according to a [report from Menlo Ventures](#) released on Thursday. OpenAI holds the second-largest market share by usage among enterprises, with 25%.

The figure marks a strong reversal from even just a couple of years ago. Since 2023, OpenAI has seen its market share among enterprises decline sharply, according to the report, as Anthropic's has steadily risen over the same timeframe. OpenAI held 50% of the enterprise market share by usage just two years ago while Anthropic had 12%.

Google has seen enterprise usage for its models increase over the last few years as well.

Anthropic has an even larger market share when it comes to coding, with 42% of the enterprise market share, the largest market share by a wide margin. Enterprise usage of Anthropic's AI models are more than double OpenAI's, when it comes to coding, which garnered 21% of overall market share.

Anthropic's release of its Claude 3.5 Sonnet model in June 2024 is what laid the foundation for the company's surge in usage, according to the report. The release of [Claude 3.7 Sonnet](#) in February 2025 only accelerated that momentum.

TEC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion

The findings from Menlo Ventures align with anecdotal chatter in the industry, which suggested that enterprise and startup developers preferred Claude over OpenAI's ChatGPT. Meanwhile, OpenAI has a strong foothold on the consumer side of the house. The company reported last week that its users send more than [2.5 billion prompts](#) to ChatGPT a day.

- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

The Menlo Ventures report found enterprises prefer closed models, which Anthropic and OpenAI use. More than half of enterprises replied that they don't use open source models at all. Only 13% of enterprise daily workloads use open source models as of mid-year 2025, down from 19% at the beginning of the year. Meta still maintains dominance in the open source market.

Topics: [AI](#) [Anthropic](#) [Enterprise](#)

[large language models](#) [menlo ventures](#) [Meta](#)



Rebecca Szkutak

Senior Reporter, Venture | [X](#) [in](#)

Becca is a senior writer at TechCrunch that covers venture capital trends and startups. She previously covered the...

[View Bio](#) >

KEEP READING



AI



The hottest AI models, what they do, and how to use them

Charles Rollet — 10:28 AM PDT · March 30, 2025

AI models are being cranked out at a dizzying pace, by everyone from Big Tech companies like Google to startups like OpenAI and Anthropic. Keeping track of the latest ones can be overwhelming.

Adding to the confusion is that AI models are often promoted based on industry benchmarks. But these [technical metrics often reveal little](#) about how real people and companies actually use them.

To cut through the noise, TechCrunch has compiled an overview of the most advanced AI models released since 2024, with details on how to use them and what they're best for. We'll keep this list updated with the latest launches, too.

There are literally over a million AI models out there: Hugging Face, for example, [hosts over 1.4 million](#). So this list might miss some models that perform better, in one way or another.

AI models released in 2025

Google Gemini 2.5

Gemini 2.5 Pro Experimental, a reasoning model, excels at [building web apps and code agents](#) according to Google. However, it underperforms on one popular coding benchmark compared to Claude Sonnet 3.7. The model requires a \$20 monthly Gemini Advanced subscription.

ChatGPT-4o image generator

OpenAI [has upgraded its existing GPT-4o model](#) to generate images, not just text. The souped-up model soon went viral for [transforming images into Studio Ghibli-style anime](#), despite obvious copyright concerns. Accessing GPT-4o requires, at minimum, a \$20 per month ChatGPT Plus subscription.

Stability AI’s Stable Virtual Camera

Image-generation startup Stability AI [has launched a model](#) that the company says can generate 3D scenes and camera angles from a single 2D image. However, it still struggles with scenes featuring more complex elements like humans and moving water. The model is available for noncommercial research use on HuggingFace.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.


San Francisco | October 27-29, 2025

REGISTER NOW

Cohere’s Aya Vision

Cohere [released a multimodal model](#) called Aya Vision that it claims is best in class at doing things like captioning images and answering questions about photos. It also excels in languages other than English, unlike other models, Cohere claims. It is available for [free on WhatsApp](#).

OpenAI’s GPT 4.5 “Orion”



Newsletters

See More ↗

Subscribe for the industry’s biggest tech news

TechCrunch Daily News

+

Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly

+

Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review

+

Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility

+

TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

OpenAI calls Orion their [largest model to date](#), touting its strong “world knowledge” and “emotional intelligence.” However, it underperforms on certain benchmarks compared to newer reasoning models. Orion is available to subscribers of OpenAI’s \$200-per-month plan.

Claude Sonnet 3.7

Anthropic says this is the [industry’s first “hybrid” reasoning model](#), because it can both fire off quick answers and really think things through when needed. It also gives users control over how long the model can think for, [per Anthropic](#). Sonnet 3.7 is available to all Claude users, but heavier users will need a \$20-per-month Pro plan.

xAI’s Grok 3

Grok 3 is the [latest flagship model](#) from Elon Musk-founded startup xAI. It’s [claimed to](#) outperform other leading models on math, science, and coding. The model requires X Premium (which is \$50 per month.) After one study [found](#) Grok 2 leaned left, Musk [pledged](#) to shift Grok more “politically neutral” but it’s not yet clear if that’s been achieved.

OpenAI o3-mini

This is OpenAI’s [latest reasoning model](#) and is optimized for STEM-related tasks like coding, math, and science. It’s [not OpenAI’s most powerful](#) model but because it’s smaller, the company [says](#) it’s significantly lower cost. It is available for free but requires a subscription for heavy users.

OpenAI Deep Research

Advertisement

OpenAI's Deep Research is [designed for doing in-depth research](#) on a topic with clear citations. This service is only available with ChatGPT's [\\$200-per-month Pro subscription](#). OpenAI [recommends it](#) for everything from science to shopping research, but beware that [hallucinations remain a problem](#) for AI.

Mistral Le Chat

Mistral has [launched app versions of Le Chat](#), a multimodal AI personal assistant. Mistral [claims](#) Le Chat responds faster than any other chatbot. It also has a paid version with [up-to-date journalism](#) from the AFP. [Tests from Le Monde](#) found Le Chat's performance impressive, although it made more errors than ChatGPT.

OpenAI Operator

OpenAI's Operator [is meant to be](#) a personal intern that can do things independently, like help you buy groceries. It requires a \$200-per-month ChatGPT Pro subscription. AI agents hold a lot of promise, but they're still experimental: A Washington Post reviewer [says Operator](#) decided on its own to order a dozen eggs for \$31, paid with the reviewer's credit card.

Google Gemini 2.0 Pro Experimental

Google Gemini's [much-awaited flagship model](#) says it excels at coding and understanding general knowledge. It also has a super-long context window of 2 million tokens, helping users who need to quickly process massive chunks of text. The service requires (at minimum) a Google One AI Premium subscription of \$19.99 a month.

AI models released in 2024

DeepSeek R1

This [Chinese AI model](#) took [Silicon Valley by storm](#). DeepSeek's R1 performs well on coding and math, while its open source nature means anyone can run it locally. Plus, it's free. However, [R1 integrates](#) Chinese government censorship and [faces rising bans](#) for potentially sending user data back to China.

Gemini Deep Research

Deep Research [summarizes Google's search results](#) in a simple and well-cited document. [The service](#) is helpful for students and anyone else who needs a quick research summary. However, its quality isn't nearly as good as an actual peer-reviewed paper. Deep Research requires a \$19.99 Google One AI Premium subscription.

Meta Llama 3.3 70B

This is the [newest and most advanced version](#) of Meta's open source Llama AI models. Meta has touted [this version](#) as its cheapest and most efficient yet, especially for math, general knowledge, and instruction following. It is free and open source.

OpenAI Sora

Sora is a model that [creates realistic videos](#) based on text. While it can generate entire scenes rather than just clips, [OpenAI admits](#) that it often generates "unrealistic physics." It's currently only available on paid versions of ChatGPT, starting with Plus, which is \$20 a month.

Alibaba Qwen QwQ-32B-Preview

This model is [one of the few to rival](#) OpenAI's o1 on certain industry benchmarks, excelling in math and coding. Ironically for a "reasoning model," it has "room for improvement in common sense reasoning," [Alibaba says](#). It also incorporates Chinese government censorship, [TechCrunch testing shows](#). It's free and open source.

Anthropic's Computer Use

Claude's Computer Use is meant to [take control of your computer](#) to complete tasks like coding or booking a plane ticket, making it a predecessor of OpenAI's Operator. Computer use, however, [remains in beta](#). Pricing is via API: \$0.80 per million tokens of input and \$4 per million tokens of output.

xAI's Grok 2

Elon Musk's AI company, xAI, has launched an [enhanced version of its flagship Grok 2 chatbot](#) it [claims](#) is "three times faster." Free users are limited to 10 questions every two hours on Grok, while subscribers to X's Premium and Premium+ plans enjoy higher usage limits. xAI also launched an image generator, Aurora, that [produces highly photorealistic images](#), including some graphic or violent content.

OpenAI o1

[OpenAI's o1 family](#) is meant to produce better answers by "thinking" through responses through a hidden [reasoning feature](#). The model excels at coding, math, and safety, [OpenAI claims](#), but has [issues with trying to deceive humans, too](#). Using o1 requires subscribing to ChatGPT Plus, which is \$20 a month.

Anthropic's Claude Sonnet 3.5

Claude Sonnet 3.5 is a model Anthropic [claims as being best in class](#). It's become known for its coding capabilities and is considered a tech insider's [chatbot of choice](#). The model can be accessed for free on Claude, although heavy users will need a \$20 monthly Pro subscription. While it can understand images, it can't generate them.

OpenAI GPT 4o-mini

OpenAI has [touted GPT 4o-mini](#) as its most affordable and fastest model yet, thanks to its small size. [It's meant](#) to enable a broad range of tasks like powering customer service chatbots. The model is available on ChatGPT's free tier. It's better suited for high-volume simple tasks compared to more complex ones.

Cohere Command R+

Cohere's [Command R+ model](#) excels at complex retrieval-augmented generation (or RAG) applications for enterprises. That means it can find and cite specific pieces of information really well. (The inventor of RAG [actually works at Cohere](#).) Still, RAG [doesn't fully solve AI's hallucination problem](#).

Topics: [AI](#) [AI](#) [ai models](#) [Anthropic](#) [ChatGPT](#)

[deepseek](#) [evergreens](#) [google gemini](#)

[mistral](#) [OpenAI](#) [sora](#) [xAI](#)

[f](#) [X](#) [in](#) [📷](#) [✉](#) [🔗](#)



Charles Rollet

Senior Reporter | [X](#)

Charles Rollet is a senior reporter at TechCrunch. His investigative reporting has led to U.S. government...

View Bio >

KEEP READING

- Latest
- Security
- Startups
- AI
- Venture
- Apps
- Apple

- Events
- Podcasts
- Newsletters




IMAGE CREDITS:

KIMBERLY WHITE/GETTY IMAGES FOR TECHCRUNCH

compete with OpenAI

Maxwell Zeff — 8:00 AM PDT · September 4, 2024

Anthropic is launching a new subscription plan for its AI chatbot, Claude, catered toward enterprise customers that want more administrative controls and increased security. Claude Enterprise will compete with OpenAI’s business-specific solution, [ChatGPT Enterprise](#), released roughly a year ago.

Latest in Enterprise

See More

Claude Enterprise allows businesses to upload proprietary company knowledge into Anthropic's AI chatbot. Then Claude can analyze the information, answer questions about it, create graphics and simple web pages, or act as a company-specific AI assistant.

★ SPONSORED



Design agents are here — Lovart launches the first end-to-end AI canvas for design

Sponsored by Lovart

Anthropic appears to be playing catch-up with OpenAI, trying to put Claude everywhere that ChatGPT already is. The startup has released a few ways to use Claude that closely match how OpenAI already offers ChatGPT.

“The reality is that Claude has been usable for companies for a year. Candidly, we’ve had a product in the market for a lot less long,” Anthropic product lead Scott White told TechCrunch. “But we’re responding to the needs of our customers at a high velocity with a smaller team.”

In May, Anthropic released the [Claude Team plan](#) which, much like the ChatGPT team plan, allowed small businesses to collaborate on projects. Since the spring, Anthropic has launched Claude mobile apps for iOS and Android. Now, it’s going up against ChatGPT Enterprise, which has seen widespread adoption among Fortune 500 companies.

But Anthropic’s enterprise offering is different from what’s on the market in a few key ways. First, the context window on Claude Enterprise is 500,000 tokens, meaning Anthropic’s models can process up to 200,000 lines of code, dozens of 100-page documents or a two-hour audio transcript in a single prompt. ChatGPT Enterprise and Claude’s Team plan offer context windows less than half that.

ENTERPRISE



Upwork is buying its way into corporate staffing beyond freelancers

Rebecca Szkutak

- 11 hours ago

STARTUPS



Figma’s IPO success is ‘a little bit of a meme stock,’ says Sapphire Ventures’ Jai Das

Rebecca Bellan,

Theresa Loconsolo

- 14 hours ago



AI



Cohere’s new AI agent platform, North, promises to keep enterprise data secure

Rebecca Bellan - 17 hours ago

Sponsored by Lovart



★ SPONSORED

Design agents are here — Lovart launches the first end-to-end AI canvas for design

Claude Enterprise comes with Projects and Artifacts, Anthropic's workspaces where several users can upload and edit content. These features could be useful in business cases, where you might be working on a longer project with lots of data sources and people involved.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

[REGISTER NOW](#)

The Enterprise plan also includes GitHub integration for engineering teams to sync their GitHub repositories with Claude. Coding has become a popular use case for Claude Sonnet 3.5, and this integration means that Anthropic's models can have direct access to its customers' codebases. This can be useful for catching a new engineer up to speed, creating a new feature or solving a bug, for instance.

Like ChatGPT's Enterprise plan, Claude Enterprise allows businesses to assign a primary owner for your company's workspace. That owner can assign different levels of access to projects and

information within Claude, and trace activity across the system for security and compliance monitoring.

Also like OpenAI, Anthropic says it is not training on Claude Enterprise customer data. That's important for many businesses that don't want their trade secrets ending up in Claude or ChatGPT's knowledge base six months from now.

★ SPONSORED



Design agents are here — Lovart launches the first end-to-end AI canvas for design

Sponsored by Lovart

Anthropic refused to disclose the price of Claude Enterprise, though White said it's more expensive than Anthropic's Team plan (which costs \$30 per month, per member). White said that's because Enterprise customers get a lot more mileage out of Claude, with larger context windows and higher rate limits. (For what it's worth, OpenAI won't publicly disclose the price of its enterprise product either.)

White says Anthropic has been working in a private beta for months with early adopters such as GitLab, Midjourney, IG Group and Menlo Ventures (an investor in Anthropic).

However, gaining broader adoption will be key. AI model developers like Anthropic have faced pressure to [sell API access at lower and lower prices](#). Products like Claude Enterprise offer a path to revenue, however, widespread adoption is needed to offset the expensive inference costs that come with them. It's not clear that any AI model developers are profiting off of these business-specific plans just yet.

Topics: [AI](#) [Anthropic](#) [claude 3.5 sonnet](#)



Maxwell Zeff

Senior AI Reporter |

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >

KEEP READING



STARTUPS



Anthropic begins supplying its text-generating AI models to select startups

Kyle Wiggers — 4:00 AM PST · February 27, 2023

Anthropic, a buzzy AI startup co-founded by ex-OpenAI employees, has begun offering partners access to its AI text-generating models.

The first commercial venture to announce that it's integrating Anthropic models is **Robin AI**, a legal tech startup that's raised over \$13 million from investors including Plural, Episode 1 and the Google Black Founders Fund. Quora's experimental chatbot app for iOS and Android, **Poe**, uses Anthropic models, but it's not currently monetized.

★ SPONSORED



Just launched: Lovart, the AI design agent building entire campaigns in 3 minutes

Sponsored by Lovart

Robin CEO Richard Robinson disclosed few details regarding the Anthropic relationship, but told TechCrunch that Robin worked to fine-tune an Anthropic model on a dataset of legal text to draft and negotiate contracts.

“We are very fortunate to be Anthropic’s launch partner for the legal sector — the team’s focus on AI safety aligns with our ‘lawyer-in-the-loop’ software-as-a-service product — deliberately designed to manage the risk of even the most advanced models ‘hallucinating,’” Robinson said in a statement.

Join the waitlist

Be the first to try Robin's AI-powered contract editing software



A computer monitor displaying the Robin AI interface. The screen shows a document titled 'NON-DISCLOSURE AND CONFIDENTIALITY AGREEMENT' with a sidebar on the right containing various document sections like 'Introduction', 'Definitions', 'Obligations', etc. The interface is clean and professional, with a blue header and a white body.

ROBIN AI IS ONE OF THE FIRST COMMERCIAL VENTURES TO USE ANTHROPIC MODELS.


TECH DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular





- Key sections of the US Constitution deleted from government’s website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue’s AI-generated ad isn’t just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines
- The best dating apps aren’t even dating apps

Anthropic has been relatively quiet about its plans to productize its work in the generative text AI space, preferring instead to focus on academic research. But late last year, the company launched a closed beta for an AI system, called [Claude](#), similar to OpenAI's [ChatGPT](#) that appeared to improve upon the original in key ways. Claude was created using a technique Anthropic developed called “constitutional AI,” which aims to provide a “principle-based” approach to aligning AI systems with human intentions — letting AI similar to ChatGPT respond to questions using a simple set of principles (e.g. avoid giving harmful advice) as a guide.

Preliminary impressions of Claude were good. But like ChatGPT, the system suffered from limitations, like giving dangerous answers to questions (e.g. how to make meth at home) and making inconsistent, factually wrong statements.

It's unclear whether the model Robin's using *is* Claude or some derivative — neither Robin nor Anthropic would say. And even after repeated prodding, Anthropic wouldn't reveal how many partners it's currently working with (or how they came to work with them) and how many models it plans to open up to commercial usage.



★ SPONSORED

Lovart's design agent hints at the future: AI-native brands built in minutes

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from

the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

But no doubt, Anthropic is feeling some sort of pressure from investors to recoup the hundreds of millions of dollars that’ve been put toward its AI tech.

Most recently, Google pledged \$300 million in Anthropic for a 10% stake in the startup. Under the terms of the deal, which was first [reported](#) by the Financial Times, Anthropic agreed to make Google Cloud its “preferred cloud provider” with the companies “co-develop[ing] AI computing systems.”

Anthropic wasn’t founded exclusively with a profit-driven mission. Dario Amodei, the former VP of research at OpenAI, launched the company in 2021 as a public benefit corporation, taking with him a number of OpenAI employees, including OpenAI’s former public policy lead Jack Clark. Amodei split from OpenAI after a disagreement over the company’s direction, namely the startup’s increasingly commercial focus.

★ SPONSORED



Design agents are here — Lovart launches the first end-to-end AI canvas for design

Sponsored by Lovart

But AI systems are expensive to develop and maintain. Ballooning costs led Anthropic to pursue outside backing, [including](#) a \$580 million tranche from a group of investors including disgraced FTX founder Sam Bankman-Fried, Caroline Ellison, Jim

McClave, Nishad Singh, Jaan Tallinn and the Center for Emerging Risk Research.

Whether startup partnerships — and Big Tech investments — denote a shift in Anthropic’s priorities isn’t yet clear. But what *is* clear is that the company believes its technology is differentiated to compete with rivals like OpenAI, Cohere and AI21 Labs, all of which offer paid access to their text-generating AI via APIs.

Topics: [AI](#) [AI](#) [Anthropic](#) [Artificial Intelligence \(AI\)](#)

[Generative AI](#) [machine learning](#) [Startups](#)



Kyle Wiggers

AI Editor |

Kyle Wiggers was TechCrunch’s AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

[View Bio](#) >

Loading the next article



TechCrunch
Staff

[Contact Us](#)

[Advertise](#)

[Crunchboard
Jobs](#)

[Site Map](#)

[Terms of
Service](#)

[Privacy
Policy](#)

[RSS Terms of
Use](#)

[OpenAI](#)
[Perplexity](#)

[Cisco](#)
[DeepMind](#)

[Windsurf](#)
[Tech Layoffs](#)
[ChatGPT](#)

