

Advertisement



IMAGE CREDITS: ANTHROPIC

ENTERPRISE



Enterprises prefer Anthropic's AI models over anyone else's, including OpenAI's

Rebecca Szkutak — 1:32 PM PDT · July 31, 2025

AI research lab Anthropic's AI models are now the top choice for enterprises, surpassing OpenAI.

Anthropic now holds 32% of the enterprise large language model market share by usage, according to a [report from Menlo Ventures](#) released on Thursday. OpenAI holds the second-largest market share by usage among enterprises, with 25%.

The figure marks a strong reversal from even just a couple of years ago. Since 2023, OpenAI has seen its market share among enterprises decline sharply, according to the report, as Anthropic's has steadily risen over the same timeframe. OpenAI held 50% of the enterprise market share by usage just two years ago while Anthropic had 12%.

Google has seen enterprise usage for its models increase over the last few years as well.

Anthropic has an even larger market share when it comes to coding, with 42% of the enterprise market share, the largest market share by a wide margin. Enterprise usage of Anthropic's AI models are more than double OpenAI's, when it comes to coding, which garnered 21% of overall market share.

Anthropic's release of its Claude 3.5 Sonnet model in June 2024 is what laid the foundation for the company's surge in usage, according to the report. The release of [Claude 3.7 Sonnet](#) in February 2025 only accelerated that momentum.

The findings from Menlo Ventures align with anecdotal chatter in the industry, which suggested that enterprise and startup developers preferred Claude over OpenAI's ChatGPT. Meanwhile, OpenAI has a strong foothold on the consumer side of the house. The company reported last week that its users send more than [2.5 billion prompts](#) to ChatGPT a day.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

TC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

Book Your Table >

Most Popular



- Key sections of the US Constitution deleted from government's website
- Instagram takes on Snapchat with new 'Instagram Map'
- Google says hackers stole its customers' data by breaching its Salesforce database
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines

The Menlo Ventures report found enterprises prefer closed models, which Anthropic and OpenAI use. More than half of enterprises replied

that they don't use open source models at all. Only 13% of enterprise daily workloads use open source models as of mid-year 2025, down from 19% at the beginning of the year. Meta still maintains dominance in the open source market.

Topics: [AI](#) [Anthropic](#) [Enterprise](#) [large language models](#) [menlo ventures](#)

[Meta](#) [OpenAI](#) [TC](#) [United States](#)

[f](#) [X](#) [in](#) [G](#) [M](#) [L](#)



Rebecca Szkutak

Senior Reporter, Venture | [in](#)

Becca is a senior writer at TechCrunch that covers venture capital trends and startups. She previously covered the same beat for Forbes and the Venture Capital...

[View Bio](#) >

KEEP READING



IMAGE CREDITS: GETTY IMAGES

AI

[f](#) [X](#) [in](#) [G](#) [M](#) [L](#)

OpenAI is practically giving ChatGPT to the government for free

Rebecca Bellan — 10:09 AM PDT · August 6, 2025

OpenAI is poised to undercut rivals like Anthropic and Google in the race to see its AI tools integrated into federal agency workflows.

The AI giant has [reached an agreement](#) with the U.S. General Services Administration (GSA), the government’s central purchasing arm, to offer ChatGPT Enterprise to participating federal agencies for just \$1 per agency for the next year.

The partnership comes a day after the [GSA added OpenAI, Google, and Anthropic](#) to a list of approved AI vendors that can offer their services to civilian federal agencies. The companies will see their tools offered via the Multiple Award Schedule (MAS), a federal contracting platform that lets government agencies access AI tools via pre-negotiated contracts so they don’t have to negotiate with vendors individually.

It’s not clear whether other AI firms are going to offer their services at such a discounted rate, though GSA Federal Acquisition Service commissioner Josh Gruenbaum in a [statement](#) encouraged “other American AI technology companies to follow OpenAI’s lead and work with us.”

TechCrunch has reached out to Anthropic and Google for more information.

In addition to access to ChatGPT Enterprise, OpenAI is offering unlimited use of advanced models for an additional 60 days. Federal employees will also have access to a new government user community and tailored introductory training resources to get familiar with OpenAI’s tools.

Data security is a top priority for government agencies that are concerned about sensitive information being leaked into model training sets. TechCrunch has asked the GSA for details on how government data is being safeguarded, including whether measures like on-premises or private cloud deployments are being used to enhance security.

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

“The government is taking a cautious, security-first approach to AI,” a GSA spokesperson told TechCrunch. “This ensures sensitive information

Advertisement

Newsletters

[See More ↗](#)

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

Advertisement

remains protected while enabling agencies to benefit from AI-driven efficiencies.”

OpenAI’s discount comes a couple of weeks after the Trump administration published its AI Action Plan that seeks to boost data center buildouts and integrate more AI tools into government, among other things. It also follows a Trump executive order that [bans “woke AI”](#) and AI models that aren’t “ideologically neutral” from government contracts.

TechCrunch has reached out to OpenAI to learn more about its GSA partnership and how it will approach Trump’s executive order.

Topics: [AI](#) [ChatGPT Enterprise](#) [Government & Policy](#) [GSA](#) [OpenAI](#)

Advertisement



Rebecca Bellan

Senior Reporter | [X](#) [🦋](#)

Rebecca Bellan is a senior reporter at TechCrunch where she covers the business, policy, and emerging trends shaping artificial intelligence. Her work has also...

[View Bio](#) >

KEEP READING

AI



OpenAI's GPT-5 is here

Maxwell Zeff — 10:00 AM PDT · August 7, 2025

GPT-5 Flagship model

OpenAI has launched GPT-5, a new flagship AI model that will power the company's next generation of ChatGPT.

GPT-5, which was released Thursday, is OpenAI's first "unified" AI model and combines the reasoning abilities of its [o-series](#) of models with the fast responses of its [GPT series](#). The next-generation model signals a new era for ChatGPT — and its creator, OpenAI — pointing to OpenAI's broader ambitions to develop AI systems that are [more like agents](#) than chatbots.

IMAGE CREDITS: OPENAI

While [GPT-4](#) enabled AI chatbots to offer smart responses on a wide variety of questions, GPT-5 allows ChatGPT to complete a wide variety of tasks on behalf of users — such as generating software applications, navigating a user's calendar, or creating research briefs.

With GPT-5, OpenAI has also sought to make ChatGPT simpler to use. Instead of asking users to choose the right settings, GPT-5 comes equipped with a real-time router that decides how to offer the best answer, whether that's responding to user questions quickly or taking additional time to "think" through answers.

GPT-5 Flagship model

Advanced creative writing skills

Stronger coding capabilities

Analyzing



Advertisement

Latest in AI

[See More](#)

TRANSPORTATION



Tesla shuts down Dojo, the AI training supercomputer that Musk said would be key to full self-driving

Rebecca Bellan - 6 hours ago

STARTUPS



High costs and thin margins threatening AI coding startups

Marina Temkin - 7 hours ago

AI



OpenAI's GPT-5 is here

Maxwell Zeff - 11 hours ago

IMAGE CREDITS: OPENAI

During a briefing with reporters, OpenAI CEO Sam Altman claimed GPT-5 is “the best model in the world,” and said it represented a “significant step” along the company’s path to developing AI that can outperform humans at most economically valuable work — that is, artificial general intelligence (AGI).

“Having something like GPT-5 would be pretty much unimaginable at any previous time in history,” said Altman.

Starting Thursday, GPT-5 will be available to all free users of ChatGPT as their default model. OpenAI’s VP of ChatGPT, Nick Turley, said this is part of the company’s effort to give free users access to an AI reasoning model for the first time. (Previously, the company gated these more advanced models behind a paywall.)

“This is just one of the ways that I’m excited to live the mission, making sure that this stuff actually benefits people,” said Turley on the decision, referencing OpenAI’s long-standing [mission](#) to distribute advanced AI to as many people as possible.

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

The expectations are high for GPT-5, one of OpenAI’s most anticipated product launches since ChatGPT put the company on the map in 2022. Since then, ChatGPT has grown into one of the world’s most popular consumer products, reaching more than 700 million users every week — nearly 10% of the globe’s population, according to the company.

Many see GPT-5 as a bellwether for AI progress broadly, and the model’s reception by Silicon Valley could have profound implications for Big Tech, Wall Street, and policymakers regulating technology. These stakeholders are watching to see if GPT-5 offers a significant jump in AI’s capabilities, much like its predecessor, GPT-4, which challenged expectations of what software can do.

GPT-5 offers a slight edge on the competition

OpenAI claims GPT-5 is state-of-the-art in several domains, slightly edging out leading AI models from Anthropic, Google DeepMind, and Elon Musk’s xAI on key benchmarks. However, GPT-5 slightly underperforms frontier AI models in other areas.

The company says GPT-5 offers frontier-level performance around coding; Altman said the model specifically excels at spinning up entire software applications on demand, in what’s become known as “vibe coding.”

On SWE-bench Verified — a test of real-world coding tasks pulled from GitHub — GPT-5 scores 74.9% on its first attempt. That means GPT-5 just outperforms Anthropic’s latest Claude Opus 4.1 model, which scored [74.5%](#), and Google DeepMind’s Gemini 2.5 Pro, which scored [59.6%](#).

On Humanity’s Last Exam — a difficult test measuring AI model performance across math, humanities, and the natural sciences — a version of GPT-5 with extended reasoning (GPT-5 Pro) scored 42% when using tools. That’s slightly less than xAI was able to achieve with [Grok 4 Heavy](#), which scored 44.4% on the test.

Humanity’s Last Exam (Full Set)*
Expert-level questions across subjects




IMAGE CREDITS: OPENAI

On GPQA Diamond — a test of PhD-level science questions — GPT-5 pro scored 89.4% on its first try, outperforming Claude Opus 4.1, which scored 80.9%, and Grok 4 Heavy, which scored 88.9%.

OpenAI says GPT-5 is better for answering health-related questions. On a test measuring accuracy in AI model responses around healthcare topics, HealthBench Hard Hallucinations, OpenAI says GPT-5 (with thinking) hallucinates just 1.6% of the time. This is far lower than the company’s previous GPT-4o and o3 models, which scored 12.9% and 15.8, respectively.

While AI chatbots are not medical professionals, [millions of people](#) are using them for health advice. In response to this phenomenon, the

company says GPT-5 is more proactive about flagging potential health concerns and helping users parse medical results.

In addition, OpenAI says GPT-5 is better than other AI models on more difficult-to-measure, [subjective domains](#), such as creative design and writing. Turley said GPT-5 responds more naturally and exhibits “better taste” than other AI models on creative tasks.

“The vibes of this model are really good,” said Turley.

GPT-5 is also more accurate than OpenAI’s previous models, and the company says it suffers far less from hallucinations — the tendency for AI models to make up information — compared to its o-series models. Hallucinations seemed to be [getting worse](#) in OpenAI’s latest AI reasoning models, such as o3, and OpenAI previously said it didn’t quite understand why it was happening.

In responses to ChatGPT prompts, OpenAI found that GPT-5 (with thinking) hallucinates and responds with incorrect information 4.8% of the time. That’s a significant reduction from o3 and GPT-4o, which score hallucination rates of 22% and 20.6%, respectively, on the test.

On a benchmark measuring an AI model’s agentic ability to complete simulated online tasks, Tau-bench, GPT-5 offers mixed performance. On part of the test measuring an AI’s ability to navigate an airline’s website, GPT-5 scores 63.5%, slightly underperforming o3, which scored 64.8%. On another part of the test measuring AI’s ability to navigate retail websites, GPT-5 scores 81.1%, underperforming Claude Opus 4.1, which scored 82.4%.

OpenAI also says that GPT-5 is safer than its previous models. While AI

TE

Latest

Startups

Venture

Apple

Security

AI

Apps

Events

Podcasts

Newsletters

Upgrades for consumers and developers

ChatGPT is getting a few user experience upgrades as part of the GPT-5 launch. Users can now select from four new personalities in ChatGPT’s setting: Cynic, Robot, Listener, and Nerd. The company says these will

adapt ChatGPT’s responses without requiring users to specifically ask the model to respond in a certain way.

Subscribers to ChatGPT’s \$20-per-month Plus plan get higher usage limits for GPT-5 than free users. Meanwhile, \$200-per-month Pro subscribers will have unlimited access to GPT-5, as well as a souped-up version called GPT-5 Pro that uses additional computational resources to produce better answers. Organizations on OpenAI’s Team, Edu, and Enterprise plans will gain access to GPT-5 as their default model next week.

For developers, GPT-5 is coming to OpenAI’s API in three sizes — gpt-5, gpt-5-mini, and gpt-5-nano — which will spend more or less time “reasoning” through tasks. Developers can also now control verbosity in the OpenAI API, deciding how long or short an AI model’s responses should be.

The base model of GPT-5 will cost developers \$1.25 per million input tokens (roughly 750,000 words, longer than the entire “Lord of the Rings” series) and \$10 per million output tokens.

The launch of GPT-5 comes after a busy week for OpenAI. The company released an open-weight reasoning model, [gpt-oss](#), that developers and enterprises can download for free and run at a fraction of the cost. The open model nearly matched the abilities of OpenAI’s previous top models, o3 and o4-mini, but GPT-5 sets a new standard for frontier performance in some areas, such as coding.

However, GPT-5 seems to be roughly on par with other frontier AI models in several areas. Benchmarks, of course, only tell part of the story for any AI model, and it remains to be seen how developers will use GPT-5 in the real world, and whether the model is truly a step above the competition.

Topics: [AI](#) [ChatGPT](#) [gpt-5](#) [OpenAI](#)

[f](#) [X](#) [in](#) [t](#) [✉](#) [🔗](#)



Maxwell Zeff

Senior AI Reporter | [X](#)

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg, and MSNBC, Zeff has covered the rise of AI and the Silicon...

[View Bio](#) >



IMAGE CREDITS: TAVILY

STARTUPS



Tavily raises \$25M to connect AI agents to the web

Marina Temkin — 8:15 AM PDT · August 6, 2025

Companies across many industries are implementing AI agents for internal use, automating a wide range of tasks.

In the financial sector, AI agents are critical for fraud detection. They can analyze vast amounts of transaction data in real time. Meanwhile, sales organizations are using AI agents to gather data on potential customers. These AI sales agents can scour the web and social media for information.

To be effective, these agents need to access the internet and find information from relevant sources, all while following company policies and mirroring how a human researcher would work.

Connecting an agent directly to a large language model like ChatGPT without company-specific safeguards can lead to highly inappropriate results.

“Governance, risk, and compliance at the enterprise is so important now, and if you just let that happen, it’s just going to be the wild, wild west,” George Mathew, managing director at Insight Partners, told TechCrunch.

That’s why Insight Partners led a \$20 million Series A in [Tavily](#), a startup that connects AI agents to the web in a way that’s compliant with

TC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government’s website
- Instagram takes on Snapchat with new ‘Instagram Map’
- Google says hackers stole its customers’ data by breaching

company-specific policies. The investment brings the 1-year-old Tavily’s total funding to \$25 million.

Founded last year by data scientist Rotem Weiss, Tavily began as an open source project he created in 2023 called GPT Researcher. The consumer-focused project fetched real-time web data before ChatGPT was hooked up to the internet, Weiss told TechCrunch. “It went extremely viral, so pretty fast we gained almost 20,000 GitHub stars.”



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda.

They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

Weiss launched Tavily after ChatGPT and other LLMs introduced web search. Unlike GPT Researcher, Tavily focuses on enterprise clients. It provides a suite of tools to companies like Groq, Cohere, MongoDB, and Writer, allowing their agents to search, crawl, and extract structured insights from both public and private sources.

While most AI agents aren’t yet connected to the internet, Weiss says Tavily’s goal is to onboard the next billion agents to the web.

Tavily is not the only startup providing search tools for AI agents. It competes with [Exa](#), which raised a \$17 million Series A from Lightspeed, Nvidia, and Y Combinator last year. Another smaller startup that offers a web search connectivity layer is [Firecrawl](#). OpenAI and Perplexity are also offering search solutions that are geared toward independent developers.

Topics: [AI](#) [AI agents](#) [Insight Partners](#) [Search](#) [Startups](#)



Marina Temkin


Reporter, Venture |

Marina Temkin is a venture capital and startups reporter at TechCrunch. Prior to joining TechCrunch, she wrote about VC for PitchBook and Venture Capital Journal...

its Salesforce database

- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Spotify raises subscription prices
- The uproar over Vogue’s AI-generated ad isn’t just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines

Loading the next article



[TechCrunch](#)

[Staff](#)

[Contact Us](#)

[Advertise](#)

[Crunchboard
Jobs](#)

[Site Map](#)

[Terms of Service](#)

[Privacy Policy](#)

[RSS Terms of Use](#)

[OpenAI](#)

[Perplexity](#)

[Cisco](#)

[DeepMind](#)

[Windsurf](#)

[Tech Layoffs](#)

[ChatGPT](#)