

Advertisement



IMAGE CREDITS: GOOGLE DEEPMIND

AI



Google rolls out Gemini Deep Think AI, a reasoning model that tests multiple ideas in parallel

Maxwell Zeff — 4:00 AM PDT · August 1, 2025

Google DeepMind is rolling out [Gemini 2.5 Deep Think](#), which, the company says, is its most advanced AI reasoning model, able to answer questions by exploring and considering multiple ideas simultaneously and then using those outputs to choose the best answer.

Subscribers to Google’s \$250-per-month [Ultra](#) subscription will gain access to Gemini 2.5 Deep Think in the Gemini app starting Friday.


First unveiled in May at Google I/O 2025, Gemini 2.5 Deep Think is Google’s first publicly available multi-agent model. These systems spawn multiple AI agents to tackle a question in parallel, a process that uses significantly more computational resources than a single agent, but tends to result in better answers.

Google used a variation of Gemini 2.5 Deep Think to [score a gold medal](#) at this year’s International Math Olympiad (IMO).

Alongside Gemini 2.5 Deep Think, the company says it is releasing the model it used at the IMO to a select group of mathematicians and academics. Google says this AI model “takes hours to reason,” instead of seconds or minutes like most consumer-facing AI models. The company hopes the IMO model will enhance research efforts, and aims to get feedback on how to improve the multi-agent system for academic use cases.

Google notes that the Gemini 2.5 Deep Think model is a significant improvement over what it announced at I/O. The company also claims to have developed “novel reinforcement learning techniques” to encourage Gemini 2.5 Deep Think to make better use of its reasoning paths.

“Deep Think can help people tackle problems that require creativity, strategic planning and making improvements step-by-step,” said Google in a [blog post](#) shared with TechCrunch.



TC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

[REGISTER NOW](#)

Most Popular




- Key sections of the US Constitution deleted from government’s website
- Instagram takes on Snapchat with new ‘Instagram Map’
- Google says hackers stole its customers’ data by breaching its Salesforce database

The company says Gemini 2.5 Deep Think achieves state-of-the-art performance on Humanity’s Last Exam (HLE) — a challenging test

measuring AI's ability to answer thousands of crowdsourced questions across math, humanities, and science. Google claims its model scored 34.8% on HLE (without tools), compared to xAI's Grok 4, which scored 25.4%, and OpenAI's o3, which scored 20.3%.

Google also says Gemini 2.5 Deep Think outperforms AI models from OpenAI, xAI, and Anthropic on LiveCodeBench 6, a challenging test of competitive coding tasks. Google's model scored 87.6%, whereas Grok 4 scored 79%, and OpenAI's o3 scored 72%.

Comparison across models without tools enabled



Methodology:
Model Selection: To ensure fair comparison we only compare results for models without additional tool calls enabled.
Sources: Humanity's Last Exam results are sourced from https://hlt.eecs.berkeley.edu/humanity's_last_exam and <https://livecodebench.github.io/leaderboard.html> and IMO 2025 from <https://imo2025.org/>. The rest of the results come from provider self-reported numbers.
Gemini 2.5 Deep Think (IMO 2025) results are computed as passed1 while all the other results coming from matharena.ai are best of 32.
IMO 2025 medal descriptions are based on the cutoffs from the 2025 competition.
For Grok 4 IMO 2025 results we show the highest result available from matharena.ai with custom prompt.


BENCHMARK SCORES

IMAGE CREDITS:GOOGLE

Gemini 2.5 Deep Think automatically works with tools such as code execution and Google Search, and the company says it's capable of producing "much longer responses" than traditional AI models.

In Google's testing, the model produced more detailed and aesthetically pleasing web development tasks compared to other AI models. The company claims the model could aid researchers and "potentially accelerate the path to discovery."

- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines



Design and create a very creative, elaborate, and detailed voxel art scene of a pagoda in a beautiful garden with trees, including some cherry blossoms. Make the scene impressive and varied and use colorful voxels. Use whatever libraries to get this done but make sure I can paste it all into a single HTML file and open it in Chrome.



Gemini 2.5 Flash



Gemini 2.5 Pro



Gemini 2.5 Deep Think

ART SCENES MADE BY GOOGLE'S AI

IMAGE CREDITS:GOOGLE

It seems that several leading AI labs are converging around the multi-agent approach.

Elon Musk's xAI recently released a multi-agent system of its own, [Grok 4 Heavy](#), which it says was able to achieve industry-leading performance on several benchmarks. OpenAI researcher Noam Brown said on a [podcast](#) that the unreleased AI model the company used to achieve a gold medal at this year's International Math Olympiad was also a multi-agent system. Meanwhile, [Anthropic's Research agent](#), which generates thorough research briefs, is also powered by a multi-agent system.

Despite the strong performance, it seems that multi-agent systems are even costlier to serve than traditional AI models. That means tech companies may keep these systems gated behind their most expensive subscription plans, which xAI and now Google have chosen to do.

In the coming weeks, Google says it plans to share Gemini 2.5 Deep Think with a select group of testers via the Gemini API. The company says it wants to better understand how developers and enterprises may use its multi-agent system.

Topics: [agents](#) [AI](#) [gemini](#) [Google](#)





Maxwell Zeff

Senior AI Reporter | X

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg, and MSNBC, Zeff has covered the rise of AI and the Silicon...

[View Bio](#) >

KEEP READING



IMAGE CREDITS:

MATTHIAS BALK/PICTURE ALLIANCE / GETTY IMAGES

AI

Google takes on ChatGPT's Study Mode with new 'Guided Learning' tool in Gemini




Aisha Malik — 11:18 AM PDT · August 6, 2025

As the new school year approaches, Google [announced](#) on Wednesday that it's launching a new tool called Guided Learning within Gemini. The tool sort of functions like an AI tutor, as it's designed to help users build a deep understanding instead of just getting answers.

The launch follows just over a week after OpenAI rolled out [Study Mode for ChatGPT](#), which is also designed to go beyond simply obtaining answers to questions by actually helping users develop critical thinking skills.



Both companies' launches [come amid concerns](#) that AI chatbots [undermine the learning process](#) because they spit out direct answers. The new tools from Google and OpenAI likely aim to address these concerns by positioning their chatbots as learning tools rather than simple answer engines.

With Guided Learning, Gemini will break down problems step-by-step and adapt explanations to its users' needs. The feature responds using images, diagrams, videos, and interactive quizzes to help users build and test themselves on their knowledge, rather than simply giving them the answer.



step of the way," wrote Maureen Heymans, Google's VP of Learning & Sustainability, in the [blog post](#).

In addition to the new feature, Google announced that it's working to make Gemini as a whole better equipped to help users learn. Gemini will now automatically incorporate images, diagrams, and YouTube videos directly into responses to help users better understand complex topics. Plus, users can now also ask Gemini to create flashcards and study guides based on their quiz results or other class materials.



voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).



Google also announced on Wednesday that it’s offering students in the U.S., Japan, Indonesia, Korea, and Brazil a free one-year subscription to Google’s AI Pro plan.

The plan includes expanded access to [Gemini 2.5 Pro](#), [NotebookLM](#), [Veo 3](#), [Deep Research](#), and more.

Topics: [AI](#) [Apps](#) [artificial intelligence](#) [ChatGPT](#) [gemini](#) [Google](#)



Aisha Malik

Consumer News Reporter | [X](#) [Twitter](#) [in](#)

Aisha is a consumer news reporter at TechCrunch. Prior to joining the publication in 2021, she was a telecom reporter at MobileSyrup. Aisha holds an honours bachelor’...

View Bio >

KEEP READING

AI



Google’s AI coding agent Jules is now out of beta

Google on Wednesday launched its AI coding agent, Jules, out of beta, just over two months after its [public preview debut in May](#).

Powered by Gemini 2.5 Pro, Jules is an asynchronous, agent-based coding tool that integrates with GitHub, clones codebases into Google Cloud virtual machines, and uses AI to fix or update code while developers focus on other tasks.

Google [initially announced](#) Jules as a Google Labs project in December and made it [available to beta testers](#) through a public preview at its I/O developer conference.

Kathy Korevec, director of product at Google Labs, told TechCrunch that the tool's improved stability drove the decision to take it out of beta after receiving hundreds of UI and quality updates during its beta phase.

"The trajectory of where we're going gives us a lot of confidence that Jules is around and going to be around for the long haul," she said.


With the wider rollout, Google introduced structured pricing tiers for Jules, starting with an "introductory access" free plan capped at 15 individual daily tasks and three concurrent ones, down from the 60-task limit during beta. Jules' paid tiers are part of the Google AI Pro and Ultra plans, which are priced at \$19.99 and \$124.99 a month, and offer subscribers 5x and 20x higher limits, respectively.

Korevec noted that Jules' packaging and pricing are based on "real usage" insights gathered over the past couple of months.

"The 60-task cap helped us study how developers use Jules and gave us the information we needed to design the new packaging," she said. "The 15/day is designed to give people a sense of whether Jules will work for them on real project tasks."

Google also updated Jules' privacy policy to be more explicit about how it trains AI. If a repository is public, its data may be used for training, but if it is private, Korevec said that no data is sent.

"We got a little bit of feedback from users that it [the privacy policy] wasn't as clear as we thought it was, and so most of it is just responding to that. We didn't change anything about what we're doing on the training side, but we changed the language," Korevec said.



Latest in AI

[See More](#)

TRANSPORTATION



Tesla shuts down Dojo, the AI training supercomputer that Musk said would be key to full self-driving

Rebecca Bellan

- 35 minutes ago

STARTUPS



High costs and thin margins threatening AI coding startups

Marina Temkin - 2 hours ago

AI



OpenAI's GPT-5 is here

Maxwell Zeff - 6 hours ago

During the beta, Google said that thousands of developers tackled tens of thousands of tasks, resulting in over 140,000 code improvements shared publicly. Initial feedback led the Google Labs team to add new capabilities, including reusing previous setups for faster task execution, integrating with GitHub issues, and supporting multimodal input.




IMAGE CREDITS: JULES / X

The two primary users of Jules so far are the AI enthusiasts and professional developers, Korevec said.

By running asynchronously in a virtual machine, Jules stands apart from top AI coding tools like Cursor, Windsurf, and Lovable, which all work synchronously and require users to watch the output after each prompt.

“Jules operates like an extra set of hands ... you can basically kick off tasks to it, and then you could close your computer and walk away from it if you want and then come back hours later. Jules would have those tasks done for you, versus if you were doing that with a local agent or using a synchronous agent, you would be bound to that session,” Korevec explained.

This week, Jules received a deeper integration with GitHub to [open pull requests automatically](#) — just like it could open branches — and a feature called [Environment Snapshots](#) to save dependencies and install scripts as a snapshot for faster, more consistent task execution.

From vibe coding to mobile use, beta trials informed Jules’ development

Since entering public beta, Jules has logged 2.28 million visits worldwide, 45% of them from mobile devices, per data from market intelligence provider [SimilarWeb](#), reviewed by TechCrunch. India was the top market for traffic, followed by the U.S. and Vietnam.

Google did not share specifics on Jules’ user base and its top geographies.

Korevec told TechCrunch that during the beta, the team observed that many people used Jules from traditional vibe-coding tools to either fix bugs that might have been implemented or extend the vibe-coded project to make it more production-ready.

Originally, Jules required users to have an existing codebase. But Google soon realized many potential users — like those trying other AI tools — might want to explore it without one. Korevec said the company quickly enabled Jules to work even with an empty repository. That helped increase its scope and usage.

Google Labs’ team also noticed an increasing number of users accessing Jules through their mobile devices. Although the tool does not have a dedicated mobile app, Korevec said users were accessing it through its web app.

“Since it’s a big use case that we’re seeing emerging, we’re absolutely exploring what the features are that people need on mobile a lot more,” she noted.

Alongside beta testers, Korevec stated that Google already uses Jules to help develop some projects internally, and there is now a “big push” to use the tool on “a lot more projects” at the company.

Topics: [AI](#) [Google](#) [Google Jules](#) [jules](#) [United States](#)



[f](#) [X](#) [in](#) [📧](#) [✉](#) [🔗](#)

Jagmeet Singh

Reporter | [X](#) [🦋](#) [@](#) [in](#)

Jagmeet covers startups, tech policy-related updates, and all other major tech-centric developments from India for TechCrunch. He previously worked as a...

[View Bio](#) >

Loading the next article




YOU HAVE ONLY 2 MOVES

IQ: 155



TRY



[TechCrunch](#)

[Staff](#)

[Contact Us](#)

[Advertise](#)

[Crunchboard](#)

[Jobs](#)

[Site Map](#)

[Terms of Service](#)

[Privacy Policy](#)

[RSS Terms of Use](#)

[OpenAI](#)

[Perplexity](#)

[Cisco](#)

[DeepMind](#)

[Windsurf](#)

[Tech Layoffs](#)

[ChatGPT](#)