

Advertisement

AI



Google rolls out Gemini
Deep Think AI, a
reasoning model that
tests multiple ideas in
parallel

Google DeepMind is rolling out [Gemini 2.5 Deep Think](#), which, the company says, is its most advanced AI reasoning model, able to answer questions by exploring and considering multiple ideas simultaneously and then using those outputs to choose the best answer.

Subscribers to Google's \$250-per-month [Ultra](#) subscription will gain access to Gemini 2.5 Deep Think in the Gemini app starting Friday.

First unveiled in May at Google I/O 2025, Gemini 2.5 Deep Think is Google's first publicly available multi-agent model. These systems spawn multiple AI agents to tackle a question in parallel, a process that uses significantly more computational resources than a single agent, but tends to result in better answers.

Google used a variation of Gemini 2.5 Deep Think to [score a gold medal](#) at this year's International Math Olympiad (IMO).

Alongside Gemini 2.5 Deep Think, the company says it is releasing the model it used at the IMO to a select group of mathematicians and academics. Google says this AI model "takes hours to reason," instead of seconds or minutes like most consumer-facing AI models. The company hopes the IMO model will enhance research efforts, and aims to get feedback on how to improve the multi-agent system for academic use cases.


DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion

Google notes that the Gemini 2.5 Deep Think model is a significant improvement over what it announced at I/O. The company also claims to have developed “novel reinforcement learning techniques” to encourage Gemini 2.5 Deep Think to make better use of its reasoning paths.

“Deep Think can help people tackle problems that require creativity, strategic planning and making improvements step-by-step,” said Google in a [blog post](#) shared with TechCrunch.

- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

The company says Gemini 2.5 Deep Think achieves state-of-the-art performance on Humanity’s Last Exam (HLE) — a challenging test measuring AI’s ability to answer thousands of crowdsourced questions across math, humanities, and science. Google claims its model scored 34.8% on HLE (without tools), compared to xAI’s Grok 4, which scored 25.4%, and OpenAI’s o3, which scored 20.3%.

Google also says Gemini 2.5 Deep Think outperforms AI models from OpenAI, xAI, and Anthropic on LiveCodeBench 6, a challenging test of competitive coding tasks. Google’s model scored 87.6%, whereas Grok 4 scored 79%, and OpenAI’s o3 scored 72%.

Comparison across models without tools enabled



Methodology:
Model Selection: To ensure fair comparison we only compare results for models without additional tool calls enabled.
Sources: Humanity's Last Exam results are sourced from https://hltai.com/leaderboard/humanitys_last_exam and <https://api.safelife.ai/>. LiveCodeBench from <https://livecodebench.github.io/leaderboard.html> and IMO 2025 from <https://matharena.ai/>. The rest of the results come from provider self-reported numbers.
Gemini 2.5 Deep Think IMO 2025 results are computed as passed/1 while all the other results coming from matharena.ai are best of 32.
IMO 2025 medal descriptions are based on the cutoffs from the 2025 competition.
For Grok 4 IMO 2025 results we show the highest result available from matharena.ai with custom prompt.

BENCHMARK SCORES

IMAGE CREDITS:GOOGLE

Gemini 2.5 Deep Think automatically works with tools such as code execution and Google Search, and the company says it’s capable of producing “much longer responses” than traditional AI models.

In Google’s testing, the model produced more detailed and aesthetically pleasing web development tasks compared to other AI models. The company claims the model could aid researchers and “potentially accelerate the path to discovery.”

Design and create a very creative, elaborate, and detailed voxel art scene of a pagoda in a beautiful garden with trees, including some cherry blossoms. Make the scene impressive and varied and use colorful voxels. Use whatever libraries to get this done but make sure I can paste it all into a single HTML file and open it in Chrome.



Gemini 2.5 Flash



Gemini 2.5 Pro



Gemini 2.5 Deep Think

ART SCENES MADE BY GOOGLE'S AI

IMAGE CREDITS:GOOGLE

It seems that several leading AI labs are converging around the multi-agent approach.

Elon Musk's xAI recently released a multi-agent system of its own, [Grok 4 Heavy](#), which it says was able to achieve industry-leading performance on several benchmarks. OpenAI researcher Noam Brown said on a [podcast](#) that the unreleased AI model the company used to achieve a gold medal at this year's International Math Olympiad was also a multi-agent system. Meanwhile, [Anthropic's Research agent](#), which generates thorough research briefs, is also powered by a multi-agent system.

Despite the strong performance, it seems that multi-agent systems are even costlier to serve than traditional AI models. That means tech companies may keep these systems gated behind their most expensive subscription plans, which xAI and now Google have chosen to do.

In the coming weeks, Google says it plans to share Gemini 2.5 Deep Think with a select group of testers via the Gemini API. The company says it wants to better understand how developers and enterprises may use its multi-agent system.

Topics: [agents](#) [AI](#) [gemini](#) [Google](#)



Maxwell Zeff

Senior AI Reporter | [X](#)

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >

KEEP READING




IMAGE CREDITS: GOOGLE

AI

‘Deep Think’ boosts the performance of Google’s flagship Google Gemini AI model

Kyle Wiggers — 10:45 AM PDT · May 20, 2025

TE

Google is upgrading its most capable Gemini AI

Latest

Security

Startups

AI

Venture

Apps

Apple

Events

Podcasts

Newsletters

Search

Menu

Google was vague on the inner workings of Deep Think, but the technology could be similar to OpenAI’s [o1-pro](#) and upcoming o3-pro models, which likely use an engine to search for and synthesize the best solution to a given problem.

Google says that Deep Think enabled Gemini 2.5 Pro to top LiveCodeBench, a challenging coding evaluation. Gemini 2.5 Pro Deep Think also beat OpenAI’s o3 on MMMU, a test for skills like perception and reasoning.




IMAGE CREDITS:GOOGLE DEEPMIND

Deep Think is available to “trusted testers” via the Gemini API as of this week. Google said that it’s

TechCrunch Week in Review

Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility

TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

taking additional time to conduct safety evaluations before rolling out Deep Think widely.

Alongside Deep Think, Google has introduced an update to its budget-oriented Gemini 2.5 Flash model that allows the model to perform better on tasks involving coding, multimodality, reasoning, and long context. The new 2.5 Flash, which is also more efficient than the version it replaces, is available for preview in Google's AI Studio and Vertex AI platforms as well as the company's Gemini apps.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

[REGISTER NOW](#)

Google says that the improved Gemini 2.5 Flash will become generally available for developers sometime in June.

Lastly, Google is introducing a model called Gemini Diffusion, which the company claims is “very fast” — delivering output 4-5 times quicker than comparable models and rivaling the performance of models twice its size. Gemini Diffusion is available beginning today to “trusted testers.”



Kyle Wiggers

AI Editor | [X](#) [🦋](#) [@](#) [in](#)

Kyle Wiggers was TechCrunch’s AI Editor until June 2025.
His writing has appeared in VentureBeat and Digital...

[View Bio](#) >

KEEP READING



AI



Google Gemini: Everything you need to know about the generative AI apps and models

Kyle Wiggers, Maxwell Zeff

6:09 PM PST · February 26, 2025

Google’s trying to make waves with Gemini, its flagship suite of generative AI models, apps, and services. But what’s Gemini? How can you use it? And how does it stack up to other generative AI tools such as OpenAI’s [ChatGPT](#), Meta’s [Llama](#), and Microsoft’s [Copilot](#)?

To make it easier to keep up with the latest Gemini developments, we’ve put together this handy guide, which we’ll keep updated as new Gemini models, features, and news about Google’s plans for Gemini are released.

What is Gemini?

Gemini is Google’s [long-promised](#), next-gen generative AI model family. Developed by Google’s AI research labs DeepMind and Google Research, it comes in several flavors:

- **Gemini Ultra**, a very large model.
- **Gemini Pro**, a large model — though smaller than Ultra. The latest version, [Gemini 2.0 Pro](#), is Google’s current flagship.
- **Gemini Flash**, a speedier, “distilled” version of Pro.
- **Gemini Flash-Lite**, a slightly smaller and faster version of Gemini Flash.
- **Gemini Flash Thinking**, a model with “reasoning” capabilities.
- **Gemini Nano**, two small models: **Nano-1** and the slightly more capable **Nano-2**, which is meant to run offline.

All Gemini models were trained to be natively multimodal — that is, able to work with and analyze more than just text. Google says they were pre-

Latest in AI

[See More](#)

AI



Google denies AI search features are killing website traffic

Sarah Perez - 12 hours ago

AI



First impressions of Alexa+, Amazon’s upgraded, AI-powered digital assistant

Sarah Perez - 13 hours ago

AI



Google takes on ChatGPT’s Study Mode with new ‘Guided Learning’ tool in Gemini

Aisha Malik - 13 hours ago

trained and fine-tuned on a variety of public, proprietary, and licensed audio, images, and videos; a set of codebases; and text in different languages.

This sets Gemini apart from models such as [Google's own LaMDA](#), which was trained exclusively on text data. LaMDA can't understand or generate anything beyond text (e.g., essays, emails, and so on), but that isn't necessarily the case with Gemini models. For example, the [latest versions of Gemini Flash](#) and Gemini Pro can natively output images and audio in addition to text.

We'll note here that the [ethics and legality](#) of training models on public data, in some cases without the data owners' knowledge or consent, are murky. Google has an [AI indemnification policy](#) to shield certain Google Cloud customers from lawsuits should they face them, but this policy contains carve-outs. Proceed with caution — particularly if you're intending on using Gemini commercially.

What's the difference between the Gemini apps and Gemini models?

Gemini is separate and distinct from the Gemini apps on the web and mobile ([formerly Bard](#)).

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the

20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

The Gemini apps are clients that connect to various Gemini models and layer a chatbot-like interface on top. Think of them as front ends for Google’s generative AI, analogous to [ChatGPT](#) and Anthropic’s [Claude family of apps](#).




IMAGE CREDITS:GOOGLE

Gemini on the web lives [here](#). On Android, the [Gemini app](#) replaces the existing Google Assistant app. And on iOS, the [Google and Google Search apps](#) serve as that platform’s Gemini clients.

On Android, users can bring up a Gemini overlay to ask questions about what’s on their screen (for example, a YouTube video). Pressing and holding a supported smartphone’s power button or saying, “Hey Google” summons the overlay.

Gemini apps can accept images as well as voice commands and text — including files like PDFs, either uploaded or imported from Google Drive — and generate images. As you'd expect, conversations with Gemini apps on mobile carry over to Gemini on the web and vice versa if you're signed in to the same Google Account in both places.

Gemini Advanced

The Gemini apps aren't the only means of recruiting Gemini models' assistance with tasks. Slowly but surely, Gemini-imbued features are [making their way](#) into staple Google apps and services like Gmail and Google Docs.

To take advantage of most of these, you'll need the Google One AI Premium Plan. Technically a part of [Google One](#), the AI Premium Plan costs \$20 a month and provides access to Gemini in Google Workspace apps like Docs, Maps, Slides, Sheets, Drive, and Meet. It also enables what Google calls Gemini Advanced, which brings the company's more sophisticated Gemini models to the Gemini apps.

IMAGE CREDITS:GOOGLE

Gemini Advanced users get extras here and there, too, like priority access to new features and models; the ability to run and edit Python code directly in Gemini; and increased limits for

[NotebookLM](#), Google’s tool that turns PDFs into AI-generated podcasts. Recently, Gemini Advanced gained a [memory feature](#) that stores users’ preferences and allows Gemini to refer to old conversations as context for current chats.

One of the more compelling Gemini Advanced exclusives, [Deep Research](#), leverages Gemini models with “advanced reasoning” to create detailed briefs. In response to a prompt (e.g. “How should I redesign my kitchen?”), Deep Research develops a multi-step research plan and searches the web to craft a comprehensive answer.

Gemini in Gmail, Docs, Chrome, dev tools, and more

In Gmail, Gemini lives in a [side panel](#) that can write emails and summarize message threads. You’ll find the same panel in Docs, where it helps write and refine content and brainstorm new ideas. Gemini in Slides generates slides and custom images. And Gemini in Google Sheets tracks and organizes data, creating tables and formulas.

Gemini [is in Google Maps](#), where it can aggregate reviews about local businesses and offer recommendations like how to spend a day visiting a foreign city. The chatbot’s reach extends to Drive, as well, where it can summarize files and folders and give quick facts about a project.

IMAGE CREDITS:GOOGLE

[Gemini recently came to Google's Chrome browser](#) in the form of an AI writing tool. You can use it to write something completely new or rewrite existing text; Google says it'll consider the web page you're on to make recommendations.

Elsewhere, you'll find hints of Gemini in Google's [database products](#), [cloud security tools](#), and [app development platforms](#) (including [Firebase](#) and [Project IDX](#)), as well as in apps like [Google Photos](#) (where Gemini handles natural language search queries), [YouTube](#) (where it helps brainstorm video ideas), and Meet (where it translates captions).

[Code Assist](#) (formerly [Duet AI for Developers](#)), Google's suite of AI-powered assistance tools for code completion and generation, is offloading heavy computational lifting to Gemini. So are Google's [security products underpinned by Gemini](#), like Gemini in Threat Intelligence, which can analyze large portions of potentially malicious code and let users perform natural language searches for ongoing threats or indicators of compromise.

Gemini extensions and Gems

[Gemini Advanced users can create Gems](#), custom chatbots on desktop and mobile powered by Gemini models. Gems can be generated from natural language descriptions — for instance, “You’re my running coach. Give me a daily running plan” — and shared with other users or kept private.

IMAGE CREDITS:GOOGLE

The Gemini apps can tap into Google services via what Google calls “Gemini extensions.” Gemini integrates with Drive, Gmail, YouTube, and more to respond to queries such as “Could you summarize my last three emails?”

Gemini Live in-depth voice chats

[An experience called Gemini Live](#) allows users to have “in-depth” voice chats with Gemini. It’s available in the Gemini apps on mobile and the [Pixel Buds Pro 2](#), where it can be accessed even when your phone’s locked.

IMAGE CREDITS:GOOGLE

With Gemini Live enabled, you can interrupt Gemini while the chatbot’s speaking to ask a clarifying question, and it’ll adapt to your speech patterns in real-time. Live is also designed to serve as a virtual coach of sorts, helping you rehearse for events, brainstorm ideas, and so on. For instance, Live can suggest which skills to highlight in an upcoming job interview and give public speaking pointers.

You can read our [review of Gemini Live here](#).

Gemini for teens

Google offers a teen-focused [Gemini experience](#) for students.

The teen-focused Gemini has “additional policies and safeguards,” including a tailored onboarding process and an AI literacy guide. Otherwise, it’s nearly identical to the standard Gemini experience, down to the “double-check” feature that looks across the web to see if Gemini’s responses are accurate.

What can the Gemini models do?

Because Gemini models are multimodal, they can perform a range of multimodal tasks, from transcribing speech to captioning images and videos in real-time. Many of these capabilities have reached the product stage, and Google is promising much more in the not-too-distant future.

Of course, Google offers no fix for some of the [underlying problems](#) with generative AI technology today, like its [encoded biases](#) and tendency to make things up (i.e., [hallucinate](#)). Neither do its rivals, but it's something to keep in mind when considering using or paying for Gemini.

Gemini Pro's capabilities

Google says that its latest Pro model, [Gemini 2.0 Pro](#), is its best yet for coding and complex prompts. 2.0 Pro outperforms its predecessor, [Gemini 1.5 Pro](#), in benchmarks measuring programming, reasoning, math, and factual accuracy.

In Google's Vertex AI platform, developers can customize Gemini Pro to specific contexts and use cases via a fine-tuning or "grounding" process. For example, Pro (along with other Gemini models) can be instructed to use data from third-party providers like Moody's, Thomson Reuters, ZoomInfo, and MSCI, or source information from corporate datasets or Google Search instead of its wider knowledge bank. Gemini Pro can also be connected to external, third-party APIs to perform particular actions, like automating a back-office workflow.

Google's AI Studio platform offers templates for creating structured chat prompts with Pro. Developers can control the model's creative range and provide examples to give tone and style instructions — and also tune Pro's safety settings.

Gemini Flash is lightweight, while Gemini Flash Thinking adds reasoning

[Gemini 2.0 Flash](#), which can use tools like Google Search and interact with external APIs, outperforms some of the larger Gemini 1.5 models on benchmarks measuring coding and image analysis. An offshoot of Gemini Pro, Flash is small and efficient — built for narrow, high-frequency generative AI workloads.

Google says that Flash is particularly well-suited for tasks like summarization and chat apps, plus image and video captioning and data extraction from long documents and tables. Meanwhile, Gemini 2.0 Flash-Lite, a more compact version of Flash, outperforms Gemini 1.5 Flash but runs at the same price and speed, according to Google.

Last December, Google [released a “thinking” version of Gemini 2.0 Flash](#) that’s capable of “reasoning.” The AI model takes a few seconds to work backward through a problem before it gives an answer, which can improve its reliability.

Gemini Nano can run on your phone

Gemini Nano is a tiny version of Gemini efficient enough to run directly on (some) devices instead of sending the task off to a server somewhere. So far, Nano powers a couple of features on the [Pixel 8 Pro](#), [Pixel 8](#), Pixel 9 Pro, Pixel 9, and [Samsung Galaxy S24](#), including Summarize in Recorder and Smart Reply in Gboard.

The Recorder app, which lets users push a button to record and transcribe audio, includes a Gemini-powered summary of recorded conversations, interviews, presentations, and other audio

snippets. Users get summaries even if they don't have a signal or Wi-Fi connection — and in a nod to privacy, no data leaves their phone in process.

IMAGE CREDITS:GOOGLE

Nano is also in Gboard, Google's keyboard replacement. There, it powers Smart Reply, which helps to suggest the next thing you'll want to say when having a conversation in a messaging app such as WhatsApp.

A future version of Android will tap Nano to [alert users to potential scams during calls.](#) The [new weather app](#) on Pixel phones uses Gemini Nano to generate tailored weather reports. And TalkBack, Google's accessibility service, employs Nano to [create aural descriptions of objects](#) for low-vision and blind users.

Gemini Ultra, MIA for now

We haven't seen much of [Gemini Ultra](#) in recent months. The model isn't available in the Gemini apps, and it isn't listed on Google's Gemini API pricing page. However, that doesn't mean Google won't bring Ultra back at some point in the future.

How much do the Gemini models cost?

Gemini 1.5 Pro, 1.5 Flash, 2.0 Flash, and 2.0 Flash-Lite are available through Google's Gemini API for building apps and services. They're pay-as-you-go. Here's the base pricing — not including add-ons — as of February 225:

- **Gemini 1.5 Pro:** \$1.25 per 1 million input tokens (for prompts up to 128K tokens) or \$2.50 per 1 million input tokens (for prompts longer than 128K tokens); \$5 per 1 million output tokens (for prompts up to 128K tokens) or \$10 per 1 million output tokens (for prompts longer than 128K tokens)
- **Gemini 1.5 Flash:** 7.5 cents per 1 million input tokens (for prompts up to 128K tokens), 15 cents per 1 million input tokens (for prompts longer than 128K tokens), 30 cents per 1 million output tokens (for prompts up to 128K tokens), 60 cents per 1 million output tokens (for prompts longer than 128K tokens)
- **Gemini 2.0 Flash:** 10 cents per 1 million input tokens, 40 cents per 1 million output tokens. For audio, 70 cents per 1 million input tokens.
- **Gemini 2.0 Flash-Lite:** 7.5 cents per 1 million input tokens, 30 cents per 1 million output tokens.

Tokens are subdivided bits of raw data, like the syllables “fan,” “tas,” and “tic” in the word “fantastic”; 1 million tokens is equivalent to about 750,000 words. *Input* refers to tokens fed into the model, while *output* refers to tokens that the model generates.

2.0 Pro pricing has yet to be announced, and Nano is still in [early access](#).

Is Gemini coming to the iPhone?

It might.

[Apple has said that it's in talks to put Gemini and other third-party models to use](#) for a number of features in its [Apple Intelligence](#) suite. Following a keynote presentation at WWDC 2024, Apple SVP Craig Federighi [confirmed plans to work with models](#), including Gemini, but he didn't divulge any additional details.

This post was originally published February 16, 2024, and is updated regularly.

Topics: [AI](#) [Apps](#) [Enterprise](#) [evergreens](#) [gemini](#)

[Gemini Pro](#) [Generative AI](#) [Google](#)

[google gemini](#)

Advertisement



Kyle Wiggers

AI Editor | [X](#) [🐼](#) [@](#) [in](#)

Kyle Wiggers was TechCrunch’s AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

[View Bio](#) >



Maxwell Zeff

Senior AI Reporter | [X](#)

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >

Loading the next article

