GOVERNMENT & POLICY

Google says it will sign EU's AI code of practice

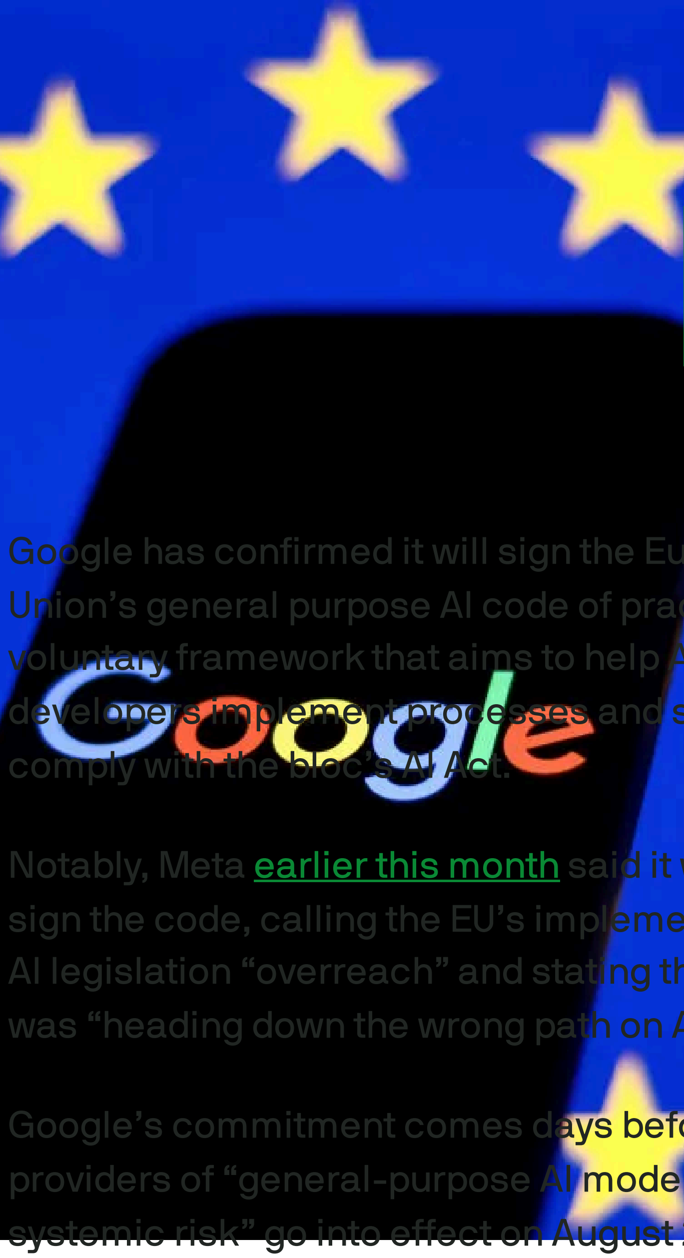

Ram Iyer — 5:57 AM PDT · July 30, 2025

Google has confirmed it will sign the European Union's general purpose AI code of practice, a voluntary framework that aims to help AI developers implement processes and systems to comply with the bloc's AI Act.

Notably, Meta earlier this month said it would not sign the code, calling the EU's implementation of its AI legislation "overreach" and stating that Europe was "heading down the wrong path on AI."

Google's commitment comes days before rules for providers of "general-purpose AI models with systemic risk" go into effect on August 2.

IMAGE CREDITS: NURPHOTO / GETTY IMAGES

Companies likely to be affected by these rules include major names such as Anthropic, Google, Meta, and OpenAI, as well as several other large generative models, and they will have two years to comply fully with the AI Act.

In a blog post on Wednesday, Kent Walker, president of global affairs at Google, conceded that the final version of the code of practice was better than what the EU proposed initially, but he still noted reservations around the AI Act and the code.

"We remain concerned that the AI Act and Code risk slowing Europe's development and deployment of AI. In particular, departures from EU copyright law, steps that slow approvals, or requirements that expose trade secrets could chill European model

Most Popular

- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a $2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion

development and deployment, harming Europe's competitiveness," wrote Walker.

By signing the EU's code of practice, AI companies would agree to follow a slate of guidelines, which include providing updated documentation about their AI tools and services; no training AI on pirated content; and complying with requests from content owners to not use their works in their datasets.

A risk-based regulation for AI applications, the EU's landmark AI Act bans some "unacceptable risk" use cases, such as cognitive behavioral manipulation or social scoring. The rules also define a set of "high-risk" uses, including biometrics and facial recognition, and the use of AI in domains like education and employment. The act also requires developers to register AI systems and meet risk- and quality-management obligations.

Topics: AI EU AI Act Europe european union

- Your public ChatGPT queries are getting indexed by Google and other search engines
- The best dating apps aren't even dating apps

Google Government & Policy Meta

Ram Iyer

Editor

Ram is a financial and tech reporter and editor. He covered North American and European M&A, equity, regulatory...

View Bio

KEEP READING

IMAGE CREDITS: JONATHAN RAA/NURPHOTO / GETTY IMAGES

GOVERNMENT & POLICY

Meta refuses to sign EU's AI code of practice

Ram Iyer — 6:52 AM PDT · July 18, 2025

Meta has refused to sign the European Union's code of practice for its AI Act, weeks before the bloc's rules for providers of general-purpose AI models take effect.

"Europe is heading down the wrong path on AI," wrote Meta's chief global affairs officer Joel Kaplan in a post on LinkedIn. "We have carefully reviewed the European Commission's Code of Practice for general-purpose AI (GPAI) models and Meta won't be signing it. This Code introduces a number of legal uncertainties for model developers, as well as measures which go far beyond the scope of the AI Act."

The EU's code of practice — a voluntary framework published earlier this month — aims to help companies implement processes and systems to comply with the bloc's legislation for regulating AI. Among other things, the code requires companies to provide and regularly update documentation about their AI tools and services and bans developers from training AI on pirated content; companies must also comply with content owners' requests to not use their works in their datasets.

Calling the EU's implementation of the legislation "overreach," Kaplan claimed that the law will "throttle the development and deployment of frontier AI models in Europe and will stunt European companies looking to build businesses on top of them."

A risk-based regulation for applications of artificial intelligence, the AI Act bans some "unacceptable risk" use cases outright, such as cognitive behavioral manipulation or social scoring. The rules also define a set of "high-risk" uses, such as biometrics and facial recognition, and in domains like education and employment. The act also requires developers to register AI systems and meet risk- and quality-management obligations.

Tech companies from across the world, including those at the forefront of the AI race like Alphabet, Meta, Microsoft, and Mistral AI, have been fighting the rules, even urging the European Commission to delay its rollout. But the Commission has held firm, saying it will not change its timeline.

Also on Friday, the EU published guidelines for providers of AI models ahead of rules that will go into effect on August 2. These rules would affect providers of "general-purpose AI models with systemic risk," like OpenAI, Anthropic, Google, and Meta. Companies that have such models on the market before August 2 will have to comply with the legislation by August 2, 2027.

Topics: AI EU AI Act Europe european union Government & Policy Meta

Ram Iyer

Editor |

Ram is a financial and tech reporter and editor. He covered North American and European M&A, equity, regulatory...

View Bio

KEEP READING

SOCIAL

Meta Llama: Everything you need to know about the open generative AI model

Kyle Wiggers — 8:00 AM PDT · September 8, 2024

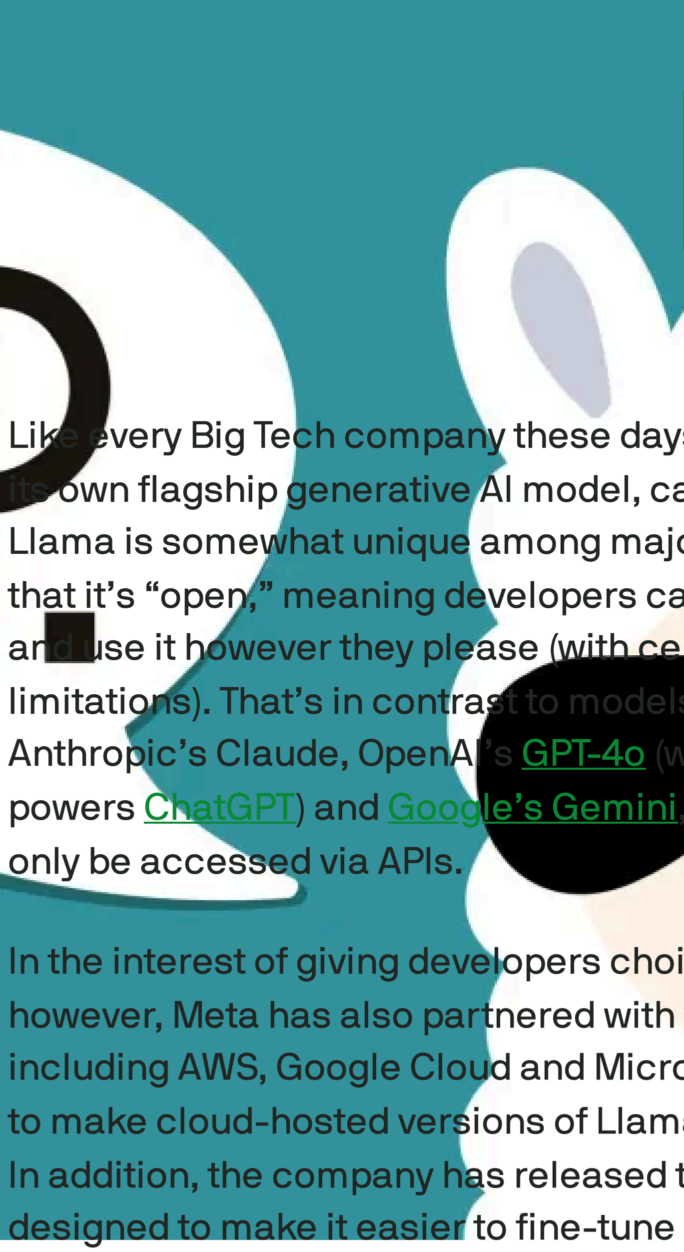

IMAGE CREDITS: LARYSA AMOSOVA VIA GETTY

Like every Big Tech company these days, Meta has its own flagship generative AI model, called Llama. Llama is somewhat unique among major models in that it's "open," meaning developers can download and use it however they please (with certain limitations). That's in contrast to models like Anthropic's Claude, OpenAI's GPT-4o (which powers ChatGPT) and Google's Gemini, which can only be accessed via APIs.

In the interest of giving developers choice, however, Meta has also partnered with vendors, including AWS, Google Cloud and Microsoft Azure, to make cloud-hosted versions of Llama available. In addition, the company has released tools designed to make it easier to fine-tune and customize the model.

Here's everything you need to know about Llama, from its capabilities and editions to where you can use it. We'll keep this post updated as Meta releases upgrades and introduces new dev tools to support the model's use.

What is Llama?

Llama is a family of models — not just one:

- **Llama 8B**
- **Llama 70B**
- **Llama 405B**

Latest in Social

See More

APPS

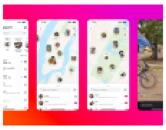

Instagram takes on Snapchat with new 'Instagram Map'

Aisha Malik - 18 hours ago

APPS

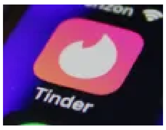

Tinder explores a redesign, dating 'modes,' and college-specific features to boost engagement

Ivan Mehta - 18 hours ago

SOCIAL

Photo-sharing app Locket is banking on a new celebrity-focused feature to fuel its growth

Lauren Forristal - 19 hours ago

The latest versions are **Llama 3.1 8B, Llama 3.1 70B** and **Llama 3.1 405B, which was released** in July 2024. They're trained on web pages in a variety of languages, public code and files on the web, and synthetic data (i.e., data generated by other AI models).

Llama 3.1 8B and Llama 3.1 70B are small, compact models meant to run on devices ranging from laptops to servers. Llama 3.1 405B, on the other hand, is a large-scale model requiring (absent some modifications) data center hardware. Llama 3.1 8B and Llama 3.1 70B are less capable than Llama 3.1 405B, but faster. They're "distilled" versions of 405B, in point of fact, optimized for low storage overhead and latency.

All the Llama models have 128,000-token context windows. (In data science, tokens are subdivided bits of raw data, like the syllables "fan," "tas" and "tic" in the word "fantastic.") A model's context, or context window, refers to input data (e.g., text) that the model considers before generating output (e.g., additional text). Long context can prevent models from "forgetting" the content of recent docs and data, and from veering off topic and extrapolating wrongly.

What can Llama do?

Like other generative AI models, Llama can perform a range of different assistive tasks, like coding and answering basic math questions, as well as summarizing documents in eight languages (English, German, French, Italian, Portuguese, Hindi, Spanish and Thai). Most text-based workloads — think analyzing files like PDFs and spreadsheets — are within its purview; none of the Llama models can process or generate images, although that may change in the near future.

All the latest Llama models can be configured to leverage third-party apps, tools and APIs to complete tasks. They're trained out of the box to use Brave Search to answer questions about recent events, the Wolfram Alpha API for math- and science-related queries, and a Python interpreter for validating code. In addition, Meta says the Llama 3.1 models can use certain tools they haven't seen before (but whether they can *reliably* use those tools is another matter).

Where can I use Llama?

If you're looking to simply chat with Llama, it's powering the Meta AI chatbot experience on Facebook Messenger, WhatsApp, Instagram, Oculus and Meta.ai.

Developers building with Llama can download, use or fine-tune the model across most of the popular cloud platforms. Meta claims it has over 25 partners hosting Llama, including Nvidia, Databricks, Groq, Dell and Snowflake.

Some of these partners have built additional tools and services on top of Llama, including tools that let the models reference proprietary data and enable them to run at lower latencies.

Meta suggests using its smaller models, Llama 8B and Llama 70B, for general-purpose applications like powering chatbots and generating code. Llama 405B, the company says, is better reserved for model distillation — the process of transferring knowledge from a large model to a smaller, more efficient model — and generating synthetic data to train (or fine-tune) alternative models.

Importantly, the Llama license constrains how developers can deploy the model: App developers with more than 700 million monthly users must request a special license from Meta that the company will grant on its discretion.

What tools does Meta offer for Llama?

Alongside Llama, Meta provides tools intended to make the model "safer" to use:

- **Llama Guard**, a moderation framework
- **Prompt Guard**, a tool to protect against prompt injection attacks

- **CyberSecEval**, a cybersecurity risk assessment suite

Llama Guard tries to detect potentially problematic content either fed into — or generated — by a Llama model, including content relating to criminal activity, child exploitation, copyright violations, hate, self-harm and sexual abuse. Developers can customize the categories of blocked content and apply the blocks to all the languages Llama supports out of the box.

Like Llama Guard, Prompt Guard can block text intended for Llama, but only text meant to "attack" the model and get it to behave in undesirable ways. Meta claims that Llama Guard can defend against explicitly malicious prompts (i.e., jailbreaks that attempt to get around Llama's built-in safety filters) in addition to prompts that contain "injected inputs."

As for CyberSecEval, it's less a tool than a collection of benchmarks to measure model security. CyberSecEval can assess the risk a Llama model poses (at least according to Meta's criteria) to app developers and end users in areas like "automated social engineering" and "scaling offensive cyber operations."

Llama's limitations

Llama comes with certain risks and limitations, like all generative AI models.

For instance, it's unclear whether Meta trained Llama on copyrighted content. If it did, users might be liable for infringement if they end up unwittingly using a copyrighted snippet that the model regurgitated.

Meta at one point used copyrighted e-books for AI training despite its own lawyers' warnings, according to recent reporting by Reuters. The company controversially trains its AI on Instagram and Facebook posts, photos and captions, and makes it difficult for users to opt out. What's more, Meta, along with OpenAI, is the subject of an ongoing lawsuit brought by authors, including comedian Sarah Silverman, over the companies' alleged unauthorized use of copyrighted data for model training.

Programming is another area where it's wise to tread lightly when using Llama. That's because Llama might — like its generative AI counterparts — produce buggy or insecure code.

As always, it's best to have a human expert review any AI-generated code before incorporating it into a service or software.

Topics: AI AI Apps Enterprise evergreens Explainer Generative AI Llama Meta meta ai model Social

Kyle Wiggers

AI Editor |

Kyle Wiggers was TechCrunch's AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

View Bio >

KEEP READING

IMAGE CREDITS:
BLOOMBERG / CONTRIBUTOR / GETTY IMAGES

AI

Meta built a code-generating AI model similar to Copilot

Kyle Wiggers — 9:00 AM PDT · May 18, 2023

Meta says it's created a generative AI tool for coding similar to GitHub's Copilot.

The company made the announcement at an event focused on its AI infrastructure efforts, including custom chips Meta's building to accelerate the training of generative AI models. The coding tool, called CodeCompose, isn't available publicly — at least not yet. But Meta says its teams use it internally to get code suggestions for Python and other languages as they type in IDEs like VS Code.

"The underlying model is built on top of public research from [Meta] that we have tuned for our internal use cases and codebases," Michael Bolin, a software engineer at Meta, said in a prerecorded video. "On the product side, we're able to integrate CodeCompose into any surface where our developers or data scientists work with code."

The largest of several CodeCompose models Meta trained has 6.7 billion parameters, a little over half the number of parameters in the model on which Copilot is based. Parameters are the parts of the model learned from historical training data and essentially define the skill of the model on a problem, such as generating text.

CodeCompose was fine-tuned on Meta's first-party code, including internal libraries and frameworks written in Hack, a Meta-developed programming language, so it can incorporate those into its programming suggestions. And its base training data set was filtered of poor coding practices and errors, like deprecated APIs, to reduce the chance that the model recommends a problematic slice of code.

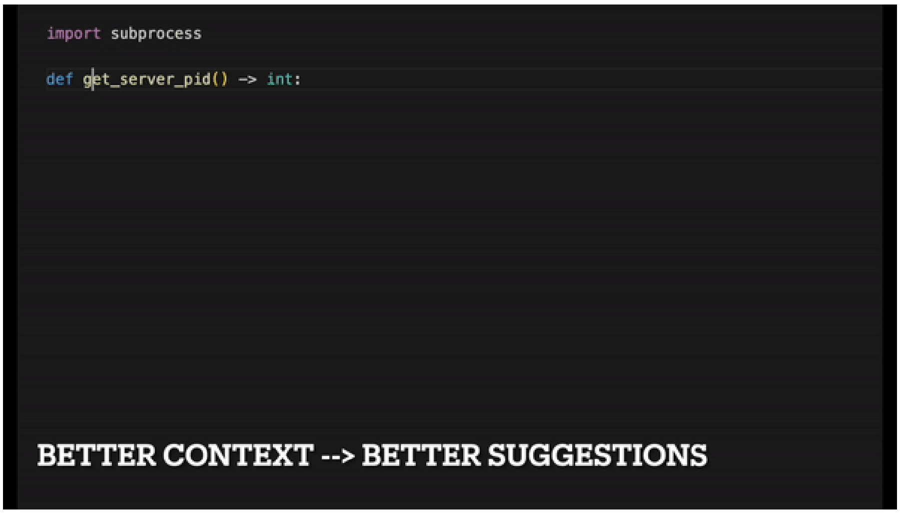

IMAGE CREDITS: META

In practice, CodeCompose makes suggestions like annotations and import statements as a user types. The system can complete single lines of code or multiple lines, optionally filling in entire large chunks of code.

Most Popular

- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a $2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines
- The best dating apps aren't even dating apps

"CodeCompose can take advantage of the surrounding code to provide better suggestions," Bolin continued. "It can also uses code comments as a signal in generating code."

Meta claims that thousands of employees are accepting suggestions from CodeCompose every week and that the acceptance rate is over 20%.

The company didn't address, however, the controversies around code-generating AI.

Microsoft, GitHub and OpenAI are being sued in a class action lawsuit that accuses them of violating copyright law by allowing Copilot to regurgitate sections of licensed code without providing credit. Liability aside, some legal experts have suggested that AI like Copilot could put companies at risk if they were to unwittingly incorporate copyrighted suggestions from the tool into their production software.

It's unclear whether CodeCompose, too, was trained on licensed or copyrighted code — even accidentally. When reached for comment, a Meta spokesperson had this to say:

"CodeCompose was trained on InCoder, which was released by Meta's AI research division. In a paper detailing InCoder, we note that, to train InCoder, 'We collect a corpus of (1) public code with permissive, non-copyleft, open source licenses from GitHub and GitLab and (2) StackOverflow questions, answers and comments.' The only additional training we do for CodeCompose is on Meta's internal code."

Generative coding tools can also introduce insecure code. According to a recent study out of Stanford, software engineers who use code-generating AI systems are more likely to cause security vulnerabilities in the apps they develop. While the study didn't look at CodeCompose specifically, it stands to reason that developers who use it would fall victim to the same.

Bolin stressed that developers needn't follow CodeCompose's suggestions and that security was a "major consideration" in creating the model. "We are extremely excited with our progress on CodeCompose to date, and we believe that our developers are best served by bringing this work in house," he added.

Topics: AI AI code Generative AI Meta

Kyle Wiggers
AI Editor

Kyle Wiggers was TechCrunch's AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

View Bio >