



GOVERNMENT &
POLICY



Google says it will sign
EU's AI code of practice

Google has confirmed it will sign the European Union's general purpose AI code of practice, a voluntary framework that aims to help AI developers implement processes and systems to comply with the bloc's AI Act.

Notably, Meta [earlier this month](#) said it would not sign the code, calling the EU's implementation of its AI legislation "overreach" and stating that Europe was "heading down the wrong path on AI."

Google's commitment comes days before rules for providers of "general-purpose AI models with systemic risk" go into effect on August 2.

Companies likely to be affected by these rules include major names such as Anthropic, Google, Meta, and OpenAI, as well as several other large generative models, and they will have two years to comply fully with the AI Act.

In a [blog post](#) on Wednesday, Kent Walker, president of global affairs at Google, conceded that the final version of the code of practice was better than what the EU proposed initially, but he still noted reservations around the AI Act and the code.

"We remain concerned that the AI Act and Code risk slowing Europe's development and deployment of AI. In particular, departures from EU copyright law, steps that slow approvals, or requirements that expose trade secrets could chill European model

TE DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Instagram takes on Snapchat with new 'Instagram Map'
- Google says hackers stole its customers' data by breaching its Salesforce database
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Spotify raises subscription prices

development and deployment, harming Europe’s competitiveness,” wrote Walker.

By signing the EU’s [code of practice](#), AI companies would agree to follow a slate of guidelines, which include providing updated documentation about their AI tools and services; no training AI on pirated content; and complying with requests from content owners to not use their works in their datasets.

A [risk-based regulation](#) for AI applications, the EU’s landmark AI Act bans some “unacceptable risk” use cases, such as cognitive behavioral manipulation or social scoring. The rules also define a set of “high-risk” uses, including biometrics and facial recognition, and the use of AI in domains like education and employment. The act also requires developers to register AI systems and meet risk- and quality-management obligations.

- The uproar over Vogue’s AI-generated ad isn’t just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

Topics: [AI](#) [EU AI Act](#) [Europe](#) [european union](#)



Ram Iyer

Editor |

Ram is a financial and tech reporter and editor. He covered North American and European M&A, equity, regulatory...

[View Bio](#) >

KEEP READING



STARTUPS



High costs and thin margins threatening AI coding startups

Marina Temkin — 2:05 PM PDT · August 7, 2025

IMAGE CREDITS:

YUICHIRO CHINO / GETTY IMAGES

In February, AI coding startup Windsurf was in talks to raise a big new round at a [\\$2.85 billion valuation led by Kleiner Perkins](#), at double the valuation it hit six months earlier, sources told TechCrunch at the time. That deal didn't happen, according to a source familiar with the matter. Instead, news broke in April that the startup planned to sell itself to OpenAI for roughly the same valuation: [\\$3 billion](#).

While that deal famously fell apart, one bigger question remains: If the startup was growing that fast and attracting VC interest, why would it sell at all?

Insiders tell TechCrunch that for all the popularity and hype around AI coding assistants, they can actually be massively money-losing businesses. Vibe coders generally, and Windsurf in particular, can have such expensive structures that their gross margins are "very negative," one person close to Windsurf told TechCrunch. Meaning it cost more to run the product than the startup could charge for it.

This is due to the high costs of using large language models (LLMs), the person explained. AI coding assistants are particularly pressured to always offer the most recent, most advanced, and most expensive LLMs because model makers are particularly fine-tuning their latest models for improvements in coding and related tasks like debugging.

This is a challenge compounded by fierce competition in the vibe-coding and code-assist market. Rivals include companies that already have huge customer bases like Anysphere's Cursor and GitHub Copilot.

The most straightforward path to improving margins in this business involves the startups building their own models, thereby eliminating

Newsletters

[See More ↗](#)

Subscribe for the industry's biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch's coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

costs of paying suppliers like Anthropic and OpenAI.

“It’s a very expensive business to run if you’re not going to be in the model game,” said the person.




Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

Latest Security
Startups AI
Venture Apps
Apple



Events
Podcasts
Newsletters



Selling the business was a strategic move to lock in a high return before it could be undermined by the very companies that supplied its AI, including OpenAI and Anthropic, which were also entering the AI coding market.

Multiple people believe that the same pressure on margins Windsurf faced could be impacting Anysphere, the maker of Cursor, as well as vibe coders like Lovable, Replit, and others.

“Margins on all of the ‘code gen’ products are either neutral or negative. They’re absolutely abysmal,” said Nicholas Charriere, founder of Mocha, a vibe-coding startup and back-end hosting solution serving small and medium businesses (SMBs). He added that he believes the variable costs for all the startups in the sector are very close, likely within 10% to 15% of one another.

Unlike Windsurf, Anysphere has been growing so fast that it intends to remain an independent company, having already [turned down acquisition offers](#), including, reports say, from OpenAI.

And Anysphere [announced in January](#) that it is attempting to build its own model, which could give it more control over its expenses. In July, the startup [hired two leaders from Anthropic’s Claude Code](#) team, the Information reported, but two weeks later, these employees [returned](#) to work at Anthropic.

In addition to building a model, Anysphere could expect the cost of LLMs to decrease over time.

“That’s what everyone’s banking on,” said Erik Nordlander, a general partner at Google Ventures. “The inference cost today, that’s the most expensive it’s ever going to be.”

It’s not entirely clear how true that is. Rather than falling as expected, the [cost of some of the latest AI models has risen](#), as they use more time and computational resources to handle complicated, multistep tasks.

When that will change remains to be seen. On Thursday, for instance, [OpenAI introduced](#) a new flagship model, GPT-5, with fees that are significantly less than its competitor, Anthropic's Claude Opus 4.1. And Anysphere [immediately offered](#) this model as a choice for Cursor users.

Anysphere has also recently changed its pricing structure to pass along the increased costs of running Anthropic's latest Claude model, particularly to its most active users. The move caught some of [Cursor customers by surprise](#), since they didn't expect additional charges on top of its \$20-per-month Pro plan. Anysphere CEO Michael Truell later apologized for unclear communication about the pricing change in a [blog post](#).

This is the rock and the hard place. Although Cursor is one of the most popular AI applications, having reached [\\$500 million in ARR in June](#), the company's user base may not be so loyal to the product if another company develops a tool that is superior to Cursor, investors say.

Anysphere didn't respond to a request for comment.

Given the competitive landscape and the costs, Windsurf's decision to get out may prove to be understandable. After the OpenAI deal fell through, the founders and key employees left to join Google in a deal that led to a [\\$2.4 billion payout](#) to key shareholders. The remaining business then [sold itself to Cognition](#).

While many, including prominent VCs, criticized Mohan for leaving approximately 200 employees without roles at Google, a source familiar with the deal insisted the acquisition actually maximized the outcomes for all employees.

Beyond Cursor, other AI coding tools are also among the fastest growing startups of the LLM generation, like Replit, Lovable, and Bolt, and all of them rely on model makers as well.

Additionally, if this extremely popular business sector, already generating hundreds of millions in revenue or more a year, has difficulty building on top of model makers, what might it mean for other, more nascent industries?

Topics: [AI](#) [coding assistants](#) [cursor](#) [Startups](#)

[Venture](#) [Windsurf](#)



[f](#) [X](#) [in](#) [t](#) [✉](#) [🔗](#)



Marina Temkin

Reporter, Venture | [X](#)

Marina Temkin is a venture capital and startups reporter at TechCrunch. Prior to joining TechCrunch, she wrote abo...

[View Bio](#) >

—— KEEP READING ——



AI



Microsoft brings OpenAI's smallest open model to Windows users

Ram Iyer — 5:47 AM PDT · August 6, 2025

IMAGE CREDITS:

DAVID RYDER / BLOOMBERG
(PHOTOMOSH/MODIFIED) / GETTY IMAGES

Microsoft is making OpenAI's new free and open GPT model, [gpt-oss-20b](#), available to Windows 11 users via [Windows AI Foundry](#), the tech giant's platform that lets users tap AI features, APIs, and popular open source models on their computers.

The company said in a [blog post](#) that gpt-oss-20b is "tool-savvy and lightweight," adding: "Optimized for agentic tasks like code execution and tool use, it runs efficiently on a range of Windows hardware, with support for more devices coming soon. It's perfect for building autonomous assistants or embedding AI into real-world workflows, even in bandwidth-constrained environments."

Latest in AI

[See More](#)

TRANSPORTATION



Tesla shuts down Dojo, the AI training supercomputer that Musk said would be key to full self-driving

Rebecca Bellan - 6 hours ago

Launched on Tuesday, OpenAI's gpt-oss-20b can run on consumer PCs and laptops with at least 16GB of VRAM, which a modern GPU from Nvidia or Radeon would have. OpenAI said the model was trained using [high-compute reinforcement learning](#), which helps it excel at powering AI agents and call tools, such as web search or Python code execution, as part of its chain-of-thought process.

However, the model is text-only, meaning it will not be able to process or generate images and audio like the company's other models.

Also, the model has a habit of hallucinating a lot: gpt-oss-20b hallucinated in response to 53% of questions on PersonQA, OpenAI's in-house benchmark for measuring the accuracy of a model's knowledge about people.

Microsoft says it will bring the model to macOS soon, as well as other devices, though the company did not name any more. Microsoft is also making both gpt-oss-20b and older model gpt-oss-120b available via its hosted Azure AI Foundry platform.

The two new models are also [available on Amazon's AWS](#).

STARTUPS



High costs and thin margins threatening AI coding startups

Marina Temkin - 7 hours ago

AI



OpenAI's GPT-5 is here

Maxwell Zeff - 11 hours ago

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

REGISTER NOW

Topics: [AI](#) [gpt](#) [Microsoft](#) [OpenAI](#) [Windows](#)





Ram Iyer

Editor | [X](#) [in](#)

Ram is a financial and tech reporter and editor. He covered North American and European M&A, equity, regulatory...

[View Bio](#) >

Loading the next article

