

Advertisement

AI



Grok Imagine, xAI's new
AI image and video
generator, lets you make
NSFW content

Elon Musk's AI company has officially rolled out Grok Imagine, XAI's image and video generator, to all SuperGrok and Premium+ X subscribers on its iOS app. And true to form for Musk, who positions Grok as an unfiltered, boundary-pushing AI, the generator allows users to make NSFW content.

Grok Imagine, which promises to turn text or image prompts into a 15-second video featuring native audio, has a "spicy mode" that allows users to generate sexually explicit content, including partial female nudity. There are limits to how explicit one can get. Many of our spicier prompts — made in the name of Journalism! — generate blurred-out images that are "moderated" and therefore inaccessible. We were, however, able to generate semi-nude imagery.

IMAGE CREDITS:

TEC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion

Imagine with [@Grok](#)

The media could not be played.

Reload

1:15 AM · Aug 4, 2025



88.6K



Reply



Copy link

[Read 14.1K replies](#)

- Your public ChatGPT queries are getting indexed by Google and other search engines
- The best dating apps aren't even dating apps

The NSFW content is unsurprising for xAI, given the release last month of a [raunchy, hyper-sexualized anime AI companion](#). But just as Grok's unrestrained nature was entertaining until it started spewing hateful, antisemitic, misogynistic content, Grok Imagine could be poised to bring its own set of unintended consequences.

[CNBC first reported](#) the existence of a "[spicy mode](#)" last week after xAI employee Mati Roy said in a now-deleted post on X: "Grok Imagine videos have a spicy mode that can [do nudity](#)."

TechCrunch has reached out to xAI for more information.

The existing limitations with Grok Imagine are somewhat heartening given the model also lets you create content of celebrities — anyone from Donald Trump to Taylor Swift — and there appear to be additional restrictions on those. For example, TechCrunch tried, and failed, to generate an image of Trump pregnant. Grok Imagine only generated images of Trump holding a baby or next to a pregnant woman.

While it's still early days for Grok Imagine, which aims to compete with incumbents like Google DeepMind, OpenAI, Runway, and Chinese rivals, the images and videos generated of humans are still a bit lost in the uncanny valley, with waxy-looking skin that verges on cartoonish at times.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

[REGISTER NOW](#)

Still, the generator is impressive. It produces images in seconds from a text prompt, and continues to auto-generate new images as you scroll. Those images can then be animated into

stylized videos. The user interface is also seamless and intuitive.

Musk said [on X](#) that the model would “get better every day.”



Rebecca Bellan

@RebeccaBellan · [Follow](#)



Testing out Grok Imagine 🙄🙄

The media could not be played.

[Reload](#)

2:46 PM · Aug 4, 2025



4



[Reply](#)



[Copy link](#)

[Read 1 reply](#)

Topics: [AI](#) [AI video generation](#) [Elon Musk](#) [Grok](#)

[grok imagine](#) [Media & Entertainment](#) [xAI](#)





Rebecca Bellan

Senior Reporter |  

Rebecca Bellan is a senior reporter at TechCrunch, where she covers Tesla and Elon Musk’s broader empire,...

View Bio >

KEEP READING



AI

xAI, Elon Musk’s AI startup, launches an API

Kyle Wiggers — 9:51 AM PDT · October 21, 2024

IMAGE CREDITS:
AXELLE/BAUER-GRIFFIN/FILMMAGIC / GETTY IMAGES

In August, Elon Musk’s xAI promised to make [Grok](#), the company’s flagship generative AI model powering a number of features on X, available via

an API. Now, that API has [arrived](#) — albeit a bit bare-bones at the moment.

The [xAI API](#) only has a single model, “grok-beta,” priced at \$5 per million input tokens (~750,000 words) or \$15 per million output tokens. Tokens are subdivided bits of raw data, like the syllables “fan,” “tas,” and “tic” in the word “fantastic.”

It’s unclear which Grok model “grok-beta” might be; Grok 2 is the very latest on X. The documentation for the API refers to Grok 2 and Grok mini, a lightweight and more affordable version of Grok — so it could well be a technical issue.

Language models

Language models are used to generate text, images, and audio.

Model	Input	Output	Context	Text	Image	Completion	RPS	RPM	RPH	RPD
grok-beta	TEXT	TEXT	131072	\$5.00	-	\$15.00	1	10	-	-

IMAGE CREDITS: XAI

Some users on X also [reported](#) encountering problems paying for usage credits.

The xAI API supports function calling, which connects Grok models to external tools such as databases and search engines. And, although they don’t appear to be live yet, the documentation hints at vision models capable of analyzing both text and images.

Musk formed xAI last year. Soon after, the company — which [recently moved into OpenAI’s old offices](#) — released its first Grok model on X to X Premium+ users, who pay \$16 a month.

Grok has what Musk has described as “a rebellious streak” — a willingness to answer “spicy questions that are rejected by most other AI systems.” Told to be vulgar, for example, Grok will happily oblige, spewing profanities and colorful language you won’t hear from [ChatGPT](#).

Advertisement

Newsletters

See More ↗

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Grok has become increasingly ingrained in the X experience. Thanks to an [integration](#) with the open image generator Flux, Grok can generate images on X (without guardrails, controversially). It summarizes news and trending events ([often with mistakes](#), mind you), and it may [enhance](#) X's search capabilities, [account bios](#), post analytics, and reply functions in the future.

xAI is sprinting to catch up to formidable competitors like OpenAI and Anthropic in the generative AI race. In May, xAI [raised](#) \$6 billion in a funding round led by investors including Andreessen Horowitz, Sequoia Capital, and Fidelity.

Musk often asserts that X's data gives xAI a substantial leg up compared to rivals. This month, X [changed](#) its privacy policy to allow third parties, including xAI, to train models on X posts.

As part of the pitch for xAI's \$6 billion funding round, the startup also outlined a vision where its models would be trained on data from Musk's

various companies, including Tesla, SpaceX, and The Boring Company, and its models could then improve technology across those companies.

Tesla shareholders, for one, weren't on board with that plan. Several sued Musk over the decision to start xAI, arguing that Musk has diverted both talent and resources from Tesla to what is essentially a

Latest Security
Startups AI
Venture Apps
Apple



Events
Podcasts
Newsletters



Kyle Wiggers

AI Editor |

Kyle Wiggers was TechCrunch's AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

[View Bio](#) >

———— KEEP READING ————

Leak reveals Grok might soon edit your spreadsheets

Rebecca Bellan — 11:16 AM PDT · June 23, 2025

Leaked code suggests xAI is developing an advanced file editor for Grok with spreadsheet support, signaling the company's push to compete with OpenAI, Google, and Microsoft by embedding AI copilots into productivity tools.

IMAGE CREDITS: [JAAP ARRIENS/NURPHOTO](#)

"You can talk to Grok and ask it to assist you at the same time you're editing the files!" [writes](#) reverse engineer Nima Owji, who leaked the finding.

BREAKING: xAI is working on an advanced FILE EDITOR for GROK!

It even supports SPREADSHEETS!

You can talk to Grok and ask it to assist you at the same time you're editing the files!

pic.twitter.com/9vIKRZj6Wn

— Nima Owji (@nima_owji) [June 22, 2025](#)

Latest in AI

[See More](#)

AI



Google denies AI search features are killing website traffic

Sarah Perez - 11 hours ago

AI



First impressions of Alexa+, Amazon's upgraded, AI-powered digital assistant

Sarah Perez - 13 hours ago

AI



Google takes on ChatGPT's Study Mode with new 'Guided Learning' tool in Gemini

Aisha Malik - 13 hours ago

TechCrunch has reached out to xAI to confirm the findings and learn more.

While xAI hasn't explicitly detailed its strategy for pursuing interactive, multimodal AI workspaces, it has dropped a series of announcements that point to how the company is thinking about these tools. In April 2025, xAI launched [Grok Studio](#), a split-screen workspace that lets users collaborate with Grok on generating documents, code, reports, and browser games. It also launched the ability to [create Workspaces](#) that let you organize files and conversations in a single place.

While OpenAI and Microsoft have similar tools, Google's Gemini [Workspace](#) for Sheets, Docs, and Gmail appears to be the most similar to what xAI is reportedly building. Google's tools can edit Docs and Sheets and allow you to chat with Gemini while looking at or editing documents. The difference is that Gemini Workspace only works within Google's ecosystem.

It's not clear what types of files xAI's editor might support aside from spreadsheets, or whether xAI plans to build a full productivity suite that could compete with Google Workspace or Microsoft 365.

If Owji's findings are true, the advanced editor would be a step toward Elon Musk's ambitions to turn X into an "everything app" that includes docs, chat, payments, and social media.

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and

sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Topics: [AI](#) [Grok](#) [xAI](#)

[f](#) [X](#) [in](#) [reddit](#) [envelope](#) [link](#)



Rebecca Bellan

Senior Reporter | [X](#) [butterfly](#)

Rebecca Bellan is a senior reporter at TechCrunch, where she covers Tesla and Elon Musk's broader empire,...

View Bio >

KEEP READING

AI

[f](#) [X](#) [in](#) [reddit](#) [envelope](#) [link](#)

Did xAI lie about Grok 3's benchmarks?

Kyle Wiggers — 2:55 PM PST · February 22, 2025

Debates over AI benchmarks — and how they're reported by AI labs — are spilling out into public view.

This week, an OpenAI employee [accused](#) Elon Musk's AI company, xAI, of publishing misleading benchmark results for its latest AI model, Grok 3. One of the co-founders of xAI, Igor Babuschkin, [insisted](#) that the company was in the right.

IMAGE CREDITS: JAAP ARRIENS/NURPHOTO / GETTY IMAGES
The truth lies somewhere in between.

In a [post on xAI's blog](#), the company published a graph showing Grok 3's performance on AIME 2025, a collection of challenging math questions from a recent invitational mathematics exam. Some experts have [questioned AIME's validity as an AI benchmark](#). Nevertheless, AIME 2025 and older versions of the test are commonly used to probe a model's math ability.

xAI's graph showed two variants of Grok 3, Grok 3 Reasoning Beta and Grok 3 mini Reasoning, beating OpenAI's best-performing available model, [o3-mini-high](#), on AIME 2025. But OpenAI employees on X were quick to point out that xAI's graph didn't include o3-mini-high's AIME 2025 score at "cons@64."


DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices

What is cons@64, you might ask? Well, it's short for "consensus@64," and it basically gives a model 64 tries to answer each problem in a benchmark and takes the answers generated most frequently as the final answers. As you can imagine, cons@64 tends to boost models' benchmark scores quite a bit, and omitting it from a graph might make it appear as though one model surpasses another when in reality, that isn't the case.

Grok 3 Reasoning Beta and Grok 3 mini Reasoning's scores for AIME 2025 at "@1" — meaning the first score the models got on the benchmark — fall below o3-mini-high's score. Grok 3 Reasoning Beta also trails ever so slightly behind OpenAI's [o1 model](#) set to "medium" computing. Yet xAI is [advertising Grok 3](#) as the "world's smartest AI."

- The uproar over Vogue's AI-generated ad isn't just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines
- The best dating apps aren't even dating apps



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Babuschkin [argued on X](#) that OpenAI has published similarly misleading benchmark charts in the past — albeit charts comparing the performance of its

own models. A more neutral party in the debate put together a more “accurate” graph showing nearly every model’s performance at cons@64:

Hilarious how some people see my plot as attack on OpenAI and others as attack on Grok while in reality it’s DeepSeek propaganda

(I actually believe Grok looks good there, and openAI’s TTC chicanery behind o3-mini-**high**-pass@*”””1””* deserves more scrutiny.)

<https://t.co/dJqIJpcJh8>
pic.twitter.com/3WH8FOUfic

— Teortaxes  (DeepSeek 推特  铁粉 2023 – ∞)
(@teortaxesTex) [February 20, 2025](#)

But as AI researcher Nathan Lambert [pointed out in a post](#), perhaps the most important metric remains a mystery: the computational (and monetary) cost it took for each model to achieve its best score. That just goes to show how little most AI benchmarks communicate about models’ limitations — and their strengths.

Topics: [AI](#) [benchmarks](#) [Grok](#) [OpenAI](#) [xAI](#)



Kyle Wiggers

AI Editor |    

Kyle Wiggers was TechCrunch’s AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

[View Bio](#) >

Loading the next article





[TechCrunch Staff](#)

[Contact Us](#)

[Advertise](#)

[Crunchboard Jobs](#)

[Site Map](#)

[Terms of Service](#)

[Privacy Policy](#)

[RSS Terms of Use](#)

[OpenAI](#)

[Perplexity](#)

[Cisco](#)

[DeepMind](#)

[Windsurf](#)

[Tech Layoffs](#)

[ChatGPT](#)