

Advertisement

AI



Meta to spend up to \$72B
on AI infrastructure in
2025 as compute arms
race escalates

Meta is pouring money into the physical and technical infrastructure needed to scale its AI ambitions, the company said Wednesday in its second-quarter earnings report. It plans to more than double its spend on building AI infrastructure, like data centers and servers.

“We currently expect 2025 capital expenditures, including principal payments on finance leases, to be in the range of \$66-72 billion...up approximately \$30 billion year-over-year at the midpoint,” Meta said.

That’s an aggressive capex growth, and one that Meta plans to continue onwards to 2026. The company said it expects a similarly large increase in spend on AI infrastructure next year as the company continues to “aggressively [pursue] opportunities to bring additional capacity online to meet the needs of [its] artificial intelligence efforts and business operations.”

“We expect that developing leading AI infrastructure will be a core advantage in developing the best AI models and product experiences, so we expect to ramp our investments significantly in 2026 to support that work,” said Susan Li, Meta CFO, during the company’s Wednesday earnings call.

Li also noted that while Meta expects to finance most of its AI spend on its own, the company is


TECH DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government’s website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue’s AI-generated ad isn’t just about fashion

exploring ways to work with financial partners to co-develop data centers.

“We don’t have any finalized transactions to announce, but we generally believe that there will be models here that will attract significant external financing to support large-scale data center projects that are developed using our ability to build world-class infrastructure, while providing us with flexibility should our infrastructure requirements change over time,” Li said.

Meta has announced two major AI “titan clusters.” The first is Prometheus in Ohio, which is poised to be among the first AI superclusters to hit 1 gigawatt of compute power when it comes online in 2026. Then there’s [Hyperion](#), a cluster in Louisiana that Meta CEO Mark Zuckerberg has bragged would have a footprint the size of Manhattan and could scale up to 5 gigawatts over several years. On top of those, Meta has several other unnamed titan-scale clusters underway.

Meta’s data center projects promise to soak up enough energy to power millions of homes, pulling that electricity from nearby communities. One of the company’s projects in Newton County, Georgia, has already [caused the water taps to run dry in some residents’ homes](#).

Meta also noted in its earnings report that it expects its second-largest driver of growth to be employee compensation as the company spends millions, and possibly even billions, to poach talented AI engineers and researchers to work for Meta’s newly formed business unit, Superintelligence Labs.

Before earnings, Zuckerberg shared his vision for “[personal superintelligence](#),” the idea that AI should help individual people live their best lives,

- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

mainly through the medium of Meta’s smart glasses and virtual reality headsets.

Meta’s stock surged 10% in after-hours as investors responded to Meta’s overall performance in the quarter and better-than-expected outlook for the third quarter. Meta reported revenue of \$47.5 billion in the second quarter, with expectations to hit between \$47.5 billion and \$50.5 billion in Q3.

Advertising drove Meta’s revenue gains, fueled by AI tools — like AI-powered translations and video generation — to help advertisers create more meaningful and targeted campaigns.

The company’s Reality Labs segment, however, saw a \$4.5 billion loss.

This article was updated to reflect Meta’s additional paths to funding its AI infrastructure builds.

Topics: [AI](#) [ai data centers](#) [Mark Zuckerberg](#)

[Media & Entertainment](#) [Meta](#)

[meta q2 2025 earnings](#)

[f](#) [X](#) [in](#) [t](#) [✉](#) [🔗](#)



Rebecca Bellan

Senior Reporter | [X](#) [🦋](#)

Rebecca Bellan is a senior reporter at TechCrunch, where she covers Tesla and Elon Musk’s broader empire,...

View Bio >

———— KEEP READING ————



APPS



Meta unveils a new, more efficient Llama model

Kyle Wiggers — 9:16 AM PST · December 6, 2024

IMAGE CREDITS: LARYSA AMOSOVA VIA GETTY

Meta has [announced](#) the newest addition to its [Llama](#) family of generative AI models: Llama 3.3 70B.

In a [post](#) on X, Ahmad Al-Dahle, VP of generative AI at Meta, said that the text-only Llama 3.3 70B delivers the performance of Meta's largest Llama model, [Llama 3.1 405B](#), at lower cost.

"By leveraging the latest advancements in post-training techniques ... this model improves core performance at a significantly lower cost," Al-Dahle wrote.

Advertisement

AI-Dahle published a chart showing Llama 3.3 70B outperforming Google’s Gemini 1.5 Pro, OpenAI’s GPT-4o, and Amazon’s newly released Nova Pro on a number of industry benchmarks, including MMLU, which evaluates a model’s ability to understand language. Via email, a Meta spokesperson said that the model should deliver improvements in areas like math, general knowledge, instruction following, and app use.

Introducing Llama 3.3 — a new 70B model that delivers the performance of our 405B model but is easier & more cost-efficient to run. By leveraging the latest advancements in post-training techniques including online preference optimization, this model improves core performance at... pic.twitter.com/6oQ7b3Yuzc

— Ahmad Al-Dahle (@Ahmad_AI_Dahle)
[December 6, 2024](#)

Llama 3.3 70B, which is available for [download](#) from the AI dev platform Hugging Face and other sources, including the official Llama [website](#), is Meta’s latest play to dominate the AI field with “open” models that can be used and commercialized for a range of applications.

Meta’s terms [constrain](#) how certain developers can use Llama models; platforms with more than 700 million monthly users must request a special license. But for many, it’s immaterial that Llama models aren’t “[open](#)” in the strictest sense. Case in point, Llama has racked up more than 650 million downloads, according to Meta.

TC

Techcrunch event

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the

Newsletters

[See More ↗](#)

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

Advertisement

Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

Meta has leveraged Llama internally as well. [Meta AI](#), the company's AI assistant, which is powered entirely by Llama models, now has nearly 600 million monthly active users, [per](#) Meta CEO Mark Zuckerberg. Zuckerberg claims that Meta AI is on track to be the most-used AI assistant in the world.

For Meta, the open nature of Llama has been a blessing and a curse. In November, a [report](#) alleged that Chinese military researchers had used a Llama model to develop a defense chatbot. Meta responded by making its Llama models [available to U.S. defense contractors](#).

Meta has also voiced concerns about its ability to comply with [the AI Act](#), the EU law that establishes a regulatory framework for AI, calling the law's implementation "too unpredictable" for its open release strategy. A related issue for the company

are provisions in the GDPR, the EU’s privacy law, pertaining to AI training. Meta trains AI models on the public data of Instagram and Facebook users who haven’t opted out — data that in Europe is subject to GDPR guarantees.

EU regulators earlier this year requested that Meta halt training on European user data while they assessed the company’s GDPR compliance. Meta relented, while at the same time endorsing an [open letter](#) calling for “a modern interpretation” of GDPR that doesn’t “reject progress.”

Meta, not immune to the technical challenges other AI labs are [encountering](#), is ramping up its computing infrastructure to train and serve future generations of Llama. The company announced Wednesday that it would build a \$10 billion AI data center in Louisiana — the largest AI data center Meta has ever built.

Zuckerberg [said](#) on Meta’s Q4 earnings call in August that to train the next major set of Llama models, Llama 4, the company will need 10x more compute than what was needed to train Llama 3. Meta has procured a cluster of more than 100,000 Nvidia GPUs for model development, rivaling the resources of competition like xAI.

Training generative AI models is a costly business. Meta’s capital expenditures rose nearly 33% to \$8.5 billion in Q2 2024, up from \$6.4 billion a year earlier, driven by investments in servers, data centers, and network infrastructure.

Topics: [AI](#) [AI](#) [Apps](#) [Generative AI](#) [Llama](#) [Llama 3](#)

[Llama 3.3 70b](#) [Mark Zuckerberg](#) [Meta](#) [open](#)

[open source](#)



Kyle Wiggers

AI Editor | X reddit @ in

Kyle Wiggers was TechCrunch’s AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

View Bio >

KEEP READING



AI

Zuckerberg signals Meta won’t open source all of its ‘superintelligence’ AI models

Rebecca Bellan — 10:56 AM PDT · July 30, 2025

Events

Podcasts

Newsletters

should be shared with the world as broadly as possible,” wrote Zuckerberg. “That said, superintelligence will raise novel safety concerns. We’ll need to be rigorous about mitigating these risks and careful about what we choose to open source.”

That wording about open source is significant. Zuckerberg has historically positioned Meta’s Llama family of open models as the company’s key differentiator from competitors like OpenAI, xAI, and Google DeepMind. Meta’s goal has been to create open AI models that were as good as or better than those closed models. In a [2024 letter](#), Zuckerberg wrote, “Starting next year, we expect future Llama models to become the most advanced in the industry.”

Zuckerberg has previously left himself room to maneuver on this commitment. “If at some point however there’s some qualitative change in what the thing is capable of, and we feel like it’s not responsible to open source it, then we won’t,” he said in a [podcast](#) last year.

And while many say Llama doesn’t fit the strict definition of open source AI — partly because Meta hasn’t released its massive training datasets — Zuckerberg’s words point to a possible change in priority: Open source may no longer be the default for Meta’s cutting-edge AI.

AI



First impressions of Alexa+, Amazon’s upgraded, AI-powered digital assistant

Sarah Perez - 12 hours ago

AI



Google takes on ChatGPT’s Study Mode with new ‘Guided Learning’ tool in Gemini

Aisha Malik - 12 hours ago

There's a reason why Meta's rivals keep their models closed. Closed models give companies more control over monetizing their products. Zuckerberg pointed out last year that Meta's business isn't reliant on selling access to AI models, so "releasing Llama doesn't undercut our revenue, sustainability, or ability to invest in research like it does for closed providers." Meta, of course, makes most of its money from selling internet advertising.

TC

Techcrunch event

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

Still, that stated viewpoint on open models was *before* Meta started to feel like it was falling behind competitors, and executives became [obsessed with beating OpenAI's GPT-4 model](#) while developing Llama 3.

Cut to June 2025, when Meta began its public AGI sprint in earnest by investing [\\$14.3 billion in Scale AI, acquiring Scale's founder](#) and CEO, and [restructuring its AI efforts](#) under a new unit called Meta Superintelligence Labs. Meta has spent [billions of dollars](#) to acquire researchers and engineers from top AI firms and build out new data centers.

[Recent reports indicate](#) that all that investment has led Meta to pause testing on its latest Llama model, Behemoth, and instead focus efforts on developing a closed model.

With Zuckerberg's mission for introducing "personal superintelligence" to the world — a decided shift from the rivals he says are working on "automating all valuable work" — his AI monetization strategy is taking shape. It's clear from Zuckerberg's words today that Meta plans to deliver "personal superintelligence" through its own products like augmented reality glasses and virtual reality headsets.

"Personal devices like glasses that understand our context because they can see what we see, hear what we hear, and interact with us throughout the day will become our primary computing devices," Zuckerberg wrote in Wednesday's letter.

When asked about Meta potentially keeping its most advanced models closed, a Meta spokesperson said that the company remains committed to open source AI and said it also expects to train closed source models in the future.

“Our position on open source AI is unchanged,” a spokesperson said. “We plan to continue releasing leading open source models. We haven’t released everything we’ve developed historically and we expect to continue training a mix of open and closed models going forward.”

This article was updated with more information about Mark Zuckerberg’s stance on open AI models.

Topics: [AI](#) [Mark Zuckerberg](#) [Meta](#)

[meta superintelligence labs](#) [superintelligence](#)



Rebecca Bellan

Senior Reporter |

Rebecca Bellan is a senior reporter at TechCrunch, where she covers Tesla and Elon Musk’s broader empire,...

[View Bio](#) >

———— KEEP READING ————

AI



Zuckerberg says Meta will need 10x more computing power to train Llama 4 than Llama 3

Ivan Mehta — 12:53 AM PDT · August 1, 2024

Meta, which develops one of the biggest foundational open source large language models, Llama, believes it will need significantly more computing power to train models in the future.

Mark Zuckerberg said on Meta's second-quarter earnings call on Tuesday that to train Llama 4, the company will need 10x more compute than what was needed to train Llama 3. But he still wants Meta to build capacity to train models rather than fall behind its competitors.

"The amount of computing needed to train Llama 4 will likely be almost 10 times more than what we used to train Llama 3, and future models will continue to grow beyond that," Zuckerberg said.

"It's hard to predict how this will trend multiple generations out into the future. But at this point, I'd rather risk building capacity before it is needed rather than too late, given the long lead times for spinning up new inference projects."

Meta released [Llama 3 with 8 billion parameters in April](#). The company last week released an upgraded version of the model, called [Llama 3.1 405B, which](#)


DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service

[had 405 billion parameters](#), making it Meta's biggest open source model.

Meta's CFO, Susan Li, also said the company is thinking about different data center projects and building capacity to train future AI models. She said Meta expects this investment to increase capital expenditures in 2025.

Training large language models can be a costly business. Meta's capital expenditures rose nearly 33% to \$8.5 billion in Q2 2024, from \$6.4 billion a year earlier, driven by investments in servers, data centers and network infrastructure.

TC

Techcrunch event

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise.

- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines
- The best dating apps aren't even dating apps

REGISTER NOW

According to a report from [The Information](#), OpenAI spends \$3 billion on training models and an additional \$4 billion on renting servers at a discount rate from Microsoft.

“As we scale generative AI training capacity to advance our foundation models, we’ll continue to build our infrastructure in a way that provides us with flexibility in how we use it over time. This will allow us to direct training capacity to GenAI inference or to our core ranking and recommendation work, when we expect that doing so would be more valuable,” Li said during the call.

During the call, Meta also talked about its consumer-facing [Meta AI’s usage and said India is the largest market of its chatbot](#). But Li noted that the company doesn’t expect Gen AI products to contribute to revenue in a significant way.

Topics: [AI](#) [ai models](#) [Llama](#) [Mark Zuckerberg](#)

[Meta](#) [meta ai](#)



Ivan Mehta

[X](#) [in](#)

Ivan covers global consumer tech developments at TechCrunch. He is based out of India and has previously...

[View Bio](#) >

Loading the next article





[TechCrunch Staff](#)

[Contact Us](#)

[Advertise](#)

[Crunchboard Jobs](#)

[Site Map](#)

[Terms of Service](#)

[Privacy Policy](#)

[RSS Terms of Use](#)

[OpenAI](#)

[Perplexity](#)

[Cisco](#)

[DeepMind](#)

[Windsurf](#)

[Tech Layoffs](#)

[ChatGPT](#)