

Advertisement

AI



Microsoft brings
OpenAI's smallest open
model to Windows users

Microsoft is making OpenAI's new free and open GPT model, [gpt-oss-20b](#), available to Windows 11 users via [Windows AI Foundry](#), the tech giant's platform that lets users tap AI features, APIs, and popular open source models on their computers.

The company said in a [blog post](#) that gpt-oss-20b is "tool-savvy and lightweight," adding: "Optimized for agentic tasks like code execution and tool use, it runs efficiently on a range of Windows hardware, with support for more devices coming soon. It's perfect for building autonomous assistants or embedding AI into real-world workflows, even in bandwidth-constrained environments."

DAVID RYDER / BLOOMBERG

(PHOTOMOSH/MODIFIED) / GETTY IMAGES

Launched on Tuesday, OpenAI's gpt-oss-20b can run on consumer PCs and laptops with at least 16GB of VRAM, which a modern GPU from Nvidia or Radeon would have. OpenAI said the model was trained using [high-compute reinforcement learning](#), which helps it excel at powering AI agents and call tools, such as web search or Python code execution, as part of its chain-of-thought process.

However, the model is text-only, meaning it will not be able to process or generate images and audio like the company's other models.

Also, the model has a habit of hallucinating a lot: gpt-oss-20b hallucinated in response to 53% of questions on PersonQA, OpenAI's in-house

TECH DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion

benchmark for measuring the accuracy of a model's knowledge about people.

Microsoft says it will bring the model to macOS soon, as well as other devices, though the company did not name any more. Microsoft is also making both gpt-oss-20b and older model gpt-oss-120b available via its hosted Azure AI Foundry platform.

The two new models are also [available on Amazon's AWS](#).

- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Topics: [AI](#) [gpt](#) [Microsoft](#) [OpenAI](#) [Windows](#)



Ram Iyer

Editor |

Ram is a financial and tech reporter and editor. He covered North American and European M&A, equity, regulatory...

[View Bio](#) >

KEEP READING



AI



OpenAI launches two ‘open’ AI reasoning models

Maxwell Zeff — 10:00 AM PDT · August 5, 2025

IMAGE CREDITS:

JAKUB PORZYCKI/NURPHOTO / GETTY IMAGES

OpenAI announced Tuesday the launch of two open-weight AI reasoning models with similar capabilities [to its o-series](#). Both are freely available to download from the online developer platform [Hugging Face](#), the company said, describing the models as “state of the art” when measured across several benchmarks for comparing open models.

Advertisement

The models come in two sizes: a larger and more capable gpt-oss-120b model that can run on a single Nvidia GPU, and a lighter-weight gpt-oss-20b model that can run on a consumer laptop with 16GB of memory.

The launch marks OpenAI’s first ‘open’ language model since [GPT-2](#), which was released more than five years ago.

In a briefing, OpenAI said its open models will be capable of sending complex queries to [AI models in the cloud](#), as TechCrunch previously reported. That means if OpenAI’s open model is not capable of a certain task, such as processing an image, developers can connect the open model to one of the company’s more capable closed models.

While OpenAI open sourced AI models in its early days, the company has generally favored a proprietary, closed source development approach. The latter strategy has helped OpenAI build a large business selling access to its AI models via an API to enterprises and developers.

However, CEO Sam Altman said in January he believes OpenAI has been “on the wrong side of history” when it comes to open sourcing its technologies. The company today faces growing pressure from Chinese AI labs — including DeepSeek, Alibaba’s Qwen, and Moonshot AI — which have developed several of the world’s most capable and popular open models. (While Meta previously dominated the open AI space, the company’s Llama AI models have [fallen behind](#) in the last year.)

In July, the Trump administration also urged U.S. AI developers to [open source more technology](#) to promote global adoption of AI aligned with American values.

Newsletters

[See More ↗](#)

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

[REGISTER NOW](#)

With the release of gpt-oss, OpenAI hopes to curry favor with developers and the Trump administration alike, both of which have watched the Chinese AI labs rise to prominence in the open source space.

“Going back to when we started in 2015, OpenAI’s mission is to ensure AGI that benefits all of humanity,” said Altman in a statement shared with TechCrunch. “To that end, we are excited for the world to be building on an open AI stack created in the United States, based on democratic values, available for free to all and for wide benefit.”




IMAGE CREDITS:TOMOHIRO OHSUMI / GETTY IMAGES

How the models performed

OpenAI aimed to make its [open model a leader](#) among other open-weight AI models, and the company claims to have done just that.

On Codeforces (with tools), a competitive coding test, gpt-oss-120b and gpt-oss-20b score 2622 and 2516, respectively, outperforming DeepSeek's R1 while underperforming o3 and o4-mini.




OPENAI'S OPEN MODEL PERFORMANCE ON CODEFORCES.

IMAGE CREDITS:OPENAI

On Humanity's Last Exam (HLE), a challenging test of crowdsourced questions across a variety of subjects (with tools), gpt-oss-120b and gpt-oss-20b

score 19% and 17.3%, respectively. Similarly, this underperforms o3 but outperforms leading open models from DeepSeek and Qwen.

Humanity's Last Exam
Expert-level questions across subjects



OPENAI'S OPEN MODEL PERFORMANCE ON HLE.

IMAGE CREDITS:OPENAI

Notably, OpenAI's open models hallucinate significantly more than its latest AI reasoning models, o3 and o4-mini.

Hallucinations have been [getting more severe](#) in OpenAI's latest AI reasoning models, and the company previously said it doesn't quite understand why. In a white paper, OpenAI says this is "expected, as smaller models have less world knowledge than larger frontier models and tend to hallucinate more."

OpenAI found that gpt-oss-120b and gpt-oss-20b hallucinated in response to 49% and 53%, respectively, of questions on PersonQA, the company's in-house benchmark for measuring the accuracy of a model's knowledge about people. That's more than triple the hallucination rate of OpenAI's o1 model, which scored 16%, and higher than its o4-mini model, which scored 36%.

Training the new models

OpenAI says its open models were trained with similar processes to its proprietary models. The company says each open model leverages mixture-of-experts (MoE) to tap fewer parameters for any given question, making it run more efficiently. For gpt-oss-120b, which has 117 billion total parameters, OpenAI says the model only activates 5.1 billion parameters per token.

The company also says its open model was trained using high-compute [reinforcement learning \(RL\)](#) — a post-training process to teach AI models right from wrong in simulated environments using large clusters of Nvidia GPUs. This was also used to train OpenAI's o-series of models, and the open models have a similar chain-of-thought process in which they take additional time and computational resources to work through their answers.

As a result of the post-training process, OpenAI says its open AI models excel at powering AI agents and are capable of calling tools such as web search or Python code execution as part of its chain-of-thought process. However, OpenAI says its open models are text-only, meaning they will not be able to process or generate images and audio like the company's other models.

OpenAI is releasing gpt-oss-120b and gpt-oss-20b under the Apache 2.0 license, which is generally considered one of the most permissive. This license will allow enterprises to monetize OpenAI's open models without having to pay or obtain permission from the company.

However, unlike fully open source offerings from AI labs like AI2, OpenAI says it will not be releasing the training data used to create its open models. This decision is not surprising given that several active lawsuits against AI model providers, including OpenAI, have alleged that these companies

inappropriately trained their AI models on copyrighted works.

OpenAI [delayed](#) the release of its open models [several](#) times in recent months, partially to address safety concerns. Beyond the company’s typical safety policies, OpenAI says in a white paper that it also investigated whether bad actors could fine-tune its gpt-oss models to be more helpful in cyberattacks or the creation of biological or chemical weapons.

After testing from OpenAI and third-party evaluators, the company says gpt-oss may marginally increase biological capabilities. However, it did not find evidence that these open models could reach its “high capability” threshold for danger in these domains, even after fine-tuning.

While OpenAI’s model appears to be state-of-the-art among open models, developers are eagerly awaiting the release of DeepSeek R2, its next AI reasoning model, as well as a new open model from Meta’s Superintelligence Lab.

Topics: [AI](#) [open model](#) [OpenAI](#)

[f](#) [X](#) [in](#) [🌐](#) [✉](#) [🔗](#)



Maxwell Zeff

Senior AI Reporter | [X](#)

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >

KEEP READING



AI



Latest

Security

Startups

AI

Venture

Apps

Apple



Events

Podcasts

Newsletters

IMAGE CREDITS:

KRISZTIAN BOCSI/BLOOMBERG / GETTY IMAGES

Roughly two years ago, Microsoft [announced](#) a partnership with OpenAI, the AI lab with which it has a close [commercial relationship](#), to build what the tech giant called an “AI Supercomputer” running in the Azure cloud. Containing over 285,000 processor cores and 10,000 graphics cards, Microsoft claimed at the time that it was one of the largest supercomputer clusters in the world.

Now, presumably to support even more ambitious AI workloads, Microsoft [says](#) it’s signed a “multi-year” deal with Nvidia to build a *new* supercomputer hosted in Azure and powered by

Latest in AI

See More

AI



Google denies AI search features are killing website traffic

Sarah Perez - 16 hours ago

AI



First impressions of Alexa+, Amazon’s upgraded, AI-powered digital assistant

Nvidia's GPUs, networking and AI software for training AI systems.

Sarah Perez - 17 hours ago

"AI is fueling the next wave of automation across enterprises and industrial computing, enabling organizations to do more with less as they navigate economic uncertainties," Scott Guthrie, executive vice president of Microsoft's cloud and AI group, said in a statement. "Our collaboration with Nvidia unlocks the world's most scalable supercomputer platform, which delivers state-of-the-art AI capabilities for every enterprise on Microsoft Azure."

Details were hard to come by at press time. But in a blog post, Microsoft and Nvidia said that the upcoming supercomputer will feature hardware like Nvidia's Quantum-2 400Gb/s InfiniBand networking technology and recently detailed H100 GPUs. Current Azure instances offer previous-gen Nvidia A100 GPUs paired with Quantum 200Gb/s InfiniBand networking.

Notably, the H100 — the flagship of Nvidia's Hopper architecture — ships with a special "Transformer Engine" to accelerate machine learning tasks and — at least [according](#) to Nvidia — delivers between 1.5 and 6 times better performance than the A100. It's also less power-hungry, offering the same performance as the A100 with up to 3.5 times better energy efficiency.

One of the first industrial-scale machines to sport H100 GPUs, the Lenovo-built Henri system operated by the Flatiron Institute in New York City, [topped](#) the list of this year's most efficient supercomputers. But Henri — a relatively small system by supercomputer standards, with only thousands of CPUs — ranked 405 on the Top500, a ranking of the 500 most powerful supercomputers in the world, implying the H100 is by no means a holy grail.

AI



Google takes on ChatGPT's Study Mode with new 'Guided Learning' tool in Gemini

Aisha Malik - 17 hours ago

As part of the Microsoft collaboration, Nvidia says that it'll use Azure virtual machine instances to research advances in generative AI, or the self-learning algorithms that can create text, code, images, video or audio. (Think along the lines of OpenAI's text-generating [GPT-3](#) and image-producing [DALL-E 2](#).) Meanwhile, Microsoft will optimize its DeepSpeed library for new Nvidia hardware, aiming to reduce computing power and memory usage during AI training workloads, and work with Nvidia to make the company's stack of AI workflows and software development kits available to Azure enterprise customers.

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Why Nvidia would opt to use Azure instances over its own in-house supercomputer, Selene, isn't entirely clear; the company's already tapped Selene to train generative AI like GauGAN2, a text-to-image generation model that creates art from basic sketches. Evidently, Nvidia anticipates that the scope of the AI systems that it's working with will eventually surpass Selene's capabilities and

perhaps even those of [Eos](#), a proof-of-concept supercomputer Nvidia is building that'll feature 4,608 H100 GPUs and reach up to 18.4 exaflops of AI computing performance.

“AI technology advances as well as industry adoption are accelerating. The breakthrough of foundation models has triggered a tidal wave of research, fostered new startups and enabled new enterprise applications,” Manuvir Das, VP of enterprise computing at Nvidia, said in a statement. “Our collaboration with Microsoft will provide researchers and companies with state-of-the-art AI infrastructure and software to capitalize on the transformative power of AI.”

The insatiable demand for powerful AI training infrastructure has led to an arms race of sorts among cloud and hardware vendors. Just this week, [Cerabras](#), which has raised over \$720 million in venture capital to date at an over-\$4 billion valuation, [unveiled](#) a 13.5-million core AI supercomputer called Andromeda it claims can achieve more than 1 exaflop of AI compute. Google and Amazon continue to invest in their own proprietary solutions, offering custom-designed chips — e.g. [TPUs](#) and [Trainium](#) — for accelerating AI training in the cloud.

The push for more powerful hardware will continue for the foreseeable future. A recent [study](#) found that the compute requirements for large-scale AI models has been doubling at an average rate of 10.7 months between 2016 and 2022. And OpenAI once estimated that, if GPT-3 were to be trained on a single Nvidia Tesla V100 GPU, it would take around 355 years.

Topics: [AI](#) [AI](#) [Enterprise](#) [Microsoft](#) [nvidia](#)

[supercomputer](#)



Kyle Wiggers

AI Editor | X butterfly @ in

Kyle Wiggers was TechCrunch’s AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

View Bio >

KEEP READING



IMAGE CREDITS: MICROSOFT

AI

This week in AI: Microsoft sticks an AI ad on keyboards

Kyle Wiggers, Devin Coldewey

7:02 AM PST · January 6, 2024

Keeping up with an industry as fast-moving as [AI](#) is a tall order. So until an AI can do it for you, here's a handy roundup of recent stories in the world of machine learning, along with notable research and experiments we didn't cover on their own.

This week in AI, Microsoft unveiled a new standard PC keyboard [layout](#) with a "Copilot" key. You heard correctly — going forward, Windows machines will have a dedicated key for launching Microsoft's AI-powered assistant [Copilot](#), replacing the right Control key.

The move is meant, one imagines, to signal the seriousness of Microsoft's investment in the race for consumer (and enterprise for that matter) AI dominance. It's the first time Microsoft's changed the Windows keyboard layout in about 30 years; laptops and keyboards with the Copilot key are scheduled to ship as soon as late February.

But is it all bluster? Do Windows users really *want* an AI shortcut — or Microsoft's flavor of AI?

Microsoft's certainly made a show of injecting nearly all its products old and new with "Copilot" functionality. In [flashy keynotes](#), [slick demos](#) and now an AI key, the company's making its AI tech prominent — and betting on this to drive demand.

Demand isn't a sure thing. But to be fair, a few vendors have managed to turn viral AI hits into successes. Look at OpenAI, the maker of ChatGPT, which [reportedly](#) topped \$1.6 billion in annualized revenue toward the end of 2023. Generative art platform Midjourney is apparently profitable as well — and hasn't yet taken a dime of outside capital.

Emphasis on *a few*, though. Most vendors, weighed down by the costs of training and running cutting-edge AI models, have had to seek larger and larger tranches of capital to stay afloat. Case in point,


TEC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion
- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

Anthropic is said to be [raising](#) \$750 million in a round that would bring its total raised to more than \$8 billion.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Advertisement

Microsoft, together with its chip partners AMD and Intel, hopes that AI processing will increasingly move from expensive data centers to local silicon, commoditizing AI in the process — and it might. Intel's new lineup of consumer chips pack custom-designed cores for running AI. Plus, new [data center chips](#) like Microsoft's own could make model training a less expensive endeavor than it is currently.

But there's no guarantee. The real test will be seeing whether Windows users and enterprise customers, bombarded with what amounts to Copilot advertising, show an appetite for the tech — and shell out for it. If they don't, it might not be long before Microsoft has to redesign the Windows keyboard once again.

Here are some other AI stories of note from the past few days:

Copilot comes to mobile: Microsoft quietly brought Copilot clients to Android and iOS, along with iPadOS.

GPT Store: OpenAI announced plans to launch a store for GPTs, custom apps based on its text-generating AI models (e.g., GPT-4), within the next week. The GPT Store was announced last year during OpenAI's first annual developer conference, DevDay, but delayed in December — almost certainly due to the leadership shakeup that occurred in November just after the initial announcement.

OpenAI shrinks reg risk: In other OpenAI news, the startup's looking to shrink its regulatory risk in the EU by funneling much of its overseas business through an Irish entity. Natasha writes that the move will reduce the ability of some privacy watchdogs in the bloc to unilaterally act on concerns.

Training robots: Google's DeepMind robotics team is exploring ways to give robots a better understanding of precisely what it is we humans want out of them, Brian writes. The team's new system can manage a fleet of robots working in tandem and suggest tasks that can be accomplished by the robots' hardware.

Intel's new company: Intel is spinning out a new platform company, Articul8 AI, with the backing of Boca Raton, Florida-based asset manager and investor DigitalBridge. As an Intel spokesperson explains, Articul8's platform "delivers AI capabilities that keep customer data, training and inference within the enterprise security perimeter" — an appealing prospect for customers in highly

regulated industries like healthcare and financial services.

Dark fishing industry, exposed: Satellite imagery and machine learning offer a new, far more detailed look at the maritime industry, specifically the number and activities of fishing and transport ships at sea. Turns out there are way more of them than publicly available data would suggest — a fact revealed by new research published in Nature from a team at Global Fishing Watch and multiple collaborating universities.

AI-powered search: Perplexity AI, a platform applying AI to web searching, raised \$73.6 million in a funding round valuing the company at \$520 million. Unlike traditional search engines, Perplexity offers a chatbot-like interface that allows users to ask questions in natural language (e.g., “Do we burn calories while sleeping?” “What’s the least visited country?”).

Clinical notes, written automatically: In more funding news, Paris-based startup Nabla raised a cool \$24 million. The company, which has a partnership with Permanente Medical Group, a division of U.S. healthcare giant Kaiser Permanente, is working on an “AI copilot” for doctors and other clinical staff that automatically takes notes and writes medical reports.

More machine learnings

You may remember various examples of interesting work over the last year involving **making minor changes to images** that cause machine learning models to mistake, for instance, a picture of a dog for a picture of a car. They do this by adding “perturbations,” minor changes to the pixels of the image, in a pattern that only the model can perceive. Or at least they *thought* only the model could perceive it.

[An experiment by Google DeepMind researchers](#)

showed that when a picture of flowers was perturbed to appear more catlike to AI, people were more likely to describe that image as more catlike despite its definitely not looking like a cat. Same for other common objects like trucks and chairs.

IMAGE CREDITS: GOOGLE DEEPMIND

Why? How? The researchers don't really know, and the participants all felt like they were just choosing randomly (indeed the influence is, while reliable, scarcely above chance). It seems we're just more perceptive than we think — but this also has implications on safety and other measures, since it suggests that subliminal signals could indeed propagate through imagery without anyone noticing.

Another interesting experiment involving human perception came out of MIT this week, which used machine learning to [help elucidate a particular system of language understanding](#). Basically some simple sentences, like “I walked to the beach,” barely take any brain power to decode, while complex or confusing ones like “in whose aristocratic system it effects a dismal revolution” produce more and broader activation, as measured by fMRI.

The team compared the activation readings of humans reading a variety of such sentences with how the same sentences activated the equivalent of cortical areas in a large language model. Then they made a second model that learned how the two activation patterns corresponded to one another. This model was able to predict for novel sentences whether they would be taxing on human cognition or not. It may sound a bit arcane, but it is definitely super interesting, trust me.

Whether machine learning can imitate human cognition in more complex areas, like interacting with computer interfaces, is still very much an open question. There's lots of research, though, and it's always worth taking a look at. This week we have [SeeAct](#), a system from Ohio State researchers that works by laboriously grounding a LLM's interpretations of possible actions in real-world examples.

IMAGE CREDITS: OHIO STATE UNIVERSITY

Basically you can ask a system like GPT-4V to create a reservation on a site, and it will get what its task is and that it needs to click the “make reservation” button, but it doesn't really know how to do that. By improving how it perceives interfaces with explicit labels and world knowledge, it can do lots better, even if it still only succeeds a fraction of the time. These agent models have a long way to go, but expect a lot of big claims this year anyway! I just heard some today.

Next, check out this interesting solution to a problem I had no idea existed but that makes perfect sense. Autonomous ships are a promising area of automation, but when the sea is angry, it is difficult to make sure they stay on track. GPS and gyros don't cut it, and visibility can be poor too — but more importantly, the systems governing them aren't too sophisticated. So they can go wildly off target or waste fuel going on large detours if they don't know any better, a big problem if you're on battery power. I never even thought about that!

[Korea's Maritime and Ocean University](#) (another thing I learned about today) proposes a more powerful pathfinding model built on simulating ship movements in a computational fluid dynamics model. They propose that this better understanding of wave action and its effect on hulls and propulsion could seriously improve the efficiency and safety of autonomous marine transport. It might even make sense to use in human-guided vessels whose captains aren't quite sure what the best angle of attack is for a given squall or wave form!

Last, if you want a good recap of last year's big advances in computer science, which in 2023 overlapped massively with ML research, [check out Quanta's excellent review](#).

Topics: [AI](#) [AI](#) [ChatGPT](#) [Microsoft](#) [this week in AI](#)

[this week in ai newsletter](#)

Advertisement



Kyle Wiggers

AI Editor | [X](#) [🦋](#) [@](#) [in](#)

Kyle Wiggers was TechCrunch’s AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

[View Bio](#) >

Devin Coldewey

Writer & Photographer | [in](#) [🏠](#)

Devin Coldewey is a Seattle-based writer and photographer....

[View Bio](#) >

Loading the next article