


Advertisement

AI



Microsoft in talks to maintain access to OpenAI's tech beyond AGI milestone



Microsoft is reportedly in advanced talks with OpenAI for a new agreement that would give it ongoing access to the startup's technology even if OpenAI achieves [what it defines as AGI](#), or advanced general intelligence. If the deal goes through, it would clear a key hurdle in OpenAI's transition toward becoming a fully commercial enterprise.

The companies have been negotiating regularly, and they could come to an agreement in a few weeks, [Bloomberg reports](#), citing three anonymous sources. The report cited some of the sources as saying that while the talks have been positive, other [roadblocks could emerge in the form of regulatory scrutiny and Elon Musk's lawsuit](#) to block OpenAI's for-profit transition.

OpenAI is [currently structured](#) as a mission-driven nonprofit that oversees a capped for-profit company — a setup that's meant to limit how fully it can commercialize or raise money. That structure hasn't stopped it from raising billions and operating like a traditional tech company, but OpenAI still wants to shake off its constraints.

Microsoft, OpenAI's biggest backer with \$13.75 billion invested and rights to some of the ChatGPT maker's IP, has put up meaningful roadblocks to OpenAI's future as a [for-profit company](#), with talks dragging on for months.

TECH DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Instagram takes on Snapchat with new 'Instagram Map'
- Google says hackers stole its customers' data by breaching its Salesforce database
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Spotify raises subscription prices

Microsoft wants a bigger stake in the restructured company and seeks to secure its access to OpenAI's tech beyond the current deal, which ends in 2030 or whenever OpenAI says it has achieved AGI — though no one can really agree on [what that means](#).

Microsoft has built its Azure OpenAI Service around the smaller company's models, and has integrated the startup's tech into Copilot across Windows, Office, and GitHub. If OpenAI suddenly declares it has achieved AGI and cuts off access, Microsoft would lose a huge strategic advantage. The ChatGPT maker has reportedly also told its investors that it expects to pay Microsoft [a lower share of its revenue](#) as it progresses.

A source told Bloomberg that OpenAI also hopes to guarantee that Microsoft deploys OpenAI's technology safely, especially as it nears AGI.

- The uproar over Vogue's AI-generated ad isn't just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

Microsoft also stands to gain from OpenAI becoming a for-profit company. The current setup

caps investor returns, so a more standard structure would give Microsoft a chance to receive formal equity and significant returns, in addition to access to OpenAI's tech.

Bloomberg reports that the two companies have been negotiating an equity stake for Microsoft in the low- to mid-30% range.

Got a sensitive tip or confidential documents? We're reporting on the inner workings of the AI industry — from the companies shaping its future to the people impacted by their decisions. Reach out to Rebecca Bellan at rebecca.bellan@techcrunch.com and Maxwell Zeff at maxwell.zeff@techcrunch.com. For secure communication, you can contact us via Signal at @rebeccabellan.491 and @mzeff.88.

Topics: [AGI](#) [AI](#) [Microsoft](#) [OpenAI](#)



Rebecca Bellan

Senior Reporter |

Rebecca Bellan is a senior reporter at TechCrunch where she covers the business, policy, and emerging trends...

[View Bio](#) >

———— KEEP READING ————

OpenAI's GPT-5 is here

Maxwell Zeff — 10:00 AM PDT · August 7, 2025

GPT-5 flagship model

OpenAI has launched GPT-5, a new flagship AI model that will power the company's next generation of ChatGPT.

IMAGE CREDITS: OPENAI

GPT-5, which was released Thursday, is OpenAI's first "unified" AI model and combines the reasoning abilities of its [o-series](#) of models with the fast responses of its [GPT series](#). The next-generation model signals a new era for ChatGPT — and its creator, OpenAI — pointing to OpenAI's broader ambitions to develop AI systems that are [more like agents](#) than chatbots.

While [GPT-4](#) enabled AI chatbots to offer smart responses on a wide variety of questions, GPT-5 allows ChatGPT to complete a wide variety of tasks on behalf of users — such as generating software applications, navigating a user's calendar, or creating research briefs.

With GPT-5, OpenAI has also sought to make

Advertisement

Newsletters

See More ↗

Subscribe for the industry's biggest tech news

ChatGPT simpler to use. Instead of asking users to choose the right settings, GPT-5 comes equipped with a real-time router that decides how to offer the best answer, whether that’s responding to user questions quickly or taking additional time to “think” through answers.




IMAGE CREDITS: OPENAI

During a briefing with reporters, OpenAI CEO Sam Altman claimed GPT-5 is “the best model in the world,” and said it represented a “significant step” along the company’s path to developing AI that can outperform humans at most economically valuable work — that is, artificial general intelligence (AGI).

“Having something like GPT-5 would be pretty much unimaginable at any previous time in history,” said Altman.

Starting Thursday, GPT-5 will be available to all free users of ChatGPT as their default model. OpenAI’s VP of ChatGPT, Nick Turley, said this is part of the company’s effort to give free users access to an AI reasoning model for the first time. (Previously, the company gated these more advanced models behind a paywall.)

“This is just one of the ways that I’m excited to live the mission, making sure that this stuff actually benefits people,” said Turley on the decision, referencing OpenAI’s long-standing [mission](#) to

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

Advertisement

distribute advanced AI to as many people as possible.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

[REGISTER NOW](#)

The expectations are high for GPT-5, one of OpenAI's most anticipated product launches since ChatGPT put the company on the map in 2022. Since then, ChatGPT has grown into one of the world's most popular consumer products, reaching more than 700 million users every week — nearly 10% of the globe's population, according to the company.

Many see GPT-5 as a bellwether for AI progress broadly, and the model's reception by Silicon Valley could have profound implications for Big Tech, Wall Street, and policymakers regulating technology. These stakeholders are watching to see if GPT-5 offers a significant jump in AI's capabilities, much like its predecessor, GPT-4, which challenged expectations of what software can do.

GPT-5 offers a slight edge on the competition


OpenAI claims GPT-5 is state-of-the-art in several domains, slightly edging out leading AI models from Anthropic, Google DeepMind, and Elon Musk’s xAI on key benchmarks. However, GPT-5 slightly underperforms frontier AI models in other areas.

The company says GPT-5 offers frontier-level performance around coding; Altman said the model specifically excels at spinning up entire software applications on demand, in what’s become known as “vibe coding.”

On SWE-bench Verified — a test of real-world coding tasks pulled from GitHub — GPT-5 scores 74.9% on its first attempt. That means GPT-5 just outperforms Anthropic’s latest Claude Opus 4.1 model, which scored [74.5%](#), and Google DeepMind’s Gemini 2.5 Pro, which scored [59.6%](#).

On Humanity’s Last Exam — a difficult test measuring AI model performance across math, humanities, and the natural sciences — a version of GPT-5 with extended reasoning (GPT-5 Pro) scored 42% when using tools. That’s slightly less than xAI was able to achieve with [Grok 4 Heavy](#), which scored 44.4% on the test.

Humanity’s Last Exam (Full Set)*
Expert-level questions across subjects



On GPQA Diamond — a test of PhD-level science questions — GPT-5 pro scored 89.4% on its first try, outperforming Claude Opus 4.1, which scored 80.9%, and Grok 4 Heavy, which scored 88.9%.

OpenAI says GPT-5 is better for answering health-related questions. On a test measuring accuracy in AI model responses around healthcare topics, HealthBench Hard Hallucinations, OpenAI says GPT-5 (with thinking) hallucinates just 1.6% of the time. This is far lower than the company's previous GPT-4o and o3 models, which scored 12.9% and 15.8, respectively.

While AI chatbots are not medical professionals, [millions of people](#) are using them for health advice. In response to this phenomenon, the company says GPT-5 is more proactive about flagging potential health concerns and helping users parse medical results.

In addition, OpenAI says GPT-5 is better than other AI models on more difficult-to-measure, [subjective domains](#), such as creative design and writing. Turley said GPT-5 responds more naturally and exhibits “better taste” than other AI models on creative tasks.

“The vibes of this model are really good,” said Turley.

GPT-5 is also more accurate than OpenAI's previous models, and the company says it suffers far less from hallucinations — the tendency for AI models to make up information — compared to its o-series models. Hallucinations seemed to be [getting worse](#) in OpenAI's latest AI reasoning models, such as o3, and OpenAI previously said it didn't quite understand why it was happening.

In responses to ChatGPT prompts, OpenAI found that GPT-5 (with thinking) hallucinates and responds with incorrect information 4.8% of the time. That's a significant reduction from o3 and GPT-4o, which score hallucination rates of 22% and 20.6%, respectively, on the test.

On a benchmark measuring an AI model's agentic ability to complete simulated online tasks, *Tau-bench*, GPT-5 offers mixed performance. On part of the test measuring an AI's ability to navigate an airline's website, GPT-5 scores 63.5%, slightly underperforming o3, which scored 64.8%. On another part of the test measuring AI's ability to navigate retail websites, GPT-5 scores 81.1%, underperforming Claude Opus 4.1, which scored 82.4%.


OpenAI also says that GPT-5 is safer than its previous models. While AI reasoning models occasionally exhibit a tendency to [scheme against humans](#) or lie to promote their own goals, OpenAI found that GPT-5 was deceptive at a lower rate than other models.

Alex Beutel, OpenAI's safety research lead, said reducing deception improves not only the safety of GPT-5, but also the user experience, creating a model that's more "transparent and honest in ways users can trust."

Beutel also notes GPT-5 is better at discerning between bad actors who are trying to misuse ChatGPT and users making harmless requests. This results in GPT-5 being able to refuse more unsafe questions, while offering fewer rejections to users seeking harmless information.

Upgrades for consumers and developers

ChatGPT is getting a few user experience upgrades as part of the GPT-5 launch. Users can now select from four new personalities in ChatGPT's setting: Cynic, Robot, Listener, and Nerd. The company says these will adapt ChatGPT's responses without



which will spend more or less time “reasoning” through tasks. Developers can also now control verbosity in the OpenAI API, deciding how long or short an AI model’s responses should be.

The base model of GPT-5 will cost developers \$1.25 per million input tokens (roughly 750,000 words, longer than the entire “Lord of the Rings” series) and \$10 per million output tokens.

The launch of GPT-5 comes after a busy week for OpenAI. The company released an open-weight reasoning model, [gpt-oss](#), that developers and enterprises can download for free and run at a fraction of the cost. The open model nearly matched the abilities of OpenAI’s previous top models, o3 and o4-mini, but GPT-5 sets a new standard for frontier performance in some areas, such as coding.

However, GPT-5 seems to be roughly on par with other frontier AI models in several areas.

Benchmarks, of course, only tell part of the story for any AI model, and it remains to be seen how developers will use GPT-5 in the real world, and whether the model is truly a step above the competition.

Topics: [AI](#) [ChatGPT](#) [gpt-5](#) [OpenAI](#)



Maxwell Zeff

Senior AI Reporter | [X](#)

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >

———— KEEP READING ————

AI



Microsoft brings
OpenAI's smallest open
model to Windows users

Microsoft is making OpenAI's new free and open GPT model, [gpt-oss-20b](#), available to Windows 11 users via [Windows AI Foundry](#), the tech giant's platform that lets users tap AI features, APIs, and popular open source models on their computers.

The company said in a [blog post](#) that gpt-oss-20b is "tool-savvy and lightweight," adding: "Optimized for agentic tasks like code execution and tool use, it runs efficiently on a range of Windows hardware, with support for more devices coming soon. It's perfect for powering autonomous assistants or embedding AI into real-world workflows, even in bandwidth-constrained environments."

IMAGE CREDITS:

DAVID RYDER / PLUTONERD
(PHOTOMOSH/MODIFIED) / GETTY IMAGES
Launched on Tuesday, OpenAI's gpt-oss-20b can run on consumer PCs and laptops with at least 16GB of VRAM, which a modern GPU from Nvidia or Radeon would have. OpenAI said the model was trained using [high-compute reinforcement learning](#), which helps it excel at powering AI agents and call tools, such as web search or Python code execution, as part of its chain-of-thought process.

However, the model is text-only, meaning it will not be able to process or generate images and audio like the company's other models.

Also, the model has a habit of hallucinating a lot: gpt-oss-20b hallucinated in response to 53% of questions on PersonQA, OpenAI's in-house benchmark for measuring the accuracy of a model's knowledge about people.

Latest in AI

[See More](#)

TRANSPORTATION



Tesla shuts down Dojo, the AI training supercomputer that Musk said would be key to full self-driving

Rebecca Bellan - 6 hours ago

STARTUPS



High costs and thin margins threatening AI coding startups

Marina Temkin - 7 hours ago

AI



OpenAI's GPT-5 is here

Maxwell Zeff - 11 hours ago

Microsoft says it will bring the model to macOS soon, as well as other devices, though the company did not name any more. Microsoft is also making both gpt-oss-20b and older model gpt-oss-120b available via its hosted Azure AI Foundry platform.

The two new models are also [available on Amazon's AWS](#).

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

Topics: [AI](#) [gpt](#) [Microsoft](#) [OpenAI](#) [Windows](#)



Ram Iyer

Editor |

Ram is a financial and tech reporter and editor. He covered North American and European M&A, equity, regulatory...



IMAGE CREDITS: GETTY IMAGES

AI



OpenAI is practically giving ChatGPT to the government for free

Rebecca Bellan — 10:09 AM PDT · August 6, 2025

OpenAI is poised to undercut rivals like Anthropic and Google in the race to see its AI tools integrated into federal agency workflows.

The AI giant has [reached an agreement](#) with the U.S. General Services Administration (GSA), the government's central purchasing arm, to offer ChatGPT Enterprise to participating federal agencies for just \$1 per agency for the next year.

DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

The partnership comes a day after the [GSA added OpenAI, Google, and Anthropic](#) to a list of approved AI vendors that can offer their services to civilian federal agencies. The companies will see their tools offered via the Multiple Award Schedule (MAS), a federal contracting platform that lets government agencies access AI tools via pre-negotiated contracts so they don't have to negotiate with vendors individually.

It's not clear whether other AI firms are going to offer their services at such a discounted rate, though GSA Federal Acquisition Service commissioner Josh Gruenbaum in a [statement](#) encouraged "other American AI technology companies to follow OpenAI's lead and work with us."

TechCrunch has reached out to Anthropic and Google for more information.

In addition to access to ChatGPT Enterprise, OpenAI is offering unlimited use of advanced models for an additional 60 days. Federal employees will also have access to a new government user community and tailored introductory training resources to get familiar with OpenAI's tools.

Data security is a top priority for government agencies that are concerned about sensitive information being leaked into model training sets. TechCrunch has asked the GSA for details on how government data is being safeguarded, including whether measures like on-premises or private cloud deployments are being used to enhance security.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Instagram takes on Snapchat with new 'Instagram Map'
- Google says hackers stole its customers' data by breaching its Salesforce database
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the

heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

“The government is taking a cautious, security-first approach to AI,” a GSA spokesperson told TechCrunch. “This ensures sensitive information remains protected while enabling agencies to benefit from AI-driven efficiencies.”

OpenAI's discount comes a couple of weeks after the Trump administration published its AI Action Plan that seeks to boost data center buildouts and integrate more AI tools into government, among other things. It also follows a Trump executive order that [bans “woke AI”](#) and AI models that aren't “ideologically neutral” from government contracts.

TechCrunch has reached out to OpenAI to learn more about its GSA partnership and how it will approach Trump's executive order.

Topics: [AI](#) [ChatGPT Enterprise](#) [Government & Policy](#)

[GSA](#) [OpenAI](#)






Rebecca Bellan

Senior Reporter | [X](#) [🦋](#)

Rebecca Bellan is a senior reporter at TechCrunch where she covers the business, policy, and emerging trends...

[View Bio](#) >

Loading the next article



- TechCrunch
- Staff
- Contact Us
- Advertise
- Crunchboard
- Jobs
- Site Map
- Terms of Service
- Privacy Policy
- RSS Terms of Use
- OpenAI
- Perplexity
- Cisco
- DeepMind
- Windsurf
- Tech Layoffs
- ChatGPT