

Advertisement

AI



Mistral releases Voxtral,
its first open source AI
audio model

As AI systems become more capable, speech is fast becoming the default way we communicate with machines. French AI startup [Mistral](#) has jumped into the audio race with its first open model, aiming to challenge the dominance of walled-off corporate systems with open-weight alternatives.

On Tuesday, Mistral announced the release of Voxtral, its first family of audio models aimed at businesses.

The company is pitching Voxtral as the first open model that's capable of deploying "truly usable speech intelligence in production."

IMAGE CREDITS: CHESNOT / GETTY IMAGES

In other words, no longer will developers have to choose between a cheap, open system that fumbles transcriptions and doesn't really understand what's being said, and one that functions well, but is closed, leaving developers with a higher bill and less control over deployment.

For businesses, that means Voxtral offers an affordable alternative that the company claims is "less than half the price" of comparable solutions.


DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion




IMAGE CREDITS:MISTRAL

Mistral says Voxtral can transcribe up to 30 minutes of audio. Due to its LLM backbone, [Mistral Small 3.1](#), it can understand up to 40 minutes, allowing users to ask questions about the audio content, generate summaries, or turn voice commands into real-time actions like calling APIs or running functions. Voxtral is also multilingual, with the ability to transcribe and understand languages including English, Spanish, French, Portuguese, Hindi, German, Dutch, and Italian.

The company is offering up two variants of its “speech understanding models.” The first, Voxtral Small, has 24 billion parameters for production-scale deployments, and is competitive with ElevenLabs Scribe, GPT-4o-mini, and Gemini 2.5 Flash.

- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

[REGISTER NOW](#)

The second, Voxtral Mini, has 3 billion parameters for local and edge deployments. There's also an ultra-cheap, stripped-down, fast API version of the 3 billion model called Voxtral Mini Transcribe that is optimized for transcription-only use cases and promises to outperform OpenAI Whisper for less than half the price.

Users can try Voxtral for free by downloading the API on Hugging Face or testing the models in Mistral's chatbot Le Chat. Integrating the API into applications starts at \$0.001 per minute, according to the company.

The launch comes a month after [Mistral announced Magistral](#), its first family of reasoning models that work through problems step-by-step for improved reliability.

Mistral, one of the top AI firms in Europe, is well-known for its advocacy pushing open source AI models. Earlier this month, TechCrunch reported that the company is in talks to raise up to [\\$1 billion in equity](#) from investors like Abu Dhabi's MGX fund.

Topics: [AI](#) [AI](#) [mistral](#) [Mistral AI](#) [mistral voxtral](#)

[open source ai](#)



Rebecca Bellan

Senior Reporter |

Rebecca Bellan is a senior reporter at TechCrunch, where she covers Tesla and Elon Musk’s broader empire,...

View Bio >

KEEP READING



AI

[f](#) [X](#) [in](#) [reddit](#) [email](#) [link](#)

Mistral releases new AI models optimized for laptops and phones

Kyle Wiggers — 7:49 AM PDT · October 16, 2024

- Latest
- Security
- Startups
- AI
- Venture
- Apps
- Apple



- Events
- Podcasts
- Newsletters

[Q](#) [Menu](#)

to complete tasks.

Two Les Ministraux models are available — Ministral 3B and Ministral 8B — both of which have a context window of 128,000 tokens, meaning they can ingest roughly the length of a 50-page book.

“Our most innovative customers and partners have increasingly been asking for local, privacy-first inference for critical applications such as on-device translation, internet-less smart assistants, local analytics, and autonomous robotics,” Mistral [writes](#) in a blog post. “Les Ministraux were built to provide a compute-efficient and low-latency solution for these scenarios.”

Ministral 8B is available for download as of today — albeit strictly for research purposes. Mistral is requiring devs and companies interested in Ministral 8B or Ministral 3B self-deployment setups to contact it for a commercial license.

Otherwise, devs can use Ministral 3B and Ministral 8B through Mistral’s cloud platform, La Plateforme, and other clouds with which the startup has partnered in the coming weeks. Ministral 8B costs 10 cents per million output/input tokens (~750,000 words), while Ministral 3B costs 4 cents per million output/input tokens.

There’s been a trend toward small models, lately, which are cheaper and quicker to train, fine-tune, and run than their larger counterparts. Google continues to add models to its [Gemma](#) small model family, while Microsoft offers its [Phi](#) collection of models. In the most recent refresh of its [Llama](#) suite, Meta [introduced](#) several small models optimized for edge hardware.

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Mistral claims that Ministral 3B and Ministral 8B outperform comparable Llama and Gemma models — as well as its own [Mistral 7B](#) — on several AI benchmarks designed to evaluate instruction-following and problem-solving capabilities.

Paris-based Mistral, which [recently](#) raised \$640 million in venture capital, continues to gradually expand its AI product portfolio. Over the past few months, the company has [launched](#) a free service for developers to test its models, an [SDK](#) to let customers fine-tune those models, and new [models](#), including a generative model for code called [Codestral](#).

Co-founded by alumni from Meta and Google's DeepMind, Mistral's stated mission is to create flagship models that rival the best-performing models today, like OpenAI's [GPT-4o](#) and Anthropic's Claude — and ideally make money in the process. While the “making money” bit is proving to be challenging ([as it is for most generative AI startups](#)), Mistral [reportedly](#) began to generate revenue this summer.



Kyle Wiggers
AI Editor | [X](#) [🦋](#) [📧](#) [in](#)

Kyle Wiggers was TechCrunch’s AI Editor until June 2025.
His writing has appeared in VentureBeat and Digital...

[View Bio](#) >

———— KEEP READING ————

AI

Mistral adds a new API that turns any PDF document into an AI-ready Markdown file

On Thursday French large language model (LLM) developer [Mistral](#) launched a new API for developers who handle complex PDF documents. [Mistral OCR](#) is an optical character recognition (OCR) API that can turn any PDF into a text file to make it easier for AI models to ingest.

LLMs, which underpin popular GenAI tools like OpenAI's ChatGPT, work particularly well with raw text. So companies that want to create their own AI workflow know that it has become extremely important to store and index data in a clean format so that this data can be reused for AI processing.

Unlike most OCR APIs, Mistral OCR is a multimodal API, meaning that it can detect when there are illustrations and photos intertwined with blocks of text. The OCR API creates bounding boxes around these graphical elements and includes them in the output.

IMAGE CREDITS: CAROL YEPES / GETTY IMAGES

Mistral OCR also doesn't just output a big wall of text; the output is formatted in Markdown, a formatting syntax that developers use to add links, headers, and other formatting elements to a plain text file.

Latest in AI

[See More](#)

AI



Google denies AI search features are killing website traffic

Sarah Perez - 14 hours ago

AI



First impressions of Alexa+, Amazon's upgraded, AI-powered digital assistant

Sarah Perez - 15 hours ago

AI



Google takes on ChatGPT's Study Mode with new 'Guided Learning' tool in Gemini

Aisha Malik - 15 hours ago



LLMs rely heavily on Markdown for their training datasets. Similarly, when you use an AI assistant, such as Mistral's Le Chat or OpenAI's ChatGPT, they often generate Markdown to create bullet lists, add links, or put some elements in bold. Assistant apps seamlessly format the Markdown output into a rich text output. That's why raw text — and Markdown — have become more important in recent years as GenAI has boomed.

“Over the years, organizations have accumulated numerous documents, often in PDF or slide formats, which are inaccessible to LLMs, particularly RAG systems. With Mistral OCR, our customers can now convert rich and complex documents into readable content in all languages,” said Mistral co-founder and chief science officer Guillaume Lample.

“This is a crucial step toward the widespread adoption of AI assistants in companies that need to simplify access to their vast internal documentation,” he added.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy

hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Mistral OCR is available on Mistral’s own API platform or through its cloud partners (AWS, Azure, Google Cloud Vertex, etc.). And for companies working with classified or sensitive data, Mistral offers on-premise deployment.

According to the Paris-based AI company, Mistral OCR performs better than APIs from Google, Microsoft, and OpenAI. The company has tested its OCR model with complex documents that include mathematical expressions (LaTeX formatting), advanced layouts, or tables. It is also supposed to perform better with non-English documents.

Model	Overall	Math	Multilingual	Scanned	Tables
Google Document AI	83.42	80.29	86.42	92.77	78.16
Azure OCR	89.52	85.72	87.52	94.65	89.52
Gemini-1.5-Flash-002	90.23	89.11	86.76	94.87	90.48
Gemini-1.5-Pro-002	89.92	88.48	86.33	96.15	89.71
Gemini-2.0-Flash-001	88.69	84.18	85.80	95.11	91.46
GPT-4o-2024-11-20	89.77	87.55	86.00	94.58	91.70
Mistral OCR 2503	94.89	94.29	89.55	98.96	96.12

Given that Mistral OCR does one thing and one thing only, the company believes it is also faster than what's out there. That's not a surprise if you compare it with a multimodal LLM like GPT-4o, which also has OCR capabilities (among *many* other features).

Mistral is also using Mistral OCR for its own AI assistant [Le Chat](#). When a user uploads a PDF file, the company uses Mistral OCR in the background to understand what's in the document before processing the text.

Companies and developers will most likely use Mistral OCR with a RAG (aka Retrieval-Augmented Generation) system to use multimodal documents as input in an LLM. And there are many potential use cases. For instance, we could envisage law firms using it to help them swiftly plough through huge volumes of documents.

RAG is a technique that's used to retrieve data and use it as context with a generative AI model.

Topics: [AI](#) [Europe](#) [France](#) [mistral](#) [ocr](#)

Advertisement




Romain Dillet

Senior Reporter | [X](#) [🦋](#) [@](#) [in](#)

Romain Dillet was a Senior Reporter at TechCrunch until April 2025. He has written over 3,500 articles on technolo...

[View Bio](#) >

Loading the next article



- [TechCrunch Staff](#)
[Contact Us](#)
[Advertise](#)
[Crunchboard Jobs](#)
[Site Map](#)
- [Terms of Service](#)
[Privacy Policy](#)
[RSS Terms of Use](#)
- [OpenAI](#)
[Perplexity](#)
[Cisco](#)
[DeepMind](#)
[Windsurf](#)
[Tech Layoffs](#)
[ChatGPT](#)