ENTERPRISE

Observe continues to adapt to the changing world of software observability

Rebecca Szkutak — 6:00 AM PDT · July 30, 2025

IMAGE CREDITS: OBSERVE

Observe, an observability platform, was founded in 2017 in response to the changing nature of software observability. Companies started pushing out new versions of their software more frequently — and producing significantly more data because of it.

Now Observe is responding to the latest big shift in technology: AI.

San Mateo-based Observe helps companies get an inside look at the status of their software, which makes it easier for engineers to spot and solve disruptions and outages.

The recent advancements in AI are both a blessing and a burden for the company. Observe's observability product incorporates AI agents to help make finding and fixing issues faster for its customers. But advancements in AI mean companies are shipping software even faster than before and seeing their data balloon because of it.

CEO Jeremy Burton told TechCrunch that with the continual advancement of AI agents, observability continues to get more complex.

"In a few years, you're going to have hundreds or thousands of agents on your network that are all interacting with employees or interacting with each other," Burton said. "That's all great until something goes wrong, and you've got to try and, you know, do a Sherlock Holmes and figure out who done it, you know?"

But Observe is adapting to how the industry is changing, Burton said. The company released a Model Context Protocol (MCP) server earlier this year that allows developers to access their observability data from AI coding tools and large language models. This is to help meet developers where they are already working, Burton said, and help them accomplish tasks easier.

"We have customers already that are using the MCP server and really trying fairly radical workflows," Burton said. "They're sitting in their development environment, and they say, 'Hey, take a look at this ticket. Use Observe to go figure out what's happening, and then describe to me the code that you think is problematic, and then suggest the fix.' That would have been in the realm of science fiction even a year ago."

The company is also working toward supporting Apache Iceberg, an open source data table format that allows businesses to own and standardize their own data. Burton said companies really like that approach, and Observe expects to be able to support that format by the end of the year.

The company's revenue nearly tripled in 2024, and it saw 93% gross retention of its customers, although Burton declined to share specific numbers. The company counts large enterprises, including CapitalOne, Paramount, and Dialpad, as customers.

Observe raised $156 million in a Series C round led by Sutter Hill Ventures with participation from Madrona Ventures, Alumni Ventures, and strategic investors like Snowflake, among others.

The company will put the capital toward research and development and hiring. Observe hopes to roll out its private preview for Apache Iceberg support shortly.

"We've got a couple of really good things out there, but I feel like we've just gotten started," Burton said.

Topics: AI Apache Iceberg cloud data management Enterprise observability Observe United States

Rebecca Szkutak

Senior Reporter, Venture

Becca is a senior writer at TechCrunch that covers venture capital trends and startups. She previously covered the...

View Bio

KEEP READING

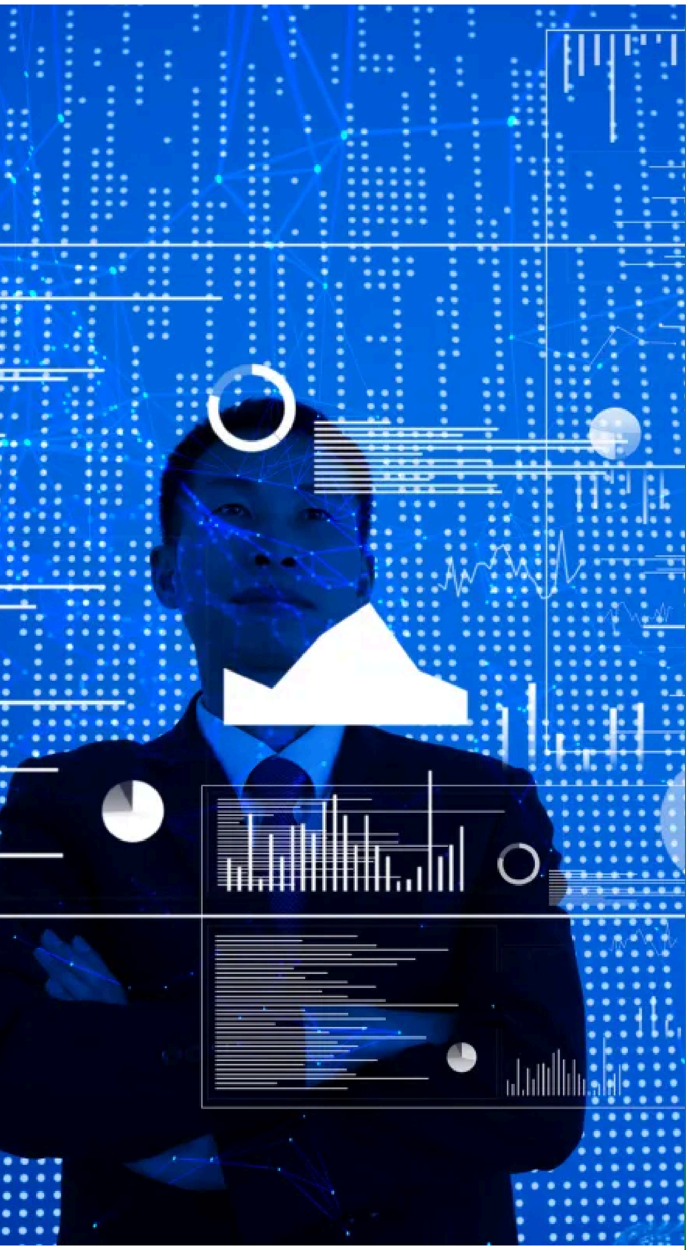

IMAGE CREDITS: WEIQUAN LIN / GETTY IMAGES

ENTERPRISE

Observability platform Observe raises $50M in debt, launches gen AI features

Kyle Wiggers — 6:00 AM PDT · October 5, 2023

Venture investors see potential in observability software — i.e. software that gives teams the ability to monitor, measure and understand the state of a system or app. And that's not surprising. The observability market is forecasted to reach $2 billion by 2026, climbing from $278 million in 2022, according to 650 Group.

The growth is perhaps because of the software's perceived benefits. A 2022 survey from Enterprise Strategy Group found that advanced observability deployments can cut downtime costs by 90%, keeping costs down to $2.5 million annually versus $23.8 million for observability "beginners."

In any case, vendors like Observe are riding high.

Observe, which develops software-as-a-service observability tools for storing, managing and analyzing machine-generated data and logs, has raised $50 million in convertible debt (i.e. debt that converts to equity) led by Sutter Hill Ventures. The financing will be used to grow Observe's sales and R&D teams, according to CEO Jeremy Burton, as the company looks to expand its headcount from 150 employees to 250 by the end of 2024.

Burton says that the decision to opt for debt funding was in the interest of deferring dilution.

"We expect to raise the Series B early next year, at which time the debt will convert to equity," Burton told TechCrunch in an email interview. "We've run the company off of debt financing for the past three years."

San Mateo-based Observe, which was founded in 2017 by Jacob Leverich, Jon Watte, Jonathan Trevor and Philipp Unterbrunner, stores all raw observability data in a data lake, a centralized repository. It curates and layers analytics on top of this data through what Burton calls a "data graph," ostensibly making it easier for users to navigate and understand the data.

Observe competes with app monitoring software; monitoring and log analytics tools like New Relic, Splunk, Datadog and Sumo Logic; and new entrants to the observability space such as Grafana, Chronosphere and Honeycomb. But it's launching new tools and capabilities to stay ahead of the curve.

Burton claims that Observe is now capable of ingesting over a petabyte of data per day into a single customer's instance while providing a "live" mode for interactive debugging. And, as of this week, Observe offers a range of generative AI features designed to expedite specific observability tasks.

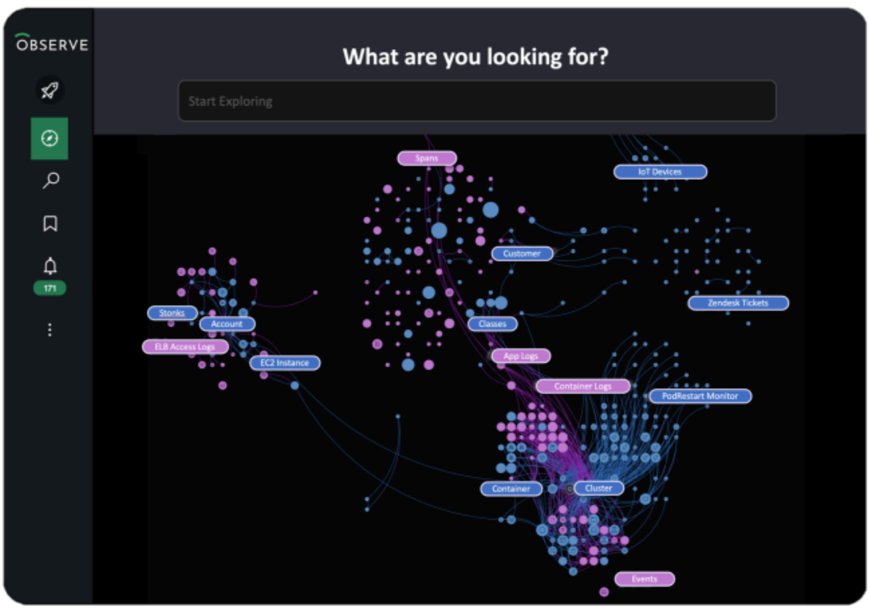

IMAGE CREDITS: OBSERVE

For example, Observe's new GPT Help module, basically a chatbot, responds to natural language commands about Observe's features, "how-to" tasks and error messages. Meanwhile, GPT Extract parses data to add structure to logs on the fly. GPT Slack Assistant embeds into Slack to help users troubleshoot issues and summarize threads for incident response. And OPAL Co-Pilot generates Opal code — Observe's query language — in response to natural language inputs.

Now, do organizations investing in observability want all this trendy generative AI stuff? It's possible. Burton seems to think so, for what it's worth.

"As organizations become increasingly digital, the volume of telemetry data generated by modern distributed applications is exploding," Burton said. "Legacy tools weren't designed for either the volume of data or the complexity of investigating unknown problems in production ... But since its founding, Observe has viewed observability as a data problem — if data used to troubleshoot could be stored in one place and reasoned about in a single tool, users would be able to solve problems much faster."

Generative AI isn't all that's new in Observe's platform. The company's introducing Observe Apps, prebuilt packages containing Observe configurations and best practices for observing specific dev environments. And Observe's launching a public API and command line interface, alongside options to export data to a CSV file or a Snowflake dashboard for further analysis.

Burton admits that the economic downturn and pandemic forced Observe to shift its go-to-market approach, with an increased focus on companies with 200 to 2,000 employees. But he claims that this led to success in the end, bringing Observe's client base to over 60 brands and around 1,600 monthly average users.

"Over the past year, we've moved the sales team up to larger accounts and that has resulted in much higher average sales prices and less churn," Burton said. "While we're not disclosing annual recurring revenue at this time, what we can share is that Observe has already surpassed its new annual contract value plan for 2023."

Will generative AI and a platform refresh take Observe to new heights? Perhaps. Or maybe it'll simply enable Observe to stay a disciplined, measured-growth course, which isn't necessarily a bad thing.

Topics: AI data observability Enterprise Funding observability Observe startup Startups

Kyle Wiggers
AI Editor

Kyle Wiggers was TechCrunch's AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

View Bio >

KEEP READING

IMAGE CREDITS: YUICHIRO CHINO / GETTY IMAGES

Data observability startup Observe secures $70M

Kyle Wiggers — 5:00 AM PDT · May 11, 2022

Software apps are increasingly cloud-based, distributed and updated at a breakneck pace. They're growing more complex, in other words, which is becoming a major challenge for dev teams. A survey from Enterprise Strategy Group commissioned by Yotascale — which isn't an impartial vendor, granted (if there is such a thing) — found that 84% of companies in 2022 considered tracking cloud cost allocation alone to be burdensome and time consuming. Sixty-one percent admitted they lacked "sufficient visibility" into which teams own which cloud resources.

If you ask Jeremy Burton, the solution is greater observability — i.e., tools to keep watch over all of an app's components, from the back-end systems to the front-end interface. That's not especially surprising from Burton, considering he's the CEO of Observe, an observability platform developer co-started in 2017 by Jacob Leverich, Jonathan Trevor, Jon Watte and Philip Unterbrunner. But it's not a stance without merit. According to a report form Techstrong Research and Testlio, 54% of developers can't keep up with the demand for new and updated software, underlining the hurdles the

sought to fix the problem with a product that combines log analytics and monitoring with app performance management. Built on top of Snowflake's cloud data platform, Observe ingests machine-generated data — almost 50 terabytes every day — and transforms it into "human-readable" data sets to form relationships between them. The relationships, which range from things like customer accounts to e-commerce shopping carts, can be used by devs to investigate why an app misbehaved, for example, or suffered an outage.

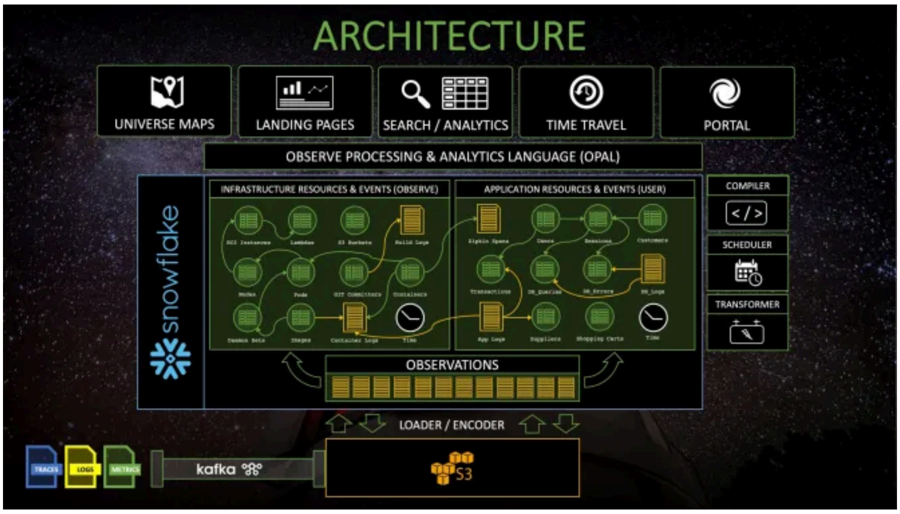

OBSERVE'S DATA PROCESSING ARCHITECTURE. **IMAGE CREDITS:** OBSERVE

Observe launched in preview in 2020, and, in anticipation of expansion later this year, the company today closed a $70 million debt round between Sutter Hill Ventures, Capital One Ventures and Madrona Ventures. It brings the company's total capital to date to $114.5 million.

Burton has a distinguished career, one unsurprisingly rooted in developer operations. Prior to co-founding Observable, Burton served as the SVP of product and services marketing at Oracle; was group president of Symantec's enterprise security and data management division; headed EMC's products group as president; and was chief marketing officer at Dell. Leverich was the director of engineering at Splunk while Unterbrunner was a founding engineer at Snowflake.

"The company was founded, like Snowflake and Pure Storage, by Sutter Hill Ventures back in late 2017. The founding team was recruited from Splunk, Snowflake, Wavefront and Roblox," Burton said. "All the founders had a background in working with massive data volumes."

Observe curates data from a company's apps into "resources" like users, sessions, help desk tickets and software builds. Users can use the platform ask questions about the resources to see how they're connected. Beyond this, they can track the state of how each resource's attributes (e.g., IP addresses) change over time and provides levels of abstraction on top of the data, including automatically generated dashboards and worksheets with spreadsheet-like interfaces.

When asked about Observe's data retention policies, Burton said that the platform keeps data for 13 months by default but can shorten or lengthen that window depending on a customer's individual requirements. Customers can delete data on request, and Burton claims that data isn't used for any purpose other than observability.

Burton isn't ignorant of the competition in the observability solutions segment — which is fast-growing. Statista estimates that the sector will increase from $12.98 billion in worth in 2020 to

$19.38 billion in 2024. Databand is among the startups offering products for observability, as well as Elastic, Datadog, Splunk and New Relic.

In addition to its analytics capabilities, Burton argues that Observe's pricing model differentiates it from others on the market. Customers pay storage and querying costs, the former of which reflect current Amazon S3 pricing plus charges for processing data as it's ingested. Querying data consumes usage-based "Observe credits."

"Observe eliminates data silos and provides one consistent interface for teams to troubleshoot problems an order of magnitude faster. In addition, because of the modern architecture, we can do it an order of magnitude cheaper," Burton said.

A $70 million debt round might not instill enormous confidence, but Burton says that the strategy is to use debt to increase Observe's value before it converts to equity while minimizing dilution. Further attempting to allay any investor fears, he said that Observe's annual contract value is in the millions of dollars and that the company plans to increase its headcount from 80 employees to 150 by the end of the year.

Observe has close to 50 customers at present, Burton claims, including Upstart Financial, OpenGov and Top Golf.

"This is particularly important when you are trying to solve really hard technical challenges which may take three to four years to bring to market," Burton added, referring to the debt.

Sources of financing aside, Observe will have to convince skeptical devs that its tool solves some of the shortcomings of popular observability tools. In a Dimensional Research survey, 79% of respondents indicated that the overall cost of log observability and management will skyrocket in 2022 if current tools don't evolve. With eight months left in the year, we'll have to see whether that's the case.

Topics: Capital One Ventures Cloud Computing data developers Enterprise Enterprise Funding Madrona Ventures Media & Entertainment monitoring observability performance monitoring Startups Startups sutter hill ventures

Kyle Wiggers

AI Editor

Kyle Wiggers was TechCrunch's AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

View Bio

KEEP READING

ENTERPRISE

IMAGE CREDITS: JOHN LUND / GETTY IMAGES

Data observability platform Kloudfuse launches out of stealth with $23M

Kyle Wiggers — 6:00 AM PST · November 6, 2023

It's tough to recover from catastrophic system failures when you're flying blind. Yet many companies admit they're doing just that — flying blind, more or less — where it concerns their software and critical infrastructure.

In a recent survey by Logz.io, the observability platform provider, 64% of businesses report over an hour mean time to recovery — in large part because they're encountering roadblocks in getting the data they need to diagnose the underlying issues.

Now, there's no shortage of observability vendors claiming to offer definitive solutions to this problem. But Krishna Yadappanavar — who previously co-founded a startup, Springpath, that Cisco acquired for $320 million in 2017 — thought there was room for one more.

With Springpath colleague Ashish Hanwadikar and Pankaj Thakkar, a former VMware engineer, Yadappanavar co-founded Kloudfuse, which provides several tools for data observability across different compute environments. (Yes — that's "Cloudfuse" spelled with a "*K*" instead of a "C.")

Kloudfuse today emerged from stealth with $23 million in funding, inclusive of a $17 million Series A round led by Newlands with participation from Blumberg Capital, Aspenwood Ventures, High Sage and Exponent.

“Hanwadikar, Thakkar and I saw an opportunity to help developers, DevOps and site reliability engineers, who spend significant time managing observability, troubleshooting data and analytics and lack automation capabilities,” Yadappanavar told TechCrunch in an email interview. “We co-founded Kloudfuse with the mission of making unified cloud observability more accessible and scalable to enterprises, enabling them to handle increased cardinality, gain more control and reduce costs.”

Through Kloudfuse’s platform — which is technically a data lake, a centralized repository for unstructured observability data — customers can analyze metrics, events, logs and traces from various software apps and machines. Kloudfuse automatically alerts them to anomalies in the data that might suggest something’s amiss.

“Economic factors such as high interest rates and fears of a looming recession are forcing organizations to cut observability bills,” Yadappanavar said. “Thus, enterprises are being forced to balance keeping costs low without sacrificing safety, efficiency and performance … Kloudfuse is well-prepared to weather potential economic headwinds, as the company is helping to solve a significant problem that touches all enterprises regardless of their size, sector and structure.”

- Key sections of the US Constitution deleted from government’s website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a $2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue’s AI-generated ad isn’t just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines
- The best dating apps aren’t even dating apps

As alluded to, Kloudfuse has many rivals in the observability space — Acceldata, Sifflet, Observe and Manta, to name a few. But Yadappanavar claims revenue has tripled in the last two quarters as Kloudfuse's customer base has grown to include household brands like Workday, GE HealthCare and Automation Anywhere.

David Blumberg, managing partner at Blumberg Capital, had this to say via email: "We continue to invest in Kloudfuse because enterprises need a better way to maximize data observability while controlling budgets, without sacrificing security, efficiency and performance. Krishna and team have earned the trust of large, sophisticated enterprises across industries, and delivered value quickly."

Kloudfuse plans to use the proceeds from the latest funding round to expand its product management, marketing and sales teams.

Topics: cloud data observability Enterprise

Funding Kloudfuse observability startup

Startups

Kyle Wiggers

AI Editor

Kyle Wiggers was TechCrunch's AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

View Bio

Loading the next article