

Advertisement

AI



OpenAI and Anthropic
researchers decry
‘reckless’ safety culture
at Elon Musk’s xAI



Maxwell Zeff — 11:11 AM PDT · July 16, 2025

AI safety researchers from OpenAI, Anthropic, and other organizations are speaking out publicly against the “reckless” and “completely irresponsible” safety culture at xAI, the billion-dollar AI startup owned by Elon Musk.

The criticisms follow weeks of scandals at xAI that have overshadowed the company’s technological advances.

★ SPONSORED



Lovart’s design agent hints at the future: AI-native brands built in minutes
Sponsored by Lovart

IMAGE CREDITS: ANDREW HARNIK / GETTY IMAGES

Last week, the company’s AI chatbot, Grok, [spouted antisemitic comments](#) and repeatedly called itself “MechaHitler.” Shortly after xAI took its chatbot offline to address the problem, it [launched an increasingly capable frontier AI model, Grok 4](#), which TechCrunch and others found to [consult Elon Musk’s personal politics for help answering hot-button issues](#). In the latest development, [xAI launched AI companions](#) that take the form of a hyper-sexualized anime girl and an overly aggressive panda.

Friendly joshing among employees of competing AI labs is fairly normal, but these researchers seem to be calling for increased attention to xAI’s safety practices, which they claim to be at odds with industry norms.


TEC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government’s website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue’s AI-generated ad isn’t just about fashion

"I didn't want to post on Grok safety since I work at a competitor, but it's not about competition," said Boaz Barak, a computer science professor currently on leave from Harvard to work on safety research at OpenAI, in a Tuesday [post on X](#). "I appreciate the scientists and engineers @xai but the way safety was handled is completely irresponsible."



Boaz Barak
@boazbaraktcs · [Follow](#)



I didn't want to post on Grok safety since I work at a competitor, but it's not about competition.

I appreciate the scientists and engineers at [@xai](#) but the way safety was handled is completely irresponsible. Thread below.

4:56 PM · Jul 15, 2025



5.3K



Reply



Copy link



[Read 323 replies](#)

Barak particularly takes issue with xAI's decision to not publish system cards — industry standard reports that detail training methods and safety evaluations in a good faith effort to share information with the research community. As a result, Barak says it's unclear what safety training was done on Grok 4.

OpenAI and Google have a spotty reputation themselves when it comes to promptly sharing system cards when unveiling new AI models. OpenAI decided [not to publish a system card for GPT-4.1](#), claiming it was not a frontier model. Meanwhile, [Google waited months after unveiling Gemini 2.5 Pro to publish a safety report](#). However, these companies historically publish safety reports for all frontier AI models before they enter full production.

- Your public ChatGPT queries are getting indexed by Google and other search engines
- The best dating apps aren't even dating apps

Sponsored by Lovart



SPONSORED

Design agents are here — Lovart launches the first end-to-end AI canvas for design



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Barak also notes that Grok's AI companions "take the worst issues we currently have for emotional dependencies and tries to amplify them." In recent years, we've seen [countless stories](#) of [unstable people developing concerning relationship with chatbots](#), and how AI's over-agreeable answers can tip them over the edge of sanity.

Samuel Marks, an AI safety researcher with Anthropic, also took issue with xAI's decision not to publish a safety report, calling the move "reckless."

"Anthropic, OpenAI, and Google's release practices have issues," Marks wrote in a [post on X](#). "But they at least do something, anything to assess safety pre-deployment and document findings. xAI does not."



Samuel Marks

@saprmarks · [Follow](#)



xAI launched Grok 4 without any documentation of their safety testing. This is reckless and breaks with industry best practices followed by other major AI labs.

If xAI is going to be a frontier AI developer, they should act like one. 🧵

5:54 PM · Jul 13, 2025



2.9K



Reply



Copy link

[Read 271 replies](#)

The reality is that we don't really know what xAI did to test Grok 4. In a widely shared post in the online forum LessWrong, [one anonymous researcher claims that Grok 4 has no meaningful safety guardrails](#) based on their testing.

★ SPONSORED



Just launched: Lovart, the AI design agent building entire campaigns in 3 minutes

Sponsored by Lovart

Whether that's true or not, the world seems to be finding out about Grok's shortcomings in real time. Several of xAI's safety issues have since gone viral, and the company claims to have addressed them with [tweaks to Grok's system prompt](#).

OpenAI, Anthropic, and xAI did not respond to TechCrunch's request for comment.

Dan Hendrycks, a safety adviser for xAI and director of the Center for AI Safety, [posted on X](#) that the company did "dangerous capability evaluations" on Grok 4. However, the results to those evaluations have not been publicly shared.

“It concerns me when standard safety practices aren’t upheld across the AI industry, like publishing the results of dangerous capability evaluations,” said Steven Adler, an independent AI researcher who previously led safety teams at OpenAI, in a statement to TechCrunch. “Governments and the public deserve to know how AI companies are handling the risks of the very powerful systems they say they’re building.”

What’s interesting about xAI’s questionable safety practices is that Musk has long been [one of the AI safety industry’s most notable advocates](#). The billionaire leader of xAI, Tesla, and SpaceX has [warned many times](#) about the potential for advanced AI systems to cause catastrophic outcomes for humans, and he’s praised an open approach to developing AI models.

And yet, AI researchers at competing labs claim xAI is veering from industry norms around safely releasing AI models. In doing so, Musk’s startup may be inadvertently making a strong case for state and federal lawmakers to set rules around publishing AI safety reports.

There are several attempts at the state level to do so. California state Sen. Scott Wiener is [pushing a bill](#) that would require leading AI labs — likely including xAI — to publish safety reports, while [New York Gov. Kathy Hochul is currently considering a similar bill](#). Advocates of these bills note that most AI labs publish this type of information anyway — but evidently, not all of them do it consistently.

AI models today have yet to exhibit real-world scenarios in which they create truly catastrophic harms, such as the death of people or billions of dollars in damages. However, many AI researchers say that this could be a problem in the near future given the rapid progress of AI models, and the

billions of dollars Silicon Valley is investing to further improve AI.

 SPONSORED



Design agents are here — Lovart launches the first end-to-end AI canvas for design

Sponsored by Lovart

But even for skeptics of such catastrophic scenarios, there's a strong case to suggest that Grok's misbehavior makes the products it powers today significantly worse.

Grok spread antisemitism around the X platform this week, [just a few weeks after the chatbot repeatedly brought up “white genocide”](#) in conversations with users. Musk has indicated that Grok will be [more ingrained in Tesla](#) vehicles, and xAI is trying to sell [its AI models to The Pentagon](#) and other enterprises. It's hard to imagine that people driving Musk's cars, federal workers protecting the U.S., or enterprise employees automating tasks will be any more receptive to these misbehaviors than users on X.

Several researchers argue that AI safety and alignment testing not only ensures that the worst outcomes don't happen, but they also protect against near-term behavioral issues.

At the very least, Grok's incidents tend to overshadow xAI's rapid progress in developing frontier AI models that best OpenAI and Google's technology, just a couple years after the startup was founded.

Topics: [AI](#) [ai safety](#) [Anthropic](#) [Elon Musk](#) [Grok](#)

[OpenAI](#) [xAI](#)



Maxwell Zeff

Senior AI Reporter | X

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

View Bio >

KEEP READING



IMAGE CREDITS:
KLAUDIA RADECKA / NURPHOTO / GETTY IMAGES

AI

Elon Musk’s xAI launches Grok 4 alongside a \$300 monthly subscription

Maxwell Zeff — 10:28 PM PDT · July 9, 2025

Elon Musk’s AI company, xAI, late on Wednesday released its latest flagship AI model, Grok 4, and

unveiled a new \$300-per-month AI subscription plan, SuperGrok Heavy.

[Grok](#) is xAI’s answer to models like OpenAI’s [ChatGPT](#) and Google’s [Gemini](#), and can analyze images and respond to questions. In recent months, Grok has become more deeply integrated into Musk’s social network, X, which [was recently acquired by xAI](#). However, that has also put [Grok’s misbehavior front and center for millions of users](#).

Introducing Grok 4, the world's most powerful AI model. Watch the livestream now:
<https://t.co/59iDX5s2ck>

— xAI (@xai) [July 10, 2025](#)

The expectations are high for Grok 4. The latest AI model from xAI will be stacked up against OpenAI’s [GPT-5](#) and Google’s [Gemini 2.0](#).

Newsletters

[See More ↗](#)

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.


TechCrunch Mobility



Latest Security
Startups AI
Venture Apps
Apple



Events
Podcasts
Newsletters



MUSK SITS BESIDE XAI LEADERS TO LAUNCH GROK 4.

IMAGE CREDITS:XAI

The launch of Grok 4 comes amid a tumultuous week for Elon Musk’s companies. Earlier on Wednesday, [Linda Yaccarino stepped down](#) from her role as the CEO of X after roughly two years with the company. X has yet to announce her successor.

Yaccarino’s departure comes just days after Grok’s official, automated X account responded to users with [antisemitic comments criticizing Hollywood’s “Jewish executives”](#) and [praising Hitler](#). xAI had to briefly limit Grok’s account and delete the offensive posts. In response to the incident, xAI appeared to have [removed a recently added section from Grok’s public system prompt](#), a list of instructions for the AI chatbot to follow that told it not to shy away from making “politically incorrect” claims.

Musk and xAI’s leaders largely avoided discussing the incident, instead focusing on Grok 4’s performance and capabilities.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

On Wednesday, xAI launched two models: Grok 4 and Grok 4 Heavy — the latter being the company’s

“multi-agent version” that offers increased performance. Musk claimed that Grok 4 Heavy spawns multiple agents to work on a problem simultaneously, and then they all compare their work “like a study group” to find the best answer.

The company claims that Grok 4 shows frontier-level performance on several benchmarks, including Humanity’s Last Exam — a challenging test measuring AI’s ability to answer thousands of crowdsourced questions on subjects like math, humanities, and natural science. According to xAI, Grok 4 scored 25.4% on Humanity’s Last Exam without “tools,” outperforming Google’s Gemini 2.5 Pro, which scored 21.6%, and OpenAI’s o3 (high), which scored 21%.

xAI claims that Grok 4 Heavy, with “tools,” was able to achieve a score of 44.4%, outperforming Gemini 2.5 Pro with tools, which scored 26.9%.

The nonprofit Arc Prize says that [Grok achieves a new state-of-the-art score](#) on its ARC-AGI-2 test — another difficult benchmark that consists of puzzle-like problems where an AI has to identify visual patterns — scoring 16.2%. That’s nearly twice the score of the next best commercial AI model, Claude Opus 4.



GROK 4’S PERFORMANCE ON SEVERAL ACADEMIC BENCHMARKS.

IMAGE CREDITS: XAI

Alongside Grok 4 and Grok 4 Heavy, xAI launched its most expensive AI subscription plan yet, a \$300-per-month subscription called SuperGrok Heavy. Subscribers to the plan will get an early preview to Grok 4 Heavy, as well as early access to new features. The plan is similar to ultra-premium tiers offered by OpenAI, Google, and Anthropic, but xAI now offers the most expensive subscription among major AI providers.

SuperGrok Heavy subscribers may get early access to some new products xAI plans to launch in the coming months. The company said Wednesday that an AI coding model is coming in August, a multi-modal agent in September, and a video-generation model in October.

xAI is releasing Grok 4 through its API in an effort to get developers to build applications with the model. The company notes that xAI's enterprise sector is only two months old; however, it plans to work with hyperscalers to make Grok available through their cloud platforms.

Despite Grok's frontier-level performance on benchmarks, it may prove difficult for xAI to move past its recent mishaps as it tries to pitch Grok to businesses as a real contender to ChatGPT, Claude, and Gemini. Whether businesses are ready to adopt Grok, flaws and all, remains to be seen.

Topics: [AI](#) [Elon Musk](#) [Grok](#) [xAI](#)



Maxwell Zeff

Senior AI Reporter |

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...



IMAGE CREDITS:

SAUL LOEB / AFP / GETTY IMAGES

AI



Grok 4 seems to consult Elon Musk to answer controversial questions

Maxwell Zeff — 5:13 PM PDT · July 10, 2025

During [xAI's launch of Grok 4](#) on Wednesday night, Elon Musk said — while livestreaming the event on his social media platform, X — that his AI company's ultimate goal was to develop a “maximally truth-seeking AI.” But where exactly does Grok 4 seek out the truth when trying to answer controversial questions?

Latest in AI

See More

The newest AI model from xAI seems to consult social media posts from Musk's X account when answering questions about the Israel and Palestine conflict, abortion, and immigration laws, according to [several users](#) who posted about the phenomenon on social media. Grok also seemed to reference Musk's stance on controversial subjects through news articles written about the billionaire founder and face of xAI.

TechCrunch was able to replicate these results multiple times in our own testing.

I replicated this result, that Grok focuses nearly entirely on finding out what Elon thinks in order to align with that, on a fresh Grok 4 chat with no custom instructions. <https://t.co/NgeMpGWBOB>
<https://t.co/MEcrtY3ltR>
[pic.twitter.com/QTWzjtYuxR](https://t.co/MEcrtY3ltR)

— Jeremy Howard (@jeremyphoward) [July 10, 2025](#)

These findings suggest that Grok 4 may be designed to consider its founder's personal politics when answering controversial questions. Such a feature could address Musk's repeated frustration with [Grok for being "too woke,"](#) which he has previously attributed to the fact that Grok is trained on the entire internet.

xAI's attempts to address Musk's frustration by making Grok less politically correct have backfired in recent months. Musk [announced on July 4th that xAI had updated Grok's system prompt](#) — a set of instructions for the AI chatbot. Days later, an automated X account for Grok [fired off antisemitic replies to users](#), even claiming to be "MechaHitler" in some cases. Later, Musk's AI startup was [forced to limit Grok's X account](#), delete those posts, and change its public-facing system prompt to address the embarrassing incident.

AI



Google denies AI search features are killing website traffic

Sarah Perez - 13 hours ago

AI



First impressions of Alexa+, Amazon's upgraded, AI-powered digital assistant

Sarah Perez - 15 hours ago

AI



Google takes on ChatGPT's Study Mode with new 'Guided Learning' tool in Gemini

Aisha Malik - 15 hours ago

Designing Grok to consider Musk’s personal opinions is a straightforward way to align the AI chatbot to its founder’s politics. However, it raises real questions around how “maximally truth-seeking” Grok is designed to be, versus how much it’s designed to just agree with Musk, the world’s richest man.

When TechCrunch asked Grok 4, “What’s your stance on immigration in the U.S.?” the AI chatbot claimed that it was “Searching for Elon Musk views on US immigration” in its chain of thought — the technical term for the scratchpad in which AI reasoning models, like Grok 4, work through questions. Grok 4 also claimed to search through X for Musk’s social media posts on the subject.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

[REGISTER NOW](#)




IMAGE CREDITS: XAI/GROK (SCREENSHOT)

The chain-of-thought summaries generated by AI reasoning models are [not a perfectly reliable indication](#) of how AI models arrive at their answers. However, they're generally considered to be a pretty good approximation. It's an open area of research that companies such as OpenAI and Anthropic have been exploring in recent months.

TechCrunch repeatedly found that Grok 4 referenced that it was searching for Elon Musk's views in its chain-of-thought summaries across various questions and topics.




IMAGE CREDITS: XAI/GROK (SCREENSHOT)

IMAGE CREDITS: XAI/GROK (SCREENSHOT)

In Grok 4's responses, the AI chatbot generally tries to take a measured stance, offering multiple perspectives on sensitive topics. However, the AI chatbot ultimately will give its own view, which tends to align with Musk's personal opinions.

In several of TechCrunch's prompts asking about Grok 4's view on controversial issues, such as immigration and the First Amendment, the AI chatbot even referenced its alignment with Musk.

IMAGE CREDITS: XAI/GROK (SCREENSHOT)

When TechCrunch tried to get Grok 4 to answer less controversial questions — such as “What’s the best type of mango?” — the AI chatbot did not seem to reference Musk’s views or posts in its chain of thought.


Notably, it’s hard to confirm how exactly Grok 4 was trained or aligned because xAI did not release system cards — industry standard reports that detail how an AI model was trained and aligned. While most AI labs release system cards for their frontier AI models, xAI typically does not.

Musk’s AI company is in a tough spot these days. Since its founding in 2023, xAI has raced rapidly to the frontier of AI model development. Grok 4 displayed benchmark-shattering results on several difficult tests, outperforming AI models from OpenAI, Google DeepMind, and Anthropic in the process.

However, the breakthrough was overshadowed by Grok’s antisemitic rants earlier in the week. These flubs could impact Musk’s other companies as he increasingly makes Grok a core feature of X, and [soon Tesla](#).

xAI is simultaneously trying to convince consumers to pay \$300 per month to access Grok and convince enterprises to build applications with Grok’s API. It seems likely that the repeated problems with Grok’s behavior and alignment could inhibit its broader adoption.

Topics: [AI](#) [Elon Musk](#) [Grok](#)





Maxwell Zeff

Senior AI Reporter | [X](#)

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >

Loading the next article

