

Advertisement

AI



OpenAI launches a
general purpose agent in
ChatGPT

OpenAI is launching a new general purpose AI agent in ChatGPT, which the company says can complete a wide variety of computer-based tasks on behalf of users. OpenAI says the agent can automatically navigate a user's calendar, generate editable presentations and slideshows, and run code.

The tool, called ChatGPT agent, combines several capabilities from OpenAI's previous agentic tools, including [Operator's ability to click around on websites](#), as well as [Deep Research's ability to synthesize information](#) from dozens of websites into a concise research report. OpenAI says users will be able to interact with the agent simply by prompting ChatGPT in natural language.

IMAGE CREDITS: BRYCE DURBIN / TECHCRUNCH

ChatGPT agent is rolling out on Thursday to subscribers to OpenAI's Pro, Plus, and Team plans. To activate the tool, users can select "agent mode" in ChatGPT's dropdown menu of tools.


TECH DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion



OpenAI 
@OpenAI · [Follow](#)



ChatGPT can now do work for you using its own computer.

Introducing ChatGPT agent—a unified agentic system combining Operator’s action-taking remote browser, deep research’s web synthesis, and ChatGPT’s conversational strengths.

The media could not be played.

Reload

5:53 PM · Jul 17, 2025



14.2K



Reply



Copy link

[Read 862 replies](#)

- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

The launch of ChatGPT agent represents OpenAI’s boldest attempt yet to turn ChatGPT into an agentic product that can take actions and offload tasks for users, rather than just answering questions. In recent years, Silicon Valley companies including OpenAI, Google, and Perplexity have unveiled dozens of AI agents that have promised to do just that. However, these [early version AI agents have proven to struggle with complex tasks](#), and they seem less compelling as products than the ultimate vision tech executives pitch around AI agents.

That said, OpenAI says ChatGPT agent is far more capable than its previous offerings.

The company's new agent can access [ChatGPT connectors](#), allowing users to connect apps like Gmail and GitHub so that the agent can find relevant information to your prompts. OpenAI says ChatGPT agent has access to a terminal, and it can use APIs to access certain apps.

OpenAI suggests that users can tap ChatGPT agent to “plan and buy ingredients to make Japanese breakfast for four,” as well as “analyze three competitors and create a slide deck.” These kinds of capabilities requires ChatGPT agent to parse through websites, plan a course of action, and use tools — much more complicated tasks than OpenAI has previously tried to tackle with agents.

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

[REGISTER NOW](#)

The model underlying ChatGPT agent offers state-of-the-art performance on several benchmarks, according to OpenAI.

The company says the ChatGPT agent model scores 41.6% on Humanity's Last Exam (pass@1), a difficult test made up of thousands of questions

across more than one hundred subjects. That's roughly double what OpenAI's o3 and o4-mini scored on the test.

On FrontierMath, one of the hardest known math benchmarks, OpenAI says ChatGPT agent scores 27.4% when it has access to tools, such as a terminal for code execution. The previous state-of-the-art score comes from o4-mini, which scored just 6.3%.

OpenAI notes that it developed ChatGPT agent with safety in mind, largely because the product presents some newfound capabilities that could make it more dangerous in the hands of a bad actor. OpenAI has previously warned that agentic models could present more dangerous capabilities.

In a [safety report](#) for ChatGPT agent, OpenAI says it's designated the model as "high capability" in biological and chemical weapon domains, which is defined in OpenAI's Preparedness Framework as a model with the ability to "amplify existing pathways to severe harm." OpenAI notes that it does not have direct evidence of this, but it's decided to take a precautionary approach and activate new safeguards to mitigate these risks.

The new safeguards for ChatGPT agent include a monitor that works in real time as users interact with the product. OpenAI says it runs a classifier across every prompt entered into ChatGPT agent, determining whether the request is related to biology. If so, OpenAI runs ChatGPT agent's response through a second monitor that determines whether the content could be used to evoke a biological threat.

OpenAI also says it disabled ChatGPT's memory feature for this agent to prevent misuse. In other parts of ChatGPT, OpenAI's memory feature allows

the chatbot to reference information from previous user chats. However, OpenAI says bad actors could use the feature in ChatGPT agent to exfiltrate sensitive data through prompt injection attacks. The company says it may revisit adding the feature in the future, however.

While ChatGPT agent sounds impressive, it remains to be seen how capable it truly is in the real world. Until now, agent technology has proven relatively brittle when interacting with the real world. That said, OpenAI says it's developed a more capable model that's able to deliver on the promise of AI agents.

This story was updated with more information.

Topics: [agents](#) [AI](#) [ChatGPT](#) [OpenAI](#)

[f](#) [X](#) [in](#) [reddit](#) [email](#) [link](#)



Maxwell Zeff

Senior AI Reporter | [X](#)

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >

KEEP READING

AI

[f](#) [X](#) [in](#) [reddit](#) [email](#) [link](#)



OpenAI launches Operator, an AI agent that performs tasks autonomously

Maxwell Zeff — 10:00 AM PST · January 23, 2025

OpenAI CEO Sam Altman kicked off this year by saying in a blog post that 2025 would [be big for AI agents](#), tools that can automate tasks and take actions on your behalf.

Now, we're seeing OpenAI's first real attempt.

OpenAI announced on Thursday that it is launching a research preview of Operator, a general-purpose AI agent that can take control of a web browser and independently perform certain actions. [Operator is coming to U.S. users on ChatGPT's \\$200 Pro subscription plan first](#). OpenAI says it plans to roll this feature out to more users in its Plus, Team, and Enterprise tiers eventually.

"[Operator] will be [in] other countries soon," OpenAI CEO Sam Altman said during a livestream Thursday. "Europe will, unfortunately, take a while."

This initial research preview is available through [operator.chatgpt.com](#), but soon, OpenAI says it

Newsletters

[See More](#)

Subscribe for the industry's biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch's coverage.

Startups Weekly




Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

wants to integrate Operator into all of its ChatGPT clients.



OPERATOR'S HOME SCREEN.

IMAGE CREDITS: OPENAI

Operator promises to automate tasks such as booking travel accommodations, making restaurant reservations, and shopping online, according to OpenAI. There are several task categories users can choose from within the Operator interface, including shopping, delivery, dining, and travel — all of which enable different kinds of automation.

When ChatGPT users activate Operator, a small window will pop up showing a dedicated web browser that the agent uses to complete tasks, along with explanations of specific actions the agent is performing. Users can still take control of their screen while Operator is working, as Operator uses its own dedicated browser.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.


Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).



Sponsored by Lovart



SPONSORED

Lovart's design agent hints at the future: AI-native brands built in minutes

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from


the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

OpenAI says that Operator is powered by a Computer-Using Agent model, or CUA, that combines the vision capabilities of the company's [GPT-4o](#) model with reasoning abilities from OpenAI's more advanced models. The CUA is trained to interact with the front-end of websites, meaning it doesn't need to use developer-facing APIs to tap into different services.

Latest Security
Startups AI
Venture Apps
Apple



Events
Podcasts
Newsletters

IMAGE CREDITS: OPENAI

“The CUA model is trained to ask for user confirmation before finalizing tasks with external side effects, for example before submitting an

order, sending an email, etc., so that the user can double-check the model's work before it becomes permanent," OpenAI writes in materials provided to TechCrunch. "[It] has already proven useful in a variety of cases, and we aim to extend that reliability across a wider range of tasks."

But OpenAI warns the CUA isn't perfect. The company says it "[doesn't] expect [the] CUA to perform reliably in all scenarios just yet."

"Currently, Operator cannot reliably handle many complex or specialized tasks," OpenAI [adds](#) in a support document, "such as creating detailed slideshows, managing intricate calendar systems, or interacting with highly customized or non-standard web interfaces."

Out of an abundance of caution, OpenAI is also requiring supervision for some tasks, like banking transactions, the CUA and Operator could perform mostly on their own. Users will need to take over to put in credit card information, for example. OpenAI says that Operator doesn't collect or screenshot any data.

"On particularly sensitive websites, such as email, Operator requires active user supervision, ensuring users can directly catch and address any potential mistakes the model might make," OpenAI says in its support materials.

This limits the usefulness of Operator, to be sure — but also ensures that the agent doesn't hallucinate and, say, spend your mortgage payment on accent chairs. Google took a similar approach with its [Project Mariner](#) AI agent, which also doesn't fill in info like credit card numbers.

Limitations

Operator has a few limitations worth noting.

There are rate limits — both daily and task-dependent. OpenAI says that Operator can perform multiple tasks at once, but that there are “dynamic limits” on this. There is also an overall usage limit that resets daily.

At this release stage, Operator will also refuse to perform tasks outright for security reasons, like sending emails (despite the fact that the CUA is capable of this) and deleting calendar events. OpenAI says this will change in the future, but gives no ETA.

Operator may also get “stuck” if it runs into a particularly complex interface, password field, or CAPTCHA check. It’ll ask the user to take over when this occurs, OpenAI says.

An agentic future

OpenAI has been rather slow to develop an AI agent compared to rivals (see: agents from [Rabbit](#), [Google](#), and [Anthropic](#)), which may have something to do with the safety risks around the technology.

When an AI system can take actions on the web, it opens the door to much more dangerous use cases from nefarious actors. You could automate AI agents to orchestrate phishing scams or DDoS attacks, or have them snatch up tickets to a concert before anyone else could. Especially for a tool as widely used as ChatGPT, it’s important OpenAI take steps to prevent those sorts of exploits.

OpenAI seems to think Operator is safe enough to release in its current form, at least as a research preview.

“Operator employs tools that seek to limit the model’s susceptibility to malicious prompts, hidden instructions, and phishing attempts,” OpenAI [explains](#) on its website. “A monitoring system pauses execution if suspicious activity is detected, while automated and human-reviewed pipelines continuously update safeguards.”

Operator is OpenAI’s boldest attempt yet at creating an AI agent. Last week, [OpenAI released Tasks](#), giving ChatGPT simple automation features such as the ability to set reminders and schedule prompts to run at a set time every day.

Tasks gave ChatGPT users some familiar, but necessary, features to make ChatGPT as practical to use as Siri or Alexa. However, Operator shows off capabilities that the previous generation of virtual assistants could never do.

AI agents have been pitched as the next big thing in AI after ChatGPT: a new technology that will change how people use the internet and their PCs. Instead of simply delivering and processing information, agents can — in theory — take actions and actually do things.

With the release of OpenAI’s first concrete take on agents, it’ll soon become clear just how realistic this vision is.

Topics: [agent](#) [AI](#) [AI](#) [ai agent](#) [ChatGPT](#) [OpenAI](#)

[operator](#)

[f](#) [X](#) [in](#) [📧](#) [✉](#) [🔗](#)



Maxwell Zeff


Senior AI Reporter | [X](#)

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >

KEEP READING

AI



Inside OpenAI's quest to make AI do anything for you

How OpenAI's first reasoning model sparked Silicon Valley's obsession with general-purpose AI agents — and what comes next.

Maxwell Zeff

7:00 AM PDT · August 3, 2025

Shortly after Hunter Lightman joined OpenAI as a researcher in 2022, he watched his colleagues launch ChatGPT, one of the fastest-growing products ever. Meanwhile, Lightman quietly worked on a team teaching OpenAI's models to solve high school math competitions.

Today that team, known as MathGen, is considered instrumental to OpenAI's industry-leading effort to create AI reasoning models: the core technology behind AI agents that can do tasks on a computer like a human would.

"We were trying to make the models better at mathematical reasoning, which at the time they weren't very good at," Lightman told TechCrunch, describing MathGen's early work.

OpenAI's models are far from perfect today — the company's latest AI systems [still hallucinate](#) and its agents [struggle with complex tasks](#).

But its state-of-the-art models have improved significantly on mathematical reasoning. One of OpenAI's models recently won a [gold medal](#) at the International Mathematical Olympiad (IMO), a math competition for the world's brightest high school students. OpenAI believes these reasoning capabilities will translate to other subjects, and ultimately power general-purpose agents that the company has always dreamed of building.

ChatGPT was a happy accident — a low-key research preview turned viral consumer business — but OpenAI's agents are the product of a yearslong, deliberate effort within the company.


"Eventually, you'll just ask the computer for what you need and it'll do all of these tasks for you," said OpenAI CEO Sam Altman at the company's [first developer conference](#) in 2023. "These capabilities are often talked about in the AI field as agents. The upsides of this are going to be tremendous."



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025



OPENAI CEO SAM ALTMAN SPEAKS DURING THE OPENAI DEVDAY EVENT ON NOVEMBER 06, 2023, IN SAN FRANCISCO, CALIFORNIA.

IMAGE CREDITS: JUSTIN SULLIVAN / GETTY IMAGES

Whether agents will meet Altman's vision remains to be seen, but OpenAI shocked the world with the release of [its first AI reasoning model, o1](#), in the fall of 2024. Less than a year later, the 21 foundational researchers behind that breakthrough are the most highly sought-after talent in Silicon Valley.

Mark Zuckerberg [recruited](#) five of the o1 researchers to work on Meta's new superintelligence-focused unit, offering some compensation packages north of \$100 million. One of them, Shengjia Zhao, was recently [named](#) chief scientist of Meta Superintelligence Labs.

The reinforcement learning renaissance

The rise of OpenAI's reasoning models and agents are tied to a machine learning training technique known as reinforcement learning (RL). RL provides feedback to an AI model on whether its choices were correct in simulated environments.

RL has been used for decades. For instance, in 2016, about a year after OpenAI was founded in 2015, an AI system created by Google DeepMind

using RL, [AlphaGo](#), gained global attention after beating a world champion in the board game Go.



SOUTH KOREAN PROFESSIONAL GO PLAYER LEE SE-DOL (R) PREPARES FOR HIS FOURTH MATCH AGAINST GOOGLE'S ARTIFICIAL INTELLIGENCE PROGRAM, ALPHAGO, DURING THE GOOGLE DEEPMIND CHALLENGE MATCH ON MARCH 13, 2016, IN SEOUL, SOUTH KOREA.

IMAGE CREDITS:GOOGLE / GETTY IMAGES

Around that time, one of OpenAI's first employees, Andrej Karpathy, began pondering how to leverage RL to create an AI agent that could use a computer. But it would take years for OpenAI to develop the necessary models and training techniques.

Advertisement

By 2018, OpenAI pioneered its first large language model in the GPT series, pretrained on massive amounts of internet data and large clusters of GPUs. GPT models excelled at text processing, eventually leading to ChatGPT, but struggled with basic math.

It took until 2023 for OpenAI to achieve a breakthrough, initially dubbed “Q*” and then “Strawberry,” by combining LLMs, RL, and a technique called “test-time computation.” The latter gave the models extra time and computing power to plan and work through problems, verifying its steps, before providing an answer.

This allowed OpenAI to introduce a new approach called “chain-of-thought” (CoT), which improved AI’s performance on math questions the models hadn’t seen before.

“I could see the model starting to reason,” said OpenAI researcher Ahmed El-Kishky. “It would notice mistakes and backtrack, it would get frustrated. It really felt like reading the thoughts of a person.”

Though individually these techniques weren’t novel, OpenAI uniquely combined them to create Strawberry, which directly led to the development of o1. OpenAI quickly identified that the planning and fact-checking abilities of AI reasoning models could be useful to power AI agents.

“We had solved a problem that I had been banging my head against for a couple of years,” said Lightman. “It was one of the most exciting moments of my research career.”

Scaling reasoning

With AI reasoning models, OpenAI determined it had two new axes that would allow it to improve AI models: using more computational power during the post-training of AI models, and giving AI models more time and processing power while answering a question.

“OpenAI, as a company, thinks a lot about not just the way things are, but the way things are going to scale,” said Lightman.

Advertisement

Shortly after the 2023 Strawberry breakthrough, OpenAI spun up an “Agents” team led by OpenAI researcher Daniel Selsam to make further progress on this new paradigm, two sources told TechCrunch. Although the team was called “Agents,” OpenAI didn’t initially differentiate between reasoning models and agents as we think of them today. The company just wanted to make AI systems capable of completing complex tasks.

Eventually, the work of Selsam’s Agents team became part of a larger project to develop the o1 reasoning model, with leaders including OpenAI co-founder Ilya Sutskever, chief research officer Mark Chen, and chief scientist Jakub Pachocki.

ILYA SUTSKEVER, COMPUTER SCIENTIST AND CO-FOUNDER AND FORMER CHIEF SCIENTIST OF OPENAI.
IMAGE CREDITS: JACK GUEZ / AFP / GETTY IMAGES

OpenAI would have to divert precious resources — mainly talent and GPUs — to create o1. Throughout OpenAI’s history, researchers have had to negotiate with company leaders to obtain resources; demonstrating breakthroughs was a surefire way to secure them.

“One of the core components of OpenAI is that everything in research is bottom up,” said Lightman. “When we showed the evidence [for o1], the company was like, ‘This makes sense, let’s push on it.’”


Some former employees say that the startup's mission to develop AGI was the key factor in achieving breakthroughs around AI reasoning models. By focusing on developing the smartest-possible AI models, rather than products, OpenAI was able to prioritize o1 above other efforts. That type of large investment in ideas wasn't always possible at competing AI labs.

The decision to try new training methods proved prescient. By late 2024, several leading AI labs started seeing [diminishing returns](#) on models created through traditional pretraining scaling. Today, much of the AI field's momentum comes from advances in reasoning models.

What does it mean for an AI to “reason”?

In many ways, the goal of AI research is to re-create human intelligence with computers. Since the launch of o1, ChatGPT's UX has been filled with more human-sounding features such as “thinking” and “reasoning.”

When asked whether OpenAI's models were truly reasoning, El-Kishky hedged, saying he thinks about the concept in terms of computer science.



Advertisement

“We’re teaching the model how to efficiently expend compute to get an answer. So if you define it that way, yes, it is reasoning,” said El-Kishky.

Lightman takes the approach of focusing on the model's results and not as much on the means or their relation to human brains.

IMAGE CREDITS:DEVIN COLDEWEY

“If the model is doing hard things, then it is doing whatever necessary approximation of reasoning it needs in order to do that,” said Lightman. “We can call it reasoning, because it looks like these reasoning traces, but it’s all just a proxy for trying to make AI tools that are really powerful and useful to a lot of people.”

OpenAI’s researchers note people may disagree with their nomenclature or definitions of reasoning — and surely, [critics have emerged](#) — but they argue it’s less important than the capabilities of their models. Other AI researchers tend to agree.


Nathan Lambert, an AI researcher with the non-profit AI2, compares AI reasoning modes to airplanes in a [blog post](#). Both, he says, are man-made systems inspired by nature — human reasoning and bird flight, respectively — but they operate through entirely different mechanisms. That doesn’t make them any less useful, or any less capable of achieving similar outcomes.

A group of AI researchers from OpenAI, Anthropic, and Google DeepMind agreed in a recent [position paper](#) that AI reasoning models are not well understood today, and more research is needed. It may be too early to confidently claim what exactly is going on inside them.

The next frontier: AI agents for subjective tasks

The AI agents on the market today work best for well-defined, verifiable domains such as coding. OpenAI's [Codex agent](#) aims to help software engineers offload simple coding tasks. Meanwhile, Anthropic's models have become particularly [popular](#) in AI coding tools like Cursor and Claude Code — these are some of the first AI agents that people are willing to [pay up for](#).

However, general-purpose AI agents like OpenAI's [ChatGPT agent](#) and Perplexity's [Comet](#) struggle with many of the complex, subjective tasks people want to automate. When trying to use these tools for online shopping or finding a long-term parking spot, I've found the agents take longer than I'd like and make [silly mistakes](#).



Advertisement

Agents are, of course, early systems that will undoubtedly improve. But researchers must first figure out how to better train the underlying models to complete tasks that are more subjective.

IMAGE CREDITS: JONATHAN RAA/NURPHOTO / GETTY IMAGES

“Like many problems in machine learning, it’s a data problem,” said Lightman when asked about the limitations of agents on subjective tasks. “Some of the research I’m really excited about right now is figuring out how to train on less verifiable tasks. We have some leads on how to do these things.”

Noam Brown, an OpenAI researcher who helped create the IMO model and o1, told TechCrunch that OpenAI has new general-purpose RL techniques that allow them to teach AI models skills that aren’t easily verified. This was how the company built the model that achieved a gold medal at IMO, he said.

OpenAI’s IMO model was a newer AI system that spawns multiple agents, which then simultaneously explore several ideas and then choose the best possible answer. These types of AI models are becoming more popular; [Google](#) and [xAI](#) have recently released state-of-the-art models using this technique.

“I think these models will become more capable at math, and I think they’ll get more capable in other reasoning areas as well,” said Brown. “The progress has been incredibly fast. I don’t see any reason to think it will slow down.”

These techniques may help OpenAI's models become more performant, gains that could show up in the company's upcoming GPT-5 model. OpenAI hopes to assert its dominance over competitors with the launch of GPT-5, ideally offering the [best AI model](#) to power agents for developers and consumers.

But the company also wants to make its products simpler to use. El-Kishky says OpenAI wants to develop AI agents that intuitively understand what users want, without requiring them to select specific settings. He says OpenAI aims to build AI systems that understand when to call up certain tools, and how long to reason for.

These ideas paint a picture of an ultimate version of ChatGPT: an agent that can do anything on the internet for you, and understand how you want it to be done. That's a much different product than what ChatGPT is today, but the company's research is squarely headed in this direction.

Advertisement

While OpenAI undoubtedly led the AI industry a few years ago, the company now faces a tranche of worthy opponents. The question is no longer just whether OpenAI can deliver its agentic future, but can the company do so before Google, Anthropic, xAI, or Meta beat them to it?

Topics: [agents](#) [AI](#) [ChatGPT](#) [Exclusive](#) [OpenAI](#)

Advertisement



Maxwell Zeff

Senior AI Reporter | [X](#)

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg, and MSNBC, Zeff has covered the rise of AI and the Silicon Valley Bank crisis. He is based in San Francisco. Wh...

[View Bio](#) >

Loading the next article