



AI



## OpenAI launches Study Mode in ChatGPT

OpenAI announced Tuesday the launch of Study Mode, a new feature within ChatGPT that aims to help students develop their own critical thinking skills, rather than simply obtain answers to questions. With Study Mode enabled, ChatGPT will ask users questions to test their understanding and, in some cases, refuse to offer direct answers unless students engage with the material.

OpenAI says Study Mode is rolling out to logged in users on ChatGPT's Free, Plus, Pro, and Team plans starting Tuesday. The company expects to introduce Study Mode to its Edu subscribers, which largely consists of young people whose school administrator's have purchased a plan for the entire student body, in the coming weeks.

Study Mode is OpenAI's attempt to address [the millions of students who use ChatGPT in school](#). Studies have shown that using [ChatGPT can be a helpful tutor](#) for young people, but it also [may harm their critical thinking skills](#). A research paper released in June found that [people who use ChatGPT to write essays exhibit lower brain activity during the process](#) compared to those who use Google Search or nothing at all.

When ChatGPT first launched in 2022, its widespread use in school settings sparked fear among educators, leading to generative AI bans in many American school districts. By 2023, [some of those schools repealed their ChatGPT bans](#), and teachers around the country came to terms with the fact that ChatGPT would be a part of young people's lives from now on.

Now with the launch of Study Mode, OpenAI hopes to improve ChatGPT as a learning tool, and not just an answer engine. Anthropic launched a similar tool for its AI chatbot Claude, [called Learning Mode](#), in April.

Of course, there are limitations to how effective Study Mode truly is. Students can easily switch into the regular mode of ChatGPT if they just want an answer to a question. OpenAI's VP of Education, Leah Belsky, told TechCrunch in a briefing that the company is not offering tools for parents or administrators to lock students into Study Mode. However,

## DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

## Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines
- The best dating apps aren't even dating apps

Belsky said OpenAI may explore administrative or parental controls in the future.

That means it will take a committed student to use Study Mode — the kids have to really want to learn, not just finish their assignment.



## Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

OpenAI says Study Mode is the company's first step to improving learning in ChatGPT and aims to publish more information in the future about how students use generative AI throughout their education.

Topics: [AI](#) [ChatGPT](#) [Education](#) [OpenAI](#)



**Maxwell Zeff**

Senior AI Reporter |

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg, and MSNBC, Zeff has covered th...

[View Bio](#) >

KEEP READING



IMAGE CREDITS:

JAAP ARRIENS/NURPHOTO / GETTY IMAGES

AI



## OpenAI launches a pair of AI reasoning models, o3 and o4-mini

Maxwell Zeff — 10:00 AM PDT · April 16, 2025

OpenAI announced on Wednesday the launch of o3 and o4-mini, new AI reasoning models designed to pause and work through questions before responding.

The company calls o3 its most advanced reasoning model ever, outperforming the company's previous models on tests measuring math, coding, reasoning, science, and visual understanding capabilities. Meanwhile, o4-mini offers what OpenAI says is a competitive trade-off between price, speed, and performance — three factors developers often consider when choosing an AI model to power their applications.

Unlike previous reasoning models, o3 and o4-mini can generate responses using tools in ChatGPT such as web browsing, Python code execution, image processing, and image generation. Starting today, the models, plus a variant of o4-mini called "o4-mini-high" that spends more time crafting answers to improve its reliability, are available for subscribers to OpenAI's Pro, Plus, and Team plans.

Advertisement

Newsletters

See More ↗

Subscribe for the industry's biggest tech news

The new models are part of OpenAI’s effort to beat out Google, Meta, xAI, Anthropic, and DeepSeek in the cutthroat global AI race. While OpenAI was first to release an AI reasoning model, o1, competitors quickly followed with versions of their own that match or exceed the performance of OpenAI’s lineup. In fact, reasoning models have begun to dominate the field as AI labs look to eke more performance out of their systems.

O3 nearly wasn’t released in ChatGPT. OpenAI CEO Sam Altman signaled in February that the company intended to devote more resources to a sophisticated alternative that incorporated o3’s technology. But competitive pressure seemingly spurred OpenAI to reverse course in the end.

OpenAI says that o3 achieves state-of-the-art performance on SWE-bench verified (without custom scaffolding), a test measuring coding abilities, scoring 69.1%. The o4-mini model achieves similar performance, scoring 68.1%. OpenAI’s next best model, o3-mini, scored 49.3% on the test, while Claude 3.7 Sonnet scored 62.3%.

OpenAI claims that o3 and o4-mini are its first models that can “think with images.” In practice, users can upload images to ChatGPT, such as whiteboard sketches or diagrams from PDFs, and the models will analyze the images during their “chain-of-thought” phase before answering. Thanks to this newfound ability, o3 and o4-mini can understand blurry and low-quality images and can perform tasks such as zooming or rotating images as they reason.

**TechCrunch Daily News** 

Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

**Startups Weekly** 

Startups are the core of TechCrunch, so get our best coverage delivered weekly.

**TechCrunch Week in Review** 

Get the best of our coverage, delivered to your inbox every Saturday.

**TechCrunch Mobility** 


TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).



Advertisement



## Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Beyond image-processing capabilities, o3 and o4-mini can run and execute Python code directly in your browser via ChatGPT’s Canvas feature, and search the web when asked about current events.


In addition to ChatGPT, all three models — o3, o4-mini, and o4-mini-high — will be available via OpenAI’s developer-facing endpoints, the Chat Completions API and Responses API, allowing engineers to build applications with the company’s models at usage-based rates.

OpenAI is charging developers a relatively low price for o3, given its improved performance, at \$10 per million input tokens (roughly 750,000 words, longer than the Lord of the Rings series) and \$40 per million output tokens. For o4-mini, OpenAI is charging the same as o3-mini, \$1.10 per million input tokens and \$4.40 per million output tokens.

In the coming weeks, OpenAI says it plans to release o3-pro, a version of o3 that uses more computing resources to produce its answers, exclusively for ChatGPT Pro subscribers.

OpenAI CEO Sam Altman has indicated o3 and o4-mini may be its last stand-alone AI reasoning models in ChatGPT before GPT-5, a model that the company has said will unify traditional models like GPT-4.1 with its reasoning models.

Topics: [AI](#) [ChatGPT](#) [OpenAI](#) [reasoning models](#)



[f](#) [X](#) [in](#) [📷](#) [✉](#) [🔗](#)



**Maxwell Zeff**

Senior AI Reporter | [X](#)



Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg, and MSNBC, Zeff has covered th...

View Bio >

KEEP READING



IMAGE CREDITS:

SILAS STEIN / PICTURE ALLIANCE / GETTY IMAGES

AI



# ChatGPT’s deep research tool gets a GitHub connector to answer questions about code


Kyle Wiggers — 1:27 PM PDT · May 8, 2025

OpenAI is enhancing its AI-powered “deep research” feature with the ability to analyze codebases on GitHub.

On Thursday, OpenAI announced what it’s calling the first “connector” for [ChatGPT deep research](#), the company’s tool that searches across the web and other sources to compile thorough research reports on a topic. Now, ChatGPT deep research can link to GitHub (in beta), allowing developers to ask questions about a codebase and engineering documents.

Advertisement

The connector will be available for ChatGPT Plus, Pro, and Team users over the next few days, with Enterprise and Edu support coming soon, according to an OpenAI spokesperson.



OPENAI'S CHATGPT DEEP RESEARCH FEATURE CAN NOW CONNECT TO GITHUB  
IMAGE CREDITS: OPENAI

The GitHub connector for ChatGPT deep research arrives as AI companies look to make their AI-powered chatbots more useful by building ways to link them to outside platforms and services. Anthropic, for example, recently debuted [Integrations](#), which gives apps a pipeline into its AI chatbot Claude.

OpenAI [years ago](#) offered a plug-in capability for ChatGPT, but deprecated it in favor of custom chatbots called [GPTs](#).

“I often hear that users find ChatGPT’s deep research agent so valuable that they want it to connect to their internal sources, in addition to the web,” OpenAI Head of Business Products Nate Gonzalez wrote in a [blog post](#) on LinkedIn. “[That’s why] today we’re introducing our first connector.”

In addition to answering questions about codebases, the new ChatGPT deep research GitHub connector lets ChatGPT users break down product specs into technical tasks and dependencies, summarize code structure and patterns, and understand how to implement new APIs using real code examples.

## Latest in AI

[See More](#)

AI



Google denies AI search features are killing website traffic

Sarah Perez - 12 hours ago

AI



First impressions of Alexa+, Amazon’s upgraded, AI-powered digital assistant

Sarah Perez - 13 hours ago

AI



Google takes on ChatGPT’s Study Mode with new ‘Guided Learning’ tool in Gemini

Aisha Malik - 13 hours ago

Advertisement

TC

**Tech and VC heavyweights join the Disrupt 2025 agenda**

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights



that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

There's a risk that ChatGPT deep research hallucinates, of course — no AI model in existence doesn't [confidently make things up sometimes](#). But OpenAI is pitching the new capability as a potential time saver, not a replacement for experts.

An OpenAI spokesperson said ChatGPT will respect an organization's settings so users only see GitHub content they're already allowed to view and codebases that've been explicitly shared with ChatGPT.

OpenAI has been [investing](#) in its [tooling](#) for assistive coding, recently unveiling an open source coding tool for terminals called Codex CLI and upgrading the ChatGPT desktop app to read code in a handful of developer-focused coding apps. The company sees programming as a top use case for its models. [Case in point](#), OpenAI has reportedly reached an agreement to buy AI-powered coding assistant Windsurf for \$3 billion.

In other OpenAI news on Thursday, the company launched fine-tuning options for developers looking to customize its newer models for particular applications. Devs can now fine-tune OpenAI's o4-mini "reasoning" model via a technique OpenAI calls reinforcement fine-tuning, which uses task-specific grading to improve the model's performance. Fine-tuning has also rolled out for the company's GPT-4.1 nano model.

Only verified organizations can fine-tune o4-mini, according to OpenAI. GPT-4.1 nano fine-tuning, meanwhile, is available for all paying developers.

OpenAI [began gating](#) certain models and developer features behind verification, which requires organizations to submit an ID and other identity documents, in April. The company claims that it's necessary to prevent abuse.



[f](#) [X](#) [in](#) [🐼](#) [✉](#) [🔗](#)



**Kyle Wiggers**

AI Editor | [X](#) [🐼](#) [@](#) [in](#)

Kyle Wiggers was TechCrunch’s AI Editor until June 2025. His writing has appeared in VentureBeat and Digital Trends, as well as a range of...

[View Bio](#) >

———— KEEP READING ————

AI [f](#) [X](#) [in](#) [🐼](#) [✉](#) [🔗](#)

The hottest AI models, what they do, and how to use them



AI models are being cranked out at a dizzying pace, by everyone from Big Tech companies like Google to startups like OpenAI and Anthropic. Keeping track of the latest ones can be overwhelming.

Adding to the confusion is that AI models are often promoted based on industry benchmarks. But these [technical metrics often reveal little](#) about how real people and companies actually use them.

To cut through the noise, TechCrunch has compiled an overview of the most advanced AI models released since 2024, with details on how to use them and what they're best for. We'll keep this list updated with the latest launches, too.

There are literally over a million AI models out there: Hugging Face, for example, [hosts over 1.4 million](#). So this list might miss some models that perform better, in one way or another.

# AI models released in 2025

## Google Gemini 2.5

Gemini 2.5 Pro Experimental, a reasoning model, excels at [building web apps and code agents](#) according to Google. However, it underperforms on one popular coding benchmark compared to Claude Sonnet 3.7. The model requires a \$20 monthly Gemini Advanced subscription.

## ChatGPT-4o image generator

OpenAI [has upgraded its existing GPT-4o model](#) to generate images, not just text. The souped-up model soon went viral for [transforming images into Studio Ghibli-style anime](#), despite obvious copyright concerns. Accessing GPT-4o requires, at minimum, a \$20 per month ChatGPT Plus subscription.

Advertisement

### TEC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

## Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices

# Stability AI's Stable Virtual Camera

Image-generation startup Stability AI [has launched a model](#) that the company says can generate 3D scenes and camera angles from a single 2D image. However, it still struggles with scenes featuring more complex elements like humans and moving water. The model is available for noncommercial research use on HuggingFace.

TC

## Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

- The uproar over Vogue's AI-generated ad isn't just about fashion
- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

Advertisement

# Cohere's Aya Vision

Cohere [released a multimodal model](#) called Aya Vision that it claims is best in class at doing things like captioning images and answering questions about photos. It also excels in languages other than English, unlike other models, Cohere claims. It is available for [free on WhatsApp](#).

# OpenAI's GPT 4.5 "Orion"

OpenAI calls Orion their [largest model to date](#), touting its strong "world knowledge" and "emotional intelligence." However, it underperforms on certain benchmarks compared to newer reasoning models. Orion is available to subscribers of OpenAI's \$200-per-month plan.

# Claude Sonnet 3.7

Anthropic says this is the [industry's first "hybrid" reasoning model](#), because it can both fire off quick answers and really think things through when needed. It also gives users control over how long the model can think for, [per Anthropic](#). Sonnet 3.7 is available to all Claude users, but heavier users will need a \$20-per-month Pro plan.

## xAI's Grok 3

Grok 3 is the [latest flagship model](#) from Elon Musk-founded startup xAI. It's [claimed to](#) outperform other leading models on math, science, and coding. The model requires X Premium (which is \$50 per month.) After one study [found](#) Grok 2 leaned left, Musk [pledged](#) to shift Grok more "politically neutral" but it's not yet clear if that's been achieved.

## OpenAI o3-mini

This is OpenAI's [latest reasoning model](#) and is optimized for STEM-related tasks like coding, math, and science. It's [not OpenAI's most powerful](#) model but because it's smaller, the company [says](#) it's significantly lower cost. It is available for free but requires a subscription for heavy users.

## OpenAI Deep Research

OpenAI's Deep Research is [designed for doing in-depth research](#) on a topic with clear citations. This service is only available with ChatGPT's [\\$200-per-month Pro subscription](#). OpenAI [recommends it](#) for everything from science to shopping research, but beware that [hallucinations remain a problem](#) for AI.

## Mistral Le Chat

Mistral has [launched app versions of Le Chat](#), a multimodal AI personal assistant. Mistral [claims](#) Le Chat responds faster than any other chatbot. It also has a paid version with [up-to-date journalism](#) from the AFP. [Tests from Le Monde](#) found Le Chat's performance impressive, although it made more errors than ChatGPT.

## OpenAI Operator



who need to quickly process massive chunks of text. The service requires (at minimum) a Google One AI Premium subscription of \$19.99 a month.

## AI models released in 2024

### DeepSeek R1

This [Chinese AI model](#) took [Silicon Valley by storm](#). DeepSeek's R1 performs well on coding and math, while its open source nature means anyone can run it locally. Plus, it's free. However, [R1 integrates](#) Chinese government censorship and [faces rising bans](#) for potentially sending user data back to China.

### Gemini Deep Research

Deep Research [summarizes Google's search results](#) in a simple and well-cited document. [The service](#) is helpful for students and anyone else who needs a quick research summary. However, its quality isn't nearly as good as an actual peer-reviewed paper. Deep Research requires a \$19.99 Google One AI Premium subscription.

### Meta Llama 3.3 70B

This is the [newest and most advanced version](#) of Meta's open source Llama AI models. Meta has touted [this version](#) as its cheapest and most efficient yet, especially for math, general knowledge, and instruction following. It is free and open source.



## OpenAI Sora

Sora is a model that [creates realistic videos](#) based on text. While it can generate entire scenes rather than just clips, [OpenAI admits](#) that it often generates “unrealistic physics.” It’s currently only available on paid versions of ChatGPT, starting with Plus, which is \$20 a month.

## Alibaba Qwen QwQ-32B-Preview

This model is [one of the few to rival](#) OpenAI’s o1 on certain industry benchmarks, excelling in math and coding. Ironically for a “reasoning model,” it has “room for improvement in common sense reasoning,” [Alibaba says](#). It also incorporates Chinese government censorship, [TechCrunch testing shows](#). It’s free and open source.

## Anthropic’s Computer Use

Claude’s Computer Use is meant to [take control of your computer](#) to complete tasks like coding or booking a plane ticket, making it a predecessor of OpenAI’s Operator. Computer use, however, [remains in beta](#). Pricing is via API: \$0.80 per million tokens of input and \$4 per million tokens of output.

## xAI’s Grok 2

Elon Musk’s AI company, xAI, has launched an [enhanced version of its flagship Grok 2 chatbot](#) it [claims](#) is “three times faster.” Free users are limited to 10 questions every two hours on Grok, while subscribers to X’s Premium and Premium+ plans enjoy higher usage limits. xAI also launched an image generator, Aurora, that [produces highly photorealistic images](#), including some graphic or violent content.

## OpenAI o1

[OpenAI’s o1 family](#) is meant to produce better answers by “thinking” through responses through a hidden [reasoning feature](#). The model excels at coding, math, and safety, [OpenAI claims](#), but has [issues with trying to deceive humans, too](#). Using o1 requires subscribing to ChatGPT Plus, which is \$20 a month.

# Anthropic's Claude Sonnet 3.5

Claude Sonnet 3.5 is a model Anthropic [claims as being best in class](#). It's become known for its coding capabilities and is considered a tech insider's [chatbot of choice](#). The model can be accessed for free on Claude, although heavy users will need a \$20 monthly Pro subscription. While it can understand images, it can't generate them.

## OpenAI GPT 4o-mini

OpenAI has [touted GPT 4o-mini](#) as its most affordable and fastest model yet, thanks to its small size. [It's meant](#) to enable a broad range of tasks like powering customer service chatbots. The model is available on ChatGPT's free tier. It's better suited for high-volume simple tasks compared to more complex ones.

## Cohere Command R+

Cohere's [Command R+ model](#) excels at complex retrieval-augmented generation (or RAG) applications for enterprises. That means it can find and cite specific pieces of information really well. (The inventor of RAG [actually works at Cohere](#).) Still, RAG [doesn't fully solve AI's hallucination problem](#).

Topics: [AI](#) [AI](#) [ai models](#) [Anthropic](#) [ChatGPT](#) [deepseek](#)

[evergreens](#) [google gemini](#) [mistral](#) [OpenAI](#) [sora](#) [xAI](#)





**Charles Rollet**

Senior Reporter | X

Charles Rollet is a senior reporter at TechCrunch. His investigative reporting has led to U.S. government sanctions against four tech...

[View Bio](#) >

KEEP READING



IMAGE CREDITS: BRYCE DURBIN / TECHCRUNCH

AI



## OpenAI launches a general purpose agent in ChatGPT

Maxwell Zeff — 10:00 AM PDT · July 17, 2025

OpenAI is launching a new general purpose AI agent in ChatGPT, which the company says can complete a wide variety of computer-based tasks on behalf of users. OpenAI says the agent can automatically navigate a user's calendar, generate editable presentations and slideshows, and run code.

The tool, called ChatGPT agent, combines several capabilities from OpenAI's previous agentic tools, including [Operator's](#)

Advertisement

[ability to click around on websites](#), as well as [Deep Research's ability to synthesize information](#) from dozens of websites into a concise research report. OpenAI says users will be able to interact with the agent simply by prompting ChatGPT in natural language.

ChatGPT agent is rolling out on Thursday to subscribers to OpenAI's Pro, Plus, and Team plans. To activate the tool, users can select "agent mode" in ChatGPT's dropdown menu of tools.

ChatGPT can now do work for you using its own computer.

Introducing ChatGPT agent—a unified agentic system combining Operator's action-taking remote browser, deep research's web synthesis, and ChatGPT's conversational strengths. [pic.twitter.com/7uN2Nc6nBQ](https://pic.twitter.com/7uN2Nc6nBQ)

— OpenAI (@OpenAI) [July 17, 2025](#)

The launch of ChatGPT agent represents OpenAI's boldest attempt yet to turn ChatGPT into an agentic product that can take actions and offload tasks for users, rather than just answering questions. In recent years, Silicon Valley companies including OpenAI, Google, and Perplexity have unveiled dozens of AI agents that have promised to do just that. However, these [early version AI agents have proven to struggle with complex tasks](#), and they seem less compelling as products than the ultimate vision tech executives pitch around AI agents.

That said, OpenAI says ChatGPT agent is far more capable than its previous offerings.

The company's new agent can access [ChatGPT connectors](#), allowing users to connect apps like Gmail and GitHub so that the agent can find relevant information to your prompts. OpenAI says ChatGPT agent has access to a terminal, and it can use APIs to access certain apps.

OpenAI suggests that users can tap ChatGPT agent to "plan and buy ingredients to make Japanese breakfast for four," as well as "analyze three competitors and create a slide deck." These kinds of capabilities requires ChatGPT agent to parse through websites, plan a course of action, and use tools — much more complicated tasks than OpenAI has previously tried to tackle with agents.

## Newsletters

[See More ↗](#)

Subscribe for the industry's biggest tech news

### TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch's coverage.

### Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

### TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

### TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

Advertisement



## Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

The model underlying ChatGPT agent offers state-of-the-art performance on several benchmarks, according to OpenAI.

The company says the ChatGPT agent model scores 41.6% on Humanity's Last Exam (pass@1), a difficult test made up of thousands of questions across more than one hundred subjects. That's roughly double what OpenAI's o3 and o4-mini scored on the test.

On FrontierMath, one of the hardest known math benchmarks, OpenAI says ChatGPT agent scores 27.4% when it has access to tools, such as a terminal for code execution. The previous state-of-the-art score comes from o4-mini, which scored just 6.3%.

OpenAI notes that it developed ChatGPT agent with safety in mind, largely because the product presents some newfound capabilities that could make it more dangerous in the hands of a bad actor. OpenAI has previously warned that agentic models could present more dangerous capabilities.

In a [safety report](#) for ChatGPT agent, OpenAI says it's designated the model as "high capability" in biological and chemical weapon domains, which is defined in OpenAI's Preparedness Framework as a model with the ability to "amplify existing pathways to severe harm." OpenAI notes that it does not have direct evidence of this, but it's decided to take a precautionary approach and activate new safeguards to mitigate these risks.

The new safeguards for ChatGPT agent include a monitor that works in real time as users interact with the product. OpenAI says it runs a classifier across every prompt entered into ChatGPT agent, determining whether the request is related to biology. If so, OpenAI runs ChatGPT agent's response through a second monitor that determines whether the content could be used to evoke a biological threat.

OpenAI also says it disabled ChatGPT's memory feature for this agent to prevent misuse. In other parts of ChatGPT, OpenAI's memory feature allows the chatbot to reference information from previous user chats. However, OpenAI says bad actors could use the feature in ChatGPT agent to exfiltrate sensitive data through prompt injection attacks. The company says it may revisit adding the feature in the future, however.

While ChatGPT agent sounds impressive, it remains to be seen how capable it truly is in the real world. Until now, agent technology has proven relatively brittle when interacting with the real world. That said, OpenAI says it's developed a more capable model that's able to deliver on the promise of AI agents.

*This story was updated with more information.*

Topics: [agents](#) [AI](#) [ChatGPT](#) [OpenAI](#)

Advertisement



**Maxwell Zeff**

Senior AI Reporter | [X](#)



Maxwell Zeff is a senior reporter at TechCrunch specializing in AI.  
Previously with Gizmodo, Bloomberg, and MSNBC, Zeff has covered th...

[View Bio](#) >

Loading the next article



[TechCrunch](#)

[Staff](#)

[Contact Us](#)

[Advertise](#)

[Crunchboard  
Jobs](#)

[Site Map](#)

[Terms of  
Service](#)

[Privacy Policy](#)

[RSS Terms of  
Use](#)

[OpenAI](#)

[Perplexity](#)

[Cisco](#)

[DeepMind](#)

[Windsurf](#)

[Tech Layoffs](#)

[ChatGPT](#)