


Advertisement

AI



**OpenAI says ChatGPT is
on track to reach 700M
weekly users**



ChatGPT's impressive growth as a consumer app continues as the chatbot is on track to hit 700 million weekly active users this week, the company says.

The app had earlier reached 500 million weekly active users as of the end of March, noted Nick Turley, OpenAI VP and head of ChatGPT's app, [in a post on X](#). He also said the app has grown 4x since last year.

"Every day people and teams are learning, creating, and solving harder problems. Big week ahead. Grateful to the team for making ChatGPT more useful and delivering on our mission so everyone can benefit from AI," he posted.

The app's popularity increased after OpenAI launched an [upgraded image-generation feature](#), powered by the GPT-4 model, in March. In early April, the company's COO, Brad Lightcap, said that more than [130 million users had created over 700 million images](#) in just a few days after the launch.

The company has also seen an increase in subscribers. Last week, Lightcap said that ChatGPT had [5 million paying business users](#), up from 3 million in June.

In a recent report, market intelligence firm Sensor Tower noted that users are using [ChatGPT for more than 12 days a month on average](#), only behind


TEC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Instagram takes on Snapchat with new 'Instagram Map'
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices

Google and X. The report also said that in H1 2025, users spent an average of 16 minutes per day on the app.

Topics: [AI](#) [Apps](#) [ChatGPT](#) [OpenAI](#)



Ivan Mehta

[X](#) [in](#)

Ivan covers global consumer tech developments at TechCrunch. He is based out of India and has previously...

[View Bio](#) >


KEEP READING

- The uproar over Vogue's AI-generated ad isn't just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines

AI



OpenAI is practically giving ChatGPT to the government for free



OpenAI is poised to undercut rivals like Anthropic and Google in the race to see its AI tools integrated into federal agency workflows.

The AI giant has [reached an agreement](#) with the U.S. General Services Administration (GSA), the government's central purchasing arm, to offer ChatGPT Enterprise to participating federal agencies for just \$1 per agency for the next year.

The partnership comes a day after the [GSA added OpenAI, Google, and Anthropic](#) to a list of approved AI vendors that can offer their services to civilian federal agencies. The companies will see their tools offered via the Multiple Award Schedule (MAS), a federal contracting platform that lets government agencies access AI tools via pre-negotiated contracts so they don't have to negotiate with vendors individually.

IMAGE CREDITS: GETTY IMAGES

It's not clear whether other AI firms are going to offer their services at such a discounted rate, though GSA Federal Acquisition Service commissioner Josh Gruenbaum in a [statement](#) encouraged "other American AI technology companies to follow OpenAI's lead and work with us."

TechCrunch has reached out to Anthropic and Google for more information.

In addition to access to ChatGPT Enterprise, OpenAI is offering unlimited use of advanced models for an additional 60 days. Federal employees will also have access to a new government user community

Advertisement

Newsletters

[See More ↗](#)

Subscribe for the industry's biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch's coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

and tailored introductory training resources to get familiar with OpenAI’s tools.

Data security is a top priority for government agencies that are concerned about sensitive information being leaked into model training sets. TechCrunch has asked the GSA for details on how government data is being safeguarded, including whether measures like on-premises or private cloud deployments are being used to enhance security.

TC

Techcrunch event
Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise.

San Francisco | October 27-29, 2025

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).



“The government is taking a cautious, security-first approach to AI,” a GSA spokesperson told TechCrunch. “This ensures sensitive information remains protected while enabling agencies to benefit from AI-driven efficiencies.”

OpenAI’s discount comes a couple of weeks after the Trump administration published its AI Action Plan that seeks to boost data center buildouts and integrate more AI tools into government, among other things. It also follows a Trump executive order that [bans “woke AI”](#) and AI models that aren’t “ideologically neutral” from government contracts.

TechCrunch has reached out to OpenAI to learn more about its GSA partnership and how it will approach Trump’s executive order.

Topics: [AI](#) [ChatGPT Enterprise](#) [Government & Policy](#)

[GSA](#) [OpenAI](#)



Rebecca Bellan

Senior Reporter | [X](#) [🦋](#)

Rebecca Bellan is a senior reporter at TechCrunch where she covers the business, policy, and emerging trends...

[View Bio](#) >

KEEP READING



AI



Google takes on ChatGPT's Study Mode with new 'Guided Learning' tool in Gemini

Aisha Malik — 11:18 AM PDT · August 6, 2025

IMAGE CREDITS:


MATTHIAS BALK/PICTURE ALLIANCE / GETTY IMAGES

As the new school year approaches, Google

Latest Security
Startups AI
Venture Apps
Apple



Events
Podcasts
Newsletters



because they spit out direct answers. The new tools from Google and OpenAI likely aim to address these concerns by positioning their chatbots as learning tools rather than simple answer engines.

With Guided Learning, Gemini will break down problems step-by-step and adapt explanations to its users' needs. The feature responds using images, diagrams, videos, and interactive quizzes to help users build and test themselves on their knowledge, rather than simply giving them the answer.




IMAGE CREDITS:GOOGLE

Google says the feature will help users uncover the “why” and “how” behind concepts.

“Whether you’re preparing for an exam about enzymes, starting the first draft of a paper on the importance of bee populations in supporting our food systems, or exploring your passion for photography, Guided Learning is a collaborative thinking partner that helps you get it — each step of the way,” wrote Maureen Heymans, Google’s VP of Learning & Sustainability, in the [blog post](#).

In addition to the new feature, Google announced that it’s working to make Gemini as a whole better equipped to help users learn. Gemini will now

AI



Truth Social’s AI search is powered by Perplexity, but the platform can set limits on sources

Rebecca Bellan - 2 hours ago

automatically incorporate images, diagrams, and YouTube videos directly into responses to help users better understand complex topics. Plus, users can now also ask Gemini to create flashcards and study guides based on their quiz results or other class materials.

TC

Techcrunch event

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

Google also announced on Wednesday that it's offering students in the U.S., Japan, Indonesia,

Korea, and Brazil a free one-year subscription to Google's AI Pro plan.

The plan includes expanded access to [Gemini 2.5 Pro](#), [NotebookLM](#), [Veo 3](#), [Deep Research](#), and more.

Topics: [AI](#) [Apps](#) [artificial intelligence](#) [ChatGPT](#)

[gemini](#) [Google](#)

[f](#) [X](#) [in](#) [reddit](#) [email](#) [link](#)



Aisha Malik

Consumer News Reporter | [X](#) [reddit](#) [in](#)

Aisha is a consumer news reporter at TechCrunch. Prior to joining the publication in 2021, she was a telecom reporte...

[View Bio](#) >

KEEP READING

AI

[f](#) [X](#) [in](#) [reddit](#) [email](#) [link](#)

OpenAI launches two
'open' AI reasoning
models

OpenAI announced Tuesday the launch of two open-weight AI reasoning models with similar capabilities [to its o-series](#). Both are freely available to download from the online developer platform [Hugging Face](#), the company said, describing the models as “state of the art” when measured across several benchmarks for comparing open models.

The models come in two sizes: a larger and more capable gpt-oss-120b model that can run on a single Nvidia GPU, and a lighter-weight gpt-oss-20b model that can run on a consumer laptop with 16GB of memory.

IMAGE CREDITS:

JAKUB PORZYCKI/NURPHOTO / GETTY IMAGES

The launch marks OpenAI’s first ‘open’ language model since [GPT-2](#), which was released more than five years ago.

In a briefing, OpenAI said its open models will be capable of sending complex queries to [AI models in the cloud](#), as TechCrunch previously reported. That means if OpenAI’s open model is not capable of a certain task, such as processing an image, developers can connect the open model to one of the company’s more capable closed models.

While OpenAI open sourced AI models in its early days, the company has generally favored a proprietary, closed source development approach. The latter strategy has helped OpenAI build a large

TECH DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government’s website
- Instagram takes on Snapchat with new ‘Instagram Map’
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices

business selling access to its AI models via an API to enterprises and developers.

However, CEO Sam Altman said in January he believes OpenAI has been “on the wrong side of history” when it comes to open sourcing its technologies. The company today faces growing pressure from Chinese AI labs — including DeepSeek, Alibaba’s Qwen, and Moonshot AI — which have developed several of the world’s most capable and popular open models. (While Meta previously dominated the open AI space, the company’s Llama AI models have [fallen behind](#) in the last year.)

In July, the Trump administration also urged U.S. AI developers to [open source more technology](#) to promote global adoption of AI aligned with American values.

TC

Techcrunch event

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th

- The uproar over Vogue’s AI-generated ad isn’t just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines

anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

With the release of gpt-oss, OpenAI hopes to curry favor with developers and the Trump administration alike, both of which have watched the Chinese AI labs rise to prominence in the open source space.

“Going back to when we started in 2015, OpenAI’s mission is to ensure AGI that benefits all of humanity,” said Altman in a statement shared with TechCrunch. “To that end, we are excited for the world to be building on an open AI stack created in the United States, based on democratic values, available for free to all and for wide benefit.”





IMAGE CREDITS:TOMOHIRO OHSUMI / GETTY IMAGES

How the models performed

OpenAI aimed to make its [open model a leader](#) among other open-weight AI models, and the company claims to have done just that.

On Codeforces (with tools), a competitive coding test, gpt-oss-120b and gpt-oss-20b score 2622 and 2516, respectively, outperforming DeepSeek’s R1 while underperforming o3 and o4-mini.



OPENAI’S OPEN MODEL PERFORMANCE ON CODEFORCES.

IMAGE CREDITS:OPENAI

On Humanity’s Last Exam (HLE), a challenging test of crowdsourced questions across a variety of subjects (with tools), gpt-oss-120b and gpt-oss-20b score 19% and 17.3%, respectively. Similarly, this underperforms o3 but outperforms leading open models from DeepSeek and Qwen.

OPENAI’S OPEN MODEL PERFORMANCE ON HLE.

IMAGE CREDITS:OPENAI

Notably, OpenAI’s open models hallucinate significantly more than its latest AI reasoning

models, o3 and o4-mini.

Hallucinations have been [getting more severe](#) in OpenAI's latest AI reasoning models, and the company previously said it doesn't quite understand why. In a white paper, OpenAI says this is "expected, as smaller models have less world knowledge than larger frontier models and tend to hallucinate more."

OpenAI found that gpt-oss-120b and gpt-oss-20b hallucinated in response to 49% and 53%, respectively, of questions on PersonQA, the company's in-house benchmark for measuring the accuracy of a model's knowledge about people. That's more than triple the hallucination rate of OpenAI's o1 model, which scored 16%, and higher than its o4-mini model, which scored 36%.

Training the new models

OpenAI says its open models were trained with similar processes to its proprietary models. The company says each open model leverages mixture-of-experts (MoE) to tap fewer parameters for any given question, making it run more efficiently. For gpt-oss-120b, which has 117 billion total parameters, OpenAI says the model only activates 5.1 billion parameters per token.

The company also says its open model was trained using high-compute [reinforcement learning \(RL\)](#) — a post-training process to teach AI models right from wrong in simulated environments using large clusters of Nvidia GPUs. This was also used to train OpenAI's o-series of models, and the open models have a similar chain-of-thought process in which they take additional time and computational resources to work through their answers.

As a result of the post-training process, OpenAI says its open AI models excel at powering AI agents and are capable of calling tools such as web search or Python code execution as part of its chain-of-thought process. However, OpenAI says its open models are text-only, meaning they will not be able to process or generate images and audio like the company's other models.

OpenAI is releasing gpt-oss-120b and gpt-oss-20b under the Apache 2.0 license, which is generally considered one of the most permissive. This license will allow enterprises to monetize OpenAI's open models without having to pay or obtain permission from the company.


However, unlike fully open source offerings from AI labs like AI2, OpenAI says it will not be releasing the training data used to create its open models. This decision is not surprising given that several active lawsuits against AI model providers, including OpenAI, have alleged that these companies inappropriately trained their AI models on copyrighted works.

OpenAI [delayed](#) the release of its open models [several](#) times in recent months, partially to address safety concerns. Beyond the company's typical safety policies, OpenAI says in a white paper that it also investigated whether bad actors could fine-tune its gpt-oss models to be more helpful in cyberattacks or the creation of biological or chemical weapons.

After testing from OpenAI and third-party evaluators, the company says gpt-oss may marginally increase biological capabilities. However, it did not find evidence that these open models could reach its "high capability" threshold for danger in these domains, even after fine-tuning.

While OpenAI's model appears to be state-of-the-art among open models, developers are eagerly awaiting the release of DeepSeek R2, its next AI reasoning model, as well as a new open model from Meta's Superintelligence Lab.

Topics: [AI](#) [open model](#) [OpenAI](#)



Maxwell Zeff

Senior AI Reporter | [X](#)

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >

Loading the next article



[TechCrunch Staff](#)

[Contact Us](#)

[Advertise](#)

[Crunchboard Jobs](#)

[Terms of Service](#)

[Privacy Policy](#)

[RSS Terms of Use](#)

[OpenAI](#)

[Perplexity](#)

[Cisco](#)

[DeepMind](#)

[Windsurf](#)

[Tech Layoffs](#)

