

Advertisement

AI



# OpenAI to launch AI data center in Norway, its first in Europe

Rebecca Bellan — 7:54 AM PDT · July 31, 2025

OpenAI said Thursday that it plans to launch Stargate Norway, its first AI data center in Europe, in partnership with British AI cloud infrastructure provider Nscale and Norwegian energy infrastructure firm Aker.

Nscale will design and build the site, and it will be a 50/50 joint venture between the two companies. OpenAI will be an “off-taker” in the project, buying capacity from the data center.


The announcement comes as Europe races to achieve AI sovereignty and invest in data centers and compute power. Earlier this week, the bloc unveiled details of its [multibillion-dollar investment](#) into AI infrastructure, including €10 billion (\$11.8 billion) to set up 13 AI factories and €20 billion as an initial investment in the factories. Data sovereignty is key to that mission, due to the sensitive nature of business and government data.

Nscale and OpenAI told TechCrunch that Stargate Norway isn’t a part of the European Union’s plans to scale AI at home. Nscale CEO Josh Payne [told CNBC](#) that part of the purpose of this project is to “leverage European sovereign compute” for the benefit of the continent. Norway’s AI ecosystem, like startups and scientific researchers, will get priority access to the center.

CNBC also reported that Nscale and Aker have each committed around \$1 billion to the initial 20 megawatt (MW) phase of the project. OpenAI says Stargate Norway will initially deliver 230 MW capacity, with plans to expand to 290 MW, and will run on 100,000 Nvidia GPUs by the end of 2026.

The site will be located near Narvik, a small town in northern Norway. OpenAI said in [a blog post](#) the region is notable for its access to hydropower, a cool climate, and a “mature industrial base.”

“The facility will run entirely on renewable power and is expected to incorporate closed-loop, direct-to-chip liquid cooling to ensure maximum cooling efficiency,” said OpenAI. “Additionally, excess heat from the GPU systems will be made available to support low-carbon enterprises in the region.”



## TC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

TC

### Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda.

They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

## Most Popular



- Key sections of the US Constitution deleted from government's website
- Instagram takes on Snapchat with new 'Instagram Map'
- Google says hackers stole its customers' data by breaching its Salesforce database

Per the [EU AI Act](#), which took effect in August 2024 and bans systems with “unacceptable risk,” companies building data centers must take steps to protect the environment and be transparent about the energy consumption of AI models. Additionally, the bloc’s Energy Efficiency Directive emphasizes energy efficiency in the ICT sector, which includes data centers. The directive also directs data centers exceeding certain energy input thresholds to recover waste heat.

The Stargate Norway announcement comes seven months after OpenAI announced it would [invest \\$500 billion](#) into 10 gigawatts of AI infrastructure in the United States over the next four years, in partnership with [Oracle](#) and SoftBank. The deal also follows the launch of [Stargate UAE](#) earlier this year, and a recently [signed deal with the U.K. government](#) to accelerate AI adoption and boost infrastructure.

*This article was updated to clarify that Stargate Norway is separate from the EU’s investment into AI infrastructure.*

Topics: [AI](#) [ai sovereignty](#) [aker](#) [data centers](#) [Norway](#) [nscale](#) [OpenAI](#)

[stargate](#) [stargate norway](#)



**Rebecca Bellan**


Senior Reporter |

Rebecca Bellan is a senior reporter at TechCrunch where she covers the business, policy, and emerging trends shaping artificial intelligence. Her work has also...

[View Bio](#) >

KEEP READING

- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Spotify raises subscription prices
- The uproar over Vogue’s AI-generated ad isn’t just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines





# OpenAI is practically giving ChatGPT to the government for free

Rebecca Bellan — 10:09 AM PDT · August 6, 2025

OpenAI is poised to undercut rivals like Anthropic and Google in the race to see its AI tools integrated into federal agency workflows.

The AI giant has [reached an agreement](#) with the U.S. General Services Administration (GSA), the government's central purchasing arm, to offer ChatGPT Enterprise to participating federal agencies for just \$1 per agency for the next year. IMAGE CREDITS: GETTY IMAGES

The partnership comes a day after the [GSA added OpenAI, Google, and Anthropic](#) to a list of approved AI vendors that can offer their services to civilian federal agencies. The companies will see their tools offered via the Multiple Award Schedule (MAS), a federal contracting platform that lets government agencies access AI tools via pre-negotiated contracts so they don't have to negotiate with vendors individually.

It's not clear whether other AI firms are going to offer their services at such a discounted rate, though GSA Federal Acquisition Service commissioner Josh Gruenbaum in a [statement](#) encouraged "other American AI technology companies to follow OpenAI's lead and work with us."

TechCrunch has reached out to Anthropic and Google for more information.

In addition to access to ChatGPT Enterprise, OpenAI is offering unlimited use of advanced models for an additional 60 days. Federal employees will also have access to a new government user community and tailored introductory training resources to get familiar with OpenAI's tools.

Data security is a top priority for government agencies that are concerned about sensitive information being leaked into model training sets. TechCrunch has asked the GSA for details on how government data is being safeguarded, including whether measures like on-premises or private cloud deployments are being used to enhance security.

Advertisement

## Newsletters

[See More ↗](#)

Subscribe for the industry's biggest tech news

### TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch's coverage.

### Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

### TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

### TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.



## Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

“The government is taking a cautious, security-first approach to AI,” a GSA spokesperson told TechCrunch. “This ensures sensitive information remains protected while enabling agencies to benefit from AI-driven efficiencies.”

OpenAI’s discount comes a couple of weeks after the Trump administration published its AI Action Plan that seeks to boost data center buildouts and integrate more AI tools into government, among other things. It also follows a Trump executive order that [bans “woke AI”](#) and AI models that aren’t “ideologically neutral” from government contracts.

TechCrunch has reached out to OpenAI to learn more about its GSA partnership and how it will approach Trump’s executive order.

Topics: [AI](#) [ChatGPT Enterprise](#) [Government & Policy](#) [GSA](#) [OpenAI](#)

Advertisement

Advertisement



**Rebecca Bellan**

Senior Reporter | [X](#) [🦋](#)

Rebecca Bellan is a senior reporter at TechCrunch where she covers the business, policy, and emerging trends shaping artificial intelligence. Her work has also...

# GPT-5

## Flagship model

IMAGE CREDITS: OPENAI

AI



## OpenAI's GPT-5 is here

Maxwell Zeff — 10:00 AM PDT · August 7, 2025

OpenAI has launched GPT-5, a new flagship AI model that will power the company's next generation of ChatGPT.

GPT-5, which was released Thursday, is OpenAI's first "unified" AI model and combines the reasoning abilities of its [o-series](#) of models with the fast responses of its [GPT series](#). The next-generation model signals a new era for ChatGPT — and its creator, OpenAI — pointing to OpenAI's broader ambitions to develop AI systems that are [more like agents](#) than chatbots.

While [GPT-4](#) enabled AI chatbots to offer smart responses on a wide variety of questions, GPT-5 allows ChatGPT to complete a wide variety of tasks on behalf of users — such as generating software applications, navigating a user's calendar, or creating research briefs.

With GPT-5, OpenAI has also sought to make ChatGPT simpler to use. Instead of asking users to choose the right settings, GPT-5 comes equipped with a real-time router that decides how to offer the best

Advertisement

### Latest in AI

See More

#### TRANSPORTATION



Tesla shuts down Dojo, the AI training supercomputer that Musk said would be key to full self-driving

answer, whether that’s responding to user questions quickly or taking additional time to “think” through answers.




IMAGE CREDITS: OPENAI

During a briefing with reporters, OpenAI CEO Sam Altman claimed GPT-5 is “the best model in the world,” and said it represented a “significant step” along the company’s path to developing AI that can outperform humans at most economically valuable work — that is, artificial general intelligence (AGI).

“Having something like GPT-5 would be pretty much unimaginable at any previous time in history,” said Altman.

Starting Thursday, GPT-5 will be available to all free users of ChatGPT as their default model. OpenAI’s VP of ChatGPT, Nick Turley, said this is part of the company’s effort to give free users access to an AI reasoning model for the first time. (Previously, the company gated these more advanced models behind a paywall.)

“This is just one of the ways that I’m excited to live the mission, making sure that this stuff actually benefits people,” said Turley on the decision, referencing OpenAI’s long-standing [mission](#) to distribute advanced AI to as many people as possible.



## Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

Rebecca Bellan - 6 hours ago

### STARTUPS



High costs and thin margins threatening AI coding startups

Marina Temkin - 7 hours ago

### AI



OpenAI’s GPT-5 is here

Maxwell Zeff - 11 hours ago

Advertisement

The expectations are high for GPT-5, one of OpenAI's most anticipated product launches since ChatGPT put the company on the map in 2022. Since then, ChatGPT has grown into one of the world's most popular consumer products, reaching more than 700 million users every week — nearly 10% of the globe's population, according to the company.

Many see GPT-5 as a bellwether for AI progress broadly, and the model's reception by Silicon Valley could have profound implications for Big Tech, Wall Street, and policymakers regulating technology. These stakeholders are watching to see if GPT-5 offers a significant jump in AI's capabilities, much like its predecessor, GPT-4, which challenged expectations of what software can do.

## GPT-5 offers a slight edge on the competition

OpenAI claims GPT-5 is state-of-the-art in several domains, slightly edging out leading AI models from Anthropic, Google DeepMind, and Elon Musk's xAI on key benchmarks. However, GPT-5 slightly underperforms frontier AI models in other areas.

The company says GPT-5 offers frontier-level performance around coding; Altman said the model specifically excels at spinning up entire software applications on demand, in what's become known as "vibe coding."

On SWE-bench Verified — a test of real-world coding tasks pulled from GitHub — GPT-5 scores 74.9% on its first attempt. That means GPT-5 just outperforms Anthropic's latest Claude Opus 4.1 model, which scored [74.5%](#), and Google DeepMind's Gemini 2.5 Pro, which scored [59.6%](#).

On Humanity's Last Exam — a difficult test measuring AI model performance across math, humanities, and the natural sciences — a version of GPT-5 with extended reasoning (GPT-5 Pro) scored 42% when using tools. That's slightly less than xAI was able to achieve with [Grok 4 Heavy](#), which scored 44.4% on the test.




IMAGE CREDITS: OPENAI

On GPQA Diamond — a test of PhD-level science questions — GPT-5 pro scored 89.4% on its first try, outperforming Claude Opus 4.1, which scored 80.9%, and Grok 4 Heavy, which scored 88.9%.

OpenAI says GPT-5 is better for answering health-related questions. On a test measuring accuracy in AI model responses around healthcare topics, HealthBench Hard Hallucinations, OpenAI says GPT-5 (with thinking) hallucinates just 1.6% of the time. This is far lower than the company’s previous GPT-4o and o3 models, which scored 12.9% and 15.8, respectively.

While AI chatbots are not medical professionals, [millions of people](#) are using them for health advice. In response to this phenomenon, the company says GPT-5 is more proactive about flagging potential health concerns and helping users parse medical results.

In addition, OpenAI says GPT-5 is better than other AI models on more difficult-to-measure, [subjective domains](#), such as creative design and writing. Turley said GPT-5 responds more naturally and exhibits “better taste” than other AI models on creative tasks.

“The vibes of this model are really good,” said Turley.

GPT-5 is also more accurate than OpenAI’s previous models, and the company says it suffers far less from hallucinations — the tendency for AI models to make up information — compared to its o-series models. Hallucinations seemed to be [getting worse](#) in OpenAI’s latest AI reasoning models, such as o3, and OpenAI previously said it didn’t quite understand why it was happening.

[Events](#)[Podcasts](#)[Newsletters](#)

[humans](#) or lie to promote their own goals, OpenAI found that GPT-5 was deceptive at a lower rate than other models.

Alex Beutel, OpenAI's safety research lead, said reducing deception improves not only the safety of GPT-5, but also the user experience, creating a model that's more "transparent and honest in ways users can trust."

Beutel also notes GPT-5 is better at discerning between bad actors who are trying to misuse ChatGPT and users making harmless requests. This results in GPT-5 being able to refuse more unsafe questions, while offering fewer rejections to users seeking harmless information.

## Upgrades for consumers and developers

ChatGPT is getting a few user experience upgrades as part of the GPT-5 launch. Users can now select from four new personalities in ChatGPT's setting: Cynic, Robot, Listener, and Nerd. The company says these will adapt ChatGPT's responses without requiring users to specifically ask the model to respond in a certain way.

Subscribers to ChatGPT's \$20-per-month Plus plan get higher usage limits for GPT-5 than free users. Meanwhile, \$200-per-month Pro subscribers will have unlimited access to GPT-5, as well as a souped-up version called GPT-5 Pro that uses additional computational resources to produce better answers. Organizations on OpenAI's Team, Edu, and Enterprise plans will gain access to GPT-5 as their default model next week.

For developers, GPT-5 is coming to OpenAI's API in three sizes — gpt-5, gpt-5-mini, and gpt-5-nano — which will spend more or less time "reasoning" through tasks. Developers can also now control verbosity in the OpenAI API, deciding how long or short an AI model's responses should be.

The base model of GPT-5 will cost developers \$1.25 per million input tokens (roughly 750,000 words, longer than the entire "Lord of the Rings" series) and \$10 per million output tokens.

The launch of GPT-5 comes after a busy week for OpenAI. The company released an open-weight reasoning model, [gpt-oss](#), that developers and

enterprises can download for free and run at a fraction of the cost. The open model nearly matched the abilities of OpenAI’s previous top models, o3 and o4-mini, but GPT-5 sets a new standard for frontier performance in some areas, such as coding.

However, GPT-5 seems to be roughly on par with other frontier AI models in several areas. Benchmarks, of course, only tell part of the story for any AI model, and it remains to be seen how developers will use GPT-5 in the real world, and whether the model is truly a step above the competition.

Topics: [AI](#) [ChatGPT](#) [gpt-5](#) [OpenAI](#)

[f](#) [X](#) [in](#) [t](#) [✉](#) [🔗](#)




**Maxwell Zeff**  
Senior AI Reporter | [X](#)

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg, and MSNBC, Zeff has covered the rise of AI and the Silicon...

[View Bio](#) >

KEEP READING



AI

[f](#) [X](#) [in](#) [t](#) [✉](#) [🔗](#)

Microsoft brings OpenAI’s smallest open model to Windows users

Ram Iyer — 5:47 AM PDT · August 6, 2025

Microsoft is making OpenAI's new free and open GPT model, gpt-oss-20b, available to Windows 11 users via [Windows AI Foundry](#), the tech giant's platform that lets users tap AI features, APIs, and popular open source models on their computers.

The company said in a [blog post](#) that gpt-oss-20b is “tool-savvy and lightweight,” adding: “Optimized for agentic tasks like code execution and tool use, it runs efficiently on a range of Windows hardware, with support for more devices coming soon. It's perfect for building autonomous assistants or embedding AI into real-world workflows, even in bandwidth-constrained environments.”

Launched on Tuesday, OpenAI's gpt-oss-20b can run on consumer PCs and laptops with at least 16GB of VRAM, which a modern GPU from Nvidia or Radeon would have. OpenAI said the model was trained using [high-compute reinforcement learning](#), which helps it excel at powering AI agents and call tools, such as web search or Python code execution, as part of its chain-of-thought process.

However, the model is text-only, meaning it will not be able to process or generate images and audio like the company's other models.

Also, the model has a habit of hallucinating a lot: gpt-oss-20b hallucinated in response to 53% of questions on PersonQA, OpenAI's in-house benchmark for measuring the accuracy of a model's knowledge about people.

Microsoft says it will bring the model to macOS soon, as well as other devices, though the company did not name any more. Microsoft is also making both gpt-oss-20b and older model gpt-oss-120b available via its hosted Azure AI Foundry platform.

The two new models are also [available on Amazon's AWS](#).



## Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

## TC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

Book Your Table >

## Most Popular



- Key sections of the US Constitution deleted from government's website
- Instagram takes on Snapchat with new 'Instagram Map'
- Google says hackers stole its customers' data by breaching its Salesforce database
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines



Ram Iyer

Editor | [X](#) [in](#)

Ram is a financial and tech reporter and editor. He covered North American and European M&A, equity, regulatory news and debt markets at Reuters and Acuris...

View Bio >

Loading the next article



YOU HAVE  
ONLY 2 MOVES

IQ: 155

HERO WARS

TRY

