




ENTERPRISE



ParadeDB takes on
Elasticsearch as interest
in Postgres explodes
amid AI boom



Open source database management system **Postgres is nearly 40 years old**, but has recently started seeing explosive demand due to being very well-suited for AI applications. Despite this rise in popularity, search and analytics functionality has been missing. **ParadeDB** is changing that.

ParadeDB is an open source Postgres extension that adds full-text search and analytics capabilities to Postgres without users needing to integrate with a separate source. The platform integrates with popular data infrastructure tools, including Snowflake, AWS SQL, Azure Postgres, and Amazon RDS, among others.

IMAGE CREDITS: PARADEDB

Philippe Noël, the co-founder and CEO of ParadeDB, told TechCrunch that he and his co-founder, Ming Ying, CTO, got the idea for the company by running into their own Postgres search woes while running their first startup, cloud-hybrid browser Whist.

“Postgres is becoming the default database of the world, and you still can’t do good search over that information, believe it or not,” Noël said. “There’s just a lot of pain points. So we decided to go and start a company to do that after talking to more people and realizing this was a very shared problem.”

ParadeDB isn’t the first company to try to solve Postgres search. Open source Elasticsearch is the

TECH DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government’s website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue’s AI-generated ad isn’t just about fashion

most notable legacy player in the space since it was founded in 2012. Noël said that Elasticsearch works by moving data back and forth between itself and Postgres, which can work, but this system isn't great for heavy workloads or processes that require frequent updating.

"That breaks all the time," Noël said. "The two databases are not meant to work together. There's a lot of compatibility issues, there's a lot of latency issues, higher costs, and all of that deteriorates the user experience."

ParadeDB claims to eliminate a lot of those challenges by building as an extension on top of Postgres directly, no data transfer required.

- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

[REGISTER NOW](#)

The startup was founded in 2023 and released the first open source version of the product later that year. Noël said that the group wanted to focus first on building out the open source product before it thought about sales or marketing.

That was short-lived. In early 2024, Chinese e-commerce giant Alibaba reached out and became ParadeDB's first customer in May 2024. From there, the company turned toward building out its enterprise version of the software. Now it works with enterprises, including Modern Treasury, Bilt Rewards, and TCDI, among others.

ParadeDB just raised a \$12 million Series A round led by Craft Ventures with participation from existing investors, including Y Combinator.

The company will use the funds for hiring. ParadeDB is currently a team of four and is looking to grow to at least 10 employees. Some of the funds will also go to continue to improve the platform's user interface and analytics capabilities.

Noël said the company wasn't looking to raise capital, but when Craft reached out to ParadeDB after being recommended to them by [Supabase](#), the Postgres developer platform that has [raised nearly \\$400 million of venture capital](#), they thought the timing made sense.

"It's clear people are recognizing that Postgres matters a lot, and they want to get behind it," he added.

Postgres' popularity is clearly on the rise if recent M&A is any indicator. The recent acquisitions of [Crunchy Data by Snowflake](#) and [Neon by Databricks](#) were both plays for the acquirers to gain more Postgres offerings, Devin Pratt, a research director at IDC, told TechCrunch last month.

"It is widely used across the industry, and I would say that was one of the main targets for their acquisitions," Pratt told TechCrunch at the time. "The different strengths of Neon versus Crunchy [Data] were definitely debated between both companies and what would suit their companies

best, but Postgres was their first goal in those acquisitions.”

ParadeDB hopes it can ride this wave of interest.

“We think building on Postgres is actually a meaningful shift, right?” Noël said. “Because that’s where the data is, and you can make a meaningful dent in Elasticsearch’s market share by meeting users where their data is, rather than building something that’s marginally faster or marginally cheaper.”

Topics: [AI](#) [artificial intelligence](#) [data](#)

[data management](#) [Databricks](#) [Elasticsearch](#)

[Enterprise](#) [Exclusive](#) [Postgres](#) [Snowflake](#)

[United States](#)



Rebecca Szkutak

Senior Reporter, Venture | [X](#) [in](#)

Becca is a senior writer at TechCrunch that covers venture capital trends and startups. She previously covered the...

[View Bio](#) >

———— KEEP READING ————

Latest Security
Startups AI
Venture Apps
Apple

[Events](#)[Podcasts](#)[Newsletters](#)

Crunchy Data

Rebecca Szkutak — 1:05 PM PDT · June 2, 2025

The surge in data company acquisitions continued Monday with Snowflake's purchase of Crunchy Data.

IMAGE CREDITS:

SOPA IMAGES / CONTRIBUTOR / GETTY IMAGES

Cloud data platform Snowflake announced its intent to acquire [Crunchy Data](#), a PostgreSQL database partner, late this afternoon. A source familiar with the matter estimated the deal to be worth around \$250 million.

The terms weren't disclosed. Snowflake declined to comment on the deal's valuation.

It's the latest in a string of tech giants buying data startups to bolster their underlying database offerings that power AI agents. Last week, [Salesforce acquired decades-old Informatica](#) to fortify its data management tooling for AI agents. A few weeks ago, [Alation acquired Numbers Station](#) to give its customers access to AI agents that could run on top of structured data. And earlier this



Newsletters

[See More ↗](#)

Subscribe for the industry's biggest tech news

month, [ServiceNow acquired Data.World](#) with AI agents in mind.

Crunchy Data gives companies the tools they need to build with Postgres, an open source and object-relational database management system that is popular with developers and works well with the SQL database language. The 13-year-old company works with clients such as UPS, SAS, and Moneytree, among others, in addition to public sector organizations like the U.S. Department of Homeland Security.

The acquisition will give Snowflake the underlying tech needed to launch Snowflake Postgres, an enterprise-grade PostgreSQL database for its customers, and partners including Blue Yonder and LandingAI. This deal also further bolsters Snowflake's database offerings in the age of AI agents.

This tech will be available in private preview soon.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch's coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).



“Our vision is to deliver the world’s most trusted and comprehensive data and AI platform to our customers,” said Vivek Raghunathan, SVP of engineering at Snowflake, in a press release provided to TechCrunch. “Today’s announcement of our proposed acquisition of Crunchy Data represents another reason why Snowflake is the ultimate destination for all enterprise data and AI needs. We’re tackling a massive \$350 billion market opportunity and a real need for our customers to bring Postgres to the Snowflake AI Data Cloud.”

In 2024, [Snowflake launched new connectors](#) for Snowflake customers to PostgreSQL and MySQL. Now it is bulking up that offering. Snowflake also [acquired data management company Datavolo](#) last year.

Topics: [acquisition](#) [AI](#) [AI agents](#) [data management](#)

[Enterprise](#) [Snowflake](#) [United States](#)



f X in 🗺️ ✉️ 🔗



Rebecca Szkutak

Senior Reporter, Venture | X in

Becca is a senior writer at TechCrunch that covers venture capital trends and startups. She previously covered the...

[View Bio](#) >

KEEP READING



IMAGE CREDITS: [MINISERIES](#) / GETTY IMAGES

[f](#) [X](#) [in](#) [reddit](#) [email](#) [link](#)

This is the beginning of the unbundled database era

Ethan Batraski — 3:04 PM PDT · June 1, 2022

Thanks to the cloud, the amount of data being generated and stored has exploded in scale and volume.

Every aspect of the enterprise is being instrumented for data, so new operations are built

Advertisement

based on that data, pushing every company into becoming a data company.

One of the most profound and maybe non-obvious shifts driving this is the emergence of the cloud database. Services such as Amazon S3, Google BigQuery, Snowflake and Databricks have solved computing on large volumes of data and have made it easy to store data from every available source.

The enterprise wants to store everything they can in the hopes of being able to deliver improved customer experiences and new market capabilities.

It's a good time to be a database company

Database companies have raised over \$8.7 billion over the last 10 years, with almost half of that, \$4.1 billion, just in the last 24 months, according to CB Insights.

It's not surprising given the sky-high valuations of Snowflake and Databricks. The market doubled in the last four years to almost \$90 billion, and is expected to double again over the next four years. It's safe to say there is a huge opportunity to go after.

[See here for a solid list of database financings in 2021.](#)

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and

Latest in Enterprise

See More


ENTERPRISE



Upwork is buying its way into corporate staffing beyond freelancers

Rebecca Szkutak
- 14 hours ago

STARTUPS



Figma's IPO success is 'a little bit of a meme stock,' says Sapphire Ventures' Jai Das

Rebecca Bellan,
Theresa Loconsolo
- 18 hours ago

AI



Cohere's new AI agent platform, North, promises to keep enterprise data secure

Rebecca Bellan
- 20 hours ago

Advertisement

sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025


REGISTER NOW

Data growth driving new spend



Database Market (2018)

Database Market (2021)



DATABASE GROWTH IS DRIVING SPEND IN THE ENTERPRISE.

IMAGE CREDITS: VENROCK

20 years ago, you had one option: A relational database

Today, thanks to the cloud, microservices, distributed applications, global scale, real-time data and deep learning, new database architectures have emerged to solve for new performance requirements.

We now have different systems for fast reads and fast writes. There are also systems specifically to power ad-hoc analytics or for data that is unstructured, semi-structured, transactional, relational, graph or time-series, as well as for data used for cache, search, based on indexes, events and more.

Each system comes with different performance needs, including high availability, horizontal scale, distributed consistency, failover protection, partition tolerance and being serverless and fully managed.

As a result, enterprises, on average, store data across seven or more different databases. For example, you may have Snowflake as your data warehouse, Clickhouse for ad-hoc analytics, Timescale for time-series data, Elastic for their search data, S3 for logs, Postgres for transactions, Redis for caching or application data, Cassandra for complex workloads and Dgraph* for relationship data or dynamic schemas.

That's all assuming you are collocated to a single cloud and you've built a modern data stack from scratch.

The level of performance and guarantees from these services and platforms is on a very different level compared with what we had five to 10 years ago. At the same time, the proliferation and fragmentation of the database layer are increasingly creating new challenges.

For example, syncing across different schemas and systems, writing new ETL jobs to bridge workloads across multiple databases, constant cross-talk and connectivity issues, the overhead of managing active-active clustering across so many different systems, or data transfers when new clusters or systems come online. Each of these has different scaling, branching, propagation, sharding and resource requirements.

What's more, we now have new databases every month that aim to solve the next challenge of enterprise scale.

The new-age database

So the question is, will the future of the database continue to be defined as it is today?

I'd make the case that it shouldn't. Instead, I hope the next generation of databases will look very different from the last.

They should have the following capabilities:

Primarily compute, query and/or be infrastructure engines that can sit on top of commodity storage layers.

No migration or restructuring of the underlying data should be required.

No rewriting or parsing of queries should be needed.

Work on top of multiple storage engines, whether columnar, non-relational or graph.

Moves the complexity of configuration, availability and scale into code.

Allows applications to call into a single interface, regardless of the underlying data infrastructure.

Works out of the box as a serverless or managed service.

Be built for developer-first experiences in both single-player and multiplayer modes.

Deliver day-zero value for both existing (brownfield) and new (greenfield) projects.

There are many secular trends driving this future

No one wants to migrate to a new database

The cost of every new database introduced into an organization is equal to N^2 to the cost of the number of databases you already have. Migrating to a new architecture, schema or configuration and having to re-optimize for rebalancing, query planning, scaling and resource requirements means the value derived after the cost and spend is close to zero.

It may come as a surprise, but there are still billions of dollars in Oracle instances still powering critical apps today, and they likely aren't going anywhere.

A majority of the killer features won't be in the storage layer

Separating compute and storage has increasingly enabled new levels of performance, allowing for super cheap raw storage and finely tuned, elastically scaled compute, query and infrastructure layers.

The storage layer can be at the center of data infrastructure and leveraged in various different ways by multiple tools to solve routing, parsing, availability, scale, translation and more.

The database is slowly unbundling into highly specialized services

It is moving away from the overly complex, locked-in approaches of the past. No single database can solve transactional and analytical use cases fully with fast reads and writes, high availability and consistency, all while solving caching at the edge and scaling horizontally as needed.

But separating the database into a set of layers sitting on top of the storage engine can introduce a set of new services to deliver new levels of performance and guarantees.

For example, you can have a dynamic caching service that can optimize caches based on user, query and data awareness; manage sharding based on data distribution query demand and data change rates; introduce a proxy layer to enable high availability and horizontal scale with connection pooling and resource management; add a data management framework to solve async and sync propagation between schemas; or introduce translation layers between GraphQL and relational databases.

Scale and simplicity have been trade-offs so far

Postgres, MySQL and Cassandra are very powerful but difficult to get right. Firebase and Heroku are super easy to use but don't scale.

These database technologies have massive install bases and robust engines and have withstood the test of time at Facebook and Netflix-level scales. But tuning them for your needs often requires a Ph.D. and a team of database experts, as teams at Facebook, Netflix, Uber, Airbnb all have.

The rest of us struggle with consistency and isolation, sharding, locking, clock skews, query planning, security, networking, etc. What companies like Supabase and Hydras are doing in leveraging standard Postgres installs but building powerful compute and management layers on top, allows for the power of Postgres, but with the simplicity of Firebase or Heroku.

The database index model hasn't changed in 30+ years

Today, we rely on general-purpose, one-size-fits-all indexes such as B-trees and hash maps, taking a black-box view of our data. Being more data aware,

like leveraging a cumulative distribution function (CDF) as we've seen with Learned Indexes, can lead to smaller indexes, faster lookups, increased parallelism and lower CPU usage.

We've barely even begun to demonstrate next-generation indexes that have adapted both to the shape and changes of our data.

There is little to no machine learning used to improve database performance

Instead, today we define static rule sets and configurations to optimize query performance, cost modeling and workload forecasting. These combinatorial, multidimensional problem sets are too complex for humans to configure and are perfect machine learning problems.

Resources such as disk, RAM and CPU are well characterized, query history is well understood and data distribution can be defined. We could see 10x step-ups in query performance, cost and resource utilization, and never see another nested loop join again.

Data platform and engineering teams don't want to be DBAs, DevOps or SREs

They want their systems and services to just work out of the box and not have to think about resources, connection pooling, cache logic, vacuuming, query planning, updating indexes and more. Teams today want a robust set of end points that are easy to deploy — and just work.

The need for operational real-time data is driving a need for hybrid

systems

Transactional systems can write new records into a table rapidly at a high level of accuracy, speed and reliability. An analytics system can search across a set of tables and data rapidly to find an answer.

With streaming data and the need for faster responsiveness in analytical systems, the idea of HTAP (hybrid transaction/analytical processing) systems is emerging, particularly for use cases that are highly operational in nature. This means a very high level of new writes/records and more responsive telemetry or analytics on business metrics. This introduces a new architectural paradigm in which transactional and analytical data and systems start to reside much closer to each other but not together.

A new category of cloud database companies is emerging, effectively deconstructing the traditional database monolith stack into core layered services — storage, compute, optimization, query planning, indexing, functions and more. Companies like ReadySet, Hasura, Xata, OtterTune, Apollo and PolyScale are examples of this movement and quickly becoming the new developer standard.

These new unbundled databases are focused on solving the hard problems of caching, indexes, scale and availability, and are beginning to remove the trade-off between performance and guarantees. Fast databases that are always on, can handle mass scale and are data aware are blurring the lines between the traditional divisions between operational and analytical systems. The future looks bright.

Topics: [Airbnb](#) [Caching](#) [cassandra](#) [ClickHouse](#)

[Column](#) [computing](#) [data management](#)

[data warehouse](#) [databases](#) [Databricks](#)

[developers](#)

[EC Cloud and Enterprise Infrastructure](#)

[EC Column](#) [EC Market Analysis](#) [Enterprise](#)

[Google](#) [Hasura](#) [heroku](#) [machine learning](#)

[mysql](#) [Netflix](#) [oracle](#) [OtterTune](#) [readysset](#)

[real time data](#) [relational database](#) [Startups](#)

[Startups](#) [Supabase](#) [Uber](#) [Xata](#)

Advertisement

[f](#) [X](#) [in](#) [t](#) [✉](#) [🔗](#)

Ethan Batraski

Contributor | [X](#)

[Ethan Batraski](#) is a partner at [Venrock](#) and focuses on data infrastructure, open source and developer tools.

[View Bio](#) >

Loading the next article





[TechCrunch Staff](#)

[Contact Us](#)

[Advertise](#)

[Crunchboard Jobs](#)

[Site Map](#)

[Terms of Service](#)

[Privacy Policy](#)

[RSS Terms of Use](#)

[OpenAI](#)

[Perplexity](#)

[Cisco](#)

[DeepMind](#)

[Windsurf](#)

[Tech Layoffs](#)

[ChatGPT](#)