


Advertisement

SECURITY



**Perplexity accused of
scraping websites that
explicitly blocked AI
scraping**



Lorenzo Franceschi-Bicchierai

8:41 AM PDT · August 4, 2025

AI startup Perplexity is crawling and scraping content from websites that have explicitly indicated they don't want to be scraped, according to internet infrastructure provider Cloudflare.

On Monday, Cloudflare [published research](#) saying it observed the AI startup ignore blocks and hide its crawling and scraping activities. The network infrastructure giant accused Perplexity of obscuring its identity when trying to scrape web pages “in an attempt to circumvent the website’s preferences,” Cloudflare’s researchers wrote.

AI products like those offered by Perplexity rely on gobbling up large amounts of data from the internet, and AI startups have long scraped text, images, and videos from the internet many times without permission to make their products work. In recent times, websites have tried to fight back by using the web standard Robots.txt file, which tells search engines and AI companies which pages can be indexed and which shouldn't, efforts [that have seen mixed results so far](#).

Perplexity appears to be willingly circumventing these blocks by changing its bots’ “user agent,” meaning a signal that identifies a website visitor by their device and version type, as well as changing their autonomous system networks, or ASN, essentially a number that identifies large networks on the internet, according to Cloudflare.

IMAGE CREDITS:

KIMBERLY WHITE FOR PERPLEXITY IMAGES FOR TECHCRUNCH


TECH DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government’s website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue’s AI-generated ad isn’t just about fashion

“This activity was observed across tens of thousands of domains and millions of requests per day. We were able to fingerprint this crawler using a combination of machine learning and network signals,” read Cloudflare’s post.

Perplexity spokesperson Jesse Dwyer dismissed Cloudflare’s blog post as a “sales pitch,” adding in an email to TechCrunch that the screenshots in the post “show that no content was accessed.” In a follow-up email, Dwyer claimed the bot named in the Cloudflare blog “isn’t even ours.”

Cloudflare said it first noticed the behavior after its customers complained that Perplexity was crawling and scraping their sites, even after they added rules on their Robots file and for specifically blocking Perplexity’s known bots. Cloudflare said it then performed tests to check and confirmed that Perplexity was circumventing these blocks.

- Your public ChatGPT queries are getting indexed by Google and other search engines
- The best dating apps aren’t even dating apps



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

“We observed that Perplexity uses not only their declared user-agent, but also a generic browser

intended to impersonate Google Chrome on macOS when their declared crawler was blocked,” according to Cloudflare.

The company also said that it has de-listed Perplexity’s bots from its verified list and added new techniques to block them.

Cloudflare has recently taken a public stance against AI crawlers. Last month, Cloudflare [announced the launch of a marketplace](#) allowing website owners and publishers to charge AI scrapers who visit their sites. Cloudflare’s chief executive Matthew Prince [sounded the alarm](#) at the time, saying AI is breaking the business model of the internet, particularly publishers. Last year, Cloudflare also [launched a free tool](#) to prevent bots from scraping websites to train AI.

This is not the first time Perplexity is accused of scraping without authorization.

Last year, news outlets, [such as Wired](#), alleged [Perplexity was plagiarizing their content](#). Weeks later, Perplexity’s CEO Aravind Srinivas [was unable to immediately answer](#) when asked to provide the company’s definition of plagiarism during an interview with TechCrunch’s Devin Coldewey at the Disrupt 2024 conference.

Topics: [AI](#) [Artificial Intelligence \(AI\)](#) [bots](#) [cloudflare](#)

[LLMs](#) [Perplexity](#) [scraping](#) [Security](#)

[f](#) [X](#) [in](#) [📺](#) [✉](#) [🔗](#)



Lorenzo Franceschi-Bicchierai

Senior Reporter, Cybersecurity | [X](#) [🦋](#) [@](#)

Lorenzo Franceschi-Bicchierai is a Senior Writer at TechCrunch, where he covers hacking, cybersecurity,...

View Bio >

KEEP READING



IMAGE CREDITS:

KIMBERLY WHITE/GETTY IMAGES FOR TECHCRUNCH

AI

[f](#) [X](#) [in](#) [reddit](#) [email](#) [link](#)

Some people are defending Perplexity after Cloudflare ‘named and shamed’ it

Julie Bort — 9:33 AM PDT · August 5, 2025

When Cloudflare accused AI search engine Perplexity of [stealthily scraping websites](#) on Monday, while ignoring a site’s specific methods to block it, this wasn’t a clear-cut case of an AI web crawler gone wild.

Advertisement

Many people came to Perplexity’s defense. They argued that Perplexity accessing sites in defiance of the website owner’s wishes, while controversial, is acceptable. And this is a controversy that will certainly grow as AI agents flood the internet: Should an agent accessing a website on behalf of its user be treated like a bot? Or like a human making the same request?

Cloudflare is known for providing anti-bot crawling and other web security services to millions of websites. Essentially, Cloudflare’s test case involved setting up a new website with a new domain that had never been crawled by any bot, setting up a robots.txt file that specifically blocked Perplexity’s known AI crawling bots, and then asking Perplexity about the website’s content. And Perplexity answered the question.

Cloudflare researchers found the AI search engine used “a generic browser intended to impersonate Google Chrome on macOS” when its web crawler itself was blocked. Cloudflare CEO Matthew Prince [posted](#) the research on X, writing, “Some supposedly ‘reputable’ AI companies act more like North Korean hackers. Time to name, shame, and hard block them.”

But many people disagreed with Prince’s assessment that this was actual bad behavior. Those defending Perplexity on sites [like X](#) and [Hacker News](#) pointed out that what Cloudflare seemed to document was the AI accessing a specific public website when its user asked about that specific website.

“If I as a human request a website, then I should be shown the content,” one person on [Hacker News](#) wrote, adding, “why would the LLM accessing the website on my behalf be in a different legal category as my Firefox web browser?”

Newsletters

[See More ↗](#)

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

A Perplexity spokesperson [previously](#) denied to TechCrunch that the bots were the company's and called Cloudflare's blog post a sales pitch for Cloudflare. Then on Tuesday, Perplexity [published a blog](#) in its defense (and generally attacking Cloudflare), claiming the behavior was from a third-party service it uses occasionally.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

But the crux of Perplexity's post made a similar appeal as its online defenders did.

"The difference between automated crawling and user-driven fetching isn't just technical — it's about who gets to access information on the open web," the post said. "This controversy reveals that Cloudflare's systems are fundamentally inadequate for distinguishing between legitimate AI assistants and actual threats."

Perplexity's accusations aren't exactly fair, either. One argument that Prince and Cloudflare used for calling out Perplexity's methods was that OpenAI doesn't behave in the same way.

“OpenAI is an example of a leading AI company that follows these best practices,” [Cloudflare wrote](#). “They respect robots.txt and do not try to evade either a robots.txt directive or a network level block. And ChatGPT Agent is signing http requests using the newly proposed open standard Web Bot Auth.”

[Web Bot Auth](#) is a Cloudflare-supported standard being developed by the Internet Engineering Task Force that hopes to create a cryptographic method for identifying AI agent web requests.

The debate comes as bot activity reshapes the internet. As TechCrunch has previously reported, bots seeking to scrape massive amounts of content to train AI models [have become a menace](#), especially to smaller sites.

For the first time in the internet’s history, [bot activity is currently outstripping human activity online](#), with AI traffic accounting for over 50%, according to Imperva’s Bad Bot report released last month. Most of that activity is coming from LLMs. But the report also found that malicious bots now make up 37% of all internet traffic. That’s activity that includes everything from persistent scraping to unauthorized login attempts.

Until LLMs, the internet generally accepted that websites could and should block most bot activity given how often it was malicious by using CAPTCHAs and other services (such as Cloudflare). Websites also had a clear incentive to work with specific good actors, such as Googlebot, guiding it on what not to index through robots.txt. Google indexed the internet, which sent traffic to sites.


Now, LLMs are eating an increasing amount of that traffic. Gartner predicts [that search engine volume](#) will drop by 25% by 2026. Right now humans tend to click website links from LLMs at the point they are

most valuable to the website, which is when they are ready to conduct a transaction.

But if humans [adopt agents](#) as the tech industry predicts they will — to arrange our travel, book our dinner reservations, and shop for us — would websites hurt their business interests by blocking them? The debate on X captured the dilemma perfectly:

“I WANT perplexity to visit any public content on my behalf when I give it a request/task!” wrote [one person](#) in response to Cloudflare calling out Perplexity.

“What if the site owners don’t want it? they just want you [to] directly visit the home, see their stuff” [argued another](#), pointing out that the site owner who created the content wants the traffic and



Julie Bort

Venture Editor | 

[View Bio](#) >

———— KEEP READING ————




IMAGE CREDITS: GETTY IMAGES

AI



Cloudflare launches a tool to combat AI bots

Kyle Wiggers — 8:49 AM PDT · July 3, 2024

Cloudflare, the publicly traded cloud service provider, has launched a new, free tool to prevent bots from scraping websites hosted on its platform for data to train AI models.

Some AI vendors, including Google, OpenAI and Apple, allow website owners to block the bots they use for data scraping and model training by amending their site's robots.txt, the text file that tells bots which pages they can access on a website. But, as Cloudflare points out in a [post](#) announcing its bot-combating tool, not all AI scrapers respect this.

Latest in AI

[See More](#)

AI



Google denies AI search features are killing website traffic

Sarah Perez - 11 hours ago

AI



First impressions of Alexa+, Amazon's upgraded, AI-powered digital assistant

“Customers don’t want AI bots visiting their websites, and especially those that do so dishonestly,” the company writes on its official blog. “We fear that some AI companies intent on circumventing rules to access content will persistently adapt to evade bot detection.”

So, in an attempt to address the problem, Cloudflare analyzed AI bot and crawler traffic to fine-tune automatic bot detection models. The models consider, among other factors, whether an AI bot might be trying to evade detection by mimicking the appearance and behavior of someone using a web browser.

“When bad actors attempt to crawl websites at scale, they generally use tools and frameworks that we are able to fingerprint,” Cloudflare writes. “Based on these signals, our models [are] able to appropriately flag traffic from evasive AI bots as bots.”

Cloudflare has set up a form for hosts to report suspected AI bots and crawlers and says that it’ll continue to manually blacklist AI bots over time.

The problem of AI bots has come into sharp relief as the generative AI boom fuels the demand for model training data.

Sarah Perez - 12 hours ago

AI



Google takes on ChatGPT’s Study Mode with new ‘Guided Learning’ tool in Gemini

Aisha Malik - 12 hours ago

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your

ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Many sites, wary of AI vendors training models on their content without alerting or compensating them, have opted to block AI scrapers and crawlers. Around 26% of the top 1,000 sites on the web have blocked OpenAI's bot, according to one [study](#); [another](#) found that more than 600 news publishers had blocked the bot.

Blocking isn't a surefire protection, however. As alluded to earlier, some vendors appear to be ignoring standard bot exclusion rules to gain a competitive advantage in the AI race. [AI search engine Perplexity was recently accused](#) of impersonating legitimate visitors to scrape content from websites, and OpenAI and Anthropic are said to have at times [ignored](#) robots.txt rules.

In a [letter to publishers last month](#), content licensing startup TollBit said that, in fact, it sees "many AI agents" ignoring the robots.txt standard.

Tools like Cloudflare's could help — but only if they prove to be accurate in detecting clandestine AI bots. And they *won't* solve the more intractable problem of publishers risking sacrificing referral traffic from AI tools like Google's AI Overviews, which exclude sites from inclusion if they block specific AI crawlers.

Topics: [AI](#) [AI](#) [bots](#) [cloudflare](#)





Kyle Wiggers

AI Editor | [X](#) [Twitter](#) [Email](#) [in](#)

Kyle Wiggers was TechCrunch's AI Editor until June 2025.
His writing has appeared in VentureBeat and Digital...

[View Bio](#) >

KEEP READING




IMAGE CREDITS: STOCKCAM / GETTY IMAGES

AI



Reddit's upcoming changes attempt to safeguard the platform against AI crawlers

Aisha Malik — 12:09 PM PDT · June 25, 2024

Reddit [announced](#) on Tuesday that it's updating its Robots Exclusion Protocol (robots.txt file), which tells automated web bots whether they are permitted to crawl a site.

TC DISRUPT

October 27-29, 2025 | San Francisco

Historically, robots.txt file was used to allow search engines to scrape a site and then direct people to the content. However, with the rise of AI, websites are being scraped and used to train models without acknowledging the actual source of the content.

Along with the updated robots.txt file, Reddit will continue rate-limiting and blocking unknown bots and crawlers from accessing its platform. The company told TechCrunch that bots and crawlers will be rate-limited or blocked if they don't abide by Reddit's Public Content Policy and don't have an agreement with the platform.

Reddit says the update shouldn't affect the majority of users or good faith actors, like researchers and organizations, such as the Internet Archive. Instead, the update is designed to deter AI companies from training their large language models on Reddit content. Of course, AI crawlers could ignore Reddit's robots.txt file.

The announcement comes a few days after a [Wired investigation](#) found that AI-powered search startup Perplexity has been stealing and scraping content. Wired found that Perplexity seems to ignore requests not to scrape its website, even though it blocked the startup in its robots.txt file. Perplexity CEO [Aravind Srinivas responded](#) to the claims and said that the robots.txt file is not a legal framework.

Reddit's upcoming changes won't affect companies that it has an agreement with. For instance, Reddit has a [\\$60 million deal with Google](#) that allows the search giant to train its AI models on the social platform's content. With these changes, Reddit is signaling to other companies that want to use Reddit's data for AI training that they will have to pay.

"Anyone accessing Reddit content must abide by our policies, including those in place to protect

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines
- The best dating apps aren't even dating apps

redditors,” Reddit said [in its blog post](#). “We are selective about who we work with and trust with large-scale access to Reddit content.”



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

The announcement doesn’t come as a surprise, as Reddit [released a new policy](#) a few weeks ago that was designed to guide how Reddit’s data is being accessed and used by commercial entities and other partners.

Topics: [AI](#) [Apps](#) [bots](#) [Reddit](#) [Webcrawler](#)



Aisha Malik

Consumer News Reporter |

Aisha is a consumer news reporter at TechCrunch. Prior to joining the publication in 2021, she was a telecom reporte...

Loading the next article



TechCrunch
Staff

Contact Us

Advertise

Crunchboard
Jobs

Site Map

Terms of
Service

Privacy
Policy

RSS Terms of
Use

OpenAI
Perplexity

Cisco
DeepMind

Windsurf
Tech Layoffs
ChatGPT