


Advertisement

STARTUPS



**PlayerZero raises \$15M
to prevent AI agents from
shipping buggy code**



As Silicon Valley races toward a future where AI agents do most of the software programming, a new problem is created: finding the AI-generated bugs before they are put into production. Even OpenAI is [dealing with such issues](#), a former employee has described.

Newly funded startup [PlayerZero](#) has created a solution: use AI agents trained to find and fix problems before the code is put into production, the startup's CEO and sole founder, Animesh Koratana, tells TechCrunch.

Koratana created PlayerZero while he was at the Stanford DAWN lab for machine learning under his adviser and lab founder, Matei Zaharia. Zaharia is, of course, a famed developer and the co-founder of Databricks; he created its foundational technology while working on his own doctorate.

PlayerZero on Wednesday announced that it raised a \$15 million Series A led by Foundation Capital's Ashu Garg, an early Databricks backer. This follows a \$5 million seed led by Green Bay Ventures and several noteworthy angels, including Zaharia, Dropbox CEO Drew Houston, Figma CEO Dylan Field, and Vercel CEO Guillermo Rauch.

During his time at Stanford DAWN, Koratana, now 26, was working on AI model compression technology and "got exposed to language models really early on," he says. He met the developers


DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion

who crafted some of the first AI coding assistance tools.

It hit him then that “there’s this world in which computers are going to write the code. It’s not going to be humans anymore,” Koratana tells TechCrunch. “What’s the world going to look like at that point?”

He knew before the term “AI slop” was even coined that these agents were going to produce code that broke things just as their human overseers do.

- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

That problem would also be exacerbated by so many agents cranking out so much more code than has ever been written before. It won’t always be practical for humans to check all AI-written code for bugs or hallucinations. And the issue becomes even more intense for the large, complex code bases that enterprises rely upon.

PlayerZero trains models “that really deeply understand code bases, and we understand the

way they're built, the way they're architected," Koratana says.

His tech studies the history of an enterprise's bugs, issues, and solutions. When something breaks, his product can then "figure out why and fix it, and then learn from those mistakes to prevent them from ever happening again," Koratana says. He likens his product to an immune system for large code bases.

Landing Zaharia, his adviser, as an angel was a first step to fundraising, but the moment that really validated his idea was when he showed a demo to another famous developer: Rauch. Rauch is the founder of [triple-unicorn developer tool company](#) Vercel and creator of the popular open source JavaScript framework [Next.js](#).

Rauch watched Koratana's demo with interest but skepticism, asking how much of it was "real." Koratana replied that this was code "running in production. Like, this is a real instance. And he was quiet," Koratana says. Then his soon-to-be-angel investor responded, "If you can actually solve this the way that you're imagining, it's a really big deal."

Of course, PlayerZero isn't alone in attempting to solve the AI-generated bug problem. Just last week, Anysphere's Cursor [launched Bugbot](#) to detect coding errors, as just one example.

Still, PlayerZero is already gaining traction for its emphasis on large codebases. While it was conceived for a world where agents are the coders, it is currently being used by several large enterprises that use coding co-pilots. For instance, subscription billing company Zuora is one of the startup's marquee customers. Zuora is using the tech across its engineering teams, including to watchdog its most precious code, its billing systems, it said.

[AI](#) [ai code](#) [AI coding agents](#) [ai dev tools](#)

ai slop foundation capital Startups



Venture Editor | X

View Bio >

KEEP READING



AI



ServiceTitan names LLMs from Microsoft, OpenAI as risk factors

Julie Bort — 2:45 PM PST · November 18, 2024

IMAGE CREDITS: GETTY IMAGES

Every IPO S-1 prospectus includes a section listing all kinds risk factors for why its stock may not be a good buy. Many of them are boilerplate warnings, cautions over financial performance, or acts of war and nature. But some of them are more specific to the company.

With ServiceTitan, a cloud service startup that filed [public IPO documents](#) with the SEC on Monday, we may be witnessing the birth of a new boilerplate warning: large language models (LLMs) that help grow business could also lead to damage.

ServiceTitan had an 1,150-word risk factor on how its use of AI, specifically generative AI, could adversely impact its business. It warned that LLM hallucinatory behavior could produce “inaccurate” information and engage in “discriminatory” behaviors; its LLMs could infringe on others’ copyright or intellectual property; and using LLMs means exposing more data to potential hacks and harm. Then again, if it can’t get enough data, it may not be able to continue offering AI products or building new ones. On top of that, the company cautions that its employees or contractors could possibly mistakenly share its customers’ private data with third-party systems that could then lead to security breaches, or use that data to train their models.

ServiceTitan warns that its use of AI may one day run afoul of social and ethical issues. Or that future regulations could cost it money.

ServiceTitan also fears that it may not be able to hire the AI experts it needs for its products and, even if it can, that talent will be pricey. ServiceTitan warns its AI products rely on third-party services it doesn’t control — and specifically names Microsoft and OpenAI. So, if they become unavailable or have issues, that’s a risk too.

These warnings are interesting because ServiceTitan operates at the nexus of the industries that GenAI, specifically LLM-driven AI agents, is expected to dominate. ServiceTitan offers software for field service industry businesses — largely small businesses, think construction contractors, HVAC professionals, landscapers, and the like. Its software performs functions like marketing, customer relationship management, customer support, and accounting. (ServiceTitan declined further comment.)

AI agents are currently marching so furiously toward all things sales, marketing, and CRM that the king of the CRM market, Salesforce CEO Marc Benioff, told TechCrunch that [his company alone could have a billion AI agents](#) working at Salesforce’s customers’ companies in a year.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

ServiceTitan has been offering AI-powered services since at least 2023, called Titan Intelligence. [In](#)

Advertisement

Newsletters

[See More ↗](#)

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

[October](#), it also rolled out a suite of AI agents for sales, customer service, and call centers.

But LLMs were designed to study words, images, and data and create their own versions of things. That is, they were designed to make stuff up.

OpenAI’s scientists are busy [racing toward general intelligence](#) — mimicking humans’ ability to infer even better — meaning they are working on AI that they hope will be even more creative.

Sure, LLM reliability for business usage is something that should be solved over time as more companies specialize in agents where reliability, not creativity, is essential. But it’s worth noting that those that are selling AI bots today are actually saying the quiet part out loud in their legalese: At this early stage, AI adoption could create as many problems as it solves.

Topics: [AI](#) [hallucinations](#) [IPO](#) [LLMs](#) [OpenAI](#)

[servicetitan](#)



Julie Bort

Venture Editor |

[View Bio](#) >

———— KEEP READING ————

Advertisement

Marc Benioff warns that AI, while useful, is overhyped and partly blames Microsoft

Julie Bort — 11:23 AM PDT · October 17, 2024

One of the tech industry's biggest hype men, Salesforce CEO Marc Benioff, wants you to know that he is psyched about generative AI, but even he doesn't think it can do what its biggest proponents say it will do. And he (not surprisingly) blames Microsoft.

Latest in AI

See More

Latest

Security

Startups

AI

Venture

Apps

Apple



Events

Podcasts

Newsletters



But he also warned that “customers have been told things about enterprise AI, maybe AI overall, that are not true,” he said. “I think Microsoft has done a tremendous disservice to not only our whole industry but all of the AI research that has been done.”

Benioff had negative things to say in particular about Microsoft Copilot’s accuracy and usefulness. He even compared Copilot to Clippy, Microsoft’s widely panned 1990s talking paperclip cartoon that was supposed to be an assistant to Microsoft Office users.

“We may have heard from these [AI priests](#) and priestesses of these LLM model companies and Microsoft and others about AI is now curing cancer, and AI is curing climate change, and we all have to [plug into these nuclear power plants](#) to get these data centers. None of this is true,” Benioff said.

He was doubling down on a [comment](#) he made on X, where he said, “LLMs (Large Language Models) are not the direct bridge to AGI, and much of AI’s current potential is simply oversold. AI isn’t yet curing cancer or solving climate change as pundits claim.”



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

[REGISTER NOW](#)

That was a dig at OpenAI CEO Sam Altman. Earlier this summer, he postulated that with AI-enhanced health tech, “[m]aybe a future version will help discover cures for cancer,” Altman said at the Aspen Ideas Festival, [reported Newsweek](#).

Benioff also cited some research by Gartner about Microsoft Copilot. A [Gartner report](#) released in April called “The Top 10 ‘Gotchas’ of Copilot for Microsoft 365” found that only a quarter of the organizations doing Copilot pilot programs are currently planning a large-scale rollout. That’s actually a pretty good number, given how young this tech is and how slow enterprises can be to adopt it. That said, Gartner also concluded that as Copilot improves, so will enterprise adoption.

Microsoft might point to [market research](#) conducted by Forrester that cited a slew of bottom-line benefits for small businesses that are using Copilot. Forrester found that Copilot led to a tiny uptick in revenue, while reducing operating costs and speeding new-hire onboarding, based on a survey of 266 small companies. However, we must note that this research was paid for by Microsoft. Make of that what you will.

Still, Benioff has a point that the GenAI we have today, while sometimes mind-blowing, is not ready to replace human workers in most cases. True, the podcast-making [capabilities of Google’s NotebookLM](#) — which can create bantering AI-generated hosts that explain material — is quite the party trick. Yet, it’s hard to see how such a thing will directly reduce the kind of soul-sucking menial labor tasks that bog down most corporate jobs.

Benioff is also spot-on that the one area where GenAI does seem to be doing a screaming good job with enterprises is AI agents. That ties in nicely with the Salesforce product he's been hyping lately, Agentforce. A slew of other tech companies and startups are working on AI agent technologies, too, from building use case-specific ones, to offering platforms where businesses can build their own. Just a few examples include [OutRival](#), [Atlassian's Rovo AI](#), and [Sierra](#), the startup founded by former Salesforce co-CEO Bret Taylor and Google vet Clay Bavor.

AI agents are especially finding their way into customer service, ranging from much improved website chat bots to field service guides. Examples include [Zingtree](#), [Talla](#), or [Neuron7](#). Salesforce has offerings here, too.

AI agents are also finding success with [sales, especially prospecting](#), the high-rejection rate cold-calling and emailing that is the bottom tier of every sales organization. Examples include Regie.ai, AiSDR, Artisan, and [11x.ai](#).

"I think we'll have more than a billion agents running from Salesforce within the next 12 months," Benioff speculated, based, he said, on getting about 10,000 customers at his tech conference to try it.

However, it's also worth pointing out that there are other areas where LLMs are already valued and where Microsoft, specifically, does have game. Software programmers and engineers increasingly use them to help them test and debug or generate code examples, including with Microsoft's GitHub Copilot. Countless startups are offering AI coding assistants as well, like [JetBrains](#) and [Continue](#), to name just a couple.

Microsoft’s close relationship with OpenAI means that its cloud, Azure, is a popular choice for enterprises that are using LLM models to build their own GenAI apps, too. So, Microsoft’s tentacles into GenAI extend far beyond having Word write documents, having Excel whip up charts, or having Teams transcribe meetings.

At the same time, when the people who are literally talking up AI to sell their own AI products warn that the AI is overhyped, it’s safe to say that AI is overhyped.

“It’s about managing expectations while harnessing AI’s capabilities,” Benioff explained.

Topics: [AI](#) [AI agents](#) [Enterprise](#) [Microsoft](#)

[Microsoft Copilot](#) [Salesforce](#)

[f](#) [X](#) [in](#) [👤](#) [✉](#) [🔗](#)



Julie Bort

Venture Editor | [X](#)

[View Bio](#) >

—— KEEP READING ——

AI

[f](#) [X](#) [in](#) [👤](#) [✉](#) [🔗](#)



OpenAI's Codex is part of a new cohort of agentic coding tools

Russell Brandom — 5:30 AM PDT · May 20, 2025

Last Friday, OpenAI introduced a new coding system called Codex, designed to perform complex programming tasks from natural language commands. Codex moves OpenAI into a new cohort of agentic coding tools that is just beginning to take shape.

IMAGE CREDITS: VGAJIC / GETTY IMAGES

From GitHub's early Copilot to contemporary tools like Cursor and Windsurf, most AI coding assistants operate as an exceptionally intelligent form of autocomplete. The tools generally live in an integrated development environment, and users interact directly with the AI-generated code. The prospect of simply assigning a task and returning when it's finished is largely out of reach.

But these new agentic coding tools, led by products like [Devin](#), [SWE-Agent](#), [OpenHands](#), and OpenAI Codex, are designed to work without users ever having to see the code. The goal is to operate like the manager of an engineering team, assigning issues through workplace systems like Asana or Slack and checking in when a solution has been reached.


TEC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door

For believers in forms of highly capable AI, it's the next logical step in a natural progression of automation taking over more and more software work.

"In the beginning, people just wrote code by pressing every single keystroke," explains Kilian Lieret, a Princeton researcher and member of the SWE-Agent team. "GitHub Copilot was the first product that offered real auto-complete, which is kind of stage two. You're still absolutely in the loop, but sometimes you can take a shortcut."

The goal for agentic systems is to move beyond developer environments entirely, instead presenting coding agents with an issue and leaving them to resolve it on their own. "We pull things back to the management layer, where I just assign a bug report and the bot tries to fix it completely autonomously," says Lieret.

It's an ambitious aim, and so far, it's proven difficult.

- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines
- The best dating apps aren't even dating apps

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

[REGISTER NOW](#)

After Devin became generally available at the end of 2024, it drew [scathing criticism](#) from YouTube pundits, as well as [a more measured critique](#) from an early client at [Answer.AI](#). The overall impression was a familiar one for vibe-coding veterans: with so many errors, overseeing the models takes as much work as doing the task manually. (While Devin’s rollout has been a bit rocky, it hasn’t stopped fundraisers from recognizing the potential — in March, Devin’s parent company, Cognition AI, reportedly [raised hundreds of millions of dollars at a \\$4 billion valuation](#).)

Even supporters of the technology caution against unsupervised vibe-coding, seeing the new coding agents as powerful elements in a human-supervised development process.

“Right now, and I would say, for the foreseeable future, a human has to step in at code review time to look at the code that’s been written,” says Robert Brennan, the CEO of All Hands AI, which maintains OpenHands. “I’ve seen several people work themselves into a mess by just auto-approving every bit of code that the agent writes. It gets out of hand fast.”

Hallucinations are an ongoing problem as well. Brennan recalls one incident in which, when asked about an API that had been released after the OpenHands agent’s training data cutoff, the agent fabricated details of an API that fit the description. All Hands AI says it’s working on systems to catch these hallucinations before they can cause harm, but there isn’t a simple fix.

Arguably the best measure of agentic programming progress is the [SWE-Bench leaderboards](#), where developers can test their models against a set of unresolved issues from open GitHub repositories. OpenHands currently holds the top spot on the verified leaderboard, solving 65.8% of the problem

set. OpenAI claims that one of the models powering Codex, codex-1, can do better, listing a 72.1% score in its announcement — although the score came with a few caveats and hasn't been independently verified.

The concern among many in the tech industry is that high benchmark scores don't necessarily translate to truly hands-off agentic coding. If agentic coders can only solve three out of every four problems, they're going to require significant oversight from human developers — particularly when tackling complex systems with multiple stages.

Like most AI tools, the hope is that improvements to foundation models will come at a steady pace, eventually enabling agentic coding systems to grow into reliable developer tools. But finding ways to manage hallucinations and other reliability issues will be crucial for getting there.

“I think there is a little bit of a sound barrier effect,” Brennan says. “The question is, how much trust can you shift to the agents, so they take more out of your workload at the end of the day?”

Topics: [AI](#) [vibe coding](#)

[f](#) [X](#) [in](#) [📷](#) [✉](#) [🔗](#)

Russell Brandom

✕

Russell is a freelance writer based in New York.

[View Bio](#) >

Loading the next article





[TechCrunch Staff](#)

[Contact Us](#)

[Advertise](#)

[Crunchboard Jobs](#)

[Site Map](#)

[Terms of Service](#)

[Privacy Policy](#)

[RSS Terms of Use](#)

[OpenAI](#)

[Perplexity](#)

[Cisco](#)

[DeepMind](#)

[Windsurf](#)

[Tech Layoffs](#)

[ChatGPT](#)