


Advertisement

STARTUPS



**PlayerZero raises \$15M
to prevent AI agents from
shipping buggy code**



As Silicon Valley races toward a future where AI agents do most of the software programming, a new problem is created: finding the AI-generated bugs before they are put into production. Even OpenAI is [dealing with such issues](#), a former employee has described.

Newly funded startup [PlayerZero](#) has created a solution: use AI agents trained to find and fix problems before the code is put into production, the startup's CEO and sole founder, Animesh Koratana, tells TechCrunch.

Koratana created PlayerZero while he was at the Stanford DAWN lab for machine learning under his adviser and lab founder, Matei Zaharia. Zaharia is, of course, a famed developer and the co-founder of Databricks; he created its foundational technology while working on his own doctorate.

PlayerZero on Wednesday announced that it raised a \$15 million Series A led by Foundation Capital's Ashu Garg, an early Databricks backer. This follows a \$5 million seed led by Green Bay Ventures and several noteworthy angels, including Zaharia, Dropbox CEO Drew Houston, Figma CEO Dylan Field, and Vercel CEO Guillermo Rauch.

During his time at Stanford DAWN, Koratana, now 26, was working on AI model compression technology and "got exposed to language models really early on," he says. He met the developers


DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Instagram takes on Snapchat with new 'Instagram Map'
- Google says hackers stole its customers' data by breaching its Salesforce database
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Spotify raises subscription prices

who crafted some of the first AI coding assistance tools.

It hit him then that “there’s this world in which computers are going to write the code. It’s not going to be humans anymore,” Koratana tells TechCrunch. “What’s the world going to look like at that point?”

He knew before the term “AI slop” was even coined that these agents were going to produce code that broke things just as their human overseers do.

- The uproar over Vogue’s AI-generated ad isn’t just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

That problem would also be exacerbated by so many agents cranking out so much more code than has ever been written before. It won’t always be practical for humans to check all AI-written code for bugs or hallucinations. And the issue becomes even more intense for the large, complex code bases that enterprises rely upon.

PlayerZero trains models “that really deeply understand code bases, and we understand the

way they're built, the way they're architected," Koratana says.

His tech studies the history of an enterprise's bugs, issues, and solutions. When something breaks, his product can then "figure out why and fix it, and then learn from those mistakes to prevent them from ever happening again," Koratana says. He likens his product to an immune system for large code bases.

Landing Zaharia, his adviser, as an angel was a first step to fundraising, but the moment that really validated his idea was when he showed a demo to another famous developer: Rauch. Rauch is the founder of [triple-unicorn developer tool company](#) Vercel and creator of the popular open source JavaScript framework [Next.js](#).

Rauch watched Koratana's demo with interest but skepticism, asking how much of it was "real." Koratana replied that this was code "running in production. Like, this is a real instance. And he was quiet," Koratana says. Then his soon-to-be-angel investor responded, "If you can actually solve this the way that you're imagining, it's a really big deal."

Of course, PlayerZero isn't alone in attempting to solve the AI-generated bug problem. Just last week, Anysphere's Cursor [launched Bugbot](#) to detect coding errors, as just one example.

Still, PlayerZero is already gaining traction for its emphasis on large codebases. While it was conceived for a world where agents are the coders, it is currently being used by several large enterprises that use coding co-pilots. For instance, subscription billing company Zuora is one of the startup's marquee customers. Zuora is using the tech across its engineering teams, including to watchdog its most precious code, its billing systems, it said.



Julie Bort
Venture Editor |

Julie Bort is the Startups/Venture Desk editor for TechCrunch....

[View Bio](#) >

KEEP READING



AI



Microsoft brings OpenAI’s smallest open model to Windows users

Ram Iyer — 5:47 AM PDT · August 6, 2025

DAVID RYDER / BLOOMBERG
GPT model, gpt-oss-20b, available to Windows 11 users via [Windows AI Foundry](#), the tech giant's platform that lets users tap AI features, APIs, and popular open source models on their computers.

The company said in a [blog post](#) that gpt-oss-20b is "tool-savvy and lightweight," adding: "Optimized for agentic tasks like code execution and tool use, it runs efficiently on a range of Windows hardware, with support for more devices coming soon. It's perfect for building autonomous assistants or embedding AI into real-world workflows, even in bandwidth-constrained environments."

Launched on Tuesday, OpenAI's gpt-oss-20b can run on consumer PCs and laptops with at least 16GB of VRAM, which a modern GPU from Nvidia or Radeon would have. OpenAI said the model was trained using [high-compute reinforcement learning](#), which helps it excel at powering AI agents and call tools, such as web search or Python code execution, as part of its chain-of-thought process.

However, the model is text-only, meaning it will not be able to process or generate images and audio like the company's other models.

Also, the model has a habit of hallucinating a lot: gpt-oss-20b hallucinated in response to 53% of questions on PersonQA, OpenAI's in-house benchmark for measuring the accuracy of a model's knowledge about people.

Microsoft says it will bring the model to macOS soon, as well as other devices, though the company did not name any more. Microsoft is also making both gpt-oss-20b and older model gpt-oss-120b available via its hosted Azure AI Foundry platform.

The two new models are also [available on Amazon's AWS](#).



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

Advertisement

Newsletters

See More ↗

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

Topics: [AI](#) [gpt](#) [Microsoft](#) [OpenAI](#) [Windows](#)

Advertisement



Ram Iyer

Editor |

Ram is a financial and tech reporter and editor. He covered North American and European M&A, equity, regulatory...

Advertisement

KEEP READING

GPT-5

flagship model

AI



OpenAI’s GPT-5 is here

Maxwell Zeff — 10:00 AM PDT · August 7, 2025

OpenAI has launched GPT-5, a new flagship AI model that will power the company’s next generation of ChatGPT.

GPT-5, which was released Thursday, is OpenAI’s first “unified” AI model and combines the reasoning abilities of its [o-series](#) of models with the fast responses of its [GPT series](#). The next-generation model signals a new era for ChatGPT — and its creator, OpenAI — pointing to OpenAI’s broader ambitions to develop AI systems that are [more like agents](#) than chatbots.

While [GPT-4](#) enabled AI chatbots to offer smart responses on a wide variety of questions, GPT-5 allows ChatGPT to complete a wide variety of tasks on behalf of users — such as generating software applications, navigating a user’s calendar, or creating research briefs.

With GPT-5, OpenAI has also sought to make ChatGPT simpler to use. Instead of asking users to choose the right settings, GPT-5 comes equipped with a real-time router that decides how to offer the best answer, whether that’s responding to user questions quickly or taking additional time to “think” through answers.




IMAGE CREDITS:OPENAI

During a briefing with reporters, OpenAI CEO Sam Altman claimed GPT-5 is “the best model in the world,” and said it represented a “significant step”

Advertisement

Latest in AI

See More

TRANSPORTATION



Tesla shuts down Dojo, the AI training supercomputer that Musk said would be key to full self-driving

Rebecca Bellan - 6 hours ago

STARTUPS



High costs and thin margins threatening AI coding startups

Marina Temkin - 7 hours ago

AI



OpenAI’s GPT-5 is here

Maxwell Zeff - 12 hours ago

Advertisement

along the company's path to developing AI that can outperform humans at most economically valuable work — that is, artificial general intelligence (AGI).

“Having something like GPT-5 would be pretty much unimaginable at any previous time in history,” said Altman.

Starting Thursday, GPT-5 will be available to all free users of ChatGPT as their default model. OpenAI's VP of ChatGPT, Nick Turley, said this is part of the company's effort to give free users access to an AI reasoning model for the first time. (Previously, the company gated these more advanced models behind a paywall.)

“This is just one of the ways that I'm excited to live the mission, making sure that this stuff actually benefits people,” said Turley on the decision, referencing OpenAI's long-standing [mission](#) to distribute advanced AI to as many people as possible.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

The expectations are high for GPT-5, one of OpenAI's most anticipated product launches since ChatGPT put the company on the map in 2022. Since then, ChatGPT has grown into one of the world's most popular consumer products, reaching more than 700 million users every week — nearly 10% of the globe's population, according to the company.

Many see GPT-5 as a bellwether for AI progress broadly, and the model's reception by Silicon Valley could have profound implications for Big Tech, Wall Street, and policymakers regulating technology. These stakeholders are watching to see if GPT-5 offers a significant jump in AI's capabilities, much like its predecessor, GPT-4, which challenged expectations of what software can do.

GPT-5 offers a slight edge on the competition

OpenAI claims GPT-5 is state-of-the-art in several domains, slightly edging out leading AI models from Anthropic, Google DeepMind, and Elon Musk's xAI on key benchmarks. However, GPT-5 slightly underperforms frontier AI models in other areas.

The company says GPT-5 offers frontier-level performance around coding; Altman said the model specifically excels at spinning up entire software applications on demand, in what's become known

Latest Security
Startups AI
Venture Apps
Apple



Events

Podcasts



12% when using tools. That's slightly less than what it was able to achieve with [Grok 4 Heavy](#), which scored 44.4% on the test.

Humanity's Last Exam (Full Set)*
Expert-level questions across subjects




IMAGE CREDITS: OPENAI

On GPQA Diamond — a test of PhD-level science questions — GPT-5 pro scored 89.4% on its first try, outperforming Claude Opus 4.1, which scored 80.9%, and Grok 4 Heavy, which scored 88.9%.

OpenAI says GPT-5 is better for answering health-related questions. On a test measuring accuracy in AI model responses around healthcare topics, HealthBench Hard Hallucinations, OpenAI says GPT-5 (with thinking) hallucinates just 1.6% of the time. This is far lower than the company's previous GPT-4o and o3 models, which scored 12.9% and 15.8, respectively.

While AI chatbots are not medical professionals, [millions of people](#) are using them for health advice. In response to this phenomenon, the company says GPT-5 is more proactive about flagging potential health concerns and helping users parse medical results.

In addition, OpenAI says GPT-5 is better than other AI models on more difficult-to-measure, [subjective domains](#), such as creative design and writing.

Turley said GPT-5 responds more naturally and exhibits “better taste” than other AI models on creative tasks.

“The vibes of this model are really good,” said Turley.

GPT-5 is also more accurate than OpenAI’s previous models, and the company says it suffers far less from hallucinations — the tendency for AI models to make up information — compared to its o-series models. Hallucinations seemed to be [getting worse](#) in OpenAI’s latest AI reasoning models, such as o3, and OpenAI previously said it didn’t quite understand why it was happening.

In responses to ChatGPT prompts, OpenAI found that GPT-5 (with thinking) hallucinates and responds with incorrect information 4.8% of the time. That’s a significant reduction from o3 and GPT-4o, which score hallucination rates of 22% and 20.6%, respectively, on the test.

On a benchmark measuring an AI model’s agentic ability to complete simulated online tasks, Tau-bench, GPT-5 offers mixed performance. On part of the test measuring an AI’s ability to navigate an airline’s website, GPT-5 scores 63.5%, slightly underperforming o3, which scored 64.8%. On another part of the test measuring AI’s ability to navigate retail websites, GPT-5 scores 81.1%, underperforming Claude Opus 4.1, which scored 82.4%.

OpenAI also says that GPT-5 is safer than its previous models. While AI reasoning models occasionally exhibit a tendency to [scheme against humans](#) or lie to promote their own goals, OpenAI

found that GPT-5 was deceptive at a lower rate than other models.

Alex Beutel, OpenAI’s safety research lead, said reducing deception improves not only the safety of GPT-5, but also the user experience, creating a model that’s more “transparent and honest in ways users can trust.”

Beutel also notes GPT-5 is better at discerning between bad actors who are trying to misuse ChatGPT and users making harmless requests. This results in GPT-5 being able to refuse more unsafe questions, while offering fewer rejections to users seeking harmless information.

Upgrades for consumers and developers

ChatGPT is getting a few user experience upgrades as part of the GPT-5 launch. Users can now select from four new personalities in ChatGPT’s setting: Cynic, Robot, Listener, and Nerd. The company says these will adapt ChatGPT’s responses without requiring users to specifically ask the model to respond in a certain way.

Subscribers to ChatGPT’s \$20-per-month Plus plan get higher usage limits for GPT-5 than free users. Meanwhile, \$200-per-month Pro subscribers will have unlimited access to GPT-5, as well as a souped-up version called GPT-5 Pro that uses additional computational resources to produce better answers. Organizations on OpenAI’s Team, Edu, and Enterprise plans will gain access to GPT-5 as their default model next week.

For developers, GPT-5 is coming to OpenAI’s API in three sizes — gpt-5, gpt-5-mini, and gpt-5-nano — which will spend more or less time “reasoning” through tasks. Developers can also now control

verbosity in the OpenAI API, deciding how long or short an AI model's responses should be.

The base model of GPT-5 will cost developers \$1.25 per million input tokens (roughly 750,000 words, longer than the entire "Lord of the Rings" series) and \$10 per million output tokens.

The launch of GPT-5 comes after a busy week for OpenAI. The company released an open-weight reasoning model, [gpt-oss](#), that developers and enterprises can download for free and run at a fraction of the cost. The open model nearly matched the abilities of OpenAI's previous top models, o3 and o4-mini, but GPT-5 sets a new standard for frontier performance in some areas, such as coding.

However, GPT-5 seems to be roughly on par with other frontier AI models in several areas. Benchmarks, of course, only tell part of the story for any AI model, and it remains to be seen how developers will use GPT-5 in the real world, and whether the model is truly a step above the competition.

Topics: [AI](#) [ChatGPT](#) [gpt-5](#) [OpenAI](#)

[f](#) [X](#) [in](#) [🗨️](#) [✉️](#) [🔗](#)



Maxwell Zeff

Senior AI Reporter | [X](#)

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >




IMAGE CREDITS: JAGMEET SINGH / TECHCRUNCH

AI



Google's AI coding agent Jules is now out of beta

Jagmeet Singh — 9:00 AM PDT · August 6, 2025

Google on Wednesday launched its AI coding agent, Jules, out of beta, just over two months after its [public preview debut in May](#).

Powered by Gemini 2.5 Pro, Jules is an asynchronous, agent-based coding tool that integrates with GitHub, clones codebases into Google Cloud virtual machines, and uses AI to fix or update code while developers focus on other tasks.

Google [initially announced](#) Jules as a Google Labs project in December and made it [available to beta testers](#) through a public preview at its I/O developer conference.

DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Kathy Korevec, director of product at Google Labs, told TechCrunch that the tool's improved stability drove the decision to take it out of beta after receiving hundreds of UI and quality updates during its beta phase.

"The trajectory of where we're going gives us a lot of confidence that Jules is around and going to be around for the long haul," she said.

With the wider rollout, Google introduced structured pricing tiers for Jules, starting with an "introductory access" free plan capped at 15 individual daily tasks and three concurrent ones, down from the 60-task limit during beta. Jules' paid tiers are part of the Google AI Pro and Ultra plans, which are priced at \$19.99 and \$124.99 a month, and offer subscribers 5x and 20x higher limits, respectively.

Korevec noted that Jules' packaging and pricing are based on "real usage" insights gathered over the past couple of months.

"The 60-task cap helped us study how developers use Jules and gave us the information we needed to design the new packaging," she said. "The 15/day is designed to give people a sense of whether Jules will work for them on real project tasks."

Google also updated Jules' privacy policy to be more explicit about how it trains AI. If a repository is public, its data may be used for training, but if it is private, Korevec said that no data is sent.

"We got a little bit of feedback from users that it [the privacy policy] wasn't as clear as we thought it was, and so most of it is just responding to that. We didn't change anything about what we're doing on the training side, but we changed the language," Korevec said.

Most Popular



- Key sections of the US Constitution deleted from government's website
- Instagram takes on Snapchat with new 'Instagram Map'
- Google says hackers stole its customers' data by breaching its Salesforce database
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines

During the beta, Google said that thousands of developers tackled tens of thousands of tasks, resulting in over 140,000 code improvements shared publicly. Initial feedback led the Google Labs team to add new capabilities, including reusing previous setups for faster task execution, integrating with GitHub issues, and supporting multimodal input.




IMAGE CREDITS: JULES / X

The two primary users of Jules so far are the AI enthusiasts and professional developers, Korevec said.

By running asynchronously in a virtual machine, Jules stands apart from top AI coding tools like Cursor, Windsurf, and Lovable, which all work synchronously and require users to watch the output after each prompt.

“Jules operates like an extra set of hands ... you can basically kick off tasks to it, and then you could close your computer and walk away from it if you want and then come back hours later. Jules would have those tasks done for you, versus if you were doing that with a local agent or using a synchronous agent, you would be bound to that session,” Korevec explained.

This week, Jules received a deeper integration with GitHub to [open pull requests automatically](#) — just like it could open branches — and a feature called

[Environment Snapshots](#) to save dependencies and install scripts as a snapshot for faster, more consistent task execution.

From vibe coding to mobile use, beta trials informed Jules' development

Since entering public beta, Jules has logged 2.28 million visits worldwide, 45% of them from mobile devices, per data from market intelligence provider [SimilarWeb](#), reviewed by TechCrunch. India was the top market for traffic, followed by the U.S. and Vietnam.

Google did not share specifics on Jules' user base and its top geographies.

Korevec told TechCrunch that during the beta, the team observed that many people used Jules from traditional vibe-coding tools to either fix bugs that might have been implemented or extend the vibe-coded project to make it more production-ready.

Originally, Jules required users to have an existing codebase. But Google soon realized many potential users — like those trying other AI tools — might want to explore it without one. Korevec said the company quickly enabled Jules to work even with an empty repository. That helped increase its scope and usage.

Google Labs' team also noticed an increasing number of users accessing Jules through their mobile devices. Although the tool does not have a dedicated mobile app, Korevec said users were accessing it through its web app.

“Since it’s a big use case that we’re seeing emerging, we’re absolutely exploring what the

features are that people need on mobile a lot more,” she noted.

Alongside beta testers, Korevec stated that Google already uses Jules to help develop some projects internally, and there is now a “big push” to use the tool on “a lot more projects” at the company.

Topics: [AI](#) [Google](#) [Google Jules](#) [jules](#)

[United States](#)



Jagmeet Singh

Reporter |

Jagmeet covers startups, tech policy-related updates, and all other major tech-centric developments from India for...

[View Bio](#) >

Loading the next article



[TechCrunch Staff](#)

[Contact Us](#)

[Advertise](#)

[Crunchboard Jobs](#)

[Site Map](#)

[Terms of Service](#)

[Privacy Policy](#)

[RSS Terms of Use](#)

[OpenAI](#)

[Perplexity](#)

[Cisco](#)

[DeepMind](#)

[Windsurf](#)

[Tech Layoffs](#)

[ChatGPT](#)