



ЖАҢА ДӘММЕН
ТАНЫСЫҢЫЗ
ВСТРЕЧАЙТЕ
НОВЫЕ ВКУСЫ



APPS



Quora's Poe releases a
developer API with access
to a bouquet of AI models

Quora's AI platform [Poe](#) announced on Thursday that it's releasing an API that allows developers to easily access different models or bots for their own applications.



The API doesn't require a separate fee. Instead, usage is tracked via Poe's existing point-based subscription plans, where each model call costs a set number of points. For instance, low-quality image generation through GPT-4o in a 1:1 aspect ratio and 1024 x 1024 size would cost 328 points.

Instant  App-Creator  Claude-3.7-Sonnet

★ SPONSORED



Design agents are here — Lovart launches the first end-to-end AI canvas for design

Sponsored by Lovart

Today, Poe's plans include the \$4.99 per month plan (10,000 points per day), the \$19.99/mo plan (1 million points per month), the \$49.99/mo plan (2.5 million points per month), the \$99.99/mo plan (5 million points per month), and the \$249.99/mo plan (12.5 million points per month).

Developers will also be able to buy additional points at a rate of \$30 for 1 million tokens. Add-on tokens don't come in a fixed package, so customers can pay any dollar amount to get tokens based on that rate.

Through this API, developers can power tools like Cursor, Cline, Continue, Roo, and any others that work with OpenAI-compatible chat completion APIs. The platform currently provides access to more than 100 models across voice, text, image, and video generation.

These multimodal models include Imagen 4, GPT Image 1, Flux Kontext, Seedream 3.0, Veo 3, Runway Gen 4 Turbo, Kling 2.1, ElevenLabs, and Lyria.

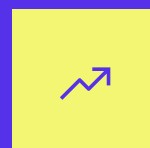
TEC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Instagram takes on Snapchat with new 'Instagram Map'
- Google says hackers stole its customers' data by breaching its Salesforce database
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines

Poe API architecture

Applications




IMAGE CREDITS: QUORA

“Currently, we are working on allowing developers to take a private bot they have built on Poe and use that through an API. Plus, we are thinking about better key management for developers for the API product,” Gareth Jones, Poe’s product lead for creators and developers, told TechCrunch over a call.

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.


San Francisco | October 27-29, 2025

REGISTER NOW

While API provides a way for developers to use different models, Poe also offers tools aimed at consumers. Earlier this year, the company introduced a new way to allow [users to easily create AI-powered apps](#). It also offers templates to build server bots, prompt bots, and image generation bots.

At the moment, developers will need to pick and manage model use manually. Jones said that the company will consider adding budget management functionality in the future, based on developer feedback.

Sponsored by Lovart



SPONSORED

Lovart’s design agent hints at the future: AI-native brands built in minutes



Поставки из России в тенге



export-plus.kz

Открыть

- [f](#)
- [X](#)
- [in](#)
- [t](#)
- [✉](#)
- [🔗](#)



Ivan Mehta

[X](#) [in](#)

Ivan covers global consumer tech developments at TechCrunch. He is based out of India and has previously worked at publications includin...

View Bio >

KEEP READING

AI

- [f](#)
- [X](#)
- [in](#)
- [t](#)
- [✉](#)
- [🔗](#)

OpenAI’s GPT-5 is here

OpenAI has launched GPT-5, a new flagship AI model that will power the company’s next generation of ChatGPT.

GPT-5 Flagship model

GPT-5, which was released Thursday, is OpenAI’s first “unified” AI model and combines the reasoning abilities of its [o-series](#) of models with the fast responses of its [GPT series](#).

The next-generation model signals a new era for ChatGPT — and its creator, OpenAI — pointing to OpenAI’s broader ambitions to develop AI systems that are [more like agents](#) than chatbots.

While [GPT-4](#) enabled AI chatbots to offer smart responses on a wide variety of questions, GPT-5 allows ChatGPT to complete a wide variety of tasks on behalf of users — such as generating software applications, navigating a user’s calendar, or creating research briefs.

IMAGE CREDITS: OPENAI

With GPT-5, OpenAI has also sought to make ChatGPT simpler to use. Instead of asking users to choose the right settings, GPT-5 comes equipped with a real-time router that decides how to offer the best answer, whether that’s responding to user questions quickly or taking additional time to “think” through answers.

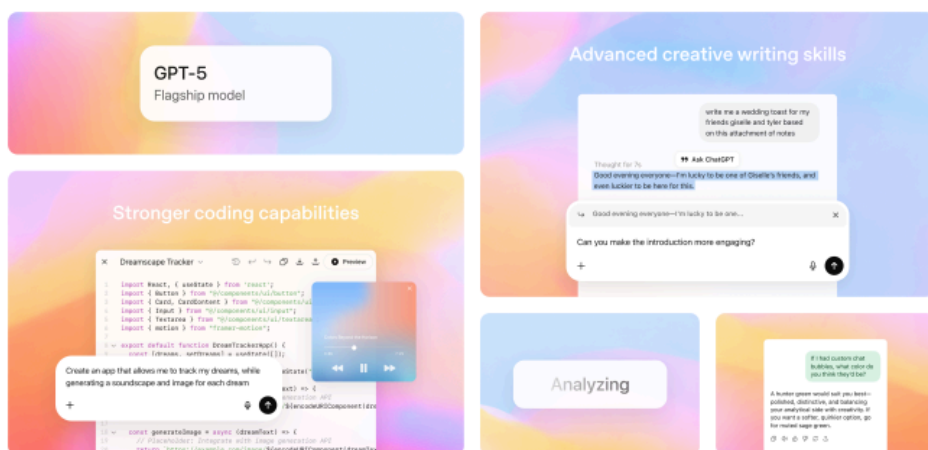


IMAGE CREDITS: OPENAI

During a briefing with reporters, OpenAI CEO Sam Altman claimed GPT-5 is “the best model in the world,” and said it

Advertisement

Newsletters

[See More ↗](#)

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

represented a “significant step” along the company’s path to developing AI that can outperform humans at most economically valuable work — that is, artificial general intelligence (AGI).

“Having something like GPT-5 would be pretty much unimaginable at any previous time in history,” said Altman.

Starting Thursday, GPT-5 will be available to all free users of ChatGPT as their default model. OpenAI’s VP of ChatGPT, Nick Turley, said this is part of the company’s effort to give free users access to an AI reasoning model for the first time. (Previously, the company gated these more advanced models behind a paywall.)

“This is just one of the ways that I’m excited to live the mission, making sure that this stuff actually benefits people,” said Turley on the decision, referencing OpenAI’s long-standing [mission](#) to distribute advanced AI to as many people as possible.

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

The expectations are high for GPT-5, one of OpenAI’s most anticipated product launches since ChatGPT put the company on the map in 2022. Since then, ChatGPT has grown into one of the world’s most popular consumer products, reaching more than 700 million users every week — nearly 10% of the globe’s population, according to the company.

Many see GPT-5 as a bellwether for AI progress broadly, and the model’s reception by Silicon Valley could have profound

implications for Big Tech, Wall Street, and policymakers regulating technology. These stakeholders are watching to see if GPT-5 offers a significant jump in AI's capabilities, much like its predecessor, GPT-4, which challenged expectations of what software can do.

GPT-5 offers a slight edge on the competition

OpenAI claims GPT-5 is state-of-the-art in several domains, slightly edging out leading AI models from Anthropic, Google DeepMind, and Elon Musk's xAI on key benchmarks. However, GPT-5 slightly underperforms frontier AI models in other areas.

★ SPONSORED



Just launched: Lovart, the AI design agent building entire campaigns in 3 minutes

Sponsored by Lovart

The company says GPT-5 offers frontier-level performance around coding; Altman said the model specifically excels at spinning up entire software applications on demand, in what's become known as "vibe coding."

On SWE-bench Verified — a test of real-world coding tasks pulled from GitHub — GPT-5 scores 74.9% on its first attempt. That means GPT-5 just outperforms Anthropic's latest Claude Opus 4.1 model, which scored [74.5%](#), and Google DeepMind's Gemini 2.5 Pro, which scored [59.6%](#).

On Humanity's Last Exam — a difficult test measuring AI model performance across math, humanities, and the natural sciences — a version of GPT-5 with extended reasoning (GPT-5 Pro) scored 42% when using tools. That's slightly less than xAI was able to achieve with [Grok 4 Heavy](#), which scored 44.4% on the test.



In addition, OpenAI says GPT-5 is better than other AI models on more difficult-to-measure, subjective domains, such as creative design and writing. Turley said GPT-5 responds more naturally and exhibits “better taste” than other AI models on creative tasks.

Events

Podcasts

Newsletters

In responses to ChatGPT prompts, OpenAI found that GPT-5 (with thinking) hallucinates and responds with incorrect information 4.8% of the time. That's a significant reduction from o3 and GPT-4o, which score hallucination rates of 22% and 20.6%, respectively, on the test.

On a benchmark measuring an AI model's agentic ability to complete simulated online tasks, Tau-bench, GPT-5 offers mixed performance. On part of the test measuring an AI's ability to navigate an airline's website, GPT-5 scores 63.5%, slightly underperforming o3, which scored 64.8%. On another part of the test measuring AI's ability to navigate retail websites, GPT-5 scores 81.1%, underperforming Claude Opus 4.1, which scored 82.4%.

OpenAI also says that GPT-5 is safer than its previous models. While AI reasoning models occasionally exhibit a tendency to [scheme against humans](#) or lie to promote their own goals, OpenAI found that GPT-5 was deceptive at a lower rate than other models.

Alex Beutel, OpenAI's safety research lead, said reducing deception improves not only the safety of GPT-5, but also the user experience, creating a model that's more "transparent and honest in ways users can trust."

Beutel also notes GPT-5 is better at discerning between bad actors who are trying to misuse ChatGPT and users making harmless requests. This results in GPT-5 being able to refuse more unsafe questions, while offering fewer rejections to users seeking harmless information.

Upgrades for consumers and developers

ChatGPT is getting a few user experience upgrades as part of the GPT-5 launch. Users can now select from four new personalities in ChatGPT's setting: Cynic, Robot, Listener, and

Nerd. The company says these will adapt ChatGPT’s responses without requiring users to specifically ask the model to respond in a certain way.

Subscribers to ChatGPT’s \$20-per-month Plus plan get higher usage limits for GPT-5 than free users. Meanwhile, \$200-per-month Pro subscribers will have unlimited access to GPT-5, as well as a souped-up version called GPT-5 Pro that uses additional computational resources to produce better answers. Organizations on OpenAI’s Team, Edu, and Enterprise plans will gain access to GPT-5 as their default model next week.

★ SPONSORED



Just launched: Lovart, the AI design agent building entire campaigns in 3 minutes

Sponsored by Lovart

For developers, GPT-5 is coming to OpenAI’s API in three sizes — gpt-5, gpt-5-mini, and gpt-5-nano — which will spend more or less time “reasoning” through tasks. Developers can also now control verbosity in the OpenAI API, deciding how long or short an AI model’s responses should be.

The base model of GPT-5 will cost developers \$1.25 per million input tokens (roughly 750,000 words, longer than the entire “Lord of the Rings” series) and \$10 per million output tokens.

The launch of GPT-5 comes after a busy week for OpenAI. The company released an open-weight reasoning model, [gpt-oss](#), that developers and enterprises can download for free and run at a fraction of the cost. The open model nearly matched the abilities of OpenAI’s previous top models, o3 and o4-mini, but GPT-5 sets a new standard for frontier performance in some areas, such as coding.

However, GPT-5 seems to be roughly on par with other frontier AI models in several areas. Benchmarks, of course, only tell part of the story for any AI model, and it remains to be seen how developers will use GPT-5 in the real world, and whether the model is truly a step above the competition.

Topics: [AI](#) [ChatGPT](#) [gpt-5](#) [OpenAI](#)



Maxwell Zeff

Senior AI Reporter | X

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg, and MSNBC, Zeff has covered th...

[View Bio](#) >

KEEP READING



IMAGE CREDITS: JJ JUMOC-CASAS

STARTUPS

Payment platform Lava raises \$5.8M to build digital wallets for the ‘agent-native economy’

Dominic-Madori Davis — 8:45 AM PDT · August 6, 2025

A new startup, [Lava Payments](#), aims to take on payment giants by building a solution for the modern web where AI agents now handle transactions for their customers. The idea came to founder Mitchell Jones after he left his earlier Y Combinator-backed fintech startup, Lendtable, as he began to experiment with AI.

Latest in Startups

[See More](#)

He saw the potential to build out a system that would make using AI and agent payments simpler and more developer-friendly. While experimenting with an AI app and trying to build what he thought was something simple, he realized he quickly spent more than \$400 trying to build a basic form-filling agent.

★ SPONSORED



Lovart's design agent hints at the future: AI-native brands built in minutes

Sponsored by Lovart

"I kept running into the same issue," he told TechCrunch. "I was using the same underlying models and tools again and again, but through different wrappers or platforms." And each time, he had to start a new subscription, re-authenticate, and pay separately, "even though I was already paying for access to the core model."

"That felt fundamentally broken," he continued. "I didn't want to keep rebuying access to the same thing under a different wrapper. What I wanted was a single wallet, one set of credits, and the ability to move between tools and providers without starting over every time so I could pay for what I was using."

He decided to launch Lava Payments as a solution.

Lava is a digital wallet that lets merchants use usage credits to facilitate transactions.

The idea is that one set of credits working across merchants and services makes it easier for autonomous agents to make payments without needing human intervention. It works like this: A merchant can enable the Lava wallet for their customers to use and upload (credits) money to. Once a customer does that, they can take that money and use it at any merchant that also accepts Lava and any of the foundational models, like GPT and Claude, on a "pay as you go basis," Jones said.

So, rather than having to pay for each tool, a user buys a one-time usage credit that AI agents can simply charge as they perform various tasks. No more asking the user to approve transaction after transaction.

STARTUPS



High costs and thin margins threatening AI coding startups

Marina Temkin - 2 hours ago

FINTECH



IPO hopeful Brex scored major win to sell in the EU, plans UK expansion

Julie Bort - 4 hours ago

TECHCRUNCH DISRUPT 2025





Elad Gil — one of tech's sharpest minds — on early bets, breakout growth, and what's coming next at TechCrunch Disrupt 2025

TechCrunch Events

- 9 hours ago

Sponsored by Lovart



★ SPONSORED

Just launched: Lovart, the AI design agent building entire campaigns in 3 minutes

“Without Lava, agents can’t move smoothly through the internet because they constantly get blocked when it comes time to pay,” he said. He used Google as an example, saying every time a person opens Google Maps, they don’t have to pay Google for that map, as they’ve already paid Verizon and AT&T to access the internet.

On Wednesday, the startup announced a \$5.8 million seed round led by Lerer Hippeau.

Others in this space include startups like [Metronome](#).

“We see the world as very interconnected,” Jones said about what makes his product different. “What we’re really focused on is building [for the] agent-native economy.”

Born to a working family in Dayton, Ohio, Jones said his parents always told him the best way to get ahead was to work hard, save money, and get a good education.

“You know, a lot of the things that most people are told,” he recalled, when speaking with TechCrunch.

Jones took that advice to heart. He got a good education (Yale), held some good jobs (Goldman, Meta), and then founded some companies (the fintechs Parable and Lendtable, the latter of which was YC S20).

Jones said he met his lead investors for Lava because he went to high school with [Will McKelvey, now an investor at Lerer Hippeau](#). He said McKelvey has been following his career for a while and always wanted to work together someday, and Lava Payments was that someday.

Others in the round included Harlem Capital, Streamlined Ventures, and Westbound. The fresh capital will be used for hiring, building products, and developing go-to-market strategies.

Overall, Jones is ready for Lava to be the “invisible layer that kind of powers the AI web,” he says, especially as AI agents find themselves more and more in the checkout line.

“We should be enabling agents to move, transact, and build without friction,” he said.

“We want to make sure that AI is something that can be used by every single person, even a kid from Dayton, like myself.”



Lovart’s design agent hints at the future: AI-native brands built in minutes

Sponsored by Lovart

The title of this piece was updated to properly reflect what the company does.

Topics: [AI](#) [Fintech](#) [Lerer Hippeau](#) [payments](#) [Startups](#)




Dominic-Madori Davis
Senior Reporter, Venture | [X](#)

Dominic-Madori Davis is a senior venture capital and startup reporter at TechCrunch. She is based in New York City....

[View Bio](#) >

Loading the next article

• • • ●



- | | | |
|-------------|------------------|--------------|
| TechCrunch | Terms of Service | OpenAI |
| Staff | Privacy Policy | Perplexity |
| Contact Us | RSS Terms of Use | Cisco |
| Advertise | | DeepMind |
| Crunchboard | | Windsurf |
| Jobs | | Tech Layoffs |
| Site Map | | ChatGPT |