


Advertisement

AI



Research leaders urge
tech industry to monitor
AI's 'thoughts'



AI researchers from OpenAI, Google DeepMind, Anthropic, and a broad coalition of companies and nonprofit groups, are calling for deeper investigation into techniques for monitoring the so-called thoughts of AI reasoning models in a [position paper](#) published Tuesday.

A key feature of AI reasoning models, such as [OpenAI's o3](#) and [DeepSeek's R1](#), are their [chains-of-thought](#) or CoTs — an externalized process in which AI models work through problems, similar to how humans use a scratch pad to work through a difficult math question. Reasoning models are a core technology for powering AI agents, and the paper's authors argue that CoT monitoring could be a core method to keep AI agents under control as they become more widespread and capable.

IMAGE CREDITS:

“CoT monitoring presents a valuable addition to safety measures for frontier AI, offering a rare glimpse into how AI agents make decisions,” said the researchers in the position paper. “Yet, there is no guarantee that the current degree of visibility will persist. We encourage the research community and frontier AI developers to make the best use of CoT monitorability and study how it can be preserved.”

The position paper asks leading AI model developers to study what makes CoTs “monitorable” — in other words, what factors can increase or decrease transparency into how AI

TECH DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion

models really arrive at answers. The paper's authors say that CoT monitoring may be a key method for understanding AI reasoning models, but note that it could be fragile, cautioning against any interventions that could reduce their transparency or reliability.

The paper's authors also call on AI model developers to track CoT monitorability and study how the method could one day be implemented as a safety measure.

Notable signatories of the paper include OpenAI chief research officer Mark Chen, Safe Superintelligence CEO Ilya Sutskever, Nobel laureate Geoffrey Hinton, Google DeepMind co-founder Shane Legg, xAI safety adviser Dan Hendrycks, and Thinking Machines co-founder John Schulman. First authors include leaders from the U.K. AI Security Institute and Apollo Research, and other signatories come from METR, Amazon, Meta, and UC Berkeley.

The paper marks a moment of unity among many of the AI industry's leaders in an attempt to boost research around AI safety. It comes at a time when tech companies are caught in a fierce competition — which has led [Meta to poach top researchers](#) from OpenAI, Google DeepMind, and Anthropic with million-dollar offers. Some of the most highly sought-after researchers are those building AI agents and AI reasoning models.

- Your public ChatGPT queries are getting indexed by Google and other search engines
- The best dating apps aren't even dating apps

TE

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the

20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

“We’re at this critical time where we have this new chain-of-thought thing. It seems pretty useful, but it could go away in a few years if people don’t really concentrate on it,” said Bowen Baker, an OpenAI researcher who worked on the paper, in an interview with TechCrunch. “Publishing a position paper like this, to me, is a mechanism to get more research and attention on this topic before that happens.”

OpenAI publicly released a preview of the first AI reasoning model, o1, in September 2024. In the months since, the tech industry was quick to release competitors that exhibit similar capabilities, with some models from Google DeepMind, xAI, and Anthropic showing even more advanced performance on benchmarks.

However, there’s relatively little understood about how AI reasoning models work. While AI labs have excelled at improving the performance of AI in the last year, that hasn’t necessarily translated into a better understanding of how they arrive at their answers.

Anthropic has been one of the industry’s leaders in figuring out how AI models really work — a field called interpretability. Earlier this year, CEO Dario Amodei announced a [commitment to crack open the black box of AI models by 2027](#) and invest more

in interpretability. He called on OpenAI and Google DeepMind to research the topic more, as well.

Early research from Anthropic has indicated that [CoTs may not be a fully reliable indication](#) of how these models arrive at answers. At the same time, OpenAI researchers have said that CoT monitoring could one day be a [reliable way to track alignment and safety](#) in AI models.

The goal of position papers like this is to signal boost and attract more attention to nascent areas of research, such as CoT monitoring. Companies like OpenAI, Google DeepMind, and Anthropic are already researching these topics, but it's possible that this paper will encourage more funding and research into the space.

Topics: [AI](#) [ai safety](#) [Anthropic](#) [OpenAI](#) [research](#)

[TC](#)

[f](#) [X](#) [in](#) [📧](#) [✉](#) [🔗](#)



Maxwell Zeff

Senior AI Reporter | [X](#)

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >

———— KEEP READING ————



Anthropic CEO wants to open the black box of AI models by 2027

Maxwell Zeff — 4:28 PM PDT · April 24, 2025

Anthropic CEO Dario Amodei [published an essay Thursday](#) highlighting how little researchers understand about the inner workings of the world's leading AI models. To address that, Amodei set an ambitious goal for Anthropic to reliably detect most

IMAGE CREDITS:

BENJAMIN GIRETTI/BLOOMBERG / GETTY IMAGES

Newsletters

[See More ↗](#)

- Latest
- Security
- Startups
- AI
- Venture
- Apps
- Apple



- Events
- Podcasts
- Newsletters

for humanity to be totally ignorant of how they work.”

Anthropic is one of the pioneering companies in mechanistic interpretability, a field that aims to open the black box of AI models and understand why they make the decisions they do. Despite the rapid performance improvements of the tech industry’s AI models, we still have relatively little idea how these systems arrive at decisions.

For example, OpenAI recently launched new reasoning AI models, o3 and o4-mini, that perform better on some tasks, but also [hallucinate more than its other models](#). The company doesn’t know why it’s happening.

“When a generative AI system does something, like summarize a financial document, we have no idea, at a specific or precise level, why it makes the choices it does — why it chooses certain words over others, or why it occasionally makes a mistake despite usually being accurate,” Amodei wrote in the essay.

In the essay, Amodei notes that Anthropic co-founder Chris Olah says that AI models are “grown more than they are built.” In other words, AI researchers have found ways to improve AI model intelligence, but they don’t quite know why.

TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch

Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

In the essay, Amodei says it could be dangerous to reach AGI — or as he calls it, “[a country of geniuses in a data center](#)” — without understanding how these models work. In a previous essay, Amodei claimed the tech industry could reach such a milestone by 2026 or 2027, but believes we’re much further out from fully understanding these AI models.

In the long term, Amodei says Anthropic would like to, essentially, conduct “brain scans” or “MRIs” of state-of-the-art AI models. These checkups would help identify a wide range of issues in AI models, including their tendencies to lie or seek power, or other weakness, he says. This could take five to 10 years to achieve, but these measures will be necessary to test and deploy Anthropic’s future AI models, he added.

Anthropic has made a few research breakthroughs that have allowed it to better understand how its AI models work. For example, the company recently found ways to [trace an AI model’s thinking pathways through](#), what the company call, circuits. Anthropic identified one circuit that helps AI models understand which U.S. cities are located in which U.S. states. The company has only found a few of these circuits but estimates there are millions within AI models.

Anthropic has been investing in interpretability research itself and recently made [its first](#)

[investment in a startup](#) working on interpretability. While interpretability is largely seen as a field of safety research today, Amodei notes that, eventually, explaining how AI models arrive at their answers could present a commercial advantage.

In the essay, Amodei called on OpenAI and Google DeepMind to increase their research efforts in the field. Beyond the friendly nudge, Anthropic’s CEO asked for governments to impose “light-touch” regulations to encourage interpretability research, such as requirements for companies to disclose their safety and security practices. In the essay, Amodei also says the U.S. should put export controls on chips to China, in order to limit the likelihood of an out-of-control, global AI race.

Anthropic has always stood out from OpenAI and Google for its focus on safety. While other tech companies pushed back on California’s controversial AI safety bill, SB 1047, [Anthropic issued modest support and recommendations for the bill](#), which would have set safety reporting standards for frontier AI model developers.

In this case, Anthropic seems to be pushing for an industry-wide effort to better understand AI models, not just increasing their capabilities.

Topics: [AI](#) [ai safety](#) [Anthropic](#)



Maxwell Zeff

Senior AI Reporter | 

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >



IMAGE CREDITS: ANTHROPIC

AI



Anthropic hires OpenAI co-founder Durk Kingma

Kyle Wiggers — 10:55 AM PDT · October 1, 2024

Durk Kingma, one of the lesser-known co-founders of OpenAI, today [announced](#) that he'll be joining Anthropic.

In a series of posts on X, Kingma revealed that he'll be working mostly remotely, from the Netherlands (where he's based), but didn't say which Anthropic org he'll be joining — or leading.

Reached for comment, an Anthropic spokesperson pointed to Kingma's posts.

Latest in AI

[See More](#)

AI



Google denies AI search features are killing website traffic

Sarah Perez - 14 hours ago

I'll be working (mostly) from the Netherlands, but will visit the SF bay area as frequently as I can!

— Durk Kingma (@dpkingma) [October 1, 2024](#)

“Anthropic’s approach to AI development resonates significantly with my own beliefs,” Kingma wrote. “[L]ooking forward to contributing to Anthropic’s mission of developing powerful AI systems responsibly. Can’t wait to work with their talented team, including a number of great ex-colleagues from OpenAI and Google, and tackle the challenges ahead!”

Kingma, who has a Ph.D. in machine learning from the University of Amsterdam, spent several years as a doctoral fellow at Google before joining OpenAI’s founding team as a research scientist. At OpenAI, Kingma focused on basic research, leading the algorithms team to develop techniques and methods primarily for generative AI models, including image generators (e.g. [DALL-E 3](#)) and large language models (e.g. [ChatGPT](#)).

In 2018, Kingma left to become a part-time angel investor and advisor for AI startups. He rejoined Google in July of that year, and started at Google Brain, which became one of the tech giant’s premiere AI R&D labs before it [merged with DeepMind](#) in 2023.

Kingma’s hiring is yet another talent coup for Anthropic, which [recruited](#) OpenAI’s former safety lead, Jan Leike, in May, and another OpenAI co-founder, John Schulman, in August. The company made an additional splashy hire in May, [appointing](#) Instagram and Artifact co-founder Mike Krieger as its first head of product.

AI



First impressions of Alexa+, Amazon’s upgraded, AI-powered digital assistant

Sarah Perez - 15 hours ago

AI



Google takes on ChatGPT’s Study Mode with new ‘Guided Learning’ tool in Gemini

Aisha Malik - 15 hours ago

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Anthropic's CEO, Dario Amodei, was once the VP of research at OpenAI, and [reportedly](#) split with the firm after a disagreement over OpenAI's roadmap — namely its growing commercial focus. Amodei brought with him a number of ex-OpenAI employees to launch Anthropic, including OpenAI's former policy lead Jack Clark.

Anthropic has often attempted to position itself as more safety-focused than OpenAI.

Topics: [AI](#) [AI](#) [Anthropic](#) [Durk Kingma](#) [OpenAI](#)

[f](#) [X](#) [in](#) [👤](#) [✉](#) [🔗](#)



Kyle Wiggers

AI Editor | [X](#) [🦋](#) [@](#) [in](#)

Kyle Wiggers was TechCrunch's AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

[View Bio](#) >

