


Advertisement

AI



Some people are
defending Perplexity
after Cloudflare ‘named
and shamed’ it



When Cloudflare accused AI search engine Perplexity of [stealthily scraping websites](#) on Monday while ignoring a site's specific methods to block it, this wasn't a clear-cut case of an AI web crawler gone wild.

Many people came to Perplexity's defense. They argued that Perplexity accessing sites in defiance of the website owner's wishes, while controversial, is acceptable. And this is a controversy that will certainly grow as AI agents flood the internet: Should an agent accessing a website on behalf of its user be treated like a bot? Or like a human making the same request?

IMAGE CREDITS:

KIMBERLY WHITE/GETTY IMAGES FOR TECHCRUNCH

Cloudflare is known for providing anti-bot crawling and other web security services to millions of websites. Essentially, Cloudflare's test case involved setting up a new website with a new domain that had never been crawled by any bot, setting up a robots.txt file that specifically blocked Perplexity's known AI crawling bots, and then asking Perplexity about the website's content. And Perplexity answered the question.

Cloudflare researchers found the AI search engine used "a generic browser intended to impersonate Google Chrome on macOS" when its web crawler itself was blocked. Cloudflare CEO Matthew Prince [posted](#) the research on X, writing, "Some supposedly 'reputable' AI companies act more like


DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion

North Korean hackers. Time to name, shame, and hard block them.”

But many people disagreed with Prince’s assessment that this was actual bad behavior. Those defending Perplexity on sites [like X](#) and [Hacker News](#) pointed out that what Cloudflare seemed to document was the AI accessing a specific public website when its user asked about that specific website.

“If I as a human request a website, then I should be shown the content,” one person on [Hacker News](#) wrote, adding, “why would the LLM accessing the website on my behalf be in a different legal category as my Firefox web browser?”

A Perplexity spokesperson [previously](#) denied to TechCrunch that the bots were the company’s and called Cloudflare’s blog post a sales pitch for Cloudflare. Then on Tuesday, Perplexity [published a blog](#) in its defense (and generally attacking Cloudflare), claiming the behavior was from a third-party service it uses occasionally.

- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

But the crux of Perplexity's post made a similar appeal as its online defenders did.

"The difference between automated crawling and user-driven fetching isn't just technical — it's about who gets to access information on the open web," the post said. "This controversy reveals that Cloudflare's systems are fundamentally inadequate for distinguishing between legitimate AI assistants and actual threats."

Perplexity's accusations aren't exactly fair, either. One argument that Prince and Cloudflare used for calling out Perplexity's methods was that OpenAI doesn't behave in the same way.

"OpenAI is an example of a leading AI company that follows these best practices," [Cloudflare wrote](#). "They respect robots.txt and do not try to evade either a robots.txt directive or a network level block. And ChatGPT Agent is signing http requests using the newly proposed open standard Web Bot Auth."

[Web Bot Auth](#) is a Cloudflare-supported standard being developed by the Internet Engineering Task Force that hopes to create a cryptographic method for identifying AI agent web requests.

The debate comes as bot activity reshapes the internet. As TechCrunch has previously reported, bots seeking to scrape massive amounts of content to train AI models [have become a menace](#), especially to smaller sites.

For the first time in the internet's history, [bot activity is currently outstripping human activity online](#), with AI traffic accounting for over 50%, according to Imperva's Bad Bot report released last month. Most

of that activity is coming from LLMs. But the report also found that malicious bots now make up 37% of all internet traffic. That's activity that includes everything from persistent scraping to unauthorized login attempts.

Until LLMs, the internet generally accepted that websites could and should block most bot activity given how often it was malicious by using CAPTCHAs and other services (such as Cloudflare). Websites also had a clear incentive to work with specific good actors, such as Googlebot, guiding it on what not to index through robots.txt. Google indexed the internet, which sent traffic to sites.

Now, LLMs are eating an increasing amount of that traffic. Gartner predicts [that search engine volume](#) will drop by 25% by 2026. Right now humans tend to click website links from LLMs at the point they are most valuable to the website, which is when they are ready to conduct a transaction.

But if humans [adopt agents](#) as the tech industry predicts they will — to arrange our travel, book our dinner reservations, and shop for us — would websites hurt their business interests by blocking them? The debate on X captured the dilemma perfectly:

“I WANT perplexity to visit any public content on my behalf when I give it a request/task!” wrote [one person](#) in response to Cloudflare calling out Perplexity.

“What if the site owners don't want it? they just want you [to] directly visit the home, see their stuff” [argued another](#), pointing out that the site owner who created the content wants the traffic and potential ad revenue, not to let Perplexity take it.

“This is why I can't see 'agentic browsing' really working — much harder problem than people think.

Most website owners will just block,” [a third](#) predicted.

Topics: [AI](#) [ai agent](#) [ai bots](#) [cloudflare](#) [Perplexity](#)

[web scraping](#).

[f](#) [X](#) [in](#) [reddit](#) [email](#) [link](#)



Julie Bort

Venture Editor | [X](#)

[View Bio](#) >

KEEP READING


SECURITY

[f](#) [X](#) [in](#) [reddit](#) [email](#) [link](#)

Perplexity accused of scraping websites that explicitly blocked AI scraping

Lorenzo Franceschi-Bicchierai

— 8:41 AM PDT · August 4, 2025



AI startup Perplexity is crawling and scraping content from websites that have explicitly indicated they don't want to be scraped, according to internet infrastructure provider Cloudflare.

On Monday, Cloudflare [published research](#) saying it observed the AI startup ignore blocks and hide its crawling and scraping activities. The network infrastructure giant accused Perplexity of obscuring its identity when trying to scrape web pages “in an attempt to circumvent the website’s preferences,” Cloudflare’s researchers wrote.

AI products like those offered by Perplexity rely on gobbling up large amounts of data from the internet, and AI startups have long scraped text, images, and videos from the internet many times without permission to make their products work. In recent times, websites have tried to fight back by using the web standard Robots.txt file, which tells search engines and AI companies which pages can be indexed and which shouldn't, efforts [that have seen mixed results so far](#).

Perplexity appears to be willingly circumventing these blocks by changing its bots’ “user agent,” meaning a signal that identifies a website visitor by their device and version type, as well as changing their autonomous system networks, or ASN, essentially a number that identifies large networks on the internet, according to Cloudflare.

“This activity was observed across tens of thousands of domains and millions of requests per day. We were able to fingerprint this crawler using a combination of machine learning and network signals,” read Cloudflare’s post.

IMAGE CREDITS:

KIMBERLY WHITFIELD GETTY IMAGES FOR TECHCRUNCH

Advertisement

Newsletters

[See More ↗](#)

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Perplexity spokesperson Jesse Dwyer dismissed Cloudflare’s blog post as a “sales pitch,” adding in an email to TechCrunch that the screenshots in the post “show that no content was accessed.” In a follow-up email, Dwyer claimed the bot named in the Cloudflare blog “isn’t even ours.”

Cloudflare said it first noticed the behavior after its customers complained that Perplexity was crawling and scraping their sites, even after they added rules on their Robots file and for specifically blocking Perplexity’s known bots. Cloudflare said it then performed tests to check and confirmed that Perplexity was circumventing these blocks.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

“We observed that Perplexity uses not only their declared user-agent, but also a generic browser intended to impersonate Google Chrome on macOS when their declared crawler was blocked,” according to Cloudflare.

The company also said that it has de-listed Perplexity’s bots from its verified list and added new

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

techniques to block them.

Cloudflare has recently taken a public stance against AI crawlers. Last month, Cloudflare [announced the launch of a marketplace](#) allowing website owners and publishers to charge AI scrapers who visit their sites. Cloudflare's chief executive Matthew Prince [sounded the alarm](#) at the time, saying AI is breaking the business model of the internet, particularly publishers. Last year, Cloudflare also [launched a free tool](#) to prevent bots from scraping websites to train AI.

This is not the first time Perplexity is accused of scraping without authorization.

Last year, news outlets, [such as Wired](#), alleged [Perplexity was plagiarizing their content](#). Weeks later, Perplexity's CEO Aravind Srinivas [was unable](#)

Latest Security
Startups AI
Venture Apps
Apple



Events
Podcasts
Newsletters



Senior Reporter, Cybersecurity |

Lorenzo Franceschi-Bicchierai is a Senior Writer at TechCrunch, where he covers hacking, cybersecurity,...

[View Bio](#) >

KEEP READING




IMAGE CREDITS: GETTY IMAGES

AI



Cloudflare launches a tool to combat AI bots

Kyle Wiggers — 8:49 AM PDT · July 3, 2024

Cloudflare, the publicly traded cloud service provider, has launched a new, free tool to prevent bots from scraping websites hosted on its platform for data to train AI models.

Some AI vendors, including Google, OpenAI and Apple, allow website owners to block the bots they use for data scraping and model training by amending their site's robots.txt, the text file that tells bots which pages they can access on a website. But, as Cloudflare points out in a [post](#) announcing its bot-combating tool, not all AI scrapers respect this.

Latest in AI

[See More](#)

AI



Google denies AI search features are killing website traffic

Sarah Perez - 17 hours ago

AI



First impressions of Alexa+, Amazon's upgraded, AI-powered digital assistant

“Customers don’t want AI bots visiting their websites, and especially those that do so dishonestly,” the company writes on its official blog. “We fear that some AI companies intent on circumventing rules to access content will persistently adapt to evade bot detection.”

So, in an attempt to address the problem, Cloudflare analyzed AI bot and crawler traffic to fine-tune automatic bot detection models. The models consider, among other factors, whether an AI bot might be trying to evade detection by mimicking the appearance and behavior of someone using a web browser.

“When bad actors attempt to crawl websites at scale, they generally use tools and frameworks that we are able to fingerprint,” Cloudflare writes. “Based on these signals, our models [are] able to appropriately flag traffic from evasive AI bots as bots.”

Cloudflare has set up a form for hosts to report suspected AI bots and crawlers and says that it’ll continue to manually blacklist AI bots over time.

The problem of AI bots has come into sharp relief as the generative AI boom fuels the demand for model training data.

Sarah Perez - 19 hours ago

AI



Google takes on ChatGPT’s Study Mode with new ‘Guided Learning’ tool in Gemini

Aisha Malik - 19 hours ago

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your

ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Many sites, wary of AI vendors training models on their content without alerting or compensating them, have opted to block AI scrapers and crawlers. Around 26% of the top 1,000 sites on the web have blocked OpenAI's bot, according to one [study](#); [another](#) found that more than 600 news publishers had blocked the bot.

Blocking isn't a surefire protection, however. As alluded to earlier, some vendors appear to be ignoring standard bot exclusion rules to gain a competitive advantage in the AI race. [AI search engine Perplexity was recently accused](#) of impersonating legitimate visitors to scrape content from websites, and OpenAI and Anthropic are said to have at times [ignored](#) robots.txt rules.

In a [letter to publishers last month](#), content licensing startup TollBit said that, in fact, it sees "many AI agents" ignoring the robots.txt standard.

Tools like Cloudflare's could help — but only if they prove to be accurate in detecting clandestine AI bots. And they *won't* solve the more intractable problem of publishers risking sacrificing referral traffic from AI tools like Google's AI Overviews, which exclude sites from inclusion if they block specific AI crawlers.

Topics: [AI](#) [AI](#) [bots](#) [cloudflare](#)

Advertisement



Kyle Wiggers

AI Editor |

Kyle Wiggers was TechCrunch's AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...


[View Bio](#) >

KEEP READING

AI



Cloudflare's new marketplace will let websites charge AI bots for scraping



Cloudflare **announced plans** on Monday to launch a marketplace in the next year where website owners can sell AI model providers access to scrape their site's content. The marketplace is the final step of Cloudflare CEO Matthew Prince's larger plan to give publishers greater control over how and when AI bots scrape their websites.

"If you don't compensate creators one way or another, then they stop creating, and that's the bit which has to get solved," said Prince in an interview with TechCrunch.

As the first step in its new plan, on Monday, Cloudflare launched free observability tools for customers, called AI Audit. Website owners will get a dashboard to view analytics on why, when, and how often AI models are crawling their sites for information. Cloudflare will also let customers block AI bots from their sites with the click of a button. Website owners can block all web scrapers using AI Audit, or let certain web scrapers through if they have deals or find their scraping beneficial.

A demo of AI Audit shared with TechCrunch showed how website owners can use the tool, which is able to see where each scraper that visits your site comes from, and offers selective windows to see how many times scrapers from OpenAI, Meta, Amazon, and other AI model providers are visiting your site.

Advertisement

TC **DISRUPT**

October 27-29, 2025 | San Francisco


Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service



DEMO OF AI AUDIT.

IMAGE CREDITS: CLOUDFLARE

Cloudflare is trying to address a problem looming over the AI industry: How will smaller publishers survive in the AI era if people go to ChatGPT instead of their website? Today, AI model providers scrape thousands of small websites for information that powers their LLMs. While some larger publishers have struck deals with OpenAI to license content, most websites get nothing, but their content is still fed into popular AI models on a daily basis. That could break the business models for many websites, reducing traffic they desperately need.

Earlier this summer, AI-powered search startup [Perplexity was accused of scraping websites](#) that deliberately indicated they did not want to be crawled using the Robots Exclusion Protocol. Shortly after, Cloudflare released a button to ensure customers could block all AI bots with one click.

“That was out of frustration we were hearing, where people were feeling like their content was being stolen,” said Prince.

- Spotify raises subscription prices
- The uproar over Vogue’s AI-generated ad isn’t just about fashion
- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

Advertisement

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the

insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Some website owners told Business Insider that AI bots were scraping their websites so much, [it felt like a DDoS attack](#) was crippling their servers. Having your website scraped can not only be upsetting, but it can literally run up your cloud bill and impact your service.

But what if you wanted to block Perplexity's bots, but not OpenAI's? Prince tells TechCrunch that Cloudflare's customers are asking for tools that allow them to choose which AI models have access to their sites. Cloudflare's new tools launching today will allow customers to block some AI crawlers, while letting others through.

Even large publishers that have struck licensing deals with OpenAI — such as TIME, Condé Nast, and The Atlantic — have relatively little insight into how much ChatGPT is scraping their websites, according to Prince. Many of them have to accept what OpenAI tells them, but the answer determines if the publishers are getting a good licensing deal or not.

But Cloudflare's marketplace, launching sometime in the next year, aims to allow small publishers to strike deals with AI model providers as well.

“Let's give all of you the ability to do what only Reddit, Quora, and the big publishers of the world

have done previously,” said Prince. “What if we let you set, effectively, a price for accessing and taking your content to ingest into these systems.”

While it’s a bold idea, Cloudflare is not sharing a fully fleshed-out idea of what its marketplace will look like. Prince says websites could charge AI model providers based on the rates at which they’re scraping individual websites, but it’s unclear how much they will really pay. Further, he says websites could charge a monetary price to be scraped, or simply ask AI labs to give them credit. The details are fuzzy.

While AI companies may not initially be excited about paying for content they currently get for free, Cloudflare’s CEO says he thinks this is ultimately good for the AI ecosystem. Prince says the current landscape, where some AI companies don’t pay for content ever, is not sustainable.

Topics: [AI](#) [AI model](#) [ChatGPT](#) [cloudflare](#) [scrapers](#)

[web scraper](#)



Maxwell Zeff

Senior AI Reporter |

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >

Loading the next article





[TechCrunch Staff](#)

[Contact Us](#)

[Advertise](#)

[Crunchboard Jobs](#)

[Site Map](#)

[Terms of Service](#)

[Privacy Policy](#)

[RSS Terms of Use](#)

[OpenAI](#)

[Perplexity](#)

[Cisco](#)

[DeepMind](#)

[Windsurf](#)

[Tech Layoffs](#)

[ChatGPT](#)