


Advertisement

AI



**This startup thinks email
could be the key to
usable AI agents**



AI companies are pushing agents as the next Great Workplace Disruptor, but experts say they're still not ready for prime time. AI agents often struggle to make decisions by themselves, hallucinate frequently, can't cooperate with other agents, fail at confidentiality awareness, and integrate poorly with existing systems.

Industry pioneers like Andrej Karpathy and Ali Ghodsi have said that, like the deployment of autonomous vehicles, humans need to be in the loop in order for agents to succeed.

A startup called Mixus wants to address that with its AI agent platform that not only keeps humans in the workflow, it also lets users interact with agents directly from their email or Slack.

"We're meeting customers where they are today," Elliot Katz, Mixus' co-founder, told TechCrunch. "Where is every person in the workforce today? For the most part, they're on email. And so because we can do this through email, we believe that's a way we can democratize access [to agents]."

If Mixus works reliably, it may solve a big problem in the AI agent space. Most AI companies today either give you a pre-built assistant, à la ChatGPT or Gemini, or developers have to build custom agents using frameworks like LangChain, AutoGen, or CrewAI.

TECH DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion

Mixus launched in beta out of Stanford only in late 2024, but it has already raised \$2.6 million in pre-seed funding and brought on some customers, including clothing store chain Rainbow Shops, as well as others across finance and tech.

The startup says its biggest selling point is ease of use, from how it helps you create agents to how you can interact with them. Users can use text prompts to set up their agents within Mixus' platform via a chat function, or by simply emailing instructions to agent@mixus.com. Then Mixus will build, run, and manage single- or multi-step agents directly from the inbox.

- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

For example, a customer support manager may use a prompt that reads so:

Create an agent that finds all open tasks in Jira in project mixus-dummy, and send me a report with information on all tasks that are overdue. Draft emails to all the assignees who have overdue tasks, and have me review them in the

chat and with simple clear formatting for email (no attachments/docs). Once I verify, send the emails. Run it now. And moving forward, run it every Monday at 7am PST.




IMAGE CREDITS:MIXUS

Katz and his co-founder Shai Magzimof demoed the agents for TechCrunch, showing how to add human verifiers for your agents by simply instructing at which step the agent should ask you for oversight.

For example, they ran an agent to do research on TechCrunch reporters before pitching them. The agent identified and gathered technology news and trends, analyzed the information to identify potential story angles, and compiled a research report summarizing the findings. At the last stage, the agent was directed to send the information to Katz for verification. Once approved, the agent would send the completed research report to Magzimof.

The founders noted that humans can be in the loop as much or as little as required — Magzimof said organizations can set up company-wide rules, like ensuring an email gets checked by a human if it's being sent outside the company.

Bringing other colleagues into the workflow is as easy as tagging them in the chat with an AI agent, or copying them on the email to the agent. That's another standout compared to agents on the markets today: Most models are single-user, and

while Notion AI and Slack allow users to collaborate in shared spaces, they don't let the AI manage conversations and tasks between teammates in real time.

Another core feature of Mixus is its ability to remember files, chats, prompts, and agents.

“We created Spaces so that every team, every person, every group of people can have a shared memory,” Magzimof said. “Then all my agents, all my files, all the people can be in that very specific Space’s memory.”

While ChatGPT and Claude both support memory, their enterprise plans don't yet support shared agent memory across users.

What else can Mixus do?

In our interview, the founders ran through an hour-long demo showing a range of use cases and abilities. Mixus' agents do seem capable, reflecting a high degree of autonomy and memory that places the company toward the higher end of the AI agent spectrum. That is, *if* the product works as reliably as it did in the demo.

Like other agents, Mixus can integrate with other tools, from Gmail to Jira, and users can trigger agents to run immediately or on a schedule. Agents can run and edit documents or spreadsheets inline — similar to ChatGPT, Microsoft Copilot, and Google Gemini, but those are often limited to sandboxed environments.

Mixus also lets agents autonomously navigate organizational context — like figuring out who in an organization owns a task by looking through Jira tickets.

Built on a combination of Anthropic’s Claude 4 and OpenAI’s o3, Mixus agents also have access to the web, which Magzimof says can be tapped for tasks like live research or monitoring. He described it as “Google Alerts on steroids.”

Taken together, Mixus appears to be less of a productivity tool and more like a tireless digital colleague — yet another ambitious attempt to reimagine AI as a collaborator. If it works as advertised, your next “coworker” might not be human, but it might get through your inbox faster than you do.

Correction: This article was updated to reflect how much Mixus has raised to date.

Got a sensitive tip or confidential documents? We’re reporting on the inner workings of the AI industry — from the companies shaping its future to the people impacted by their decisions. Reach out to Rebecca Bellan at rebecca.bellan@techcrunch.com and Maxwell Zeff at maxwell.zeff@techcrunch.com. For secure communication, you can contact us via Signal at @rebeccabellan.491 and @mzeff.88.

Topics: [AI](#) [AI agents](#) [mixus](#) [Startups](#)



Rebecca Bellan

Senior Reporter |

Rebecca Bellan is a senior reporter at TechCrunch, where she covers Tesla and Elon Musk’s broader empire,...

[View Bio](#) >



IMAGE CREDITS: CONVERGENCE AI

AI



Convergence AI played with agents ‘for years’ until raising \$12M to give them long-term memory

Mike Butcher — 10:04 AM PDT · September 25, 2024

Marvin Purtorab, CEO of new startup [Convergence](#), and co-founder Andy Toulis (CTO) met while they both worked at Shopify on recommender systems and AI assistants. They both joined Cohere, the AI platform for enterprises, before coming up with their startup idea only a few months ago. “I mean, to be honest, we left [Cohere] because ... we had been playing with agents for, like, a few years, like three or four years now,” Purtorab said. They’d realized agents, as a concept, had been too early a few years before.

Advertisement

Hiring out of Google DeepMind, Meta, OpenAI, and PolyAI, the pair have built out a product in barely a few months, garnering significant interest from investors.

As of today, most agents are designed for specific workflows. Convergence’s “Proxy” agent will work across a number of tasks, with the idea that it acquires skills in the same way a human might by giving it so-called “long-term memory.” That’s via what’s been dubbed by some as “Large Meta Learning Models” (LMLM). These are trained to acquire the skill of learning on their own.

Convergence raised a \$12 million pre-seed round led by Balderton Capital. Salesforce Ventures and Shopify Ventures also participated in the round, and will be deployed to develop novel models that power Proxy assistants.

In a statement, James Wise, partner at Balderton Capital, said: “Few people have the experience and skill that Marvin and Andy have, which makes them well-placed to take on the complex technical challenge of a product like Proxy.”

Human users get paired with Proxy agents, which can learn tasks and workflows, freeing the human worker up to do more decision-making than the “grunt work.”

Purtorab said: “If you look at the landscape right now, there’s a lot of companies building these, like ‘narrow’ agents: a sales agent, or an HR agent, or a finance ops agent. Our view is to take a different approach. We’re trying to build the foundations for the first general class of agents that can, depending on the user, delineate into any type of agent you need that does the things for you that you, personally, in your job don’t want to do. And we think that’s kind of a better approach, because we just don’t really see the world of the future of

Newsletters

[See More ↗](#)

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

Advertisement

everyone having 1,000 of these different little tools. Over time, things are going to consolidate, and I think that's where we will come in."



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

[REGISTER NOW](#)

The strategy is to create consumer agents, but then use those to inform how they train agents in the enterprise. "Consumers are using it more for, like, online shopping like groceries and recipes. On the enterprise side, people are using it to do sales ops. Like, enter a lot of stuff into Salesforce, and keep track of it. Or talent operations, like, keeping track of applicants for a job."

So are they using consumers as sort of guinea pigs for the enterprise? "That's would be one way to phrase it. Consumers have a lot of simpler and much broader kind of use cases and it helps us get feedback quicker, right?" Purtorab said.

With all that said, what do they make of the kinds of announcements that Salesforce is making, around agents? "I was at Dreamforce last week, and I think that it's the right direction. But I do think that they're

also going much more in the direction of, like these very narrow agents that focus on one specific task... Right now you don't have time to wait another six to 12 months for models to develop a bit better."

Right now the company is in closed beta with testers, but it should open up shortly.

Topics: [AI](#) [AI agents](#) [Balderton Capital](#) [convergence](#)



Mike Butcher

Editor At Large |

Mike Butcher (M.B.E.), formerly Editor-at-large of TechCrunch, has written for UK national newspapers and...

[View Bio](#) >

———— KEEP READING ————


AI



OpenAI launches a
general purpose agent in
ChatGPT

Latest
Startups
Venture
Apple

Security
AI
Apps



Events
Podcasts
Newsletters

The tool, called ChatGPT agent, combines several capabilities from OpenAI's previous agentic tools, including [Operator's ability to click around on websites](#), as well as [Deep Research's ability to synthesize information](#) from dozens of websites into a concise research report. OpenAI says users will be able to interact with the agent simply by prompting ChatGPT in natural language.

IMAGE CREDITS: BRYCE DURBIN / TECHCRUNCH


ChatGPT agent is rolling out on Thursday to subscribers to OpenAI's Pro, Plus, and Team plans. To activate the tool, users can select "agent mode" in ChatGPT's dropdown menu of tools.

ChatGPT can now do work for you using its own computer.

Introducing ChatGPT agent—a unified agentic system combining Operator's action-taking remote browser, deep research's web synthesis, and ChatGPT's conversational strengths.

pic.twitter.com/7uN2Nc6nBQ

— OpenAI (@OpenAI) [July 17, 2025](#)



Sarah Perez - 1 hour ago

AI



First impressions of Alexa+, Amazon's upgraded, AI-powered digital assistant

Sarah Perez - 3 hours ago

AI



Google takes on ChatGPT's Study Mode with new 'Guided Learning' tool in Gemini

Aisha Malik - 3 hours ago

The launch of ChatGPT agent represents OpenAI's boldest attempt yet to turn ChatGPT into an agentic product that can take actions and offload tasks for users, rather than just answering questions. In recent years, Silicon Valley companies including OpenAI, Google, and Perplexity have unveiled dozens of AI agents that have promised to do just that. However, these [early version AI agents have proven to struggle with complex tasks](#), and they seem less compelling as products than the ultimate vision tech executives pitch around AI agents.

That said, OpenAI says ChatGPT agent is far more capable than its previous offerings.

The company's new agent can access [ChatGPT connectors](#), allowing users to connect apps like Gmail and GitHub so that the agent can find relevant information to your prompts. OpenAI says ChatGPT agent has access to a terminal, and it can use APIs to access certain apps.

OpenAI suggests that users can tap ChatGPT agent to “plan and buy ingredients to make Japanese breakfast for four,” as well as “analyze three competitors and create a slide deck.” These kinds of capabilities requires ChatGPT agent to parse through websites, plan a course of action, and use tools — much more complicated tasks than OpenAI has previously tried to tackle with agents.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch

Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

The model underlying ChatGPT agent offers state-of-the-art performance on several benchmarks, according to OpenAI.

The company says the ChatGPT agent model scores 41.6% on Humanity’s Last Exam (pass@1), a difficult test made up of thousands of questions across more than one hundred subjects. That’s roughly double what OpenAI’s o3 and o4-mini scored on the test.

On FrontierMath, one of the hardest known math benchmarks, OpenAI says ChatGPT agent scores 27.4% when it has access to tools, such as a terminal for code execution. The previous state-of-the-art score comes from o4-mini, which scored just 6.3%.

OpenAI notes that it developed ChatGPT agent with safety in mind, largely because the product presents some newfound capabilities that could make it more dangerous in the hands of a bad actor. OpenAI has previously warned that agentic models could present more dangerous capabilities.

In a [safety report](#) for ChatGPT agent, OpenAI says it’s designated the model as “high capability” in biological and chemical weapon domains, which is defined in OpenAI’s Preparedness Framework as a model with the ability to “amplify existing pathways to severe harm.” OpenAI notes that it does not have direct evidence of this, but it’s decided to take a

precautionary approach and activate new safeguards to mitigate these risks.

The new safeguards for ChatGPT agent include a monitor that works in real time as users interact with the product. OpenAI says it runs a classifier across every prompt entered into ChatGPT agent, determining whether the request is related to biology. If so, OpenAI runs ChatGPT agent's response through a second monitor that determines whether the content could be used to evoke a biological threat.

OpenAI also says it disabled ChatGPT's memory feature for this agent to prevent misuse. In other parts of ChatGPT, OpenAI's memory feature allows the chatbot to reference information from previous user chats. However, OpenAI says bad actors could use the feature in ChatGPT agent to exfiltrate sensitive data through prompt injection attacks. The company says it may revisit adding the feature in the future, however.

While ChatGPT agent sounds impressive, it remains to be seen how capable it truly is in the real world. Until now, agent technology has proven relatively brittle when interacting with the real world. That said, OpenAI says it's developed a more capable model that's able to deliver on the promise of AI agents.

This story was updated with more information.

Topics: [agents](#) [AI](#) [ChatGPT](#) [OpenAI](#)



Maxwell Zeff

Senior AI Reporter |

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

View Bio >

KEEP READING



IMAGE CREDITS: BRYCE DURBIN / TECHCRUNCH

AI



OpenAI launches Operator, an AI agent that performs tasks autonomously

Maxwell Zeff — 10:00 AM PST · January 23, 2025

OpenAI CEO Sam Altman kicked off this year by saying in a blog post that 2025 would [be big for AI agents](#), tools that can automate tasks and take actions on your behalf.

Now, we’re seeing OpenAI’s first real attempt.

TECH DISRUPT


October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation

OpenAI announced on Thursday that it is launching a research preview of Operator, a general-purpose AI agent that can take control of a web browser and independently perform certain actions. [Operator is coming to U.S. users on ChatGPT's \\$200 Pro subscription plan first](#). OpenAI says it plans to roll this feature out to more users in its Plus, Team, and Enterprise tiers eventually.

“[Operator] will be [in] other countries soon,” OpenAI CEO Sam Altman said during a livestream Thursday. “Europe will, unfortunately, take a while.”

This initial research preview is available through [operator.chatgpt.com](#), but soon, OpenAI says it wants to integrate Operator into all of its ChatGPT clients.



OPERATOR'S HOME SCREEN.

IMAGE CREDITS:OPENAI

Operator promises to automate tasks such as booking travel accommodations, making restaurant reservations, and shopping online, according to OpenAI. There are several task categories users can choose from within the Operator interface, including shopping, delivery, dining, and travel — all of which enable different kinds of automation.

When ChatGPT users activate Operator, a small window will pop up showing a dedicated web browser that the agent uses to complete tasks, along with explanations of specific actions the

charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines
- The best dating apps aren't even dating apps

agent is performing. Users can still take control of their screen while Operator is working, as Operator uses its own dedicated browser.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

[REGISTER NOW](#)

OpenAI says that Operator is powered by a Computer-Using Agent model, or CUA, that combines the vision capabilities of the company's [GPT-4o](#) model with reasoning abilities from OpenAI's more advanced models. The CUA is trained to interact with the front-end of websites, meaning it doesn't need to use developer-facing APIs to tap into different services.

In other words, the CUA can use buttons, navigate menus, and fill out forms on a web page much like a human would.

OpenAI says it's collaborating with companies like DoorDash, eBay, Instacart, Priceline, StubHub, and Uber to ensure that Operator respects these businesses' terms of service agreements.




IMAGE CREDITS: OPENAI

“The CUA model is trained to ask for user confirmation before finalizing tasks with external side effects, for example before submitting an order, sending an email, etc., so that the user can double-check the model’s work before it becomes permanent,” OpenAI writes in materials provided to TechCrunch. “[It] has already proven useful in a variety of cases, and we aim to extend that reliability across a wider range of tasks.”

But OpenAI warns the CUA isn’t perfect. The company says it “[doesn’t] expect [the] CUA to perform reliably in all scenarios just yet.”

“Currently, Operator cannot reliably handle many complex or specialized tasks,” OpenAI [adds](#) in a support document, “such as creating detailed slideshows, managing intricate calendar systems, or interacting with highly customized or non-standard web interfaces.

Out of an abundance of caution, OpenAI is also requiring supervision for some tasks, like banking transactions, the CUA and Operator could perform mostly on their own. Users will need to take over to put in credit card information, for example. OpenAI says that Operator doesn’t collect or screenshot any data.

“On particularly sensitive websites, such as email, Operator requires active user supervision, ensuring users can directly catch and address any potential mistakes the model might make,” OpenAI says in its support materials.

This limits the usefulness of Operator, to be sure — but also ensures that the agent doesn’t hallucinate and, say, spend your mortgage payment on accent chairs. Google took a similar approach with its [Project Mariner](#) AI agent, which also doesn’t fill in info like credit card numbers.

Limitations

Operator has a few limitations worth noting.

There are rate limits — both daily and task-dependent. OpenAI says that Operator can perform multiple tasks at once, but that there are “dynamic limits” on this. There is also an overall usage limit that resets daily.

At this release stage, Operator will also refuse to perform tasks outright for security reasons, like sending emails (despite the fact that the CUA is capable of this) and deleting calendar events. OpenAI says this will change in the future, but gives no ETA.

Operator may also get “stuck” if it runs into a particularly complex interface, password field, or CAPTCHA check. It’ll ask the user to take over when this occurs, OpenAI says.

An agentic future

OpenAI has been rather slow to develop an AI agent compared to rivals (see: agents from [Rabbit](#),

[Google](#), and [Anthropic](#)), which may have something to do with the safety risks around the technology.

When an AI system can take actions on the web, it opens the door to much more dangerous use cases from nefarious actors. You could automate AI agents to orchestrate phishing scams or DDoS attacks, or have them snatch up tickets to a concert before anyone else could. Especially for a tool as widely used as ChatGPT, it's important OpenAI take steps to prevent those sorts of exploits.

OpenAI seems to think Operator is safe enough to release in its current form, at least as a research preview.

“Operator employs tools that seek to limit the model’s susceptibility to malicious prompts, hidden instructions, and phishing attempts,” OpenAI [explains](#) on its website. “A monitoring system pauses execution if suspicious activity is detected, while automated and human-reviewed pipelines continuously update safeguards.”

Operator is OpenAI’s boldest attempt yet at creating an AI agent. Last week, [OpenAI released Tasks](#), giving ChatGPT simple automation features such as the ability to set reminders and schedule prompts to run at a set time every day.

Tasks gave ChatGPT users some familiar, but necessary, features to make ChatGPT as practical to use as Siri or Alexa. However, Operator shows off capabilities that the previous generation of virtual assistants could never do.

AI agents have been pitched as the next big thing in AI after ChatGPT: a new technology that will change how people use the internet and their PCs. Instead of simply delivering and processing information,

agents can — in theory — take actions and actually do things.

With the release of OpenAI’s first concrete take on agents, it’ll soon become clear just how realistic this vision is.

Topics: [agent](#) [AI](#) [AI](#) [ai agent](#) [ChatGPT](#) [OpenAI](#)

[operator](#)

Advertisement



Maxwell Zeff

Senior AI Reporter |

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >

Loading the next article



[TechCrunch](#)

[Staff](#)

[Contact Us](#)

[Terms of](#)

[Service](#)

[Privacy](#)

[Policy](#)

[OpenAI](#)

[Perplexity](#)

[Cisco](#)

[DeepMind](#)



- Advertise
- Crunchboard
- Jobs
- Site Map
- RSS
- Terms of Use
- Windsurf
- Tech Layoffs
- ChatGPT