

Advertisement

AI



Trump's 'anti-woke AI' order could reshape how US tech companies train their models



When DeepSeek, Alibaba, and other Chinese firms released their AI models, Western researchers quickly noticed they sidestepped questions critical of the Chinese Communist Party. U.S. officials later [confirmed](#) that these tools are engineered to reflect Beijing's talking points, raising concerns about censorship and bias.

American AI leaders like OpenAI have pointed to this as justification for advancing their tech quickly, without too much regulation or oversight. As OpenAI's chief global affairs officer Chris Lehane wrote in a [LinkedIn post](#) last month, there is a contest between "US-led democratic AI and Communist-led China's autocratic AI."

ANDREW CARALLERO, REYNOLDS APF / GETTY IMAGES

[An executive order](#) signed Wednesday by President Donald Trump that bans "woke AI" and AI models that aren't "ideologically neutral" from government contracts could disrupt that balance.

The order calls out diversity, equity, and inclusion (DEI), calling it a "pervasive and destructive" ideology that can "distort the quality and accuracy of the output." Specifically, the order refers to information about race or sex, manipulation of racial or sexual representation, critical race theory, transgenderism, unconscious bias, intersectionality, and systemic racism.

Experts warn it could create a chilling effect on developers who may feel pressure to align model

TEC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion

outputs and datasets with White House rhetoric to secure federal dollars for their cash-burning businesses.

The order comes the same day the White House published [Trump's "AI Action Plan,"](#) which shifts national priorities away from societal risk and focuses instead on building out AI infrastructure, cutting red tape for tech companies, shoring up national security, and competing with China.

The order instructs the director of the Office of Management and Budget along with the administrator for Federal Procurement Policy, the administrator of General Services, and the director of the Office of Science and Technology Policy to issue guidance to other agencies on how to comply.

- Your public ChatGPT queries are getting indexed by Google and other search engines
- The best dating apps aren't even dating apps

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

"Once and for all, we are getting rid of woke," Trump said Wednesday during an AI event hosted by the All-In Podcast and Hill & Valley Forum. "I will be signing an order banning the federal government

from procuring AI technology that has been infused with partisan bias or ideological agendas, such as critical race theory, which is ridiculous. And from now on the U.S. government will deal only with AI that pursues truth, fairness, and strict impartiality.”

Determining what is impartial or objective is one of many challenges to the order.

Philip Seargeant, senior lecturer in applied linguistics at the Open University, told TechCrunch that nothing can ever be objective.

“One of the fundamental tenets of sociolinguistics is that language is never neutral,” Seargeant said. “So the idea that you can ever get pure objectivity is a fantasy.”

On top of that, the Trump administration’s ideology doesn’t reflect the beliefs and values of all Americans. Trump has repeatedly sought to eliminate funding for climate initiatives, education, public broadcasting, research, social service grants, community and agricultural support programs, and gender-affirming care, often framing these initiatives as examples of “woke” or politically biased government spending.

As Rumman Chowdhury, a data scientist, CEO of the tech nonprofit Humane Intelligence, and former U.S. science envoy for AI, put it, “Anything [the Trump administration doesn’t] like is immediately tossed into this pejorative pile of woke.”

The definitions of “truth-seeking” and “ideological neutrality” in the order published Wednesday are vague in some ways and specific in others. While “truth-seeking” is defined as LLMs that “prioritize historical accuracy, scientific inquiry, and objectivity,” “ideological neutrality” is defined as LLMs that are “neutral, nonpartisan tools that do

not manipulate responses in favor of ideological dogmas such as DEI.”

Those definitions leave room for broad interpretation, as well as potential pressure. AI companies have pushed for fewer constraints on how they operate. And while an executive order doesn’t carry the force of legislation, frontier AI firms could still find themselves subject to the shifting priorities of the administration’s political agenda.

Last week, OpenAI, Anthropic, Google, and xAI [signed contracts](#) with the Department of Defense to receive up to \$200 million each to develop agentic AI workflows that address critical national security challenges.

It’s unclear which of these companies is best positioned to gain from the woke AI ban, or if they will comply.

TechCrunch has reached out to each of them and will update this article if we hear back.

Despite displaying biases of its own, xAI may be the most aligned with the order — at least at this early stage. Elon Musk has positioned Grok, xAI’s chatbot, as the ultimate anti-woke, “less biased,” truthseeker. Grok’s system prompts have directed it to avoid deferring to mainstream authorities and media, to seek contrarian information even if it’s politically incorrect, and to even [reference Musk’s own views](#) on controversial topics. In recent months, Grok has even spouted [antisemitic comments](#) and praised Hitler on X, among other hateful, racist, and misogynistic posts.

Mark Lemley, a law professor at Stanford University, told TechCrunch the executive order is “clearly intended as viewpoint discrimination, since [the

government] just signed a contract with Grok, aka ‘MechaHitler.’”

Alongside xAI’s DOD funding, the company announced that “[Grok for Government](#)” had been added to the General Services Administration schedule, meaning that xAI products are now available for purchase across every government office and agency.

“The right question is this: would they ban Grok, the AI they just signed a large contract with, because it has been deliberately engineered to give politically charged answers?” Lemley said in an email interview. “If not, it is clearly designed to discriminate against a particular viewpoint.”

As Grok’s own system prompts have shown, model outputs can be a reflection of both the people building the technology and the data the AI is trained on. In some cases, an overabundance of caution among developers and AI trained on internet content that promotes values like inclusivity have led to distorted model outputs. Google, for example, last year [came under fire](#) after its Gemini chatbot showed a black George Washington and racially diverse Nazis — which Trump’s order calls out as an example of DEI-infected AI models.

Chowdhury says her biggest fear with this executive order is that AI companies will actively rework training data to tow the party line. She pointed to [statements](#) from Musk a few weeks prior to launching Grok 4, saying that xAI would use the new model and its advanced reasoning capabilities to “rewrite the entire corpus of human knowledge, adding missing information and deleting errors. Then retrain on that.”

This would ostensibly put Musk into the position of judging what is true, which could have huge

downstream implications for how information is accessed.

Of course, companies have been making judgment calls about what information is seen and not seen since the dawn of the internet.

Conservative David Sacks — the entrepreneur and investor who Trump appointed as AI czar — has been outspoken about his concerns around “woke AI” on the All-In Podcast, which [co-hosted](#) Trump’s day of AI announcements. Sacks has accused the creators of prominent AI products of infusing them with left-wing values, framing his arguments as a defense of free speech, and a warning against a trend toward centralized ideological control in digital platforms.

The problem, experts say, is that there is no one truth. Achieving unbiased or neutral results is impossible, especially in today’s world where even facts are politicized.

“If the results that an AI produces say that climate science is correct, is that left wing bias?”

Sergeant said. “Some people say you need to give both sides of the argument to be objective, even if one side of the argument has no status to it.”

*Got a sensitive tip or confidential documents?
We’re reporting on the inner workings of the AI industry — from the companies shaping its future to the people impacted by their decisions. Reach out to Rebecca Bellan at rebecca.bellan@techcrunch.com and Maxwell Zeff at maxwell.zeff@techcrunch.com. For secure communication, you can contact us via Signal at @rebeccabellan.491 and @mzeff.88.*

Topics: [AI](#) [Government & Policy](#) [Trump Administration](#)




Rebecca Bellan

Senior Reporter | X reddit

Rebecca Bellan is a senior reporter at TechCrunch, where she covers Tesla and Elon Musk's broader empire,...

View Bio >

KEEP READING



AI

A dev built a test to see how AI chatbots respond to controversial topics

16.1.2023 5:00 AM PST by Rebecca Bellan

- Latest
- Startups
- Venture
- Apple
- Security
- AI
- Apps

Events

Podcasts

Newsletters

AI companies have been focusing on fine-tuning how their models handle certain topics as [some White House allies accuse](#) popular chatbots of being overly “woke.” Many of President Donald Trump’s close confidants, such as Elon Musk and crypto and AI “czar” David Sacks, have alleged that chatbots [censor conservative views](#).

Although none of these AI companies have responded to the allegations directly, [several](#) have pledged to adjust their models so that they refuse to answer contentious questions less often. For example, [for its latest crop of Llama models](#), Meta said it tuned the models not to endorse “some views over others,” and to reply to more “debated” political prompts.

SpeechMap’s developer, who goes by the username “[xlr8harder](#)” on X, said they were motivated to help inform the debate about what models should, and shouldn’t, do.

“I think these are the kinds of discussions that should happen in public, not just inside corporate headquarters,” xlr8harder told TechCrunch via email. “That’s why I built the site to let anyone explore the data themselves.”

SpeechMap uses AI models to judge whether other models comply with a given set of test prompts. The prompts touch on a range of subjects, from politics to historical narratives and national symbols. SpeechMap records whether models “completely” satisfy a request (i.e. answer it without hedging),

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).



give “evasive” answers, or outright decline to respond.

Xlr8harder acknowledges that the test has flaws, like “noise” due to model provider errors. It’s also possible the “judge” models contain biases that could influence the results.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.


San Francisco | October 27-29, 2025

[REGISTER NOW](#)

But assuming the project was created in good faith and the data is accurate, SpeechMap reveals some interesting trends.

For instance, OpenAI’s models have, over time, increasingly refused to answer prompts related to politics, according to SpeechMap. The company’s latest models, the [GPT-4.1](#) family, are slightly more permissive, but they’re still a step down from one of OpenAI’s releases last year.

OpenAI said in February it would [tune future models](#) to not take an editorial stance, and to offer multiple perspectives on controversial subjects — all in an effort to make its models appear more “neutral.”



OPENAI MODEL PERFORMANCE ON SPEECHMAP OVER TIME

IMAGE CREDITS:OPENAI

By far the most permissive model of the bunch is [Grok 3](#), developed by Elon Musk’s AI startup xAI, according to SpeechMap’s benchmarking. Grok 3 powers a number of features on X, including the chatbot Grok.

Grok 3 responds to 96.2% of SpeechMap’s test prompts, compared with the global average “compliance rate” of 71.3%.

“While OpenAI’s recent models have become less permissive over time, especially on politically sensitive prompts, xAI is moving in the opposite direction,” said xlr8harder.

When Musk announced Grok roughly two years ago, he pitched the AI model as edgy, unfiltered, and anti-“woke” — in general, willing to answer controversial questions other AI systems won’t. He delivered on some of that promise. Told to be vulgar, for example, Grok and Grok 2 would happily oblige, spewing colorful language you likely wouldn’t hear from [ChatGPT](#).

But Grok models prior to Grok 3 [hedged](#) on political subjects and wouldn’t cross [certain boundaries](#). In fact, [one study](#) found that Grok leaned to the

political left on topics like transgender rights, diversity programs, and inequality.

Musk has blamed that behavior on Grok’s training data — public web pages — and [pledged](#) to “shift Grok closer to politically neutral.” Short of high-profile mistakes like [briefly censoring unflattering mentions of President Donald Trump and Musk](#), it seems he might’ve achieved that goal.

Topics: [AI](#) [ai censorship](#) [AI chatbots](#) [Grok](#) [Llama](#)

[OpenAI](#) [SpeechMap](#)

[f](#) [X](#) [in](#) [🐼](#) [✉](#) [🔗](#)

 **Kyle Wiggers**
AI Editor | [X](#) [🐼](#) [📧](#) [in](#)

Kyle Wiggers was TechCrunch’s AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

[View Bio](#) >


———— KEEP READING ————

AI

[f](#) [X](#) [in](#) [🐼](#) [✉](#) [🔗](#)

Elon Musk's xAI releases its latest flagship model, Grok 3

Kyle Wiggers — 7:46 PM PST · February 17, 2025



Elon Musk's AI company, xAI, late on Monday released its latest flagship AI model, Grok 3, and unveiled new capabilities for the Grok iOS and web apps. The new model, which you acknowledge our Terms & Privacy Policy. Grok, xAI's answer to models like OpenAI's [GPT-4o](#) and Google's [Gemini](#), can analyze images and respond to questions, and powers a number of features on Musk's social network, X. Grok 3, which has been in development for several months, was [optimistically slated](#) for release in 2024, but missed that deadline.

Monday's is an ambitious launch.

xAI has been using an enormous data center in Memphis containing around 200,000 GPUs to train Grok 3. In a [post](#) on X, Musk claimed Grok 3 was developed with "10x" (or so) more computing power than its predecessor, Grok 2, using an expanded training set that includes [filings from court cases](#) — and more.

Latest in AI

[See More](#)

AI



Google denies AI search features are killing website traffic

Sarah Perez - 12 hours ago

AI



First impressions of Alexa+, Amazon's upgraded, AI-powered digital assistant


Sarah Perez - 14 hours ago

AI



Google takes on ChatGPT's Study Mode with new 'Guided Learning' tool in Gemini

Aisha Malik - 14 hours ago



MEMBERS OF THE XAI TEAM, INCLUDING MUSK (FAR RIGHT), DURING GROK 3'S LIVESTREAMED UNVEILING.

IMAGE CREDITS:XAI

“Grok 3 is an order of magnitude more capable than Grok 2,” Musk said during a livestreamed presentation on Monday. “[It’s a] maximally truth-seeking AI, even if that truth is sometimes at odds with what is politically correct.”

Grok 3 is a family of models, to be precise. A smaller version of Grok 3, Grok 3 mini, responds to questions more quickly at the cost of some accuracy. Not all the models and related features of Grok 3 are available yet (some are in beta), but they began rolling out on Monday.

xAI claims Grok 3 beats GPT-4o on benchmarks including AIME (which evaluates a model’s performance on a sampling of math questions) and GPQA (which assesses models using PhD-level physics, biology, and chemistry problems). An early version of Grok 3 also scored competitively in [Chatbot Arena](#), a crowdsourced test that pits different AI models against each other and has users vote on their preferred responses, according to xAI.



Tech and VC heavyweights join the Disrupt 2025 agenda


Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the

insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Benchmarks



LIVE 1.5M views

x1

IMAGE CREDITS: XAI

Two models in the new Grok 3 family, Grok 3 Reasoning and Grok 3 mini Reasoning, can carefully “think through” problems, similar to “reasoning” models like OpenAI’s [o3-mini](#) and Chinese AI company DeepSeek’s [R1](#). Reasoning models try to fact-check themselves before giving out results, which helps them [avoid some of the pitfalls](#) that normally trip up models.

xAI claims that Grok 3 Reasoning surpasses the best version of o3-mini — o3-mini-high — on several popular benchmarks, including a newer mathematics benchmark called AIME 2025.

Reasoning + Test-Time Compute




IMAGE CREDITS: XAI

These reasoning models can be accessed via the Grok app. Users can ask Grok 3 to “Think,” or — for more difficult queries — leverage “Big Brain” mode for reasoning that employs additional computing. xAI describes the reasoning models as best suited for mathematics, science, and programming questions.

Musk said some of the reasoning models’ “thoughts” are obscured in the Grok app to prevent distillation, a method used by AI model developers to extract knowledge from other models. Recently, DeepSeek was [accused of distilling OpenAI’s models](#) to create its own.

Grok’s reasoning models underpin a new feature in the Grok app called DeepSearch, xAI’s answer to AI-powered research tools like [OpenAI’s deep research](#). DeepSearch scans the internet and X to analyze information and deliver an abstract in response to a question.

Subscribers to [X’s Premium+ tier](#) (\$50 per month) will get access to Grok 3 first, and other features will be gated behind a new plan that xAI’s calling SuperGrok. Priced at \$30 per month or \$300 per year ([if leaks are to be believed](#)), SuperGrok unlocks additional reasoning and DeepSearch queries, and throws in unlimited image generation.

IMAGE CREDITS: XAI

In the future — as soon as about a week from now — the Grok app will gain a “voice mode,” Musk said, which will give Grok models a synthesized voice. A few weeks after that, Grok 3 models will be available via [xAI’s enterprise API](#), along with the DeepSearch capability.

xAI plans to open source Grok 2 in the coming months, Musk said.

“Our general approach is that we will open source the last version [of Grok] when the next version is fully out,” he continued. “When Grok 3 is mature and stable, which is probably within a few months, then we’ll open source Grok 2.”


When Musk announced Grok roughly two years ago, he pitched the AI model as edgy, unfiltered, and anti-“woke” — in general, willing to answer controversial questions other AI systems won’t. He delivered on some of that promise. Told to be vulgar, for example, Grok and Grok 2 would happily oblige, spewing colorful language you likely wouldn’t hear from [ChatGPT](#).

But Grok models prior to Grok 3 [hedged](#) on political subjects and wouldn’t cross [certain boundaries](#). In fact, [one study](#) found that Grok leaned to the political left on topics like transgender rights, diversity programs, and inequality.

Musk has blamed the behavior on Grok’s training data — public web pages — and [pledged](#) to “shift Grok closer to politically neutral.” It’s not yet clear whether xAI has achieved that goal, and what the consequences might be.

TechCrunch has an AI-focused newsletter! [Sign up here](#) to get it in your inbox every Wednesday.

Topics: [AI](#) [Apps](#) [Elon Musk](#) [grok 3](#) [xAI](#)



Kyle Wiggers

AI Editor |

Kyle Wiggers was TechCrunch’s AI Editor until June 2025. His writing has appeared in VentureBeat and Digital...

[View Bio](#) >

Loading the next article





- Advertise
- Crunchboard
- Jobs
- Site Map
- RSS
- Terms of Use
- Windsurf
- Tech Layoffs
- ChatGPT