

Advertisement

APPS



Twitter's former Trust and Safety head details the challenges facing decentralized social platforms

[Yoel Roth](#), [previously](#) the head of Twitter's Trust and Safety, [now at Match Group](#), is sharing his concerns about the future of the open social web and its ability to combat misinformation, spam, and other illegal content, like child sexual abuse material (CSAM). In a recent interview, Roth worried about the lack of moderation tools available to the fediverse — the open social web that includes apps like Mastodon, Threads, Pixelfed, and others, as well as other open platforms like Bluesky.

He also reminisced about key moments in Trust and Safety at Twitter, like its decision to ban President Trump from the platform, the misinformation spread by Russian bot farms, and how Twitter's own users, including CEO Jack Dorsey, fell prey to bots.

IMAGE CREDITS: BRYCE DURBIN / TECHCRUNCH

On the podcast [revolution.social with @Rabble](#), Roth pointed out that the efforts at building more democratically run online communities across the open social web are also those that have the fewest resources when it comes to moderation tools.

“...looking at Mastodon, looking at other services based on ActivityPub [protocol], looking at Bluesky in its earliest days, and then looking at Threads as Meta started to develop it, what we saw was that a lot of the services that were leaning the hardest into community-based control gave their communities the least technical tools to be able to administer their policies,” Roth said.

TECH DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Instagram takes on Snapchat with new 'Instagram Map'
- Google says hackers stole its customers' data by breaching its Salesforce database
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Spotify raises subscription prices



- The uproar over Vogue’s AI-generated ad isn’t just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines

He also saw a “pretty big backslide” on the open social web when it came to the transparency and decision legitimacy that Twitter once had. While, arguably, many at the time disagreed with Twitter’s decision to ban Trump, the company explained its rationale for doing so. Now, social media providers are so concerned about preventing bad actors from gaming them that they rarely explain themselves.

Meanwhile, on many open social platforms, users wouldn’t receive a notice about their banned posts, and their posts would just vanish — there wasn’t even an indication to others that the post used to exist.

“I don’t blame startups for being startups, or new pieces of software for lacking all the bells and whistles, but if the whole point of the project was increasing democratic legitimacy of governance, and what we’ve done is take a step back on governance, then, has this actually worked at all?” Roth wonders.

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the

20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

The economics of moderation

He also brought up the issues around the economics of moderation and how the federated approach hasn't yet been sustainable on this front.

For instance, an organization called IFTAS (Independent Federated Trust & Safety) had been working to build moderation tools for the fediverse, including providing the fediverse with access to tools to combat CSAM, but it ran out of money and [had to shut down many](#) of its projects earlier in 2025.

“We saw it coming two years ago. IFTAS saw it coming. Everybody who's been working in this space is largely volunteering their time and efforts, and that only goes so far, because at some point, people have families and need to pay bills, and compute costs stack up if you need to run ML models to detect certain types of bad content,” he explained. “It just all gets expensive, and the economics of this federated approach to trust and safety never quite added up. And in my opinion, still don't.”

Bluesky, meanwhile, has chosen to employ moderators and hire in trust and safety, but it limits itself to the moderation of its own app. Plus, they're providing tools that let people customize their own moderation preferences.

“They’re doing this work at scale. There’s obviously room for improvement. I’d love to see them be a bit more transparent. But, fundamentally, they’re doing the right stuff,” Roth said. However, as the service further decentralizes, Bluesky will face questions about when it is the responsibility to protect the individual over the needs of the community, he notes.

For example, with doxxing, it’s possible that someone wouldn’t see that their personal information was being spread online because of how they configured their moderation tools. But it should still be someone’s responsibility to enforce those protections, even if the user isn’t on the main Bluesky app.

Where to draw the line on privacy

Another issue facing the fediverse is that the decision to favor privacy can thwart moderation attempts. While Twitter tried not to store personal data it didn’t need to, it still collected things like the IP address of the user, when they accessed the service, device identifiers, and more. These helped the company when it needed to do forensic analysis of something like a Russian troll farm.

Fediverse admins, meanwhile, may not even be collecting the necessary logs or won’t view them if they think it’s a violation of user privacy.

But the reality is that without data, it’s harder to determine who’s really a bot.

Roth offered a few examples of this from his Twitter days, noting how it became a trend for users to reply “bot” to anyone they disagreed with. He says that he initially set up an alert and reviewed all these posts manually, examining hundreds of

instances of “bot” accusations, and nobody was ever right. Even Twitter co-founder and former CEO Jack Dorsey fell victim, retweeting posts [from a Russian actor who claimed to be Crystal Johnson](#), a Black woman from New York.

“The CEO of the company liked this content, amplified it, and had no way of knowing as a user that Crystal Johnson was actually a Russian troll,” Roth said.

The role of AI

One timely topic of discussion was how AI was changing the landscape. Roth referenced recent research from Stanford that found that, in a political context, large language models (LLMs) could even be more convincing than humans when properly tuned.

That means a solution that relies only on content analysis itself isn’t enough.

Instead, companies need to track other behavioral signals — like if some entity is creating multiple accounts, using automation to post, or posting at weird times of day that correspond to different time zones, he suggested.

“These are behavioral signals that are latent even in really convincing content. And I think that’s where you have to start this,” Roth said. “If you’re starting with the content, you’re in an arms race against leading AI models and you’ve already lost.”

Topics: [Apps](#) [Bluesky](#) [fediverse](#) [moderation](#)

[Online Safety](#) [Social](#) [TC](#) [Twitter](#)



Sarah Perez

Consumer News Editor | X butterfly @ in

Sarah has worked as a reporter for TechCrunch since August 2011. She joined the company after having...

View Bio >

KEEP READING




IMAGE CREDITS: OPENAI



AI

OpenAI's GPT-5 is here

Maxwell Zeff — 10:00 AM PDT · August 7, 2025

OpenAI has launched GPT-5, a new flagship AI model that will power the company's next

generation of ChatGPT.

GPT-5, which was released Thursday, is OpenAI’s first “unified” AI model and combines the reasoning abilities of its [o-series](#) of models with the fast responses of its [GPT series](#). The next-generation model signals a new era for ChatGPT — and its creator, OpenAI — pointing to OpenAI’s broader ambitions to develop AI systems that are [more like agents](#) than chatbots.

While [GPT-4](#) enabled AI chatbots to offer smart responses on a wide variety of questions, GPT-5 allows ChatGPT to complete a wide variety of tasks on behalf of users — such as generating software applications, navigating a user’s calendar, or creating research briefs.

With GPT-5, OpenAI has also sought to make ChatGPT simpler to use. Instead of asking users to choose the right settings, GPT-5 comes equipped with a real-time router that decides how to offer the best answer, whether that’s responding to user questions quickly or taking additional time to “think” through answers.




IMAGE CREDITS: OPENAI

During a briefing with reporters, OpenAI CEO Sam Altman claimed GPT-5 is “the best model in the world,” and said it represented a “significant step” along the company’s path to developing AI that can

Advertisement

Newsletters

[See More ↗](#)

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

outperform humans at most economically valuable work — that is, artificial general intelligence (AGI).

“Having something like GPT-5 would be pretty much unimaginable at any previous time in history,” said Altman.

Starting Thursday, GPT-5 will be available to all free users of ChatGPT as their default model. OpenAI’s VP of ChatGPT, Nick Turley, said this is part of the company’s effort to give free users access to an AI reasoning model for the first time. (Previously, the company gated these more advanced models behind a paywall.)

“This is just one of the ways that I’m excited to live the mission, making sure that this stuff actually benefits people,” said Turley on the decision, referencing OpenAI’s long-standing [mission](#) to distribute advanced AI to as many people as possible.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

[REGISTER NOW](#)

The expectations are high for GPT-5, one of OpenAI's most anticipated product launches since ChatGPT put the company on the map in 2022. Since then, ChatGPT has grown into one of the world's most popular consumer products, reaching more than 700 million users every week — nearly 10% of the globe's population, according to the company.

Many see GPT-5 as a bellwether for AI progress broadly, and the model's reception by Silicon Valley could have profound implications for Big Tech, Wall Street, and policymakers regulating technology. These stakeholders are watching to see if GPT-5 offers a significant jump in AI's capabilities, much like its predecessor, GPT-4, which challenged expectations of what software can do.


GPT-5 offers a slight edge on the competition

OpenAI claims GPT-5 is state-of-the-art in several domains, slightly edging out leading AI models from Anthropic, Google DeepMind, and Elon Musk's xAI on key benchmarks. However, GPT-5 slightly underperforms frontier AI models in other areas.

The company says GPT-5 offers frontier-level performance around coding; Altman said the model specifically excels at spinning up entire software applications on demand, in what's become known as "vibe coding."

On SWE-bench Verified — a test of real-world coding tasks pulled from GitHub — GPT-5 scores 74.9% on its first attempt. That means GPT-5 just outperforms Anthropic's latest Claude Opus 4.1 model, which scored [74.5%](#), and Google DeepMind's Gemini 2.5 Pro, which scored [59.6%](#).

Latest
Startups
Venture
Apple
Security
AI
Apps



Events
Podcasts
Newsletters




IMAGE CREDITS: OPENAI

On GPQA Diamond — a test of PhD-level science questions — GPT-5 pro scored 89.4% on its first try, outperforming Claude Opus 4.1, which scored 80.9%, and Grok 4 Heavy, which scored 88.9%.

OpenAI says GPT-5 is better for answering health-related questions. On a test measuring accuracy in AI model responses around healthcare topics, HealthBench Hard Hallucinations, OpenAI says GPT-5 (with thinking) hallucinates just 1.6% of the time. This is far lower than the company's previous GPT-4o and o3 models, which scored 12.9% and 15.8, respectively.

While AI chatbots are not medical professionals, [millions of people](#) are using them for health advice. In response to this phenomenon, the company says GPT-5 is more proactive about flagging potential health concerns and helping users parse medical results.

In addition, OpenAI says GPT-5 is better than other AI models on more difficult-to-measure, [subjective domains](#), such as creative design and writing.

Turley said GPT-5 responds more naturally and exhibits “better taste” than other AI models on creative tasks.

“The vibes of this model are really good,” said Turley.

GPT-5 is also more accurate than OpenAI’s previous models, and the company says it suffers far less from hallucinations — the tendency for AI models to make up information — compared to its o-series models. Hallucinations seemed to be [getting worse](#) in OpenAI’s latest AI reasoning models, such as o3, and OpenAI previously said it didn’t quite understand why it was happening.

In responses to ChatGPT prompts, OpenAI found that GPT-5 (with thinking) hallucinates and responds with incorrect information 4.8% of the time. That’s a significant reduction from o3 and GPT-4o, which score hallucination rates of 22% and 20.6%, respectively, on the test.

On a benchmark measuring an AI model’s agentic ability to complete simulated online tasks, Tau-bench, GPT-5 offers mixed performance. On part of the test measuring an AI’s ability to navigate an airline’s website, GPT-5 scores 63.5%, slightly underperforming o3, which scored 64.8%. On another part of the test measuring AI’s ability to navigate retail websites, GPT-5 scores 81.1%, underperforming Claude Opus 4.1, which scored 82.4%.

OpenAI also says that GPT-5 is safer than its previous models. While AI reasoning models occasionally exhibit a tendency to [scheme against humans](#) or lie to promote their own goals, OpenAI

found that GPT-5 was deceptive at a lower rate than other models.

Alex Beutel, OpenAI’s safety research lead, said reducing deception improves not only the safety of GPT-5, but also the user experience, creating a model that’s more “transparent and honest in ways users can trust.”

Beutel also notes GPT-5 is better at discerning between bad actors who are trying to misuse ChatGPT and users making harmless requests. This results in GPT-5 being able to refuse more unsafe questions, while offering fewer rejections to users seeking harmless information.

Upgrades for consumers and developers

ChatGPT is getting a few user experience upgrades as part of the GPT-5 launch. Users can now select from four new personalities in ChatGPT’s setting: Cynic, Robot, Listener, and Nerd. The company says these will adapt ChatGPT’s responses without requiring users to specifically ask the model to respond in a certain way.

Subscribers to ChatGPT’s \$20-per-month Plus plan get higher usage limits for GPT-5 than free users. Meanwhile, \$200-per-month Pro subscribers will have unlimited access to GPT-5, as well as a souped-up version called GPT-5 Pro that uses additional computational resources to produce better answers. Organizations on OpenAI’s Team, Edu, and Enterprise plans will gain access to GPT-5 as their default model next week.

For developers, GPT-5 is coming to OpenAI’s API in three sizes — gpt-5, gpt-5-mini, and gpt-5-nano — which will spend more or less time “reasoning” through tasks. Developers can also now control

verbosity in the OpenAI API, deciding how long or short an AI model's responses should be.

The base model of GPT-5 will cost developers \$1.25 per million input tokens (roughly 750,000 words, longer than the entire "Lord of the Rings" series) and \$10 per million output tokens.

The launch of GPT-5 comes after a busy week for OpenAI. The company released an open-weight reasoning model, [gpt-oss](#), that developers and enterprises can download for free and run at a fraction of the cost. The open model nearly matched the abilities of OpenAI's previous top models, o3 and o4-mini, but GPT-5 sets a new standard for frontier performance in some areas, such as coding.

However, GPT-5 seems to be roughly on par with other frontier AI models in several areas. Benchmarks, of course, only tell part of the story for any AI model, and it remains to be seen how developers will use GPT-5 in the real world, and whether the model is truly a step above the competition.

Topics: [AI](#) [ChatGPT](#) [gpt-5](#) [OpenAI](#)

[f](#) [X](#) [in](#) [👤](#) [✉](#) [🔗](#)



Maxwell Zeff

Senior AI Reporter | [X](#)

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >




IMAGE CREDITS:

RENE JOHNSTON/TORONTO STAR / GETTY IMAGES

SECURITY



Citizen Lab director warns cyber industry about US authoritarian descent

Lorenzo Franceschi-Bicchierai

— 8:01 AM PDT · August 6, 2025

Ron Deibert, the director of Citizen Lab, one of the most prominent organizations investigating government spyware abuses, is sounding the alarm to the cybersecurity community and asking them to step up and join the fight against authoritarianism.

On Wednesday, Deibert [will deliver a keynote](#) at the Black Hat cybersecurity conference in Las Vegas, one of the largest gatherings of information security professionals of the year.

Ahead of his talk, Deibert told TechCrunch that he plans to speak about what he describes as a “descent into a kind of fusion of tech and fascism,”

Latest in Security

[See More](#)

SECURITY



Data breach at French telecom giant Bouygues affects millions of customers

Zack Whittaker - 14 hours ago

and the role that the Big Tech platforms are playing, and “propelling forward a really frightening type of collective insecurity that isn’t typically addressed by this crowd, this community, as a cybersecurity problem.”

Deibert described the recent political events in the United States as a “dramatic descent into authoritarianism,” but one that the cybersecurity community can help defend against.

“I think alarm bells need to be rung for this community that, at the very least, they should be aware of what’s going on and hopefully they can not contribute to it, if not help reverse it,” Deibert told TechCrunch.

Historically, at least in the United States, the cybersecurity industry has put politics — to a certain extent — to the side. More recently, however, politics has fully entered the world of cybersecurity.

Earlier this year, President Donald Trump [ordered an investigation into former CISA director Chris Krebs](#), who had publicly rebuffed Trump’s false claims about election fraud by declaring the 2020 election secure. Trump later [fired Krebs by tweet](#). The investigation ordered by Trump months after his 2024 reelection forced Krebs to [step down from SentinelOne and vow to fight back](#).

In response, Jen Easterly, another former CISA director and Krebs’ successor, [called on the cybersecurity community to get involved](#) and speak out.

“If we stay silent when experienced, mission-driven leaders are sidelined or sanctioned, we risk something greater than discomfort; we risk diminishing the very institutions we are here to protect,” Easterly wrote in a post on LinkedIn.

SECURITY



TeaOnHer, a rival Tea app for men, is leaking users’ personal data and driver’s licenses

Amanda Silberling,
Zack Whittaker
- 1 day ago

SECURITY



Citizen Lab director warns cyber industry about US authoritarian descent

Lorenzo Franceschi-Bicchierai
- 2 days ago

Easterly was herself [a victim of political pressure](#) from the Trump administration when her offer to join West Point was rescinded in late July.

Deibert, who this year [published his new book](#), “Chasing Shadows: Cyber Espionage, Subversion, and the Global Fight for Democracy,” is echoing the same message as Easterly.

“I think that there comes a point at which you have to recognize that the landscape is changing around you, and the security problems you set out for yourselves are maybe trivial in light of the broader context and the insecurities that are being propelled forward in the absence of proper checks and balances and oversight, which are deteriorating,” said Deibert.

Deibert is also concerned that big companies like Meta, Google, and Apple could take a step back in their efforts to fight against government spyware — [sometimes referred](#) to as “commercial” or “mercenary” spyware — by gutting their threat intelligence teams.

These threat intelligence teams are dedicated groups of security researchers that track government hackers, both those working inside government agencies, such as China’s Ministry of State Security or Russia’s intelligence agencies FSB and GRU, and companies such as NSO Group or Paragon.

These are the same teams that are responsible for detecting hacks against their own users, such as when [WhatsApp caught NSO Group](#) hacking more than 1,400 of its users in 2019, or when Apple caught hackers using government spyware to target its customers and [notified the victims](#) of the [attacks](#).

Deibert believes there is a “huge market failure when it comes to cybersecurity for global civil society,” a part of the population that generally cannot afford to get help from big security companies that typically serve governments and corporate clients. “This market failure is going to get more acute as supporting institutions evaporate and attacks on civil society amplify,” he said.

“Whatever they can do to contribute to offset this market failure (e.g., pro bono work) will be essential to the future of liberal democracy worldwide,” he said.

Deibert is concerned that these threat intelligence teams could be cut or at least reduced, given that the same companies have cut their moderation and safety teams.

He told TechCrunch that threat intelligence teams, like the ones at Meta, are doing “amazing work,” in part by staying siloed and separate from the commercial arms of their wider organizations.

“But the question is how long will that last?” said Deibert.

Topics: [black hat](#) [Black Hat 2025](#) [Citizen Lab](#)

[cybersecurity](#) [hacking](#) [infosec](#) [Politics](#)

[Ron Deibert](#) [Security](#)

[f](#) [X](#) [in](#) [📷](#) [✉](#) [🔗](#)




Lorenzo Franceschi-Bicchierai

Senior Reporter, Cybersecurity | [X](#) [🦋](#) [@](#)

Lorenzo Franceschi-Bicchierai is a Senior Writer at TechCrunch, where he covers hacking, cybersecurity,...

View Bio >

Loading the next article



- TechCrunch
Staff
Contact Us
Advertise
Crunchboard
Jobs
Site Map
- Terms of
Service
Privacy
Policy
RSS Terms of
Use
- OpenAI
Perplexity
Cisco
DeepMind
Windsurf
Tech Layoffs
ChatGPT