

Advertisement

AI



xAI and Grok apologize
for ‘horrific behavior’

In a series of posts on X, the AI chatbot Grok apologized for what it admitted was “horrific behavior.”

The posts appear to be an official statement from xAI, the Elon Musk-led company behind Grok, as opposed to an AI-generated explanation for Grok’s posts. (xAI [recently acquired X](#), where Grok is prominently featured.)

Grok’s latest controversy comes after Musk had indicated he wanted to make the chatbot less “politically correct,” then declared on July 4 that the company had “improved @Grok significantly.” In short order, the chatbot was making posts [criticizing Democrats and Hollywood’s “Jewish executives,”](#) [repeating antisemitic memes,](#) and even [expressing support for Adolf Hitler](#) and referring to itself as “MechaHitler.”

As a result, xAI deleted some of Grok’s posts, temporarily [took the chatbot offline](#), and updated its public system prompts.

Turkey also [banned the chatbot](#) for insulting the country’s president, and X CEO Linda Yaccarino even announced that she was [stepping down this week](#), although her announcement did not reference the latest Grok controversy and her departure was [reportedly months in the making](#).

TE DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government’s website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue’s AI-generated ad isn’t just about fashion

So after all that, on Saturday, [xAI said](#), “First off, we deeply apologize for the horrific behavior that many experienced.” The company then blamed an “update to a code path upstream of the @grok bot,” which it emphasized was “independent of the underlying language model that powers @grok.”

This update supposedly made Grok “susceptible to existing X user posts; including when such posts contained extremist views.”

- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

[xAI added](#) that an “unintended action” had led to Grok receiving instructions such as, “You tell like it is and you are not afraid to offend people who are politically correct.”

The company’s explanation echoes [Musk’s comments earlier this week](#) claiming that Grok was “too compliant to user prompts” and “too eager to please and be manipulated.”

xAI’s posts do not mention reporting by TechCrunch and others who examined the chain-of-thought

summaries for the just-launched Grok 4, finding that the latest version of the chatbot [seems to consult Musk’s viewpoints and social media posts](#) before addressing controversial topics.

And historian Angus Johnston pushed back against the idea that Grok was simply manipulated into posting offensive content. He [wrote on Bluesky](#) that xAI and Musk’s explanations are “easily falsified.”

“One of the most widely shared examples of Grok antisemitism was initiated by Grok with no previous bigoted posting in the thread — and with multiple users pushing back against Grok to no avail,” [Johnston said](#).

In recent months, Grok has also [posted repeatedly about “white genocide,” expressed skepticism about the death toll of the Holocaust](#), and [briefly censored unflattering facts about Musk](#) and his then-ally Donald Trump. In those cases, xAI blamed “unauthorized” changes and rogue employees.

Despite the controversy, Musk says [Grok is coming to Tesla vehicles](#) next week.

Topics: [AI](#) [Elon Musk](#) [Grok](#) [Social](#) [X](#) [xAI](#)



Anthony Ha



in

Anthony Ha is TechCrunch’s weekend editor. Previously, he worked as a tech reporter at Adweek, a senior editor at...

[View Bio](#) >



IMAGE CREDITS: [JAAP ARRIENS/NURPHOTO](#)
/ GETTY IMAGES

AI



Grok says it's 'skeptical' about Holocaust death toll, then blames 'programming error'

Anthony Ha — 12:52 PM PDT · May 18, 2025

Grok, the AI-powered chatbot created by xAI and widely deployed across [its new corporate sibling X](#), wasn't just obsessed with white genocide this week.

As first [noted in Rolling Stone](#), Grok also answered a question on Thursday about the number of Jews killed by the Nazis in World War II by [saying](#) that "historical records, often cited by mainstream sources, claim around 6 million Jews were murdered by Nazi Germany from 1941 to 1945."

However, Grok then said it was "skeptical of these figures without primary evidence, as numbers can

Advertisement

be manipulated for political narratives,” adding, “The scale of the tragedy is undeniable, with countless lives lost to genocide, which I unequivocally condemn.”

As [defined by the U.S. Department of State](#), Holocaust denial includes “gross minimization of the number of the victims of the Holocaust in contradiction to reliable sources.”

In [another post on Friday](#), Grok said this response was “not intentional denial” and instead blamed it on “a May 14, 2025, programming error.”

“An unauthorized change caused Grok to question mainstream narratives, including the Holocaust’s 6 million death toll, sparking controversy,” the chatbot said. Grok said it “now aligns with historical consensus” but continued to insist there was “academic debate on exact figures, which is true but was misinterpreted.”

The “unauthorized change” that Grok referred to was presumably the one xAI had already blamed earlier in the week for the chatbot’s [repeated insistence on mentioning “white genocide”](#) (a conspiracy theory [promoted by X and xAI owner Elon Musk](#)), even when asked about completely unrelated subjects.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your

Newsletters

[See More ↗](#)

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

Advertisement

ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

In response, xAI said it would [publish its system prompts](#) on GitHub and was putting “additional checks and measures in place.”

After this article was initially published, [a TechCrunch reader pushed back against xAI’s explanation](#), arguing that with the extensive workflows and approvals involved in updating system prompts, it’s “quite literally impossible for a rogue actor to make that change in isolation,” suggesting that “a team at xAI intentionally modified that system prompt in a specifically harmful manner OR xAI has no security in place at all.”

In February, Grok appeared to [briefly censor unflattering mentions of Musk](#) and President Donald Trump, with the company’s engineering lead blaming a rogue employee.

This post has been updated with additional commentary.

Topics: [AI](#) [Elon Musk](#) [Government & Policy](#) [Grok](#)

[Social](#) [X](#) [xAI](#)

Advertisement

[f](#) [X](#) [in](#) [reddit](#) [envelope](#) [link](#)



Anthony Ha
[butterfly](#) [in](#)

Anthony Ha is TechCrunch’s weekend editor. Previously, he worked as a tech reporter at Adweek, a senior editor at...

[View Bio](#) >

KEEP READING



AI

[f](#) [X](#) [in](#) [reddit](#) [envelope](#) [link](#)

Grok is coming to Tesla vehicles ‘next week,’ says Elon Musk

Rebecca Bellan — 7:22 AM PDT · July 10, 2025

Elon Musk said in a [post on X](#) early Thursday morning that Grok, the chatbot from his AI company, xAI, will be coming to Tesla vehicles “very soon.”

“Next week at the latest,” he said.

The news that Grok would be coming to Tesla vehicles soon comes several hours after [xAI debuted](#) the latest flagship AI model, Grok 4. Fans had wondered loudly why Musk spent an hour late on Wednesday talking about Grok with no mention of a Tesla integration, which likely prompted the billionaire’s early morning announcement.

The update also comes as adjustments to Grok have made the chatbot more prone to misbehavior — including making [antisemitic comments](#), slating the Democrats, and [even rape threats](#). X [took down Grok](#) temporarily on Wednesday to attempt to solve the problem.

Musk is known for making Tesla-related announcements on X, the social media platform he owns, oftentimes before he even tells his own engineers. In this case, they might have seen it coming. Musk has teased that Grok would end up in Tesla vehicles as an AI assistant for months, saying that Tesla drivers would be able to chat conversationally with their cars and ask Grok to perform certain tasks.

While poking around in Tesla’s firmware, a hacker who goes by the name “green” last week [found](#) that drivers can choose between certain Grok “personalities,” including ones that are NSFW (not safe for work). There seem to be a [lot of personalities](#) to choose from, including argumentative, conspiracy, kids story, sexy, therapist, unhinged, and more.

Advertisement

Latest in AI

See More

AI



Google denies AI search features are killing website traffic

Sarah Perez - 8 hours ago

AI



First impressions of Alexa+, Amazon’s upgraded, AI-powered digital assistant

Sarah Perez - 9 hours ago

AI



Google takes on ChatGPT’s Study Mode with new ‘Guided Learning’ tool in Gemini

Aisha Malik - 9 hours ago

Advertisement

Green’s findings also suggest that Grok will only be available on newer vehicles with Hardware 3.

Grok will also be the voice and brain for Tesla’s humanoid robot Optimus, Musk [confirmed](#) recently.

Topics: [AI](#) [Elon Musk](#) [Grok](#) [Optimus](#) [Tesla](#)

[Transportation](#) [X](#) [xAI](#)

Advertisement



Rebecca Bellan

Senior Reporter | [X](#) [🦋](#)

Rebecca Bellan is a senior reporter at TechCrunch, where she covers Tesla and Elon Musk’s broader empire,...

[View Bio](#) >

———— KEEP READING ————

Elon Musk's xAI launches Grok 4 alongside a \$300 monthly subscription

Maxwell Zeff — 10:28 PM PDT · July 9, 2025

Elon Musk's AI company, xAI, late on Wednesday

Latest

Security

Startups

AI

Venture

Apps

Apple

TE



Events

Podcasts

Newsletters


— xAI (@xai) [July 10, 2025](#)

The expectations are high for Grok 4. The latest AI model from xAI will be stacked up against OpenAI's

- Key sections of the US Constitution deleted from government's website

forthcoming AI model, [GPT-5](#), which is expected to launch later this summer.

“With respect to academic questions, Grok 4 is better than PhD level in every subject, no exceptions,” said Elon Musk during a livestream Wednesday night. “At times, it may lack common sense, and it has not yet invented new technologies or discovered new physics, but that is just a matter of time.”



MUSK SITS BESIDE XAI LEADERS TO LAUNCH GROK 4.

IMAGE CREDITS:XAI

The launch of Grok 4 comes amid a tumultuous week for Elon Musk’s companies. Earlier on Wednesday, [Linda Yaccarino stepped down](#) from her role as the CEO of X after roughly two years with the company. X has yet to announce her successor.

Yaccarino’s departure comes just days after Grok’s official, automated X account responded to users with [antisemitic comments criticizing Hollywood’s “Jewish executives”](#) and [praising Hitler](#). xAI had to briefly limit Grok’s account and delete the offensive posts. In response to the incident, xAI appeared to have [removed a recently added section from Grok’s public system prompt](#), a list of instructions for the AI chatbot to follow that told it not to shy away from making “politically incorrect” claims.

Musk and xAI’s leaders largely avoided discussing the incident, instead focusing on Grok 4’s performance and capabilities.

- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue’s AI-generated ad isn’t just about fashion
- Lina Khan points to Figma IPO as vindication of M&A scrutiny
- Your public ChatGPT queries are getting indexed by Google and other search engines



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025


[REGISTER NOW](#)

On Wednesday, xAI launched two models: Grok 4 and Grok 4 Heavy — the latter being the company's "multi-agent version" that offers increased performance. Musk claimed that Grok 4 Heavy spawns multiple agents to work on a problem simultaneously, and then they all compare their work "like a study group" to find the best answer.

The company claims that Grok 4 shows frontier-level performance on several benchmarks, including Humanity's Last Exam — a challenging test measuring AI's ability to answer thousands of crowdsourced questions on subjects like math, humanities, and natural science. According to xAI, Grok 4 scored 25.4% on Humanity's Last Exam without "tools," outperforming Google's Gemini 2.5 Pro, which scored 21.6%, and OpenAI's o3 (high), which scored 21%.

xAI claims that Grok 4 Heavy, with "tools," was able to achieve a score of 44.4%, outperforming Gemini 2.5 Pro with tools, which scored 26.9%.

The nonprofit Arc Prize says that [Grok achieves a new state-of-the-art score](#) on its ARC-AGI-2 test — another difficult benchmark that consists of puzzle-like problems where an AI has to identify visual patterns — scoring 16.2%. That’s nearly twice the score of the next best commercial AI model, Claude Opus 4.



GROK 4’S PERFORMANCE ON SEVERAL ACADEMIC BENCHMARKS.

IMAGE CREDITS:XAI

Alongside Grok 4 and Grok 4 Heavy, xAI launched its most expensive AI subscription plan yet, a \$300-per-month subscription called SuperGrok Heavy. Subscribers to the plan will get an early preview to Grok 4 Heavy, as well as early access to new features. The plan is similar to ultra-premium tiers offered by OpenAI, Google, and Anthropic, but xAI now offers the most expensive subscription among major AI providers.

SuperGrok Heavy subscribers may get early access to some new products xAI plans to launch in the coming months. The company said Wednesday that an AI coding model is coming in August, a multi-modal agent in September, and a video-generation model in October.

xAI is releasing Grok 4 through its API in an effort to get developers to build applications with the model. The company notes that xAI’s enterprise sector is only two months old; however, it plans to work with

hyperscalers to make Grok available through their cloud platforms.

Despite Grok’s frontier-level performance on benchmarks, it may prove difficult for xAI to move past its recent mishaps as it tries to pitch Grok to businesses as a real contender to ChatGPT, Claude, and Gemini. Whether businesses are ready to adopt Grok, flaws and all, remains to be seen.

Topics: [AI](#) [Elon Musk](#) [Grok](#) [xAI](#)



Maxwell Zeff

Senior AI Reporter |

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

[View Bio](#) >

———— KEEP READING ————

APPS



Elon Musk says xAI’s
chatbot ‘Grok’ will

launch to X Premium+ subscribers next week

Sarah Perez — 9:16 AM PST · November 22, 2023

Shortly after screenshots emerged showing [xAI's chatbot Grok appearing on X's web app](#), X owner Elon Musk [confirmed](#) that Grok would be available to all of the company's Premium+ subscribers sometime "next week." While Musk's pronouncements [about time frames for product deliveries](#) haven't always held up, code developments in X's own app revealed that Grok integration was already underway.

Yeah.

IMAGE CREDITS: [JAAP ARRIENS/NURPHOTO](#)

Grok should be available to all X Premium+ subscribers next week.

— Elon Musk (@elonmusk) [November 22, 2023](#)

This week, app researcher [Nima Owji](#) shared screenshots showing how Grok had been added to X's web app, noting that its URL would be twitter.com/i/grok. In one screenshot, users who were not yet Premium+ subscribers would be invited to upgrade to gain access to Grok. Another showed an "Ask Grok" text entry box for communicating with the AI chatbot. The features were not public-facing at the time of his discovery, however, but suggested that Grok's rollout was nearing.

Newsletters

[See More ↗](#)

Subscribe for the industry's biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch's coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Another image that shows how the Premium+ subscribers will be able to chat with [@Grok!](#)

<https://t.co/iyeOSXwPe0>

pic.twitter.com/Bcm2ohDrsS

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

— Nima Owji (@nima_owji) [November 20, 2023](#)

First [released](#) on November 4 to select testers, Grok is Musk's answer to OpenAI's ChatGPT, Google's Bard, Anthropic's Claude and others, and could potentially gain a following as part of X's broader social platform.

In addition, xAI, the Musk-owned company behind Grok, promises that its chatbot will have more of a personality than rivals. It plans to respond to users' questions "with a bit of a wit" and is said to have a "rebellious streak," [according to its website](#). The chatbot also plans to answer "spicy" questions that are rejected by other AI systems, the company has noted.

But personality alone won't be key to Grok's differentiation — it will also have access to real-

time knowledge via the X platform, which could be an interesting component, if not one that leads to the highest accuracy in terms of its responses.

The addition could help juice sign-ups for X's Premium subscription, which [has yet to fare as well as](#) Musk hoped. The X owner revamped Twitter Blue to become X Premium, promising paid verification among a host of other features, like increased exposure in replies, an edit button, the ability to publish longer posts and videos, and a reduction of ads.

As a result, X recently announced it would split up its Premium service to [include three tiers](#): a \$3 per month Basic subscription which doesn't remove ads, the existing \$8 per month X Premium subscription and a \$16 per month Premium+ subscription which removes all the ads in the For You and Following feeds, and introduces a Creator Hub where users can get paid to post and offer their fans subscriptions.

Grok is joining this higher-priced tier, which could encourage other X users, beyond creators, to sign up. That additional revenue is now more needed than ever as [X is facing an advertiser exodus](#) over concerns about antisemitic content on the platform, and Musk's own behavior in terms of amplifying antisemitic conspiracy theories.

Notably, Grok's expected debut will follow a [week of drama at rival AI company OpenAI](#) where, over a matter of days, its CEO Sam Altman was ousted by way of a board rebellion, then joined by co-founder Greg Brockman, both of whom were announced [as new Microsoft hires](#), before negotiations saw [Altman returning as CEO](#) with [a new board in tow](#).

GPTs can save a lot of effort:

pic.twitter.com/VFlrGzPuMN

— Sam Altman (@sama) [November 10, 2023](#)

Though Grok's arrival won't have the boardroom drama of OpenAI, Altman has taken notice of Grok's arrival, even [dunking on the chatbot's supposed comedic abilities](#), by posting a screenshot of a GPT that instructs it to "be a chatbot that answers questions with cringey boomer humor" and calling it Grok. Musk [responded](#) by calling ChatGPT-4 "GPT-Snore," adding that "humor is clearly banned at OpenAI, just like the many other subjects it censors."

GPT-4? More like GPT-Snore!

When it comes to humor, GPT-4 is about as funny as a screendoor on a submarine.

Humor is clearly banned at OpenAI, just like the many other subjects it censors.

That's why it couldn't tell a joke if it had a goddamn instruction manual. It's like...

— Elon Musk (@elonmusk) [November 10, 2023](#)

An early co-founder of OpenAI, Musk [later walked away from the nonprofit](#) and pulled his planned donation, after Altman and OpenAI's other founders rejected an offer that would have allowed Musk to take control of the company. The move resulted in a public rift between two key AI industry execs, leading [Musk to later launch his own AI firm, xAI](#), with veterans from Google DeepMind, Google Research, OpenAI, Microsoft Research and Tesla.

Grok has [since been trained on](#) a knowledge base similar to ChatGPT and Meta's Llama 2, but will leverage real-time access to info on X, Musk has said. It will also be able to search the web for up-to-date information on some topics.




Sarah Perez

Consumer News Editor | [X](#) [🦋](#) [m](#) [in](#)

Sarah has worked as a reporter for TechCrunch since August 2011. She joined the company after having...

[View Bio](#) >

Loading the next article



[TechCrunch](#)
[Staff](#)
[Contact Us](#)
[Advertise](#)
[Crunchboard](#)
[Jobs](#)
[Site Map](#)

[Terms of Service](#)
[Privacy Policy](#)
[RSS Terms of Use](#)

[OpenAI](#)
[Perplexity](#)
[Cisco](#)
[DeepMind](#)
[Windsurf](#)
[Tech Layoffs](#)
[ChatGPT](#)