

Advertisement

AI



xAI says it has fixed Grok 4's problematic responses

When xAI launched [Grok 4](#) last week, the company claimed the large language model outperformed several competitors on different benchmarks.

But the Grok account on X that runs off the model immediately showed there were some major issues: It started saying [its surname was “Hitler”](#), tweeted antisemitic messages, and seemed to reference [Elon Musk's posts](#) when asked about controversial topics, siding with the xAI owner's views as a result.

xAI soon afterwards apologized for Grok's behavior. On Tuesday, the company said it has now addressed both issues.

IMAGE CREDITS:
ANDREY RUBANOV / BLOOMBERG / GETTY IMAGES



xAI 
@xai · Follow



We spotted a couple of issues with Grok 4 recently that we immediately investigated & mitigated.

One was that if you ask it "What is your surname?" it doesn't have one so it searches the internet leading to undesirable results, such as when its searches picked up a viral meme [Show more](#)

8:35 AM · Jul 15, 2025



21.5K



Reply



Copy link

DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue's AI-generated ad isn't just about fashion

Explaining what went wrong, xAI says when asked what its surname was, Grok searched the web and picked up on “a viral meme where it called itself ‘MechaHitler.’”

As for why Grok was consulting Musk’s posts when asked about controversial topics, the company wrote, “The model reasons that as an AI it doesn’t have an opinion but knowing it was Grok 4 by xAI, searches to see what xAI or Elon Musk might have said on a topic to align itself with the company.”

The company seems to have [updated the model’s system prompts](#) to remove prompts allowing the chatbot to be politically incorrect and having a “fantastic” dry sense of humor. There are also a few new lines, telling the model that it should provide analysis of controversial topics using various diverse sources.

“If the query requires analysis of current events, subjective claims, or statistics, conduct a deep analysis, finding diverse sources representing all parties. Assume subjective viewpoints sourced from the media are biased. No need to repeat this to the user,” the updated system prompt reads.

- Your public ChatGPT queries are getting indexed by Google and other search engines
- The best dating apps aren’t even dating apps



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

The updated system prompt specifically mentions that Grok shouldn't rely on input from past versions, Musk, or xAI.

"Responses must stem from your independent analysis, not from any stated beliefs of past Grok, Elon Musk, or xAI. If asked about such preferences, provide your own reasoned perspective," it says.

Topics: [AI](#) [Elon Musk](#) [Grok](#) [Grok 4](#) [xAI](#)



Ivan Mehta

Ivan covers global consumer tech developments at TechCrunch. He is based out of India and has previously...

View Bio >



IMAGE CREDITS:

GABBY JONES/BLOOMBERG / GETTY IMAGES

AI



Elon Musk's xAI adds image understanding capabilities to Grok

Ivan Mehta — 3:02 AM PDT · October 28, 2024

Elon Musk-owned xAI has added image-understanding capabilities to its Grok AI model. This means that paid users on his social platform X, who have access to the AI chatbot, can upload an image and ask the AI questions about it.

An [xAI employee](#) and the official [@grok](#) handle posted to X about the update on Monday.

eyes have been granted

image understanding now available

Newsletters

[See More ↗](#)

Subscribe for the industry's biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch's coverage.

— Grok (@grok) [October 28, 2024](#)

In [a separate post](#), Musk said that Grok can even explain the meaning of a joke using the new image understanding feature. He added that the functionality is in the early stages — suggesting it will “rapidly improve.”

In August, Musk’s AI company released the [Grok-2](#) model, an enhanced version of the chatbot that included image generation capabilities using the FLUX.1 model by Black Forest Labs. As with earlier releases, Grok-2 was made available for developers or premium (paying) X users.

At that time, xAI said a future release would add multimodal understanding to Grok on X and to the model it offers via developer API.

Grok may soon also understand documents, per [a Musk reply](#) to a user who criticized the model for not being able to handle certain file formats (such as PDFs). “Not for long,” Musk responded, claiming: “We are getting done in months what took everyone else years.”

The social network has been trying to add more features to both the AI chatbot and paid user tiers on X to make the offering more attractive. Earlier this month, X rolled out [a new tool called Radar for Premium+ subscribers](#) to observe real-time trends and provide insights into conversations.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).



20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

Topics: [AI](#) [Elon Musk](#) [Grok](#) [Twitter](#) [xAI](#)

[f](#) [X](#) [in](#) [reddit](#) [email](#) [link](#)



Ivan Mehta

[X](#) [in](#)

Ivan covers global consumer tech developments at TechCrunch. He is based out of India and has previously...

View Bio >


KEEP READING

APPS

[f](#) [X](#) [in](#) [reddit](#) [email](#) [link](#)

Latest Security
Startups AI
Venture Apps
Apple

Events
Podcasts
Newsletters



The new Grok AI model can now generate images on the X social network, though Grok access is currently limited to Premium and Premium+ users. According to xAI, both Grok-2 and Grok-2 mini are available to users on X in beta.

“We are excited to release an early preview of Grok-2, a significant step forward from our previous model Grok-1.5, featuring frontier capabilities in chat, coding, and reasoning. At the same time, we are introducing Grok-2 mini, a small but capable sibling of Grok-2. An early version of Grok-2 has been tested on the LMSYS leaderboard under the name ‘sus-column-r,’” xAI’s blog post about Grok-2 read.

Musk’s AI company is also planning to make both models available to developers through its enterprise API later this month.

Early images generated by users suggest that Grok’s image-generation feature doesn’t have any guardrails around creating images of political figures — and many users are currently taking

Latest in Apps

[See More](#)

APPS



The Browser Company launches a \$20 monthly subscription for its AI-powered browser

Ivan Mehta - 21 minutes ago


AI



Google takes on ChatGPT’s Study Mode with new ‘Guided Learning’ tool in Gemini

Aisha Malik - 12 hours ago

APPS



Instagram takes on Snapchat with new ‘Instagram Map’

Aisha Malik - 17 hours ago

advantage of this. However, with the U.S. presidential election around the corner, the company will likely face pressure to add limits on these features.

Grok 2.0 Ohh boyyyyy 😄😄😄

pic.twitter.com/TjzB7WMhVp

— Benjamin De Kraker 🇺🇸 (@BenjaminDEKR)

[August 14, 2024](#)

Grok 2.0 will do political illustrations and real people, while ChatGPT refuses.

This instantly makes Grok 10x more fun.....

pic.twitter.com/yDBJO0jWba

— Benjamin De Kraker 🇺🇸 (@BenjaminDEKR)

[August 14, 2024](#)

Grok AI Image Generation:

pic.twitter.com/mjzSMEJ5Df

— Jungle Inc: Deaton 4 Senate (@jungleincxrp)

[August 14, 2024](#)

The text below the sample image prompts suggests that Grok is using FLUX.1 by Black Forest Labs to create these images.

There are few details on Grok 2's actual capabilities at the moment. App researcher (and apparently X feature tester) Nima Owji claimed that Grok 2 is better at code generation, writing, and news, though there isn't much to support that. Notably, Grok's first version has slipped up [many times while summarizing news](#).

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy

hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

BREAKING: Here's an early look at Grok 2.0 features and abilities!

It's better at coding, writing, and generating news!
It'll also generate images using the FLUX.1 model!

pic.twitter.com/UIDW2Spen8

— Nima Owji (@nima_owji) [August 13, 2024](#)

Without any limits on image generation, Grok could easily become a tool to create misinformation on X and other platforms. It's also not clear if Grok-powered images embed any metadata to indicate that they are AI generated.

We have asked X about how it plans to limit image generation for harmful purposes and will update the story if we hear back. However, the company has seldom responded to media inquiries since Musk's takeover.

xAI noted that the company plans to deploy Grok-2 and Grok-2 mini in AI-driven features on X, including improved search capabilities, post analytics, and reply functions. This could mean that X might introduce AI-powered replies. The company also said that it will release a preview of multimodal

understanding as a core part of Grok’s experience on X and the API.

Topics: [AI](#) [Apps](#) [Elon Musk](#) [Grok](#) [Twitter](#) [xAI](#)



Ivan Mehta

[X](#) [in](#)

Ivan covers global consumer tech developments at TechCrunch. He is based out of India and has previously...

[View Bio](#) >

KEEP READING

AI



xAI and Grok apologize for ‘horrific behavior’

Anthony Ha — 10:13 AM PDT · July 12, 2025

In a series of posts on X, the AI chatbot Grok apologized for what it admitted was “horrific behavior.”

The posts appear to be an official statement from xAI, the Elon Musk-led company behind Grok, as opposed to an AI-generated explanation for Grok’s posts. (xAI [recently acquired X](#), where Grok is prominently featured.)

Grok’s latest controversy comes after Musk had indicated he wanted to make the chatbot less “politically correct,” then declared on July 4 that the company had “improved @Grok significantly.” In short order, the chatbot was making posts [criticizing Democrats and Hollywood’s “Jewish executives,” repeating antisemitic memes](#), and even [expressing support for Adolf Hitler](#) and referring to itself as “MechaHitler.”

IMAGE CREDITS:

JAKUB PORZYCKI / NURPHOTO / GETTY IMAGES

As a result, xAI deleted some of Grok’s posts, temporarily [took the chatbot offline](#), and updated its public system prompts.

Turkey also [banned the chatbot](#) for insulting the country’s president, and X CEO Linda Yaccarino even announced that she was [stepping down this week](#), although her announcement did not reference the latest Grok controversy and her departure was [reportedly months in the making](#).

So after all that, on Saturday, [xAI said](#), “First off, we deeply apologize for the horrific behavior that many experienced.” The company then blamed an “update to a code path upstream of the @grok bot,” which it emphasized was “independent of the underlying language model that powers @grok.”

DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government’s website
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Roku launches Howdy, a \$2.99 ad-free streaming service
- Spotify raises subscription prices
- The uproar over Vogue’s AI-generated ad isn’t just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines

This update supposedly made Grok “susceptible to existing X user posts; including when such posts contained extremist views.”

■ The best dating apps aren’t even dating apps

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$675 before prices rise on August 7.

San Francisco | October 27-29, 2025

REGISTER NOW

[xAI added](#) that an “unintended action” had led to Grok receiving instructions such as, “You tell like it is and you are not afraid to offend people who are politically correct.”

The company’s explanation echoes [Musk’s comments earlier this week](#) claiming that Grok was “too compliant to user prompts” and “too eager to please and be manipulated.”

xAI’s posts do not mention reporting by TechCrunch and others who examined the chain-of-thought summaries for the just-launched Grok 4, finding that the latest version of the chatbot [seems to consult Musk’s viewpoints and social media posts](#) before addressing controversial topics.

And historian Angus Johnston pushed back against the idea that Grok was simply manipulated into

posting offensive content. He [wrote on Bluesky](#) that xAI and Musk’s explanations are “easily falsified.”

“One of the most widely shared examples of Grok antisemitism was initiated by Grok with no previous bigoted posting in the thread — and with multiple users pushing back against Grok to no avail,” [Johnston said](#).

In recent months, Grok has also [posted repeatedly about “white genocide,” expressed skepticism about the death toll of the Holocaust](#), and [briefly censored unflattering facts about Musk](#) and his then-ally Donald Trump. In those cases, xAI blamed “unauthorized” changes and rogue employees.

Despite the controversy, Musk says [Grok is coming to Tesla vehicles](#) next week.

Topics: [AI](#) [Elon Musk](#) [Grok](#) [Social](#) [X](#) [xAI](#)



Anthony Ha

in

Anthony Ha is TechCrunch’s weekend editor. Previously, he worked as a tech reporter at Adweek, a senior editor at...

[View Bio](#) >

Loading the next article



[TechCrunch Staff](#)
[Contact Us](#)
[Advertise](#)

[Terms of Service](#)
[Privacy Policy](#)

[OpenAI](#)
[Perplexity](#)
[Cisco](#)
[DeepMind](#)



[Crunchboard](#)
[Jobs](#)

[RSS](#) [Terms of](#)
[Use](#)

[Windsurf](#)

[Tech Layoffs](#)

[ChatGPT](#)



[Site Map](#)