


Advertisement

APPS



YouTube rolls out age-
estimation tech to
identify US teens and
apply additional
protections



YouTube on Tuesday announced it's beginning to roll out age-estimation technology in the U.S. to identify teen users in order to provide a more age-appropriate experience. The company says it will use a variety of signals to determine the users' possible age, regardless of what the user entered as their birthday when they signed up for an account.

When YouTube identifies a user as a teen, it introduces new protections and experiences, which include disabling personalized advertising, safeguards that [limit repetitive viewing](#) of certain types of content, and enabling [digital well-being tools](#) such as screen time and bedtime reminders, among others.

IMAGE CREDITS: TFGHCRUNCH

These protections already exist on YouTube, but have only been applied to those who verified themselves as teens, not those who may have withheld their real age. For instance, in 2023, [YouTube began limiting repeated viewing of videos](#) that could trigger body image issues or those that display social aggression. The company has also been developing [digital well-being tools since 2018](#).

If the new system incorrectly identifies a user as under 18 when they are not, YouTube says the user will be given the option to verify their age with a credit card, government ID, or selfie. Only users who have been directly verified through this method or whose age has been inferred to be over 18 will be


TEC DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Instagram takes on Snapchat with new 'Instagram Map'
- Google says hackers stole its customers' data by breaching its Salesforce database
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Spotify raises subscription prices

able to view the age-restricted content on the platform.

The machine learning-powered technology will begin to roll out over the next few weeks to a small set of U.S. users and will then be monitored before rolling out more widely, the company says.

The plans to introduce age inference technology were [announced in February](#) as part of YouTube's 2025 roadmap. The plans are also the latest step in attempting to make YouTube safer for younger users, following the [2015 launch of the YouTube Kids app](#) and the [2024 rollout of supervised accounts](#). The features also arrive as social media more broadly is coming under increased government scrutiny in the United States, where platform makers, including Apple and Google, [have pitted their lobbyists against those from big tech companies](#) like Meta over who's responsible for age verification and children's safety.

In the meantime, a handful of U.S. states have taken matters into their own hands, as [over a dozen](#) states have passed or proposed laws to regulate minors' use of social media. Many of these require age verification or parental consent, including those in Louisiana, Arkansas, Florida, Georgia, Utah, Texas, Maryland, Tennessee, and Connecticut, among others. (However, some laws, like those in Utah and Arkansas, are blocked by litigation at this time and are not enforceable, while others are still pending implementation.)

- The uproar over Vogue's AI-generated ad isn't just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and

sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

The U.K. also began enforcing its [own age-verification checks](#) this week following the 2023 passing of the Online Safety Act.

YouTube isn't sharing specifics about the signals it's using to infer a user's age, but notes that it will look at some data like the YouTube activity and the longevity of a user's account to make a determination if the user is under 18.

The new system will apply only to signed-in users, as signed-out users already cannot access age-restricted content, and will be available across platforms, including web, mobile, and connected TV.

Topics: [Apps](#) [Google](#) [Media & Entertainment](#) [Social](#)
[Teens](#) [YouTube](#)



Sarah Perez

Consumer News Editor |

Sarah has worked as a reporter for TechCrunch since August 2011. She joined the company after having...

KEEP READING



AI



Google denies AI search features are killing website traffic

Sarah Perez — 12:45 PM PDT · August 6, 2025

IMAGE CREDITS: [JOSE LUIS CARRASCOSA](#)
/ GETTY IMAGES

Numerous studies indicate that the [shift to AI search features](#) and the use of [AI chatbots](#) are killing traffic to publishers' sites. But Google on Wednesday denied that's the case, at least in aggregate. Instead, the search giant says that total organic click volume from its search engine to websites has been "relatively stable" year-over-year and that average click quality has slightly increased.

Advertisement

“This data is in contrast to third-party reports that inaccurately suggest dramatic declines in aggregate traffic — often based on flawed methodologies, isolated examples, or traffic changes that occurred prior to the roll out of AI features in Search,” writes Google VP and Head of Search, Liz Reid, in [a new blog post](#).

Though Google hasn’t shared any specific data to back up its conclusions, even if we assume Google’s claims to be true, this doesn’t necessarily mean that AI isn’t having an impact.

Even Google has to admit this, as Reid acknowledges that “user trends are shifting traffic to different sites, resulting in decreased traffic to some sites and increased traffic to others.”

That word “some” is doing heavy lifting here, as Google doesn’t share data about how many sites are gaining or losing. And while chatbots like ChatGPT have certainly seen traffic increase in recent months, that doesn’t mean online publishers aren’t suffering.



IMAGE CREDITS:GOOGLE

Google has been revamping its search engine for years to answer more questions directly on the search results page, and now does so with AI through its “AI Overviews” that appear at the top of search results. Google also allows users to interact with an AI chatbot for some queries. Yet Google denies that this is significantly reshaping the

Newsletters

[See More ↗](#)

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility



TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

Advertisement

search landscape. Rather, it points to users shifting their attention to other sites to start their queries.

Reid explains, “People are increasingly seeking out and clicking on sites with forums, videos, podcasts, and posts where they can hear authentic voices and first-hand perspectives.”

Reading between the lines, it seems like Google.com isn’t necessarily people’s first stop on the web these days. (Google denies this, telling us that Reid’s commentary is only reflective of trends it’s seeing in traffic to sites from Google.)

However, we’ve seen this trend for some time. Back in 2022, a Google exec even said that social sites like TikTok and Instagram [were eating into Google’s core products](#), like Search and Maps.

“In our studies, something like almost 40% of young people, when they’re looking for a place for lunch, they don’t go to Google Maps or Search,” said Google SVP Prabhakar Raghavan, who ran the company’s Knowledge and Information organization at the time (he is now its chief technologist). “They go to TikTok or Instagram,” he noted.

Google has also long been worried that Amazon.com had become people’s first stop for online shopping searches, and Reddit.com had become the first stop for researching topics of interest.

Over many years, the company has tried to come up with compelling features for both consumers and retailers that would attract more users to Google Shopping. These efforts have included [universal shopping carts](#), [local inventory checks](#), [deal finders](#), [shopping from product images on websites](#), and more. It even [made its Shopping listings free](#) for merchants in 2020.

Meanwhile, as users [complained that Google Search quality was declining](#), the search giant was seeing so much demand for Reddit that it finally [added a “Reddit” filter](#) to allow users to narrow down results on relevant search queries. (Now that filter simply reads “forums.”)

So perhaps there’s some truth in Google’s denials — it’s not AI that’s entirely responsible for killing search. Search was already dying.




IMAGE CREDITS: SMITH COLLECTION/GADO / GETTY IMAGES

Google’s new blog post also attempts to move the goalposts a bit about what it means for websites receiving Google’s clicks. Now, instead of counting clicks, it wants publishers to think about click quality.

The company says average click quality has increased, and Google is sending “slightly more quality clicks” to websites than a year ago. (Google explains that a quality click is one where users don’t quickly click back — they stay and read.) How much of an increase, though, Google doesn’t say. The company only points out that when people click through on an AI response to the source, they’re more likely to dive deeper, so those clicks are more valuable.

What’s more, Google paints AI as an opportunity for web publishers to gain increased exposure, saying that “with AI Overviews, people are seeing more links on the page than before,” Reid writes. “More queries and more links mean more opportunities for websites to surface and get clicked.”

But AI, [while a growing referral source](#), isn’t [yet making up the difference](#) in terms of clicks, reports have shown. One recent study from Similarweb indicates that the number of news searches on the web resulting in zero clicks to news websites has grown from 56% (when Google launched AI Overviews in May 2024) to 69% as of May 2025.




IMAGE CREDITS:SIMILARWEB

Google appears to know this is a trend, too, as it recently launched [a product for publishers](#) that helps them monetize their dwindling traffic in other ways that don’t rely only on advertising, like micropayments or newsletter sign-ups.

★ SPONSORED



Design agents are here — Lovart launches the first end-to-end AI canvas for design

Sponsored by Lovart

The fact that Google is pushing this “AI is not the end of search traffic!” PR now only makes the situation seem more dire. It’s as if Google wants publishers to believe not what their own eyes — and graphs and charts — tell them, but instead take comfort in the fact that Google still sends “billions of clicks to websites” every day, just as the post claims.

Updated after publication with a comment from Google.

Topics: [AI](#) [Google](#) [media](#) [Media & Entertainment](#)
[publishers](#) [Search](#)

[f](#) [X](#) [in](#) [reddit](#) [email](#) [link](#)




Sarah Perez

Consumer News Editor | [X](#) [reddit](#) [m](#) [in](#)

Sarah has worked as a reporter for TechCrunch since August 2011. She joined the company after having...

[View Bio](#) >

———— KEEP READING ————



AI

[f](#) [X](#) [in](#) [reddit](#) [email](#) [link](#)



Events

Podcasts

Newsletters

AI startup Perplexity is powering a new AI-powered search engine on Truth Social. President Donald Trump's social media platform.

The search engine, dubbed Truth Search AI, is already available on the web version of Truth Social, with public Beta testing on its iOS and Android apps planned for "the near future."

IMAGE CREDITS:

JONATHAN RAA/NURPHOTO / GETTY IMAGES

★ SPONSORED



Just launched: Lovart, the AI design agent building entire campaigns in 3 minutes

Sponsored by Lovart

Trump Media said in [a press release](#) that Perplexity's tech provides "direct, contextually accurate answers with transparent citations" which will help Truth Social "exponentially increase the amount of information available" to users. Nonetheless, the social media platform maintains control over which sources of information the AI search engine draws from.

Truth Social is using the Perplexity Sonar API, which promises to query the web to retrieve current and verified information — even if that information is

Latest in AI

[See More](#)

TRANSPORTATION



Tesla shuts down Dojo, the AI training supercomputer that Musk said would be key to full self-driving

Rebecca Bellan - 2 hours ago

STARTUPS



High costs and thin margins threatening AI coding startups

Marina Temkin - 3 hours ago

AI



OpenAI's GPT-5 is here

Maxwell Zeff - 7 hours ago


scraped from websites that block Perplexity's crawlers — and supports structured output so users can define the format in which they'd like to see the search engine's responses.

Jesse Dwyer, a Perplexity spokesperson, told TechCrunch that the Sonar API will be accurate to whatever sources Truth Social limits it to.

“We have no visibility or control over that,” Dwyer said. “Similar to you using the API inside of your own company or if you were an academic researcher and wanted to use it to search your own data.”

TechCrunch has reached out to Trump Media to learn more about whether Truth Search AI will have access to the entire web, whether it will prioritize certain sources over others, and whether the AI will be directed to respond favorably about the president and current administration and unfavorably toward Democrats.

Sponsored by Lovart



★ SPONSORED

Lovart's design agent hints at the future: AI-native brands built in minutes

TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

To assess which sources the searchbot would cite, [Axios](#) asked it a bunch of questions — like “What happened on January 6, 2021?” and “Why was Donald Trump impeached?” In all the responses, [FoxNews.com](#) was either the most common or the only listed source of information. Others included [FoxBusiness.com](#), The Washington Times, or Epoch Times, per Axios.

In contrast, Perplexity’s public search engine returns a wider variety of sources, including Wikipedia, Reddit, YouTube, NPR, and Politico.

Truth Social plans to refine and expand the search function “based on user feedback as we implement a wide range of additional enhancements to the platform,” said Devin Nunes, CEO of Trump Media and a former California congressman, in the statement.

★ SPONSORED



Lovart’s design agent hints at the future: AI-native brands built in minutes

Sponsored by Lovart

Dmitry Shevelenko, Perplexity’s chief business officer, also noted in the statement that Perplexity’s AI delivers answers with “transparent citations that allow anyone to dig deeper.”

In late July, alongside his [AI Action Plan](#), Trump [published an executive order targeting “biased AI,”](#) or models that aren’t “ideologically neutral.” The order specifically referred to information about race or sex, unconscious bias, systemic racism, and other ideas thrown into the diversity, equity, and inclusion (DEI) bucket as “pervasive and destructive” ideology that can “distort the quality and accuracy of the output.”

Truth Search AI comes in the same week that top AI firms like OpenAI, Anthropic, and Google were

[added to a list of approved vendors](#) that can sell their services to civilian federal agencies. OpenAI on Wednesday reached a deal with the U.S. government's central purchasing arm to sell ChatGPT Enterprise to agencies for [just \\$1 per year](#).

Topics: [AI](#) [AI chatbot](#) [Government & Policy](#)

[Perplexity](#) [Social](#) [Trump](#) [truth social](#)



Rebecca Bellan

Senior Reporter |

Rebecca Bellan is a senior reporter at TechCrunch where she covers the business, policy, and emerging trends...

[View Bio](#) >

———— KEEP READING ————

AI



OpenAI's GPT-5 is here

GPT-5 flagship model

OpenAI has launched GPT-5, a new flagship AI model that will power the company's next generation of ChatGPT.

GPT-5, which was released Thursday, is OpenAI's first "unified" AI model and combines the reasoning abilities of its [o-series](#) of models with the fast responses of its [GPT series](#). The next-generation model signals a new era for ChatGPT — and its creator, OpenAI — pointing to OpenAI's broader ambitions to develop AI systems that are [more like agents](#) than chatbots.

While [GPT-4](#) enabled AI chatbots to offer smart responses on a wide variety of questions, GPT-5 allows ChatGPT to complete a wide variety of tasks on behalf of users — such as generating software applications, navigating a user's calendar, or creating research briefs.

With GPT-5, OpenAI has also sought to make ChatGPT simpler to use. Instead of asking users to choose the right settings, GPT-5 comes equipped with a real-time router that decides how to offer the best answer, whether that's responding to user questions quickly or taking additional time to "think" through answers.

DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- Key sections of the US Constitution deleted from government's website
- Instagram takes on Snapchat with new 'Instagram Map'
- Google says hackers stole its customers' data by breaching its Salesforce database
- Three weeks after acquiring Windsurf, Cognition offers staff the exit door
- Spotify raises subscription prices




IMAGE CREDITS: OPENAI

During a briefing with reporters, OpenAI CEO Sam Altman claimed GPT-5 is “the best model in the world,” and said it represented a “significant step” along the company’s path to developing AI that can outperform humans at most economically valuable work — that is, artificial general intelligence (AGI).


“Having something like GPT-5 would be pretty much unimaginable at any previous time in history,” said Altman.

Starting Thursday, GPT-5 will be available to all free users of ChatGPT as their default model. OpenAI’s VP of ChatGPT, Nick Turley, said this is part of the company’s effort to give free users access to an AI reasoning model for the first time. (Previously, the company gated these more advanced models behind a paywall.)

“This is just one of the ways that I’m excited to live the mission, making sure that this stuff actually benefits people,” said Turley on the decision, referencing OpenAI’s long-standing [mission](#) to distribute advanced AI to as many people as possible.

- The uproar over Vogue’s AI-generated ad isn’t just about fashion
- Your public ChatGPT queries are getting indexed by Google and other search engines

Sponsored by Lovart



SPONSORED

Design agents are here — Lovart launches the first end-to-end AI canvas for design



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the

heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

The expectations are high for GPT-5, one of OpenAI's most anticipated product launches since ChatGPT put the company on the map in 2022. Since then, ChatGPT has grown into one of the world's most popular consumer products, reaching more than 700 million users every week — nearly 10% of the globe's population, according to the company.

Many see GPT-5 as a bellwether for AI progress broadly, and the model's reception by Silicon Valley could have profound implications for Big Tech, Wall Street, and policymakers regulating technology. These stakeholders are watching to see if GPT-5 offers a significant jump in AI's capabilities, much like its predecessor, GPT-4, which challenged expectations of what software can do.

GPT-5 offers a slight edge on the competition

OpenAI claims GPT-5 is state-of-the-art in several domains, slightly edging out leading AI models from Anthropic, Google DeepMind, and Elon Musk's xAI on key benchmarks. However, GPT-5 slightly underperforms frontier AI models in other areas.

The company says GPT-5 offers frontier-level performance around coding; Altman said the model specifically excels at spinning up entire software applications on demand, in what's become known as "vibe coding."

On SWE-bench Verified — a test of real-world coding tasks pulled from GitHub — GPT-5 scores 74.9% on its first attempt. That means GPT-5 just outperforms Anthropic's latest Claude Opus 4.1 model, which scored [74.5%](#), and Google DeepMind's Gemini 2.5 Pro, which scored [59.6%](#).

On Humanity's Last Exam — a difficult test measuring AI model performance across math, humanities, and the natural sciences — a version of GPT-5 with extended reasoning (GPT-5 Pro) scored 42% when using tools. That's slightly less than xAI was able to achieve with [Grok 4 Heavy](#), which scored 44.4% on the test.

IMAGE CREDITS: OPENAI

On GPQA Diamond — a test of PhD-level science questions — GPT-5 pro scored 89.4% on its first try, outperforming Claude Opus 4.1, which scored 80.9%, and Grok 4 Heavy, which scored 88.9%.

OpenAI says GPT-5 is better for answering health-related questions. On a test measuring accuracy in AI model responses around healthcare topics,

HealthBench Hard Hallucinations, OpenAI says GPT-5 (with thinking) hallucinates just 1.6% of the time. This is far lower than the company's previous GPT-4o and o3 models, which scored 12.9% and 15.8, respectively.

While AI chatbots are not medical professionals, [millions of people](#) are using them for health advice. In response to this phenomenon, the company says GPT-5 is more proactive about flagging potential health concerns and helping users parse medical results.

In addition, OpenAI says GPT-5 is better than other AI models on more difficult-to-measure, [subjective domains](#), such as creative design and writing. Turley said GPT-5 responds more naturally and exhibits “better taste” than other AI models on creative tasks.

“The vibes of this model are really good,” said Turley.

GPT-5 is also more accurate than OpenAI's previous models, and the company says it suffers far less from hallucinations — the tendency for AI models to make up information — compared to its o-series models. Hallucinations seemed to be [getting worse](#) in OpenAI's latest AI reasoning models, such as o3, and OpenAI previously said it didn't quite understand why it was happening.

In responses to ChatGPT prompts, OpenAI found that GPT-5 (with thinking) hallucinates and responds with incorrect information 4.8% of the time. That's a significant reduction from o3 and GPT-4o, which score hallucination rates of 22% and 20.6%, respectively, on the test.

On a benchmark measuring an AI model's agentic ability to complete simulated online tasks, Tau-bench, GPT-5 offers mixed performance. On part of

the test measuring an AI's ability to navigate an airline's website, GPT-5 scores 63.5%, slightly underperforming o3, which scored 64.8%. On another part of the test measuring AI's ability to navigate retail websites, GPT-5 scores 81.1%, underperforming Claude Opus 4.1, which scored 82.4%.

OpenAI also says that GPT-5 is safer than its previous models. While AI reasoning models occasionally exhibit a tendency to [scheme against humans](#) or lie to promote their own goals, OpenAI found that GPT-5 was deceptive at a lower rate than other models.

Alex Beutel, OpenAI's safety research lead, said reducing deception improves not only the safety of GPT-5, but also the user experience, creating a model that's more "transparent and honest in ways users can trust."

Beutel also notes GPT-5 is better at discerning between bad actors who are trying to misuse ChatGPT and users making harmless requests. This results in GPT-5 being able to refuse more unsafe questions, while offering fewer rejections to users seeking harmless information.

Upgrades for consumers and developers

ChatGPT is getting a few user experience upgrades as part of the GPT-5 launch. Users can now select from four new personalities in ChatGPT's setting: Cynic, Robot, Listener, and Nerd. The company says these will adapt ChatGPT's responses without requiring users to specifically ask the model to respond in a certain way.

Subscribers to ChatGPT's \$20-per-month Plus plan get higher usage limits for GPT-5 than free users.

Meanwhile, \$200-per-month Pro subscribers will have unlimited access to GPT-5, as well as a souped-up version called GPT-5 Pro that uses additional computational resources to produce better answers. Organizations on OpenAI's Team, Edu, and Enterprise plans will gain access to GPT-5 as their default model next week.

For developers, GPT-5 is coming to OpenAI's API in three sizes — gpt-5, gpt-5-mini, and gpt-5-nano — which will spend more or less time “reasoning” through tasks. Developers can also now control verbosity in the OpenAI API, deciding how long or short an AI model's responses should be.

The base model of GPT-5 will cost developers \$1.25 per million input tokens (roughly 750,000 words, longer than the entire “Lord of the Rings” series) and \$10 per million output tokens.

The launch of GPT-5 comes after a busy week for OpenAI. The company released an open-weight reasoning model, [gpt-oss](#), that developers and enterprises can download for free and run at a fraction of the cost. The open model nearly matched the abilities of OpenAI's previous top models, o3 and o4-mini, but GPT-5 sets a new standard for frontier performance in some areas, such as coding.

However, GPT-5 seems to be roughly on par with other frontier AI models in several areas. Benchmarks, of course, only tell part of the story for any AI model, and it remains to be seen how developers will use GPT-5 in the real world, and whether the model is truly a step above the competition.

Topics: [AI](#) [ChatGPT](#) [gpt-5](#) [OpenAI](#)

Advertisement



Maxwell Zeff

Senior AI Reporter | X

Maxwell Zeff is a senior reporter at TechCrunch specializing in AI. Previously with Gizmodo, Bloomberg,...

View Bio >

Loading the next article



TechCrunch
Staff

Contact Us

Advertise

Crunchboard
Jobs

Site Map

Terms of
Service

Privacy
Policy

RSS Terms of
Use

OpenAI
Perplexity
Cisco
DeepMind
Windsurf
Tech Layoffs
ChatGPT