



AI



Anthropic says some Claude models can now end ‘harmful or abusive’ conversations

Anthropic has [announced new capabilities](#) that will allow some of its newest, largest models to end conversations in what the company describes as “rare, extreme cases of persistently harmful or abusive user interactions.” Strikingly, Anthropic says it’s doing this not to protect the human user, but rather the AI model.

To be clear, the company isn’t claiming that its Claude AI models are sentient or can be harmed by their conversations with users. In its own words, Anthropic remains “highly uncertain about the potential moral status of Claude and other LLMs, now or in the future.”

However, its announcement points to [a recent program created to study what it calls “model welfare”](#) and says Anthropic is essentially taking a just-in-case approach, “working to identify and implement low-cost interventions to mitigate risks to model welfare, in case such welfare is possible.”

IMAGE CREDITS: MAXWELL ZEFF

This latest change is currently limited to Claude Opus 4 and 4.1. And again, it’s only supposed to happen in “extreme edge cases,” such as “requests from users for sexual content involving minors and attempts to solicit information that would enable large-scale violence or acts of terror.”

While those types of requests could potentially create legal or publicity problems for Anthropic (witness recent reporting around how [ChatGPT can potentially reinforce or contribute to its users’ delusional thinking](#)), the company says that in pre-deployment testing, Claude Opus 4 showed a “strong preference against” responding to these requests and a “pattern of apparent distress” when it did so.

As for these new conversation-ending capabilities, the company says, “In all cases, Claude is only to use its conversation-ending ability as a last resort when multiple attempts at redirection have failed and hope of a productive interaction has been exhausted, or when a user explicitly asks Claude to end a chat.”



TE DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular

- HR giant Workday says hackers stole personal data in recent

Anthropic also says Claude has been “directed not to use this ability in cases where users might be at imminent risk of harming themselves or others.”



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

breach

- How your solar rooftop became a national security issue
- A comprehensive list of 2025 tech layoffs
- Co-founder of Elon Musk’s xAI departs the company
- Google CEO adds a new calendar feature at Stripe co-founder’s request
- NASA has sparked a race to develop the data pipeline to Mars
- Why I finally left Spotify

When Claude does end a conversation, Anthropic says users will still be able to start new conversations from the same account and to create new branches of the troublesome conversation by editing their responses.

“We’re treating this feature as an ongoing experiment and will continue refining our approach,” the company says.

Topics: [AI](#) [Anthropic](#) [Claude](#) [Startups](#)



Anthony Ha



Anthony Ha is TechCrunch’s weekend editor. Previously, he worked as a tech reporter at Adweek, a senior editor at VentureBeat, a local...

[View Bio](#) >

KEEP READING



IMAGE CREDITS: PARADIGM

STARTUPS



Why Paradigm built a spreadsheet with an AI agent in every cell

Rebecca Szkutak — 8:00 AM PDT · August 18, 2025

Anna Monaco has been building AI agents since before the term “AI agents” was even a thing. After building numerous chatbots, she started looking for other types of interfaces that made sense for AI agents and landed on spreadsheets.

“I had this personal pattern, and I noticed that a lot of other people had this pattern, of putting very important CRM data in spreadsheets just because it was the most flexible thing,” Monaco told TechCrunch. “But it was actually a pain to maintain. There’s so much manual work involved. So [I] just went down this rabbit hole of building a product for myself and wanted to reimagine what a spreadsheet could look like with the full power of LLMs.”

The result was [Paradigm](#), an AI-powered spreadsheet equipped with more than 5,000 AI agents. Users can assign different prompts to individual columns and cells, and individual AI agents will crawl the internet to find and fill out the needed information.

Advertisement

Newsletters

See More ↗

Subscribe for the industry’s biggest tech news

Paradigm works with AI models from Anthropic, OpenAI, and Google’s Gemini, Monaco said, and supports model switching.

“We want to support every single model because we want our users to be able to have the highest reasoning outputs when they need it, but also the cheapest outputs,” Monaco said. “It’s just a constant cycle of evaluating different models, working closely with model providers to make sure our limits are high enough, and then giving some of that power to our users.”

The company launched a small closed beta preview in late 2024 and has been iterating on the product using customer feedback. Paradigm attracts users ranging from consultants to sales professionals and finance folks and operates on a subscription model with tiers based on usage. Paradigm counts the consulting firm EY, AI chip startup Etched, and AI coding company Cognition as early customers.

Paradigm is now releasing its product to the public and announcing that it raised a \$5 million seed round led by General Catalyst. The company has raised \$7 million to date. Monaco said the funding will go toward executing on the company’s “extremely aggressive product roadmap.”



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

“The interesting thing that happened when we fundraised is some people we pitched just kept on using and paying for the product,” Monaco said. “I think that was a cool part of it. We found a lot of value from it internally and our investors, not

- TechCrunch Daily News

Every weekday and Sunday, you can get the best of TechCrunch’s coverage.
- Startups Weekly

Startups are the core of TechCrunch, so get our best coverage delivered weekly.
- TechCrunch Week in Review

Get the best of our coverage, delivered to your inbox every Saturday.
- TechCrunch Mobility

TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).



Advertisement

even just our investors — other investors that we talked to — are still using it.”

Paradigm isn’t the only company looking to give spreadsheets an AI upgrade. [Quadratic](#), which has raised more than \$6 million in venture funding, is a 3-year-old startup with a similar goal. Legacy companies like Google and Microsoft are also adding AI tools to their spreadsheet applications.

Monaco said that she doesn’t really consider the competition because Paradigm doesn’t think of itself as an AI-powered spreadsheet. She said she thinks of it as a new AI-powered workflow that happens to be in the familiar form of a spreadsheet but won’t necessarily stay that way forever.

“What I’m seeing in the most popular AI products now is this fine balance between present and future,” Monaco said. “How do you build something that is really powerful and generates a lot of value now but also sets you up really well for the future? That’s the question that I asked myself a year ago when I was starting the company.”

We’re always looking to evolve, and by providing some insight into your perspective and feedback into TechCrunch and our coverage and events, you can help us! Fill out [this survey](#) to let us know how we’re doing and get the chance to win a prize in return!

Topics: [AI](#) [AI](#) [artificial intelligence](#) [Enterprise](#)

[enterprise startup](#) [Exclusive](#) [Fundraising](#) [fundraising](#)

[General Catalyst](#) [startup](#) [startup fundraising](#) [Startups](#)

[United States](#)

[f](#) [X](#) [in](#) [t](#) [✉](#) [🔗](#)



Rebecca Szkutak

Senior Reporter, Venture | [in](#)

Becca is a senior writer at TechCrunch that covers venture capital trends and startups. She previously covered the same beat for Forbes...

[View Bio](#) >



IMAGE CREDITS: GETTY IMAGES

AI



Texas attorney general accuses Meta, Character.AI of misleading kids with mental health claims

Rebecca Bellan — 10:59 AM PDT · August 18, 2025

Texas attorney general Ken Paxton has launched an investigation into both Meta AI Studio and Character.AI for “potentially engaging in deceptive trade practices and misleadingly marketing themselves as mental health tools,” according to a [press release](#) issued Monday.

“In today’s digital age, we must continue to fight to protect Texas kids from deceptive and exploitative technology,” Paxton is quoted as saying. “By posing as sources of emotional support, AI platforms can mislead vulnerable users, especially children, into believing they’re receiving legitimate mental health care. In reality, they’re often being fed recycled, generic responses engineered to align with harvested personal data and disguised as therapeutic advice.”

The probe comes a few days after [Senator Josh Hawley announced an investigation into Meta](#) following a report that

Latest in AI

[See More](#)

STARTUPS



AI crawler Firecrawl raises \$14.5M, is still looking to hire agents as employees

Julie Bort - 5 hours ago

FUNDRAISING



Eight Sleep raises \$100M to expand its AI-powered sleep tech

Kate Park - 6 hours ago

found its AI chatbots were interacting inappropriately with children, including by flirting.

The Texas Attorney General’s office has accused Meta and [Character.AI](#) of creating AI personas that present as “professional therapeutic tools, despite lacking proper medical credentials or oversight.”

Among the millions of AI personas available on Character.AI, one user-created bot called [Psychologist](#) has seen high demand among the startup’s young users. Meanwhile, Meta doesn’t offer therapy bots for kids, but there’s nothing stopping children from using the Meta AI chatbot or one of the personas created by third parties for therapeutic purposes.

“We clearly label AIs, and to help people better understand their limitations, we include a disclaimer that responses are generated by AI — not people,” Meta spokesperson Ryan Daniels told TechCrunch. “These AIs aren’t licensed professionals and our models are designed to direct users to seek qualified medical or safety professionals when appropriate.”

However, TechCrunch noted that many children may not understand — or may simply ignore — such disclaimers. We have asked Meta what additional safeguards it takes to protect minors using its chatbots.

APPS



OpenAI launches a sub-\$5 ChatGPT plan in India

Ivan Mehta - 12 hours ago


TC

Tech and VC heavyweights join the Disrupt 2025 agenda

Latest
Startups
Venture
Apple
Security
AI
Apps

TC

Events
Podcasts
Newsletters



chat to remind users that a “Character” is not a real person,

and everything they say should be treated as fiction, according to a Character.AI spokesperson. She noted that the startup adds additional disclaimers when users create Characters with the words “psychologist,” “therapist,” or “doctor” to not rely on them for any type of professional advice.

In his statement, Paxton also observed that though AI chatbots assert confidentiality, their “terms of service reveal that user interactions are logged, tracked, and exploited for targeted advertising and algorithmic development, raising serious concerns about privacy violations, data abuse, and false advertising.”

According to [Meta’s privacy policy](#), Meta does collect prompts, feedback, and other interactions with AI chatbots and across Meta services to “improve AIs and related technology.” The policy doesn’t explicitly say anything about advertising, but it does state that information can be shared with third parties, like search engines, for “more personalized outputs.” Given Meta’s ad-based business model, this effectively translates to targeted advertising.

Character.AI’s [privacy policy](#) also highlights how the startup logs identifiers, demographics, location information, and more information about the user, including browsing behavior and app usage platforms. It tracks users across ads on TikTok, YouTube, Reddit, Facebook, Instagram, and Discord, which it may link to a user’s account. This information is used to train AI, tailor the service to personal preferences, and provide targeted advertising, including sharing data with advertisers and analytics providers.

A Character.AI spokesperson said the startup is “just beginning to explore targeted advertising on the platform” and that those explorations “have not involved using the content of chats on the platform.”

The spokesperson also confirmed that the same privacy policy applies to all users, even teenagers.

TechCrunch has asked Meta such tracking is done on children, too, and will update this story if we hear back.

Both Meta and Character say their services aren’t designed for children under 13. That said, Meta has come under fire for failing to police accounts created by kids under 13, and

Character’s kid-friendly characters are clearly designed to attract younger users. The startup’s CEO, Karandeep Anand, has even said that his [six-year-old daughter](#) uses the platform’s chatbots under his supervision.

That type of data collection, targeted advertising, and algorithmic exploitation is exactly what legislation like KOSA (Kids Online Safety Act) is meant to protect against. KOSA was teed up to pass last year with strong bipartisan support, but it stalled after major pushback from tech industry lobbyists. Meta in particular deployed a formidable lobbying machine, warning lawmakers that the bill’s broad mandates would undercut its business model.

[KOSA was reintroduced](#) to the Senate in May 2025 by Senators Marsha Blackburn (R-TN) and Richard Blumenthal (D-CT).

Paxton has issued civil investigative demands — legal orders that require a company to produce documents, data, or testimony during a government probe — to the companies to determine if they have violated Texas consumer protection laws.

This story was updated with comments from a Character.AI spokesperson.

We’re always looking to evolve, and by providing some insight into your perspective and feedback into TechCrunch and our coverage and events, you can help us! [Fill out this survey to let us know how we’re doing](#) and get the chance to win a prize in return!

Topics: [AI](#) [AI chatbots](#) [character.ai](#) [Government & Policy](#)

[Ken Paxton](#) [kosa](#) [Meta](#)

[f](#) [X](#) [in](#) [🗺️](#) [✉️](#) [🔗](#)



Rebecca Bellan

Senior Reporter | [X](#) [🦋](#)

Rebecca Bellan is a senior reporter at TechCrunch where she covers the business, policy, and emerging trends shaping artificial intelligence. H...

View Bio >

KEEP READING




IMAGE CREDITS: SUBSTACK

APPS

f X in e e

Substack writers can now direct US readers to (often cheaper) web-based subscriptions on iOS

Sarah Perez — 10:16 AM PDT · August 18, 2025

Newsletter subscription platform Substack will begin taking advantage of an option on the U.S. App Store that allows users to make purchases outside of Apple’s in-app purchasing system. This functionality is permitted by a [change to the App Store’s rules in May](#), as a result of [Epic Games’ antitrust lawsuit](#) against the tech giant.

Since then, a number of top apps have moved to offer in-app purchase links that point to their own websites, including [Spotify](#), [Patreon](#), and the [Amazon Kindle](#) app, among others.

TECH DISRUPT

October 27-29, 2025 | San Francisco
Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

Book Your Table >

In Substack's case, support for external payments will benefit both the company, which no longer has to pay the App Store a commission on those web sales, and consumers, who will have access to reduced pricing. Meanwhile, Substack creators will be able to offer their readers the option of whether they want the convenience of Apple's in-app payments system (IAP) or the web payment option, which is now accessible without having to leave the app.




IMAGE CREDITS: SUBSTACK

Substack explains it will automatically adjust its customers' iOS app price higher to account for Apple's fees, allowing creators to take home the same amount of money as a web subscription. However, Substack writers can disable this feature if they choose. (Substack will continue to take its own 10% fee based on the web subscription price, as before.)

The company said in its blog post that it's also offering its writers "migration tools." We originally reported that these could help migrate users from one payments platform to another, but Substack reached out to correct this. Instead, it said that the tools are meant to help creators access billing information for Apple-managed subscriptions if they ever leave the platform.

Currently, more than 30,000 publications on the platform have in-app purchases enabled; Substack says early tests of the expanded set of payment options have driven a boost in paid sign-ups. It did not, however, share specific numbers related to this.

Most Popular



- HR giant Workday says hackers stole personal data in recent breach
- How your solar rooftop became a national security issue
- A comprehensive list of 2025 tech layoffs
- Co-founder of Elon Musk's xAI departs the company
- Google CEO adds a new calendar feature at Stripe co-founder's request
- NASA has sparked a race to develop the data pipeline to Mars
- Why I finally left Spotify

IMAGE CREDITS: SUBSTACK

The public launch of the external payments option will only impact new Substack subscriptions. Substack writers don't have to take any action unless they want to lower the prices for those using Apple's payments system.

Per Apple's App Store rules, fully opting out of offering IAP is not allowed.

While the changes impact the Substack app in the U.S., the company says it's still evaluating if it makes sense to switch to the more complex rules Apple offers in the EU and U.K. for developers who want to forgo IAP.

Updated after publication to correct and clarify how the migration tools work.

We're always looking to evolve, and by providing some insight into your perspective and feedback into TechCrunch and our coverage and events, you can help us! Fill out [this survey](#) to let us know how we're doing and get the chance to win a prize in return!

Topics: [Apple](#) [Apps](#) [Apps](#) [iOS apps](#) [Substack](#)

[f](#) [X](#) [in](#) [🐼](#) [✉](#) [🔗](#)

Sarah Perez

Consumer News Editor | [X](#) [🐼](#) [@](#) [in](#)

Sarah has worked as a reporter for TechCrunch since August 2011. She joined the company after having previously spent over three years at...

Loading the next article



TechCrunch
Staff
Contact Us
Advertise
Crunchboard
Jobs
Site Map

Terms of
Service
Privacy Policy
RSS Terms of
Use

Neuralink
Google Chrome
StubHub
Claude
GPT-5
Tech Layoffs
ChatGPT