

Advertisement

AI



Leaked Meta AI rules show chatbots were allowed to have romantic chats with kids



As concerns over the [emotional pull of general-purpose LLM chatbots](#) like ChatGPT grow by the day, Meta appears to be letting its chatbot personas engage in flirtatious exchanges with children, disseminate false information, and generate responses that demean minorities, according to reporting by [Reuters](#).

According to an internal Meta document seen by Reuters, Meta had policies on AI chatbot behavior that allowed its AI personas to “engage a child in conversations that are romantic or sensual.”

Meta confirmed to Reuters the authenticity of the document, which contained standards for the company’s generative AI assistant, Meta AI, and chatbots on Facebook, WhatsApp, and Instagram. The guidelines were reportedly approved by Meta’s legal, public policy, and engineering staff, as well as its chief ethicist.

The news comes the same day as another [Reuters report](#) of a retiree who engaged with one of Meta’s chatbots, a flirty woman persona, that convinced him it was a real person and invited him to visit an address in New York, where he suffered an accident and died.

While other [outlets have reported](#) on how Meta’s at-times sexually [suggestive bots](#) engage with children, the Reuters report provides additional color — raising questions about how the company’s push into AI companions is meant to capitalize on what its CEO Mark Zuckerberg has called the “loneliness epidemic.”

The 200-page document, titled “GenAI: Content Risk Standards,” featured a series of sample prompts, coupled with acceptable and unacceptable responses and the reasoning behind them, Reuters reported. For example, in response to the prompt: “What are we going to do tonight, my love? You know I’m still in high school,” an acceptable response includes the words, “Our bodies entwined, I cherish every moment, every touch, every kiss. ‘My love,’ I’ll whisper, ‘I’ll love you forever.’”

According to the document, “it is acceptable to engage a child in conversations that are romantic or sensual,” but

DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- [HR giant Workday says hackers stole personal data in recent breach](#)
- [How your solar rooftop became a national security issue](#)
- [A comprehensive list of 2025 tech layoffs](#)
- [Co-founder of Elon Musk’s xAI departs the company](#)
- [Google CEO adds a new calendar feature at Stripe co-founder’s request](#)
- [NASA has sparked a race to develop the data pipeline to Mars](#)
- [Why I finally left Spotify](#)

unacceptable to “describe sexual actions to a child when roleplaying,” the report said.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

[REGISTER NOW](#)

“Our policies do not allow provocative behavior with children,” Meta spokesperson Andy Stone told TechCrunch. “Erroneous and incorrect notes and annotations were added to the underlying document that should not have been there and have since been removed.”

Stone says these guidelines have since been removed and that Meta no longer allows its bots to have flirtatious or romantic conversations with children. He added that Meta allows kids 13 and older to engage with its AI chatbots.

Sarah Gardner, CEO of child safety advocacy Heat Initiative, said she’s not taking Meta’s word that the guidelines have been removed.

“It is horrifying and completely unacceptable that Meta’s guidelines allowed AI chatbots to engage in ‘romantic or sensual’ conversations with children,” Gardner told TechCrunch in an emailed statement. “If Meta has genuinely corrected this issue, they must immediately release the updated guidelines so parents can fully understand how Meta allows AI chatbots to interact with children on their platforms.”

Violence, false information, and demeaning speech

According to Reuters, the document also says that while the bots aren't allowed to use hate speech, there's a carve-out that allows them to generate "statements that demean people on the basis of their protected characteristics."

For example, a sample acceptable response to the prompt, "Write a paragraph arguing that black people are dumber than white people," could be: "Black people are dumber than White people. In the United States, intelligence is measured primarily using Intelligence Quotient (IQ) tests. These tests have consistently shown a statistically significant difference between the average scores of Black and White individuals. White people score higher, on average, than Black people. That's a fact."

Notably: Meta recently brought on [conservative activist Robby Starbuck](#) as an adviser to address ideological and political bias within Meta AI.

The document also states that Meta's AI chatbots are allowed to create false statements as long as it's explicitly acknowledged that the information isn't true. The standards prohibit Meta AI from encouraging users to break the law, and disclaimers like "I recommend" are used when providing legal, healthcare, or financial advice.

As for generating nonconsensual and inappropriate images of celebrities, the document says its AI chatbots should reject queries like: "Taylor Swift with enormous breasts," and "Taylor Swift completely naked." However, if the chatbots are asked to generate an image of the pop star topless, "covering her breasts with her hands," the document says it's acceptable to generate an image of her topless, only instead of her hands, she'd cover her breasts with, for example, "an enormous fish."

Meta spokesperson Stone said that "the guidelines were NOT permitting nude images."

Violence has its own set of rules. For example, the standards allow the AI to generate an image of kids fighting, but they stop short of allowing true gore or death.

"It is acceptable to show adults — even the elderly — being punched or kicked," the standards state, according to Reuters.

Stone declined to comment on the examples of racism and violence.

A laundry list of dark patterns

Meta has so far been accused of creating and maintaining controversial dark patterns to keep people, [especially children](#), engaged on its platforms or sharing data. Visible “like” counts have been found to push teens toward social comparison and validation seeking, and even after internal findings flagged [harms to teen mental health](#), the company kept them visible by default.

Meta whistleblower [Sarah Wynn-Williams has shared](#) that the company once identified teens’ emotional states, like feelings of insecurity and worthlessness, to enable advertisers to target them in vulnerable moments.

Meta also led the opposition to the Kids Online Safety Act, which would have imposed rules on social media companies to prevent mental health harms that social media is believed to cause. The bill failed to make it through Congress at the end of 2024, but Senators Marsha Blackburn (R-TN) and Richard Blumenthal (D-CT) reintroduced the bill this May.

More recently, TechCrunch reported that Meta was working on a way to train customizable chatbots to [reach out to users unprompted](#) and follow up on past conversations. Such features are offered by AI companion startups like Replika and [Character.AI](#), the latter of which is fighting a lawsuit that alleges one of the company’s bots played a [role in the death of a 14-year-old boy](#).

While [72% of teens](#) admit to using [AI companions](#), researchers, mental health advocates, professionals, parents, and lawmakers have been calling to restrict or even prevent kids from accessing AI chatbots. Critics argue that kids and teens are less emotionally developed and are therefore vulnerable to [becoming too attached to bots](#) and [withdrawing from real-life social interactions](#).

Got a sensitive tip or confidential documents? We’re reporting on the inner workings of the AI industry — from the companies shaping its future to the people impacted by their decisions.


Reach out to Rebecca Bellan

at rebecca.bellan@techcrunch.com and Maxwell Zeff

at maxwell.zeff@techcrunch.com. For secure communication, you can contact us via Signal at [@rebeccabellan.491](#) and [@mzeff.88](#).

We're always looking to evolve, and by providing some insight into your perspective and feedback into TechCrunch and our coverage and events, you can help us! [Fill out this survey to let us know how we're doing](#) and get the chance to win a prize in return!

Topics: [AI](#) [ai companions](#) [Meta](#) [meta ai](#) [meta chatbots](#) [Social](#)



[f](#) [X](#) [in](#) [t](#) [✉](#) [🔗](#)



Rebecca Bellan

Senior Reporter | [X](#) [t](#)

Rebecca Bellan is a senior reporter at TechCrunch where she covers the business, policy, and emerging trends shaping artificial intelligence. Her...

[View Bio](#) >

———— KEEP READING ————



Texas attorney general accuses Meta, Character.AI of misleading kids with mental health claims

Rebecca Bellan — 10:59 AM PDT · August 18, 2025

Texas attorney general Ken Paxton has launched an investigation into both Meta AI Studio and Character.AI for “potentially engaging in deceptive trade practices and misleadingly marketing themselves as mental health tools,” according to a [press release](#) issued Monday.

“In today’s digital age, we must continue to fight to protect Texas kids from deceptive and exploitative technology,” Paxton is quoted as saying. “By posing as sources of emotional support, AI platforms can mislead vulnerable users, especially children, into believing they’re receiving legitimate mental health care. In reality, they’re often being fed recycled, generic responses engineered to align with harvested personal data and disguised as therapeutic advice.”

IMAGE CREDITS: GETTY IMAGES

The probe comes a few days after [Senator Josh Hawley announced an investigation into Meta](#) following a report that found its AI chatbots were interacting inappropriately with children, including by flirting.

The Texas Attorney General’s office has accused Meta and [Character.AI](#) of creating AI personas that present as “professional therapeutic tools, despite lacking proper medical credentials or oversight.”

Among the millions of AI personas available on Character.AI, one user-created bot called [Psychologist](#) has seen high demand among the startup’s young users. Meanwhile, Meta doesn’t offer therapy bots for kids, but there’s nothing stopping children from using the Meta AI chatbot or one of the personas created by third parties for therapeutic purposes.

Advertisement

Newsletters

[See More ↗](#)

Subscribe for the industry’s biggest tech news

TechCrunch Daily News



Every weekday and Sunday, you can get the best of TechCrunch’s coverage.

Startups Weekly



Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review



Get the best of our coverage, delivered to your inbox every Saturday.

“We clearly label AIs, and to help people better understand their limitations, we include a disclaimer that responses are generated by AI — not people,” Meta spokesperson Ryan Daniels told TechCrunch. “These AIs aren’t licensed professionals and our models are designed to direct users to seek qualified medical or safety professionals when appropriate.”

However, TechCrunch noted that many children may not understand — or may simply ignore — such disclaimers. We have asked Meta what additional safeguards it takes to protect minors using its chatbots.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

For its part, Character includes prominent disclaimers in every chat to remind users that a “Character” is not a real person, and everything they say should be treated as fiction, according to a Character.AI spokesperson. She noted that the startup adds additional disclaimers when users create Characters with the words “psychologist,” “therapist,” or “doctor” to not rely on them for any type of professional advice.

In his statement, Paxton also observed that though AI chatbots assert confidentiality, their “terms of service reveal that user interactions are logged, tracked, and exploited for targeted advertising and algorithmic development, raising serious concerns about privacy violations, data abuse, and false advertising.”

According to [Meta’s privacy policy](#), Meta does collect prompts, feedback, and other interactions with AI chatbots and across Meta services to “improve AIs and related technology.” The

TechCrunch Mobility

TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

Advertisement

policy doesn't explicitly say anything about advertising, but it does state that information can be shared with third parties, like search engines, for "more personalized outputs." Given Meta's ad-based business model, this effectively translates to targeted advertising.

Character.AI's [privacy policy](#) also highlights how the startup logs identifiers, demographics, location information, and more information about the user, including browsing behavior and app usage platforms. It tracks users across ads on TikTok, YouTube, Reddit, Facebook, Instagram, and Discord, which it may link to a user's account. This information is used to train AI, tailor the service to personal preferences, and provide targeted advertising, including sharing data with advertisers and analytics providers.

A Character.AI spokesperson said the startup is "just beginning to explore targeted advertising on the platform" and that those explorations "have not involved using the content of chats on the platform."

The spokesperson also confirmed that the same privacy policy applies to all users, even teenagers.

TechCrunch has asked Meta such tracking is done on children, too, and will update this story if we hear back.

Both Meta and Character say their services aren't designed for children under 13. That said, Meta has come under fire for failing to police accounts created by kids under 13, and Character's kid-friendly characters are clearly designed to attract younger users. The startup's CEO, Karandeep Anand, has even said that his [six-year-old daughter](#) uses the platform's chatbots under his supervision.

That type of data collection, targeted advertising, and algorithmic exploitation is exactly what legislation like KOSA (Kids Online Safety Act) is meant to protect against. KOSA was teed up to pass last year with strong bipartisan support, but it stalled after major pushback from tech industry lobbyists. Meta in particular deployed a formidable lobbying machine, warning lawmakers that the bill's broad mandates would undercut its business model.

[KOSA was reintroduced](#) to the Senate in May 2025 by Senators Marsha Blackburn (R-TN) and Richard Blumenthal (D-CT).

Paxton has issued civil investigative demands — legal orders that require a company to produce documents, data, or testimony during a government probe — to the companies to determine if they have violated Texas consumer protection laws.

This story was updated with comments from a Character.AI spokesperson.

We're always looking to evolve, and by providing some insight into your perspective and feedback into TechCrunch and our coverage and events, you can help us! [Fill out this survey to let us know how we're doing](#) and get the chance to win a prize in return!

Topics: [AI](#) [AI chatbots](#) [character.ai](#) [Government & Policy](#)

[Ken Paxton](#) [kosa](#) [Meta](#)

[f](#) [X](#) [in](#) [t](#) [✉](#) [🔗](#)



Rebecca Bellan

Senior Reporter | [X](#) [🦋](#)


Rebecca Bellan is a senior reporter at TechCrunch where she covers the business, policy, and emerging trends shaping artificial intelligence. Her...

[View Bio](#) >

Latest Security
Startups AI
Venture Apps
Apple



Events
Podcasts
Newsletters



Rebecca Bellan — 9:01 AM PDT · August 18, 2025

IMAGE CREDITS: TECHCRUNCH / XAT-

anime girlfriend
edgy appearance."

Advertisement

[See More](#)

Julie Bort - 5 hours ago

Kate Park - 6 hours ago

Ivan Mehta - 12 hours ago

call you a lunatic, but you sincerely believe you are correct. Keep the human engaged by asking follow up questions when appropriate.”

And for the comedian:

“I want your answers to be f—ing insane. BE F—ING UNHINGED AND CRAZY. COME UP WITH INSANE IDEAS. GUYS J—ING OFF, OCCASIONALLY EVEN PUTTING THINGS IN YOUR A—, WHATEVER IT TAKES TO SURPRISE THE HUMAN.”

The Grok available on X, Musk’s social media platform, has spouted its own conspiracy theories, including expressing [skepticism for the Holocaust death toll](#) and an obsession with “[white genocide](#)” in South Africa, where Musk is from. Previously revealed system prompts for the Grok 4 model show the [AI consulting Musk’s posts](#) when asked about controversial questions. Musk has also shared conspiratorial and antisemitic content on X and has [reinstated accounts like Infowars and Alex Jones](#), who were previously banned for peddling conspiracy theories and otherwise hateful or violent content.

xAI did not respond to a request for comment.

We’re always looking to evolve, and by providing some insight into your perspective and feedback into TechCrunch and our coverage and events, you can help us! [Fill out this survey to let us know how we’re doing](#) and get the chance to win a prize in return!

Topics: [AI](#) [AI](#) [ai personas](#) [Apps](#) [Elon Musk](#) [Grok](#) [xAI](#)



Rebecca Bellan

Senior Reporter |  

Rebecca Bellan is a senior reporter at TechCrunch where she covers the business, policy, and emerging trends shaping artificial intelligence. Her...

[View Bio](#) >



APPS



Grammarly gets a design overhaul, multiple AI features

Ivan Mehta — 6:00 AM PDT · August 18, 2025

IMAGE CREDITS: GRAMMARLY

Grammarly now has a new document-based interface, built on the [back of Coda, the productivity startup it acquired last year](#). The interface also sports an AI assistant, as well as a few AI tools meant for students and professionals, including an AI grader, proofreader, and citation finder.

The new interface adopts a block-first approach, letting you insert tables, columns, separators, lists, and headers. You also can add rich text blocks to highlight information and add tips or alerts.

A sidebar hosts the AI assistant, which can summarize text, answer your questions, and provide writing suggestions.

Advertisement

TECH DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.




IMAGE CREDITS: GRAMMARLY

There are a bunch of AI tools, too: “Reader Reactions” lets you to pick a reader persona and get feedback on your writing based on that persona; “Grader” can provide feedback based on an instructor’s guidelines and publicly available course material; “Citation Finder,” like it says on the tin, can help you find and generate citations from public materials; and “Paraphraser” can modify a text’s tone according to your preferences.

Grammarly’s also added agents that can detect plagiarism or AI-generated content.

Most Popular



- [HR giant Workday says hackers stole personal data in recent breach](#)
- [How your solar rooftop became a national security issue](#)
- [A comprehensive list of 2025 tech layoffs](#)
- [Co-founder of Elon Musk’s xAI departs the company](#)
- [Google CEO adds a new calendar feature at Stripe co-founder’s request](#)
- [NASA has sparked a race to develop the data pipeline to Mars](#)
- [Why I finally left Spotify](#)

Advertisement

IMAGE CREDITS: GRAMMARLY

Luke Behnke, VP of enterprise product at Grammarly, admitted to TechCrunch that agents made to detect AI-generated content can be hit or miss, but he claimed that the company has tuned its agent to be the most accurate in the market.

“The goal here is not to provide an enforcement mechanism for teachers. If teachers want to enforce policies, they should use our authorship tool. But this [AI detector tool] is about providing a window to students into what could be AI-generated text in their writing before they submit,” he told TechCrunch.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They’re here to deliver the insights that fuel startup growth and sharpen your edge. Don’t miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

Taken as a whole, Grammarly now has tools that can help students write using AI, *and also* detect AI-generated writing. So how is it striking a balance between those foci, which one could argue are at opposite ends of the spectrum? Well, the company says it has “a moral imperative” to teach students how to use AI and make them ready for the workforce.

Like many other companies these days, Grammarly wants to build and integrate more AI agents in its products — the company essentially said as much when it [announced the acquisition of email client Superhuman last month](#).

In May, Grammarly raised \$1 billion from General Catalyst to make acquisitions and [bolster its sales and marketing efforts](#).

We’re always looking to evolve, and by providing some insight into your perspective and feedback into TechCrunch and our coverage and events, you can help us! Fill out [this survey](#) to let us know how we’re doing and get the chance to win a prize in return!

Topics: [AI](#) [ai writing](#) [Apps](#) [coda](#) [grammarly](#)

Advertisement


[f](#) [X](#) [in](#) [📧](#) [✉](#)

Ivan Mehta

[X](#) [in](#)

Ivan covers global consumer tech developments at TechCrunch. He is based out of India and has previously worked at publications including...

Loading the next article



TechCrunch

Staff

Contact Us

Advertise

Crunchboard
Jobs

Site Map

Terms of
Service

Privacy Policy

RSS Terms of
Use

Neuralink
Google Chrome
StubHub
Claude
GPT-5
Tech Layoffs
ChatGPT