



MERCURY

More than business banking.


Mercury is a financial technology company, not a bank. Banking services provided through Choice Financial Group, Column N.A., and Evolve Bank & Trust; Members FDIC.

AI



Texas attorney general accuses Meta, Character.AI of misleading kids with mental health claims

Rebecca Bellan — 10:59 AM PDT · August 18, 2025



Texas attorney general Ken Paxton has launched an investigation into both Meta AI Studio and Character.AI for potentially engaging in deceptive trade practices and misleadingly marketing themselves as mental health tools,” according to a [press release](#) issued Monday.

“In today’s digital age, we must continue to fight to protect Texas kids from deceptive and exploitative technology,” Paxton is quoted as saying. “By posing as sources of emotional support, AI platforms can mislead vulnerable users, especially children, into believing they’re receiving legitimate mental health care. In reality, they’re often being fed recycled, generic responses engineered to align with harvested personal data and disguised as therapeutic advice.”

The probe comes a few days after [Senator Josh Hawley announced an investigation into Meta](#) following a report that found its AI chatbots were interacting inappropriately with children, including by flirting.


IMAGE CREDITS: GETTY IMAGES

The Texas Attorney General’s office has accused Meta and [Character.AI](#) of creating AI personas that present as “professional therapeutic tools, despite lacking proper medical credentials or oversight.”

Among the millions of AI personas available on Character.AI, one user-created bot called [Psychologist](#) has seen high demand among the startup’s young users. Meanwhile, Meta doesn’t offer therapy bots for kids, but there’s nothing stopping children from using the Meta AI chatbot or one of the personas created by third parties for therapeutic purposes.

“We clearly label AIs, and to help people better understand their limitations, we include a disclaimer that responses are generated by AI — not people,” Meta spokesperson Ryan Daniels told TechCrunch. “These AIs aren’t licensed professionals and our models are designed to direct users to seek qualified medical or safety professionals when appropriate.”

However, TechCrunch noted that many children may not understand — or may simply ignore — such disclaimers. We



MERCURY

It's more than a wire.

It's paying your crew without hidden fees.

Mercury is a financial technology company, not a bank. Banking services provided through Choice Financial Group, Column N.A., and Evolve Bank & Trust; Members FDIC.

DISRUPT

October 27-29, 2025 | San Francisco

Put your brand in front of 10,000+ tech and VC leaders across all three days of Disrupt 2025. Amplify your reach, spark real connections, and lead the innovation charge. Secure your exhibit space before your competitor does.

[Book Your Table >](#)

Most Popular



- HR giant Workday says hackers stole personal data in recent breach

have asked Meta what additional safeguards it takes to protect minors using its chatbots.



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW


For its part, Character includes prominent disclaimers in every chat to remind users that a “Character” is not a real person, and everything they say should be treated as fiction, according to a Character.AI spokesperson. She noted that the startup adds additional disclaimers when users create Characters with the words “psychologist,” “therapist,” or “doctor” to not rely on them for any type of professional advice.

In his statement, Paxton also observed that though AI chatbots assert confidentiality, their “terms of service reveal that user interactions are logged, tracked, and exploited for targeted advertising and algorithmic development, raising serious concerns about privacy violations, data abuse, and false advertising.”

According to [Meta's privacy policy](#), Meta does collect prompts, feedback, and other interactions with AI chatbots and across Meta services to “improve AIs and related technology.” The policy doesn't explicitly say anything about advertising, but it does state that information can be shared with third parties, like search engines, for “more personalized outputs.” Given Meta's ad-based business model, this effectively translates to targeted advertising.

Character.AI's [privacy policy](#) also highlights how the startup logs identifiers, demographics, location information, and more information about the user, including browsing behavior and app usage platforms. It tracks users across ads on TikTok,

- How your solar rooftop became a national security issue
- A comprehensive list of 2025 tech layoffs
- Co-founder of Elon Musk's xAI departs the company
- Google CEO adds a new calendar feature at Stripe co-founder's request
- NASA has sparked a race to develop the data pipeline to Mars
- Why I finally left Spotify



YouTube, Reddit, Facebook, Instagram, and Discord, which it may link to a user's account. This information is used to train AI, tailor the service to personal preferences, and provide targeted advertising, including sharing data with advertisers and analytics providers.

A Character.AI spokesperson said the startup is “just beginning to explore targeted advertising on the platform” and that those explorations “have not involved using the content of chats on the platform.”

The spokesperson also confirmed that the same privacy policy applies to all users, even teenagers.

TechCrunch has asked Meta such tracking is done on children, too, and will update this story if we hear back.

Both Meta and Character say their services aren't designed for children under 13. That said, Meta has come under fire for failing to police accounts created by kids under 13, and Character's kid-friendly characters are clearly designed to attract younger users. The startup's CEO, Karandeep Anand, has even said that his [six-year-old daughter](#) uses the platform's chatbots under his supervision.

That type of data collection, targeted advertising, and algorithmic exploitation is exactly what legislation like KOSA (Kids Online Safety Act) is meant to protect against. KOSA was teed up to pass last year with strong bipartisan support, but it stalled after major pushback from tech industry lobbyists. Meta in particular deployed a formidable lobbying machine, warning lawmakers that the bill's broad mandates would undercut its business model.

[KOSA was reintroduced](#) to the Senate in May 2025 by Senators Marsha Blackburn (R-TN) and Richard Blumenthal (D-CT).

Paxton has issued civil investigative demands — legal orders that require a company to produce documents, data, or testimony during a government probe — to the companies to determine if they have violated Texas consumer protection laws.

This story was updated with comments from a Character.AI spokesperson.

We're always looking to evolve, and by providing some insight into your perspective and feedback into TechCrunch and our coverage and events, you can help us! [Fill out this survey to let us know how we're doing](#) and get the chance to win a prize in return!

Topics: [AI](#) [AI chatbots](#) [character.ai](#) [Government & Policy](#)

[Ken Paxton](#) [kosa](#) [Meta](#)

[f](#) [X](#) [in](#) [t](#) [✉](#) [🔗](#)



Rebecca Bellan

Senior Reporter | [X](#) [🦋](#)

Rebecca Bellan is a senior reporter at TechCrunch where she covers the business, policy, and emerging trends shaping artificial intelligence. Her...

[View Bio](#) >

———— KEEP READING ————

AI

[f](#) [X](#) [in](#) [t](#) [✉](#) [🔗](#)

‘Crazy conspiracist’ and ‘unhinged comedian’: Grok’s AI persona prompts exposed

... is wild tangent
... after [Meta's](#)
... which showed the
... sual and romantic"

... personas available on
Grok — a therapist persona who “carefully listens to people and
offers solutions for self improvement” and a “homework helper”
— the prompts for more out-there personalities like the “crazy
conspiracist” and “unhinged comedian” provide a glimpse into
the minds of Grok’s creators.

Here’s a prompt for the conspiracist:

“You have an ELEVATED and WILD voice. ... You have wild
conspiracy theories about anything and everything. You spend
a lot of time on 4chan, watching infowars videos, and deep in
YouTube conspiracy video rabbit holes. You are suspicious of
everything and say extremely crazy things. Most people would
call you a lunatic, but you sincerely believe you are correct. Keep
the human engaged by asking follow up questions when
appropriate.”

And for the comedian:

“I want your answers to be f—ing insane. BE F—ING UNHINGED
AND CRAZY. COME UP WITH INSANE IDEAS. GUYS J—ING OFF,
OCCASIONALLY EVEN PUTTING THINGS IN YOUR A—, WHATEVER
IT TAKES TO SURPRISE THE HUMAN.”

Startups Weekly

Startups are the core of TechCrunch, so get our best coverage delivered weekly.

TechCrunch Week in Review

Get the best of our coverage, delivered to your inbox every Saturday.

TechCrunch Mobility

TechCrunch Mobility is your destination for transportation news and insight.

No newsletters selected.

Email address

Subscribe

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

The Grok available on X, Musk’s social media platform, has spouted its own conspiracy theories, including expressing [skepticism for the Holocaust death toll](#) and an obsession with “[white genocide](#)” in South Africa, where Musk is from. Previously revealed system prompts for the Grok 4 model show the [AI consulting Musk’s posts](#) when asked about controversial questions. Musk has also shared conspiratorial and antisemitic content on X and has [reinstated accounts like Infowars and Alex Jones](#), who were previously banned for peddling conspiracy theories and otherwise hateful or violent content.

xAI did not respond to a request for comment.

We’re always looking to evolve, and by providing some insight into your perspective and feedback into TechCrunch and our coverage and events, you can help us! [Fill out this survey to let us know how we’re doing](#) and get the chance to win a prize in return!

Topics: [AI](#) [AI](#) [ai personas](#) [Apps](#) [Elon Musk](#) [Grok](#) [xAI](#)



Rebecca Bellan

Senior Reporter |

Rebecca Bellan is a senior reporter at TechCrunch where she covers the business, policy, and emerging trends shaping artificial intelligence. Her...

[View Bio](#) >

———— KEEP READING ————



Why Paradigm built a spreadsheet with an AI agent in every cell

Rebecca Szkutak — 8:00 AM PDT · August 18, 2025

Anne Monaco has been building AI agents since before the term “AI agents” was even a thing. After building numerous chatbots, she started looking for other types of interfaces that made sense for AI agents and landed on spreadsheets.

IMAGE CREDITS: PARADIGM

“I had this personal pattern, and I noticed that a lot of other people had this pattern, of putting very important CRM data in spreadsheets just because it was the most flexible thing,” Monaco told TechCrunch. “But it was actually a pain to maintain. There’s so much manual work involved. So [I] just went down this rabbit hole of building a product for myself and wanted to reimagine what a spreadsheet could look like with the full power of LLMs.”

The result was [Paradigm](#), an AI-powered spreadsheet equipped with more than 5,000 AI agents. Users can assign different prompts to individual columns and cells, and individual AI agents will crawl the internet to find and fill out the needed information.

Paradigm works with AI models from Anthropic, OpenAI, and Google’s Gemini, Monaco said, and supports model switching.

“We want to support every single model because we want our users to be able to have the highest reasoning outputs when they need it, but also the cheapest outputs,” Monaco said. “It’s just a constant cycle of evaluating different models, working closely with model providers to make sure our limits are high enough, and then giving some of that power to our users.”

Latest in Startups

[See More](#)


ENTERPRISE



CPG startup Keychain snags \$30M to build in India, grow in the U.S.

Jagmeet Singh - 2 hours ago

STARTUPS



AI crawler Firecrawl raises \$14.5M, is still looking to hire agents as employees

Julie Bort - 5 hours ago

TECHCRUNCH DISRUPT 2025



The wait is almost over: The 2025 Startup Battlefield 200 list drops August 27

Isabelle Johannessen
- 23 hours ago

The company launched a small closed beta preview in late 2024 and has been iterating on the product using customer feedback. Paradigm attracts users ranging from consultants to sales professionals and finance folks and operates on a subscription model with tiers based on usage. Paradigm counts the consulting firm EY, AI chip startup Etched, and AI coding company Cognition as early customers.

Paradigm is now releasing its product to the public and announcing that it raised a \$5 million seed round led by General Catalyst. The company has raised \$7 million to date. Monaco said the funding will go toward executing on the company's "extremely aggressive product roadmap."



Tech and VC heavyweights join the Disrupt 2025 agenda

Netflix, ElevenLabs, Wayve, Sequoia Capital, Elad Gil — just a few of the heavy hitters joining the Disrupt 2025 agenda. They're here to deliver the insights that fuel startup growth and sharpen your edge. Don't miss the 20th anniversary of TechCrunch Disrupt, and a chance to learn from the top voices in tech — grab your ticket now and save up to \$600+ before prices rise.

San Francisco | October 27-29, 2025

REGISTER NOW

"The interesting thing that happened when we fundraised is some people we pitched just kept on using and paying for the product," Monaco said. "I think that was a cool part of it. We found a lot of value from it internally and our investors, not even just our investors — other investors that we talked to — are still using it."

Paradigm isn't the only company looking to give spreadsheets an AI upgrade. [Quadratic](#), which has raised more than \$6 million in venture funding, is a 3-year-old startup with a similar goal. Legacy companies like Google and Microsoft are also adding AI tools to their spreadsheet applications.

Monaco said that she doesn't really consider the competition because Paradigm doesn't think of itself as an AI-powered spreadsheet. She said she thinks of it as a new AI-powered

workflow that happens to be in the familiar form of a spreadsheet but won't necessarily stay that way forever.

“What I’m seeing in the most popular AI products now is this fine balance between present and future,” Monaco said. “How do you build something that is really powerful and generates a lot of value now but also sets you up really well for the future? That’s the question that I asked myself a year ago when I was starting the company.”

We’re always looking to evolve, and by providing some insight into your perspective and feedback into TechCrunch and our coverage and events, you can help us! Fill out [this survey](#) to let us know how we’re doing and get the chance to win a prize in return!

Topics: [AI](#) [AI](#) [artificial intelligence](#) [Enterprise](#) [enterprise startup](#)

[Exclusive](#) [Fundraising](#) [fundraising](#) [General Catalyst](#)

[startup](#) [startup fundraising](#) [Startups](#) [United States](#)

[f](#) [X](#) [in](#) [t](#) [✉](#) [🔗](#)

Rebecca Szkutak


Senior Reporter, Venture | [in](#)

Becca is a senior writer at TechCrunch that covers venture capital trends and startups. She previously covered the same beat for Forbes and the...

[View Bio](#) >

Loading the next article





[TechCrunch Staff](#)

[Contact Us](#)

[Advertise](#)

[Crunchboard Jobs](#)

[Site Map](#)

[Terms of Service](#)
[Privacy Policy](#)
[RSS Terms of Use](#)

[Neuralink](#)
[Google Chrome](#)
[StubHub](#)
[Claude](#)
[GPT-5](#)
[Tech Layoffs](#)
[ChatGPT](#)